

The Cylc Suite Engine User Guide

7.6.0

Released Under the GNU GPL v3.0 Software License

Copyright (C) 2008-2018 NIWA

Hilary Oliver

February 7, 2018



Abstract

Cylc (“silk”) is a metascheduler¹ for cycling environmental forecasting suites containing many forecast models and associated processing tasks. *Cylc* has a novel self-organising scheduling algorithm: a pool of task proxy objects, that each know just their own inputs and outputs, negotiate dependencies so that correct scheduling emerges naturally at run time. *Cylc* does not group tasks artificially by forecast cycle² (each task has a private cycle time and is self-spawning - there is no suite-wide cycle time) and handles dependencies within and between cycles equally so that tasks from multiple cycles can run at once to the maximum possible extent. This matters in particular whenever the external driving data³ for upcoming cycles are available in advance: *cylc* suites can catch up from delays very quickly, parallel test suites can be started behind the main operation to catch up quickly, and one can likewise achieve greater throughput in historical case studies; the usual sequence of distinct forecast cycles emerges naturally if a suite catches up to real time operation. *Cylc* can easily use existing tasks and can run suites distributed across a heterogeneous network. Suites can be stopped and restarted in any state of operation, and they dynamically adapt to insertion and removal of tasks, and to delays or failures in particular tasks or in the external environment: tasks not directly affected will carry on cycling as normal while the problem is addressed, and then the affected tasks will catch up as quickly as possible. *Cylc* has comprehensive command line and graphical interfaces, including a dependency graph based suite control GUI. Other notable features include suite databases; a fast simulation mode; a structured, validated suite definition file format; dependency graph plotting; task event hooks for centralized alerting; and cryptographic suite security.

¹A metascheduler determines when dependent jobs are *ready to run* and then submits them to run by other means, usually a batch queue scheduler. The term can also refer to an aggregate view of multiple distributed resource managers, but that is not the topic of this document. We drop the “meta” prefix from here on because a metascheduler is also a type of scheduler.

²A *forecast cycle* comprises all tasks with a common *cycle time* (later referred to here as *cycle point*) i.e. the analysis time or nominal start time of a forecast model, or that of the associated forecast model(s) for other tasks.

³Forecast suites are typically driven by real time observational data or timely model fields from an external forecasting system.

Contents

1	Introduction: How Cylc Works	16
1.1	Scheduling Forecast Suites	16
1.2	EcoConnect	16
1.3	Dependence Between Tasks	16
1.3.1	Intra-cycle Dependence	16
1.3.2	Inter-Cycle Dependence	19
1.4	The Cylc Scheduling Algorithm	21
2	Cylc Screenshots	22
3	Installation	25
3.1	External Software Packages	25
3.2	Software Bundled With Cylc	26
3.3	Installing Cylc	26
3.3.1	Local User Installation	26
3.3.2	Create A Site Config File	27
3.3.3	Configure Site Environment on Job Hosts	27
3.4	Automated Tests	27
4	Cylc Terminology	27
4.1	Jobs and Tasks	27
4.2	Cycle Points	27
5	Workflows For Cycling Systems	28
5.1	Cycling Workflows	28
5.2	Parameterized Tasks as a Proxy for Cycling	28
5.3	Mixed Cycling Workflows	29
6	Global (Site, User) Configuration Files	29
7	Tutorial	29
7.1	User Config File	29
7.1.1	Configure Environment on Job Hosts	30
7.2	User Interfaces	30
7.2.1	Command Line Interface (CLI)	30
7.2.2	Graphical User Interface (GUI)	30
7.3	Suite Definitions	30
7.4	Suite Registration	30
7.5	Suite Passphrases	31
7.6	Import The Example Suites	31
7.7	Rename The Imported Tutorial Suites	31
7.8	Suite Validation	31
7.9	Hello World in Cylc	32
7.10	Editing Suites	32
7.11	Running Suites	33
7.11.1	CLI	33
7.11.2	GUI	33
7.12	Discovering Running Suites	34

7.13	Task Identifiers	34
7.14	Job Submission: How Tasks Are Executed	34
7.15	Locating Suite And Task Output	35
7.16	Remote Tasks	36
7.17	Task Triggering	37
7.17.1	Task Failure And Suicide Triggering	38
7.18	Runtime Inheritance	39
7.19	Triggering Families	39
7.20	Triggering Off Of Families	40
7.21	Suite Visualization	41
7.22	External Task Scripts	41
7.23	Cycling Tasks	43
7.23.1	ISO 8601 Date-Time Syntax	44
7.23.2	Inter-Cycle Triggers	44
7.23.3	Initial Non-Repeating (R1) Tasks	45
7.23.4	Integer Cycling	46
7.24	Jinja2	47
7.25	Task Retry On Failure	49
7.26	Other Users' Suites	49
7.27	Other Things To Try	50
8	Suite Name Registration	50
9	Suite Definition	50
9.1	Suite Definition Directories	51
9.2	Suite.rc File Overview	51
9.2.1	Syntax	51
9.2.2	Include-Files	52
9.2.2.1	Editing Temporarily Inlined Suites	52
9.2.2.2	Include-Files via Jinja2	52
9.2.3	Syntax Highlighting For Suite Definitions	52
9.2.4	Gross File Structure	53
9.2.5	Validation	53
9.3	Scheduling - Dependency Graphs	53
9.3.1	Graph String Syntax	54
9.3.2	Interpreting Graph Strings	54
9.3.2.1	Splitting Up Long Graph Lines	56
9.3.3	Graph Types	56
9.3.3.1	One-off (Non-Cycling)	56
9.3.3.2	Cycling Graphs	56
9.3.4	Graph Section Headings	56
9.3.4.1	Syntax Rules	57
9.3.4.2	Referencing The Initial And Final Cycle Points	58
9.3.4.3	Excluding Dates	58
9.3.4.4	Advanced exclusion syntax	59
9.3.4.5	How Multiple Graph Strings Combine	59
9.3.4.6	Advanced Examples	60
9.3.4.7	Advanced Starting Up	60
9.3.4.8	Integer Cycling	63

	9.3.4.8.1	Example	63
	9.3.4.8.2	Advanced Integer Cycling Syntax	65
9.3.5		Trigger Types	65
	9.3.5.1	Success Triggers	65
	9.3.5.2	Failure Triggers	65
	9.3.5.3	Start Triggers	66
	9.3.5.4	Finish Triggers	66
	9.3.5.5	Message Triggers	66
	9.3.5.6	Job Submission Triggers	67
	9.3.5.7	Conditional Triggers	67
	9.3.5.8	Suicide Triggers	67
	9.3.5.9	Family Triggers	69
	9.3.5.10	Writing Efficient Inter-Family Triggering	70
	9.3.5.11	Inter-Cycle Triggers	71
	9.3.5.12	Special Sequential Tasks	72
	9.3.5.13	Future Triggers	73
	9.3.5.14	Clock Triggers	73
	9.3.5.15	Clock-Expire Triggers	74
	9.3.5.16	External Triggers	74
9.3.6		Model Restart Dependencies	75
9.3.7		How The Graph Determines Task Instantiation	76
9.4		Runtime - Task Configuration	77
	9.4.1	Namespace Names	78
	9.4.2	Root - Runtime Defaults	78
	9.4.3	Defining Multiple Namespaces At Once	78
	9.4.4	Runtime Inheritance - Single	78
	9.4.5	Runtime Inheritance - Multiple	79
	9.4.5.1	Suite Visualization And Multiple Inheritance	80
	9.4.6	How Runtime Inheritance Works	81
	9.4.7	Task Execution Environment	81
	9.4.7.1	User Environment Variables	81
	9.4.7.2	Overriding Environment Variables	81
	9.4.7.3	Task Job Script Variables	82
	9.4.7.4	Suite Share Directories	83
	9.4.7.5	Task Work Directories	83
	9.4.7.6	Environment Variable Evaluation	83
	9.4.8	How Tasks Get Access To The Suite Directory	83
	9.4.9	Remote Task Hosting	84
	9.4.9.1	Dynamic Host Selection	84
	9.4.9.2	Remote Task Log Directories	84
9.5		Visualization	85
	9.5.1	Collapsible Families In Suite Graphs	85
9.6		Parameterized Tasks	85
	9.6.1	Parameter Expansion	86
	9.6.1.1	Zero-Padded Integer Values	87
	9.6.1.2	Parameters as Full Task Names	88
	9.6.2	Passing Parameter Values To Tasks	88
	9.6.3	Selecting Specific Parameter Values	89
	9.6.4	Selecting Partial Parameter Ranges	89

9.6.5	Parameter Offsets In The Graph	90
9.6.6	Task Families And Parameterization	90
9.6.7	Parameterized Cycling	91
9.6.7.1	Cycle Point And Parameter Offsets At Start-Up	93
9.7	Jinja2	93
9.7.1	Accessing Environment Variables With Jinja2	95
9.7.2	Custom Jinja2 Filters	96
9.7.2.1	pad	96
9.7.2.2	strftime	97
9.7.3	Associative Arrays In Jinja2	97
9.7.4	Jinja2 Default Values And Template Inputs	97
9.7.5	Jinja2 Variable Scope	99
9.7.6	Raising Exceptions	99
9.7.6.1	Raise	99
9.7.6.2	Assert	99
9.8	Omitting Tasks At Runtime	99
9.9	Naked Dummy Tasks And Strict Validation	100
10	Task Implementation	100
10.1	Task Job Scripts	100
10.2	Inlined Tasks	101
10.3	Task Messages	101
10.4	Aborting Job Scripts on Error	101
10.4.1	Custom Failure Messages	102
10.5	Avoid Detaching Processes	102
11	Task Job Submission and Management	102
11.1	Supported Job Submission Methods	103
11.1.1	background	103
11.1.2	at	103
11.1.3	loadleveler	103
11.1.4	lsf	104
11.1.5	pbs	104
11.1.6	moab	105
11.1.7	sge	105
11.1.8	slurm	106
11.1.9	Default Directives Provided	106
11.1.10	Directives Section Quirks (PBS, SGE, ...)	106
11.2	Task stdout And stderr Logs	107
11.3	Overriding The Job Submission Command	107
11.4	Job Polling	108
11.5	Job Killing	109
11.6	Execution Time Limit	109
11.6.1	Execution Time Limit and Execution Timeout	109
11.7	Custom Job Submission Methods	109
11.7.1	An Example	109
11.7.2	Where To Put Batch System Handler Modules	110
12	Running Suites	110
12.1	Suite Start-Up	111

12.1.1	Cold Start	111
12.1.2	Warm Start	111
12.1.3	Restart and Suite State Checkpoints	111
12.1.3.1	Restart From Latest Checkpoint	112
12.1.3.2	Restart From Another Checkpoint	112
12.1.3.3	Checkpointing With A Task	113
12.1.3.4	Behaviour of Tasks on Restart	113
12.2	Reloading The Suite Definition At Runtime	113
12.3	Task Job Access To Cylc	114
12.4	The Suite Contact File	114
12.5	Task Job Polling	114
12.5.1	Routine Polling	114
12.6	Tracking Task State	115
12.6.1	HTTPS Task Messaging	115
12.6.2	Ssh Task Messaging	115
12.6.3	Polling to Track Job Status	116
12.6.4	Task Communications Configuration	116
12.7	The Suite Service Directory	117
12.8	File-Reading Commands	117
12.8.1	Remote Host, Shared Home Directory	117
12.8.2	Remote Host, Different Home Directory	118
12.8.3	Same Host, Different User Account	118
12.9	Client-Server Interaction	118
12.9.1	Public Access - No Auth Files	118
12.9.2	Full Control - With Auth Files	118
12.10	GUI-to-Suite Interaction	119
12.11	Remote Control	119
12.12	Scan And Gscan	120
12.13	Task States Explained	120
12.14	What The Suite Control GUI Shows	121
12.15	Network Connection Timeouts	121
12.16	Runahead Limiting	121
12.17	Limiting Activity With Internal Queues	121
12.18	Automatic Task Retry On Failure	122
12.19	Task Event Handling	123
12.20	Handling Job Preemption	124
12.21	Manual Task Triggering and Edit-Run	125
12.22	Cylc Broadcast	125
12.23	The Meaning And Use Of Initial Cycle Point	125
12.23.1	CYLC_SUITE_INITIAL_CYCLE_POINT	126
12.24	Simulating Suite Behaviour	126
12.24.1	Proportional Simulated Run Length	126
12.24.2	Limitations Of Suite Simulation	126
12.24.3	Restarting Suites With A Different Run Mode?	127
12.25	Automated Reference Test Suites	127
12.25.1	Roll-your-own Reference Tests	127
12.26	Triggering Off Of Tasks In Other Suites	128
12.27	Suite Server Logs	129
12.28	Suite Run Databases	130

12.29 Disaster Recovery	130
13 Suite Storage, Discovery, Revision Control, and Deployment	131
13.1 Rose	132
A Suite.rc Reference	133
A.1 Top Level Items	133
A.2 [meta]	133
A.2.1 title	133
A.2.2 description	133
A.2.3 URL	133
A.2.4 group	134
A.2.5 __MANY__	134
A.3 [cylc]	134
A.3.1 required run mode	134
A.3.2 UTC mode	134
A.3.3 cycle point format	135
A.3.4 cycle point num expanded year digits	135
A.3.5 cycle point time zone	135
A.3.6 abort if any task fails	136
A.3.7 health check interval	136
A.3.8 task event mail interval	136
A.3.9 disable automatic shutdown	136
A.3.10 log resolved dependencies	137
A.3.11 [[parameters]]	137
A.3.12 [[parameter templates]]	137
A.3.13 [[events]]	137
A.3.13.1 EVENT handler	138
A.3.13.2 handlers	138
A.3.13.3 handler events	139
A.3.13.4 mail events	139
A.3.13.5 mail footer	139
A.3.13.6 mail from	139
A.3.13.7 mail smtp	139
A.3.13.8 mail to	140
A.3.13.9 timeout	140
A.3.13.10inactivity	140
A.3.13.11reset timer	140
A.3.13.12abort on stalled	140
A.3.13.13abort on timeout	141
A.3.13.14abort on inactivity	141
A.3.13.15abort if startup handler fails	141
A.3.14 [[environment]]	141
A.3.14.1 __VARIABLE__	141
A.3.15 [[reference test]]	142
A.3.15.1 suite shutdown event handler	142
A.3.15.2 required run mode	142
A.3.15.3 allow task failures	142
A.3.15.4 expected task failures	142

A.3.15.5	live mode suite timeout	143
A.3.15.6	simulation mode suite timeout	143
A.3.15.7	dummy mode suite timeout	143
A.3.16	[[authentication]]	143
A.3.16.1	public	143
A.3.17	[[simulation]]	143
A.3.17.1	disable suite event handlers	144
A.4	[scheduling]	144
A.4.1	cycling	144
A.4.2	initial cycle point	144
A.4.3	final cycle point	144
A.4.4	initial cycle point constraints	145
A.4.5	final cycle point constraints	145
A.4.6	hold after point	145
A.4.7	runahead limit	145
A.4.8	max active cycle points	145
A.4.9	spawn to max active cycle points	146
A.4.10	[[queues]]	146
A.4.10.1	[[[__QUEUE__]]]	146
A.4.10.2	limit	146
A.4.10.3	members	146
A.4.11	[[special tasks]]	147
A.4.11.1	clock-trigger	147
A.4.11.2	clock-expire	147
A.4.11.3	external-trigger	147
A.4.11.4	sequential	147
A.4.11.5	exclude at start-up	148
A.4.11.6	include at start-up	148
A.4.12	[[dependencies]]	148
A.4.12.1	graph	148
A.4.12.2	[[[__RECURRENCE__]]]	148
A.4.12.2.1	graph	149
A.5	[runtime]	149
A.5.1	[[__NAME__]]	149
A.5.1.1	inherit	150
A.5.1.2	inherit	150
A.5.1.3	init-script	150
A.5.1.4	env-script	150
A.5.1.5	err-script	150
A.5.1.6	pre-script	151
A.5.1.7	script	151
A.5.1.8	post-script	151
A.5.1.9	work sub-directory	151
A.5.1.10	[[[meta]]]	152
A.5.1.10.1	title	152
A.5.1.10.2	description	152
A.5.1.10.3	URL	152
A.5.1.10.4	__MANY__	153
A.5.1.11	[[[job]]]	153

A.5.1.11.1	batch system	153
A.5.1.11.2	execution time limit	153
A.5.1.11.3	batch submit command template	154
A.5.1.11.4	shell	154
A.5.1.11.5	submission retry delays	154
A.5.1.11.6	execution retry delays	154
A.5.1.11.7	submission polling intervals	155
A.5.1.11.8	execution polling intervals	155
A.5.1.12	[[[remote]]]	155
A.5.1.12.1	host	155
A.5.1.12.2	owner	156
A.5.1.12.3	retrieve job logs	156
A.5.1.12.4	retrieve job logs max size	156
A.5.1.12.5	retrieve job logs retry delays	157
A.5.1.12.6	suite definition directory	157
A.5.1.13	[[[events]]]	157
A.5.1.13.1	EVENT handler	158
A.5.1.13.2	submission timeout	158
A.5.1.13.3	execution timeout	159
A.5.1.13.4	reset timer	159
A.5.1.13.5	handlers	159
A.5.1.13.6	handler events	159
A.5.1.13.7	handler retry delays	159
A.5.1.13.8	mail events	159
A.5.1.13.9	mail from	160
A.5.1.13.10	mail retry delays	160
A.5.1.13.11	mail smtp	160
A.5.1.13.12	mail to	160
A.5.1.14	[[[environment]]]	160
A.5.1.14.1	__VARIABLE__	161
A.5.1.15	[[[environment filter]]]	161
A.5.1.15.1	include	161
A.5.1.15.2	exclude	161
A.5.1.16	[[[parameter environment templates]]]	162
A.5.1.16.1	__VARIABLE__	162
A.5.1.17	[[[directives]]]	162
A.5.1.17.1	__DIRECTIVE__	162
A.5.1.18	[[[outputs]]]	162
A.5.1.18.1	__OUTPUT__	163
A.5.1.19	[[[suite state polling]]]	163
A.5.1.19.1	run-dir	163
A.5.1.19.2	template	163
A.5.1.19.3	interval	163
A.5.1.19.4	max-polls	163
A.5.1.19.5	user	164
A.5.1.19.6	host	164
A.5.1.19.7	verbose	164
A.5.1.20	[[[simulation]]]	164
A.5.1.20.1	default run length	164

	A.5.1.20.2 speedup factor	164
	A.5.1.20.3 time limit buffer	164
	A.5.1.20.4 fail cycle points	165
	A.5.1.20.5 fail try 1 only	165
	A.5.1.20.6 disable task event handlers	165
A.6	[visualization]	165
	A.6.1 initial cycle point	165
	A.6.2 final cycle point	165
	A.6.3 number of cycle points	166
	A.6.4 collapsed families	166
	A.6.5 use node color for edges	166
	A.6.6 use node fillcolor for edges	166
	A.6.7 node penwidth	166
	A.6.8 edge penwidth	167
	A.6.9 use node color for labels	167
	A.6.10 default node attributes	167
	A.6.11 default edge attributes	167
	A.6.12 [[node groups]]	167
	A.6.12.1 __GROUP__	167
	A.6.13 [[node attributes]]	168
	A.6.13.1 __NAME__	168
B	Global (Site, User) Config File Reference	169
B.1	Top Level Items	169
	B.1.1 temporary directory	169
	B.1.2 process pool size	169
	B.1.3 disable interactive command prompts	169
	B.1.4 enable run directory housekeeping	169
	B.1.5 run directory rolling archive length	170
	B.1.6 task host select command timeout	170
B.2	[task messaging]	170
	B.2.1 retry interval	170
	B.2.2 maximum number of tries	170
	B.2.3 connection timeout	170
B.3	[suite logging]	171
	B.3.1 roll over at start-up	171
	B.3.2 rolling archive length	171
	B.3.3 maximum size in bytes	171
B.4	[documentation]	171
	B.4.1 [[files]]	171
	B.4.1.1 html index	171
	B.4.1.2 pdf user guide	172
	B.4.1.3 multi-page html user guide	172
	B.4.1.4 single-page html user guide	172
	B.4.2 [[urls]]	172
	B.4.2.1 internet homepage	172
	B.4.2.2 local index	172
B.5	[document viewers]	172
	B.5.1 pdf	172

B.5.2	html	173
B.6	[editors]	173
B.6.1	terminal	173
B.6.2	gui	173
B.7	[communication]	173
B.7.1	method	174
B.7.2	base port	174
B.7.3	maximum number of ports	174
B.7.4	proxies on	174
B.7.5	options	174
B.8	[monitor]	175
B.8.1	monitor	175
B.9	[hosts]	175
B.9.1	[[HOST]]	175
B.9.1.1	run directory	176
B.9.1.2	work directory	176
B.9.1.3	task communication method	176
B.9.1.4	execution polling intervals	176
B.9.1.5	submission polling intervals	176
B.9.1.6	scp command	177
B.9.1.7	ssh command	177
B.9.1.8	use login shell	177
B.9.1.9	cylc executable	177
B.9.1.10	global init-script	178
B.9.1.11	copyable environment variables	178
B.9.1.12	retrieve job logs	178
B.9.1.13	retrieve job logs command	178
B.9.1.14	retrieve job logs max size	178
B.9.1.15	retrieve job logs retry delays	178
B.9.1.16	task event handler retry delays	178
B.9.1.17	local tail command template	178
B.9.1.18	remote tail command template	179
B.9.1.19	[[[batch systems]]]	179
B.9.1.19.1	[[[[SYSTEM]]]]err tailer	179
B.9.1.19.2	[[[[SYSTEM]]]]out tailer	179
B.9.1.19.3	[[[[SYSTEM]]]]err viewer	179
B.9.1.19.4	[[[[SYSTEM]]]]out viewer	180
B.9.1.19.5	[[[[SYSTEM]]]]job name length maximum	180
B.9.1.19.6	[[[[SYSTEM]]]]execution time limit polling intervals	180
B.10	[suite host self-identification]	181
B.10.1	method	181
B.10.2	target	181
B.10.3	host	181
B.11	[suite host scanning]	181
B.11.1	hosts	181
B.12	[task events]	182
B.13	[test battery]	182
B.13.1	remote host with shared fs	182
B.13.2	remote host with shared fs	182

B.13.3	[[batch systems]]	182
B.13.3.1	[[[SYSTEM]]]	182
B.13.3.1.1	host	182
B.13.3.1.2	err viewer	182
B.13.3.1.3	out viewer	183
B.13.3.1.4	[[[[directives]]]]	183
B.14	[cylc]	183
B.14.1	UTC mode	183
B.14.2	health check interval	183
B.14.3	task event mail interval	183
B.14.4	[events]	183
B.14.4.1	handlers	184
B.14.4.2	handler events	184
B.14.4.3	startup handler	184
B.14.4.4	shutdown handler	184
B.14.4.5	mail events	184
B.14.4.6	mail footer	184
B.14.4.7	mail from	184
B.14.4.8	mail smtp	184
B.14.4.9	mail to	184
B.14.4.10	timeout handler	184
B.14.4.11	timeout	184
B.14.4.12	abort on timeout	184
B.14.4.13	stalled handler	184
B.14.4.14	abort on stalled	184
B.14.4.15	inactivity handler	184
B.14.4.16	inactivity	184
B.14.4.17	abort on inactivity	184
B.15	[authentication]	184
B.15.1	public	184
C	Gcylc GUI (cylc gui) Config File Reference	186
C.1	Top Level Items	186
C.1.1	dot icon size	186
C.1.2	initial side-by-side views	186
C.1.3	initial views	186
C.1.4	maximum update interval	186
C.1.5	sort by definition order	187
C.1.6	sort column	187
C.1.7	sort column ascending	187
C.1.8	task filter highlight color	187
C.1.9	task states to filter out	187
C.1.10	transpose dot	188
C.1.11	transpose graph	188
C.1.12	ungrouped views	188
C.1.13	use theme	188
C.1.14	window size	189
C.2	[themes]	189
C.2.1	[THEME]	189

C.2.1.1	inherit	189
C.2.1.2	defaults	189
C.2.1.3	STATE	189
D	Gscan GUI (cylc gscan) Config File Reference	191
D.1	Top Level Items	191
D.1.1	activate on startup	191
D.1.2	columns	191
D.1.3	suite listing update interval	191
D.1.4	suite status update interval	192
D.1.5	window size	192
D.1.6	hide main menubar	192
E	Remote Job Host Interaction	193
F	Command Reference	195
F.1	Command Categories	195
F.1.1	admin	195
F.1.2	all	196
F.1.3	control	198
F.1.4	discovery	198
F.1.5	hook	198
F.1.6	information	198
F.1.7	license	199
F.1.8	preparation	199
F.1.9	task	200
F.1.10	utility	200
F.2	Commands	200
F.2.1	5to6	200
F.2.2	broadcast	201
F.2.3	cat-log	203
F.2.4	cat-state	203
F.2.5	check-software	204
F.2.6	check-triggering	204
F.2.7	check-versions	204
F.2.8	checkpoint	205
F.2.9	conditions	205
F.2.10	cycle-point	206
F.2.11	diff	206
F.2.12	documentation	207
F.2.13	dump	207
F.2.14	edit	208
F.2.15	email-suite	209
F.2.16	email-task	209
F.2.17	ext-trigger	210
F.2.18	get-directory	211
F.2.19	get-gui-config	211
F.2.20	get-site-config	211
F.2.21	get-suite-config	212
F.2.22	get-suite-contact	213

F.2.23	get-suite-version	213
F.2.24	gpanel	214
F.2.25	graph	214
F.2.26	graph-diff	215
F.2.27	gscan	215
F.2.28	gui	216
F.2.29	hold	217
F.2.30	import-examples	218
F.2.31	insert	218
F.2.32	jobs-kill	219
F.2.33	jobs-poll	220
F.2.34	jobs-submit	220
F.2.35	jobscript	221
F.2.36	kill	221
F.2.37	list	222
F.2.38	ls-checkpoints	223
F.2.39	message	223
F.2.40	monitor	223
F.2.41	nudge	224
F.2.42	ping	225
F.2.43	poll	225
F.2.44	print	226
F.2.45	profile-battery	227
F.2.46	register	228
F.2.47	release	229
F.2.48	reload	230
F.2.49	remote-init	230
F.2.50	remote-tidy	231
F.2.51	remove	231
F.2.52	report-timings	232
F.2.53	reset	233
F.2.54	restart	234
F.2.55	run	235
F.2.56	scan	236
F.2.57	scp-transfer	237
F.2.58	search	237
F.2.59	set-verbosity	238
F.2.60	show	239
F.2.61	spawn	239
F.2.62	stop	240
F.2.63	submit	241
F.2.64	suite-state	242
F.2.65	test-battery	243
F.2.66	trigger	244
F.2.67	upgrade-run-dir	245
F.2.68	validate	246
F.2.69	version	246
F.2.70	view	247
F.2.71	warranty	248

G	The gcylc Graph View	248
H	Cylc README File	248
I	Cylc INSTALL File	249
J	Cylc Development History - Major Changes	250
K	Communication Method	251
L	Cylc 6 Migration Reference	251
L.1	Timeouts and Delays	251
L.2	Runahead Limit	251
L.3	Cycle Time/Cycle Point	252
L.4	Cycling	252
L.5	No Implicit Creation of Tasks by Offset Triggers	253
M	Known Issues	254
M.1	Current Known Issues	254
M.2	Notable Known Issues	254
M.2.1	Use of pipes in job scripts	254
N	GNU GENERAL PUBLIC LICENSE v3.0	254

List of Figures

1	A single cycle point dependency graph for a simple suite	17
2	A single cycle point job schedule for real time operation	17
3	What if the external driving data is available early?	18
4	Attempted overlap of consecutive single-cycle-point job schedules	18
5	The only safe multi-cycle-point job schedule?	18
6	The complete multi-cycle-point dependency graph	20
7	The optimal two-cycle-point job schedule	20
8	Comparison of job schedules after a delay	20
9	Optimal job schedule when all external data is available	21
10	The cylc task pool	22
11	gcylc graph and dot views	23
12	gcylc text view	23
13	gscan multi-suite state summary GUI	24
14	A large-ish suite graphed by cylc	24
15	The <i>tut/oneoff/fttrigger2</i> dependency and runtime inheritance graphs	42
16	The <i>tut/cycling/one</i> suite	44
17	The <i>tut/cycling/two</i> suite	45
18	The <i>tut/cycling/three</i> suite	47
19	The <i>tut/cycling/integer</i> suite	48
20	Example Suite	54
21	One-off (Non-Cycling) Tasks	56
22	Cycling Tasks	57
23	Staggered Start Suite	61
24	Restricted First Cycle Point Suite	62

25	The <code>examples/satellite</code> integer suite	64
26	Conditional Triggers	68
27	Automated failure recovery via suicide triggers	68
28	Screenshot of <code>cylc</code> graph showing one task as a “ghost node”	77
29	<i>namespaces</i> example suite graphs	85
30	Parameter expansion example.	88
31	Parameterized (top) and cycling (bottom) versions of the same workflow.	92
32	The Jinja2 ensemble example suite graph.	94
33	Jinja2 cities example suite graph.	96

1 Introduction: How Cylc Works

1.1 Scheduling Forecast Suites

Environmental forecasting suites generate forecast products from a potentially large group of interdependent scientific models and associated data processing tasks. They are constrained by availability of external driving data: typically one or more tasks will wait on real time observations and/or model data from an external system, and these will drive other downstream tasks, and so on. The dependency diagram for a single forecast cycle point in such a system is a *Directed Acyclic Graph* as shown in Figure 1 (in our terminology, a *forecast cycle point* is comprised of all tasks with a common *cycle point*, which is the nominal analysis time or start time of the forecast models in the group). In real time operation processing will consist of a series of distinct forecast cycle points that are each initiated, after a gap, by arrival of the new cycle point's external driving data.

From a job scheduling perspective task execution order in such a system must be carefully controlled in order to avoid dependency violations. Ideally, each task should be queued for execution at the instant its last prerequisite is satisfied; this is the best that can be done even if queued tasks are not able to execute immediately because of resource contention.

1.2 EcoConnect

Cylc was developed for the EcoConnect Forecasting System at NIWA (National Institute of Water and Atmospheric Research, New Zealand). EcoConnect takes real time atmospheric and stream flow observations, and operational global weather forecasts from the Met Office (UK), and uses these to drive global sea state and regional data assimilating weather models, which in turn drive regional sea state, storm surge, and catchment river models, plus tide prediction, and a large number of associated data collection, quality control, preprocessing, post-processing, product generation, and archiving tasks.⁴ The global sea state forecast runs once daily. The regional weather forecast runs four times daily but it supplies surface winds and pressure to several downstream models that run only twice daily, and precipitation accumulations to catchment river models that run on an hourly cycle assimilating real time stream flow observations and using the most recently available regional weather forecast. EcoConnect runs on heterogeneous distributed hardware, including a massively parallel supercomputer and several Linux servers.

1.3 Dependence Between Tasks

1.3.1 Intra-cycle Dependence

Most dependence between tasks applies within a single forecast cycle point. Figure 1 shows the dependency diagram for a single forecast cycle point of a simple example suite of three forecast models (*a*, *b*, and *c*) and three post processing or product generation tasks (*d*, *e* and *f*). A scheduler capable of handling this must manage, within a single forecast cycle point, multiple parallel streams of execution that branch when one task generates output for several downstream tasks, and merge when one task takes input from several upstream tasks.

⁴Future plans for EcoConnect include additional deterministic regional weather forecasts and a statistical ensemble.




Figure 1: The dependency graph for a single forecast cycle point of a simple example suite. Tasks *a*, *b*, and *c* represent forecast models, *d*, *e* and *f* are post processing or product generation tasks, and *x* represents external data that the upstream forecast model depends on.




Figure 2: The optimal job schedule for two consecutive cycle points of our example suite during real time operation, assuming that all tasks trigger off upstream tasks finishing completely. The horizontal extent of a task bar represents its execution time, and the vertical blue lines show when the external driving data becomes available.




Figure 3: If the external driving data is available in advance, can we start running the next cycle point early?




Figure 4: A naive attempt to overlap two consecutive cycle points using the single-cycle-point dependency graph. The red shaded tasks will fail because of dependency violations (or will not be able to run because of upstream dependency violations).

Figure 2 shows the optimal job schedule for two consecutive cycle points of the example suite in real time operation, given execution times represented by the horizontal extent of the task bars. There is a time gap between cycle points as the suite waits on new external driving data. Each task in the example suite happens to trigger off upstream tasks *finishing*, rather than off any intermediate output or event; this is merely a simplification that makes for clearer diagrams.

Now the question arises, what happens if the external driving data for upcoming cycle points is available in advance, as it would be after a significant delay in operations, or when running a historical case study? While the forecast model *a* appears to depend only on the external data *x* at this stage of the discussion, in fact it would typically also depend on its own previous instance for the model *background state* used in initializing the new forecast. Thus, as alluded to in Figure 3, task *a* could in principle start as soon as its predecessor has finished. Figure 4 shows, however, that starting a whole new cycle point at this point is dangerous - it results in dependency violations in half of the tasks in the example suite. In fact the situation could be




Figure 5: The best that can be done *in general* when inter-cycle dependence is ignored.

even worse than this - imagine that task *b* in the first cycle point is delayed for some reason *after* the second cycle point has been launched. Clearly we must consider handling inter-cycle dependence explicitly or else agree not to start the next cycle point early, as is illustrated in Figure 5.

1.3.2 Inter-Cycle Dependence

Forecast models typically depend on their own most recent previous forecast for background state or restart files of some kind (this is called *warm cycling*) but there can also be inter-cycle dependence between different tasks. In an atmospheric forecast analysis suite, for instance, the weather model may generate background states for observation processing and data-assimilation tasks in the next cycle point as well as for the next forecast model run. In real time operation inter-cycle dependence can be ignored because it is automatically satisfied when one cycle point finishes before the next begins. If it is not ignored it drastically complicates the dependency graph by blurring the clean boundary between cycle points. Figure 6 illustrates the problem for our simple example suite assuming minimal inter-cycle dependence: the warm cycled models (*a*, *b*, and *c*) each depend on their own previous instances.

For this reason, and because we tend to see forecasting suites in terms of their real time characteristics, other metaschedulers have ignored inter-cycle dependence and are thus restricted to running entire cycle points in sequence at all times. This does not affect normal real time operation but it can be a serious impediment when advance availability of external driving data makes it possible, in principle, to run some tasks from upcoming cycle points before the current cycle point is finished - as was suggested at the end of the previous section. This can occur, for instance, after operational delays (late arrival of external data, system maintenance, etc.) and to an even greater extent in historical case studies and parallel test suites started behind a real time operation. It can be a serious problem for suites that have little downtime between forecast cycle points and therefore take many cycle points to catch up after a delay. Without taking account of inter-cycle dependence, the best that can be done, in general, is to reduce the gap between cycle points to zero as shown in Figure 5. A limited crude overlap of the single cycle point job schedule may be possible for specific task sets but the allowable overlap may change if new tasks are added, and it is still dangerous: it amounts to running different parts of a dependent system as if they were not dependent and as such it cannot be guaranteed that some unforeseen delay in one cycle point, after the next cycle point has begun, (e.g. due to resource contention or task failures) won't result in dependency violations.

Figure 7 shows, in contrast to Figure 4, the optimal two cycle point job schedule obtained by respecting all inter-cycle dependence. This assumes no delays due to resource contention or otherwise - i.e. every task runs as soon as it is ready to run. The scheduler running this suite must be able to adapt dynamically to external conditions that impact on multi-cycle-point scheduling in the presence of inter-cycle dependence or else, again, risk bringing the system down with dependency violations.

To further illustrate the potential benefits of proper inter-cycle dependency handling, Figure 8 shows an operational delay of almost one whole cycle point in a suite with little downtime between cycle points. Above the time axis is the optimal schedule that is possible in principle when inter-cycle dependence is taken into account, and below it is the only safe schedule possible *in general* when it is ignored. In the former case, even the cycle point immediately after the delay is hardly affected, and subsequent cycle points are all on time, whilst in the latter case it takes five full cycle points to catch up to normal real time operation.




Figure 6: The complete dependency graph for the example suite, assuming the least possible inter-cycle dependence: the forecast models (*a*, *b*, and *c*) depend on their own previous instances. The dashed arrows show connections to previous and subsequent forecast cycle points.




Figure 7: The optimal two cycle job schedule when the next cycle's driving data is available in advance, possible in principle when inter-cycle dependence is handled explicitly.




Figure 8: Job schedules for the example suite after a delay of almost one whole forecast cycle point, when inter-cycle dependence is taken into account (above the time axis), and when it is not (below the time axis). The colored lines indicate the time that each cycle point is delayed, and normal “caught up” cycle points are shaded gray.




Figure 9: Job schedules for the example suite in case study mode, or after a long delay, when the external driving data are available many cycle points in advance. Above the time axis is the optimal schedule obtained when the suite is constrained only by its true dependencies, as in Figure 3, and underneath is the best that can be done, in general, when inter-cycle dependence is ignored.

Similarly, Figure 9 shows example suite job schedules for an historical case study, or when catching up after a very long delay; i.e. when the external driving data are available many cycle points in advance. Task *a*, which as the most upstream forecast model is likely to be a resource intensive atmosphere or ocean model, has no upstream dependence on co-temporal tasks and can therefore run continuously, regardless of how much downstream processing is yet to be completed in its own, or any previous, forecast cycle point (actually, task *a* does depend on co-temporal task *x* which waits on the external driving data, but that returns immediately when the data is available in advance, so the result stands). The other forecast models can also cycle continuously or with a short gap between, and some post processing tasks, which have no previous-instance dependence, can run continuously or even overlap (e.g. *e* in this case). Thus, even for this very simple example suite, tasks from three or four different cycle points can in principle run simultaneously at any given time.

In fact, if our tasks are able to trigger off internal outputs of upstream tasks (message triggers) rather than waiting on full completion, then successive instances of the forecast models could overlap as well (because model restart outputs are generally completed early in the forecast) for an even more efficient job schedule.

1.4 The Cylc Scheduling Algorithm

Cylc manages a pool of proxy objects that represent the real tasks in a suite. Task proxies know how to run the real tasks that they represent, and they receive progress messages from the tasks as they run (usually reports of completed outputs). There is no global cycling mechanism to advance the suite; instead individual task proxies have their own private cycle point and spawn their own successors when the time is right. Task proxies are self-contained - they know their own prerequisites and outputs but are not aware of the wider suite. Inter-cycle dependence is not treated as special, and the task pool can be populated with tasks with many different cycle points. The task pool is illustrated in Figure 10. *Whenever any task changes state due to completion of an output, every task checks to see if its own prerequisites have been satisfied.* In effect, cylc gets a pool of tasks to self-organize by negotiating their own dependencies so that optimal scheduling, as described in the previous section, emerges naturally at run time.




Figure 10: How cylc sees a suite, in contrast to the multi-cycle-point dependency graph of Figure 6. Task colors represent different cycle points, and the small squares and circles represent different prerequisites and outputs. A task can run when its prerequisites are satisfied by the outputs of other tasks in the pool.

2 Cylc Screenshots

2 CYLC SCREENSHOTS




Figure 11: gcylic graph and dot views.

task	state	host	Job ID	T-submit	T-start	T-finish	dT-mean	latest message
2012080800	submit-failed							
family	queued							
m1	succeeded	localhost	30237	22:01:38+12	22:01:44+12	22:01:53+12	PT9S	succeeded at 22:01:52+12
m2	succeeded	localhost	30701	22:02:11+12	22:02:16+12	22:02:18+12	PT5S	succeeded at 22:02:18+12
m3	succeeded	localhost	30230	22:01:38+12	22:01:44+12	22:01:58+12	PT8S	succeeded at 22:01:57+12
m_x	queued	*	*	*	*	*	*	*
bad	failed	localhost	30370	22:01:47+12	22:01:52+12	22:02:02+12	*	failed at 22:02:01+12
bad2	submit-failed	*	*	*	*	*	*	*
bad_good	succeeded	localhost	30373	22:01:47+12	22:01:52+12	22:02:02+12	PT8S	succeeded at 22:02:02+12
cfoo	succeeded	localhost	29968	22:00:45+12	22:00:50+12	22:01:04+12	PT14S	succeeded at 22:01:04+12
foo	succeeded	localhost	30089	22:01:06+12	22:01:11+12	22:01:16+12	PT6S	succeeded at 22:01:16+12
bar	waiting	*	*	*	*	*	*	*
2012080900	failed							
family	running							
m1	running	localhost	31067	22:02:36+12	22:02:41+12	22:02:50+12	PT9S	started at 22:02:40+12
m1	waiting							
submit-retrying	submit-retrying							
running	running							
succeeded	succeeded							
failed	failed							
ready	ready							
submitted	submitted							
submit-failed	submit-failed							
retrying	retrying							

Figure 12: gcylic text view.

Suite	Status
battery-24834.tests.QuickStart.b	■ □ □
battery-24834.tests.QuickStart.c	■ □ □
battery-24834.tests.broadcast	□ ■ ■
battery-24834.tests.combined	■ □ ■
battery-24834.tests.events.suite	■ □ □
battery-24834.tests.events.task	■ ■ ■ □
battery-24834.tests.host-select	■
battery-24834.tests.intercycle.one	■
battery-24834.tests.internal-outputs	■ ■ ■
battery-24834.tests.jobscript	■
battery-24834.tests.modes.simulation	■

Figure 13: gscan multi-suite state summary GUI.




Figure 14: A large-ish suite graphed by cylc.

3 Installation

Cylc runs on Unix variants, usually Linux, and including Apple OS X.

3.1 External Software Packages

Python `>= 2.6` is required (but not yet Python 3). Python should already be installed in your Linux system. <https://python.org>.

For Cylc's HTTPS communications layer:

- **OpenSSL** - <https://www.openssl.org/>
- **pyOpenSSL** - <http://www.pyopenssl.org/>
- **python-requests** - <http://docs.python-requests.org/>
- (**python-urllib3** - should be bundled with python-requests)

The following packages are highly recommended, but are technically optional as you can construct and run suites without dependency graph visualisation or the Cylc GUIs:

- **PyGTK** - GUI toolkit <http://www.pygtk.org>. *Note PyGTK typically comes with your system Python. It is alledgedly quite difficult to install if you need to do so for another Python version.*
- **Graphviz** - graph layout engine (tested 2.36.0): <http://www.graphviz.org>.
- **Pygraphviz** - Python Graphviz interface (tested 1.2): <http://pygraphviz.github.io/>.
To build this you may need some *devel* packages too:
 - python-devel
 - graphviz-devel

The User Guide is generated from L^AT_EXsource files by running `make` in the top level Cylc directory. The specific packages required may vary by distribution, e.g.:

- texlive
- texlive-tocloft
- texlive-framed
- texlive-preprint (for `fullpage.sty`)
- texlive-tex4ht

To generate the HTML User Guide **ImageMagick** is also needed.

In most modern Linux distributions all of the software above can be installed via the system package manager. Otherwise download packages manually and follow their native installation instructions. To check that all (non L^AT_EXpackages) are installed properly:

```
$ cylc check-software
Checking for Python >= 2.6 ... found 2.7.6 ... ok
Checking for non-Python packages:
+ Graphviz ... ok
Checking for Python packages:
+ pygraphviz ... ok
+ pygtk ... ok
```

If errors are reported then the packages concerned are either not installed or not in your Python search path. (Note that `cylc check-software` has become quite trivial as we've removed or bundled some former dependencies, but in future we intend to make it print a comprehensive list of library versions etc. to include in with bug reports.)

3.2 Software Bundled With Cylc

Cylc bundles several third party packages which do not need to be installed separately.

- **cherryypy 6.0.2** (slightly modified): a pure Python HTTP framework that we use as a web server for communication between server processes (suite server programs) and client programs (running tasks, GUIs, CLI commands). Client communication is via the Python **requests** library if available (recommended) or else pure Python via **urllib2**.
<http://www.cherryypy.org/>
<http://docs.python-requests.org/>
- **Jinja2 2.10**: a full featured template engine for Python, and its dependency **MarkupSafe 0.23**; both BSD licensed.
<http://jinja.pocoo.org/>
<http://www.pocoo.org/projects/markupsafe/>
- the **xdot** graph viewer (modified), LGPL licensed:
<https://github.com/jrfonseca/xdot.py>

3.3 Installing Cylc

Cylc releases can be downloaded from <https://cylc.github.io/cylc>.

The wrapper script `admin/cylc-wrapper` should be installed as `cylc` in the system executable search path (e.g. `/usr/local/bin/`) and modified slightly to point to a location such as `/opt` where successive Cylc releases will be unpacked side by side.

To install Cylc, unpack the release tarball in the right location, e.g. `/opt/cylc-7.4.0`, type `make` inside the release directory, and set site defaults - if necessary - in a site global config file (below).

Make a symbolic link from `cylc` to the latest installed version: `ln -s /opt/cylc-7.4.0 /opt/cylc`. This will be invoked by the central wrapper if a specific version is not requested. Otherwise, the wrapper will attempt to invoke the Cylc version specified in `$CYLC_VERSION`, e.g. `CYLC_VERSION=7.4.0`. This variable is automatically set in task job scripts to ensure that jobs use the same Cylc version as their parent suite server program. It can also be set by users, manually or in login scripts, to fix the Cylc version in their environment.

Installing subsequent releases is just a matter of unpacking the new tarballs next to the previous releases, running `make` in them, and copying in (possibly with modifications) the previous site global config file.

3.3.1 Local User Installation

It is easy to install Cylc under your own user account if you don't have root or sudo access to the system: just put the central Cylc wrapper in `$HOME/bin/` (making sure that is in your `$PATH`) and modify it to point to a directory such as `$HOME/cylc/` where you will unpack and install release tarballs. Local installation of third party dependencies like Graphviz is also possible, but that depends on the particular installation methods used and is outside of the scope of this document.

3.3.2 Create A Site Config File

Site and user global config files define some important parameters that affect all suites, some of which may need to be customized for your site. See [6](#) for how to generate an initial site file and where to install it. All legal site and user global config items are defined in [B](#).

3.3.3 Configure Site Environment on Job Hosts

If your users submit task jobs to hosts other than the hosts they use to run their suites, you should ensure that the job hosts have the correct environment for running cylc. A cylc suite generates task job scripts that normally invoke `bash -l`, i.e. it will invoke bash as a login shell to run the job script. Users and sites should ensure that their bash login profiles are able to set up the correct environment for running cylc and their task jobs.

Your site administrator may customise the environment for all task jobs by adding a `${CYLC_DIR}/conf/job-init-env` file and populate it with the appropriate contents. If customisation is still required, you can add your own `${HOME}/.cylc/job-init-env.sh` file and populate it with the appropriate contents.

- `${HOME}/.cylc/job-init-env.sh`
- `${CYLC_DIR}/conf/job-init-env.sh`

The job will attempt to source the first of these files it finds to set up its environment.

3.4 Automated Tests

The cylc test battery is primarily intended for developers to check that changes to the source code don't break existing functionality. Note that some test failures can be expected to result from suites timing out, even if nothing is wrong, if you run too many tests in parallel. See `cylc test-battery --help`.

4 Cylc Terminology

4.1 Jobs and Tasks

A *job* is a program or script that runs on a computer, and a *task* is a workflow abstraction - a node in the suite dependency graph - that represents a job.

4.2 Cycle Points

A *cycle point* is a particular date-time (or integer) point in a sequence of date-time (or integer) points. Each cylc task has a private cycle point and can advance independently to subsequent cycle points. It may sometimes be convenient, however, to refer to the “current cycle point” of a suite (or the previous or next one, etc.) with reference to a particular task, or in the sense of all tasks instances that “belong to” a particular cycle point. But keep in mind that different tasks may pass through the “current cycle point” (etc.) at different times as the suite evolves.

5 Workflows For Cycling Systems

A model run and associated processing may need to be cycled for the following reasons:

- In real time forecasting systems, a new forecast may be initiated at regular intervals when new real time data comes in.
- It may be convenient (or necessary, e.g. due to batch scheduler queue limits) to split single long model runs into many smaller chunks, each with associated pre- and post-processing workflows.

Cylc provides two ways of constructing workflows for cycling systems: *cycling workflows* and *parameterized tasks*.

5.1 Cycling Workflows

This is cylc's classic cycling mode as described in the Introduction. Each instance of a cycling job is represented by a new instance of *the same task*, with a new cycle point. The suite configuration defines patterns for extending the workflow on the fly, so it can keep running indefinitely if necessary. For example, to cycle `model.exe` on a monthly sequence we could define a single task `model`, an initial cycle point, and a monthly sequence. Cylc then generates the date-time sequence and creates a new task instance for each cycle point as it comes up. Workflow dependencies are defined generically with respect to the “current cycle point” of the tasks involved.

This is the only sensible way to run very large suites or operational suites that need to continue cycling indefinitely. The cycling is configured with standards-based ISO 8601 date-time *recurrence expressions*. Multiple cycling sequences can be used at once in the same suite. See Section 9.3.

5.2 Parameterized Tasks as a Proxy for Cycling

It is also possible to run cycling jobs with a pre-defined static workflow in which each instance of a cycling job is represented by *a different task*: as far as the abstract workflow is concerned there is no cycling. The sequence of tasks can be constructed efficiently, however, using cylc's built-in suite parameters (9.6.7) or explicit Jinja2 loops (9.7).

For example, to run `model.exe` 12 times on a monthly cycle we could loop over an integer parameter `R = 0, 1, 2, ..., 11` to define tasks `model-R0`, `model-R1`, `model-R2`, ... `model-R11`, and the parameter values could be multiplied by the interval `P1M` (one month) to get the start point for the corresponding model run.

This method is only good for smaller workflows of finite duration because every single task has to be mapped out in advance, and cylc has to be aware of all of them throughout the entire run. Additionally Cylc's *cycling workflow* capabilities (above) are more powerful, more flexible, and generally easier to use (Cylc will generate the cycle point date-times for you, for instance), so that is the recommended way to drive most cycling systems.

The primary use for parameterized tasks in cylc is to generate ensembles and other groups of related tasks at the same cycle point, not as a proxy for cycling.

5.3 Mixed Cycling Workflows

For completeness we note that parameterized cycling can be used within a cycling workflow. For example, in a daily cycling workflow long (daily) model runs could be split into four shorter runs by parameterized cycling. A simpler six-hourly cycling workflow should be considered first, however.

6 Global (Site, User) Configuration Files

Cylc site and user global configuration files contain settings that affect all suites. Some of these, such as the range of network ports used by cylc, should be set at site level,

```
# cylc site global config file
/path/to/cylc/conf/global.rc
# Deprecated path to cylc site global config file
/path/to/cylc/conf/siterc/site.rc
```

Others, such as the preferred text editor for suite definitions, can be overridden by users,

```
# cylc user global config file
~/.cylc/global.rc
# Deprecated cylc user global config file
~/.cylc/user.rc
```

The `cylc get-site-config` command retrieves current global settings consisting of cylc defaults overridden by site settings, if any, overridden by user settings, if any. If you need to generate an example user global config file filled with the default values commented out (useful as a quick reference), you can do:

```
$ cylc get-site-config | sed 's/^/#/' > ~/.cylc/global.rc.example
```

Settings that do not need to be changed should be deleted to reduce conflicts with site changes and changes in future versions of cylc. Copy the example file into `~/.cylc/global.rc` if you need to apply any changes.

Legal items, values, and system defaults are documented in (B).

7 Tutorial

This section provides a hands-on tutorial introduction to basic cylc functionality.

7.1 User Config File

Some settings affecting cylc's behaviour can be defined in site and user *global config files*. For example, to choose the text editor invoked by cylc on suite definitions:

```
# $HOME/.cylc/global.rc
[editors]
    terminal = vim
    gui = gvim -f
```

- For more on site and user global config files see 6 and B.

7.1.1 Configure Environment on Job Hosts

See [3.3.3](#) for information.

7.2 User Interfaces

You should have access to the `cylc` command line (CLI) and graphical (GUI) user interfaces once `cylc` has been installed as described in [Section 3.3](#).

7.2.1 Command Line Interface (CLI)

The command line interface is unified under a single top level `cylc` command that provides access to many sub-commands and their help documentation.

```
$ cylc help          # Top level command help.
$ cylc run --help    # Example command-specific help.
```

Command help transcripts are printed in [F](#) and are available from the GUI Help menu.

`Cylc` is *scriptable* - the error status returned by commands can be relied on.

7.2.2 Graphical User Interface (GUI)

The `cylc` GUI covers the same functionality as the CLI, but it has more sophisticated suite monitoring capability. It can start and stop suites, or connect to suites that are already running; in either case, shutting down the GUI does not affect the suite itself.

```
$ gcylc & # or:
$ cylc gui & # Single suite control GUI.
$ cylc gscan & # Multi-suite monitor GUI.
```

Clicking on a suite in `gscan`, shown in [Figure 13](#), opens a `gcylc` instance for it.

7.3 Suite Definitions

`Cylc` suites are defined by extended-INI format `suite.rc` files (the main file format extension is section nesting). These reside in *suite definition directories* that may also contain a `bin` directory and any other suite-related files.

- For more on the suite definition file format, see [9](#) and [A](#).

7.4 Suite Registration

Suite registration creates a run directory (under `~/cylc-run/` by default) and populates it with authentication files and a symbolic link to a suite definition directory. `Cylc` commands that parse suite definitions can take the file path or the suite name as input. Commands that interact with running suites have to target the suite by name.


```
# Target a suite by file path:
$ cylc validate /path/to/my/suite/suite.rc
$ cylc graph /path/to/my/suite/suite.rc

# Register a suite:
$ cylc register my.suite /path/to/my/suite/

# Target a suite by name:
$ cylc graph my.suite
$ cylc validate my.suite
$ cylc run my.suite
$ cylc stop my.suite
# etc.
```

7.5 Suite Passphrases

Registration (above) also generates a suite-specific passphrase file under `.service/` in the suite run directory. It is loaded by the suite server program at start-up and used to authenticate connections from client programs.

Possession of a suite's passphrase file gives full control over it. Without it, the information available to a client is determined by the suite's public access privilege level.

For more on connection authentication, suite passphrases, and public access, see [12.9](#).

7.6 Import The Example Suites

Run the following command to copy cylc's example suites and register them for your own use:

```
$ cylc import-examples /tmp
```

7.7 Rename The Imported Tutorial Suites

Suites can be renamed by simply renaming (i.e. moving) their run directories. Make the tutorial suite names shorter, and print their locations with `cylc print`:

```
$ mv ~/cylc-run/$(cylc --version)/examples/tutorial ~/cylc-run/tut
$ cylc print -ya tut
tut/oneoff/jinja2 | /tmp/cylc-examples/7.0.0/tutorial/oneoff/jinja2
tut/cycling/two | /tmp/cylc-examples/7.0.0/tutorial/cycling/two
tut/cycling/three | /tmp/cylc-examples/7.0.0/tutorial/cycling/three
# ...
```

See `cylc print --help` for other display options.

7.8 Suite Validation

Suite definitions can be validated to detect syntax (and other) errors:

```
# pass:
$ cylc validate tut/oneoff/basic
Valid for cylc-6.0.0
$ echo $?
0
# fail:
$ cylc validate my/bad/suite
Illegal item: [scheduling]special tasks
```

```
$ echo $?
1
```

7.9 Hello World in Cylc

```
suite: tut/oneoff/basic
```

Here's the traditional *Hello World* program rendered as a cylc suite:

```
[meta]
    title = "The cylc Hello World! suite"
[scheduling]
    [[dependencies]]
        graph = "hello"
[runtime]
    [[hello]]
        script = "sleep 10; echo Hello World!"
```

Cylc suites feature a clean separation of scheduling configuration, which determines *when* tasks are ready to run; and runtime configuration, which determines *what* to run (and *where* and *how* to run it) when a task is ready. In this example the `[scheduling]` section defines a single task called `hello` that triggers immediately when the suite starts up. When the task finishes the suite shuts down. That this is a *dependency graph* will be more obvious when more tasks are added. Under the `[runtime]` section the `script` item defines a simple inlined implementation for `hello`: it sleeps for ten seconds, then prints `Hello World!`, and exits. This ends up in a *job script* generated by cylc to encapsulate the task (below) and, thanks to some defaults designed to allow quick prototyping of new suites, it is submitted to run as a background job on the suite host. In fact cylc even provides a default task implementation that makes the entire `[runtime]` section technically optional:

```
[meta]
    title = "The minimal complete runnable cylc suite"
[scheduling]
    [[dependencies]]
        graph = "foo"
# (actually, 'title' is optional too ... and so is this comment)
```

(the resulting *dummy task* just prints out some identifying information and exits).

7.10 Editing Suites

The text editor invoked by cylc on suite definitions is determined by cylc site and user global config files, as shown above in 7.2. Check that you have renamed the tutorial examples suites as described just above and open the *Hello World* suite definition in your text editor:

```
$ cylc edit tut/oneoff/basic # in-terminal
$ cylc edit -g tut/oneoff/basic & # or GUI
```

Alternatively, start gcylc on the suite:

```
$ gcylc tut/oneoff/basic &
```

and choose *Suite* → *Edit* from the menu.


The editor will be invoked from within the suite definition directory for easy access to other suite files (in this case there are none). There are syntax highlighting control files for several text editors under `/path/to/cylc/conf/`; see in-file comments for installation instructions.

7.11 Running Suites

7.11.1 CLI

Run `tut/oneoff/basic` using the `cylc run` command. As a suite runs detailed timestamped information is written to a *suite log* and progress can be followed with `cylc`'s suite monitoring tools (below). By default a suite server program daemonizes after printing a short message so that you can exit the terminal or even log out without killing the suite:

```
$ cylc run tut/oneoff/basic
```



```

The Cylc Suite Engine [7.0.0]
Copyright (C) 2008-2018 NIWA

This program comes with ABSOLUTELY NO WARRANTY;
see 'cylc warranty'. It is free software, you
are welcome to redistribute it under certain
conditions; see 'cylc conditions'.

*** listening on nwp-1:43027 ***

To view suite server program contact information:
$ cylc get-suite-contact tut/oneoff/basic

Other ways to see if the suite is still running:
$ cylc scan -n '\btut/oneoff/basic\b' nwp-1
$ cylc ping -v --host=nwp-1 tut/oneoff/basic
$ ps h -opid,args 123456 # on nwp-1
```

If you're quick enough (this example only takes 10-15 seconds to run) the `cylc scan` command will detect the running suite:

```
$ cylc scan
tut/oneoff/basic oliverh@nwp-1:43027
```

Note you can use the `--no-detach` and `--debug` options to `cylc-run` to prevent the suite from daemonizing (i.e. to make it stay attached to your terminal until it exits).

When a task is ready `cylc` generates a *job script* to run it, by default as a background jobs on the suite host. The job process ID is captured, and job output is directed to log files in standard locations under the suite run directory.

Log file locations relative to the suite run directory look like `job/1/hello/01/` where the first digit is the *cycle point* of the task `hello` (for non-cycling tasks this is just '1'); and the final `01` is the *submit number* (so that job logs do not get overwritten if a job is resubmitted for any reason).

The suite shuts down automatically once all tasks have succeeded.

7.11.2 GUI

The `cylc` GUI can start and stop suites, or (re)connect to suites that are already running:

```
$ cylc gui tut/oneoff/basic &
```

Use the tool bar *Play* button, or the *Control* → *Run* menu item, to run the suite again. You may want to alter the suite definition slightly to make the task take longer to run. Try right-clicking on the `hello` task to view its output logs. The relative merits of the three *suite views* - dot, text, and graph - will be more apparent later when we have more tasks. Closing the GUI does not affect the suite itself.

7.12 Discovering Running Suites

Suites that are currently running can be detected with command line or GUI tools:

```
# list currently running suites and their port numbers:
$ cylc scan
tut/oneoff/basic oliverh@nwp-1:43001

# GUI summary view of running suites:
$ cylc gscan &
```

The scan GUI is shown in Figure 13; clicking on a suite in it opens gcylc.

7.13 Task Identifiers

At run time, task instances are identified by *name*, which is determined entirely by the suite definition, and a *cycle point* which is usually a date-time or an integer:

```
foo.20100808T00Z    # a task with a date-time cycle point
bar.1              # a task with an integer cycle point (could be non-cycling)
```

Non-cycling tasks usually just have the cycle point 1, but this still has to be used to target the task instance with cylc commands.

7.14 Job Submission: How Tasks Are Executed

```
suite: tut/oneoff/jobsub
```

Task *job scripts* are generated by cylc to wrap the task implementation specified in the suite definition (environment, script, etc.) in error trapping code, messaging calls to report task progress back to the suite server program, and so forth. Job scripts are written to the *suite job log directory* where they can be viewed alongside the job output logs. They can be accessed at run time by right-clicking on the task in the cylc GUI, or printed to the terminal:

```
$ cylc cat-log tut/oneoff/basic hello.1
```

This command can also print the suite log (and stdout and stderr for suites in daemon mode) and task stdout and stderr logs (see `cylc cat-log --help`). A new job script can also be generated on the fly for inspection:

```
$ cylc jobscript tut/oneoff/basic hello.1
```

Take a look at the job script generated for `hello.1` during the suite run above. The custom scripting should be clearly visible toward the bottom of the file.

The `hello` task in the first tutorial suite defaults to running as a background job on the suite host. To submit it to the Unix `at` scheduler instead, configure its job submission settings as in `tut/oneoff/jobsub`:

```
[runtime]
[[hello]]
    script = "sleep 10; echo Hello World!"
[[[job]]]
    batch system = at
```

Run the suite again after checking that `atd` is running on your system.

Cylc supports a number of different batch systems. Tasks submitted to external batch queuing systems like `at`, `PBS`, `SLURM`, `Moab`, or `LoadLeveler`, are displayed as *submitted* in the cylc GUI until they start executing.

- For more on task job scripts, see [10.1](#).
- For more on batch systems, see [11.1](#).

7.15 Locating Suite And Task Output

If the `--no-detach` option is not used, suite stdout and stderr will be directed to the suite run directory along with the time-stamped suite log file, and task job scripts and job logs (task stdout and stderr). The default suite run directory location is `$HOME/cylc-run`:

```
$ tree $HOME/cylc-run/tut/oneoff/basic/
|-- .service                # location of run time service files
|   |-- contact             # detail on how to contact the running suite
|   |-- db                  # private suite run database
|   |-- passphrase         # passphrase for client authentication
|   |-- source              # symbolic link to source directory
|   |-- ssl.cert            # SSL certificate for the suite server
|   |-- ssl.pem             # SSL private key
|-- cylc-suite.db           # back compat symlink to public suite run database
|-- share                   # suite share directory (not used in this example)
|-- work                    # task work space (sub-dirs are deleted if not used)
|   |-- 1                   # task cycle point directory (or 1)
|   |   |-- hello          # task work directory (deleted if not used)
|-- log                     # suite log directory
|   |-- db                  # public suite run database
|   |-- job                 # task job log directory
|   |   |-- 1              # task cycle point directory (or 1)
|   |   |   |-- hello      # task name
|   |   |   |   |-- 01      # task submission number
|   |   |   |   |   |-- job # task job script
|   |   |   |   |   |-- job-activity.log # task job activity log
|   |   |   |   |   |-- job.err          # task stderr log
|   |   |   |   |   |-- job.out          # task stdout log
|   |   |   |   |   |-- job.status       # task status file
|   |   |   |-- NN -> 01 # symlink to latest submission number
|   |   |-- suite        # suite server log directory
|   |   |   |-- err      # suite server stderr log (daemon mode only)
|   |   |   |-- out      # suite server stdout log (daemon mode only)
|   |   |   |-- log      # suite server event log (timestamped info)
```

The suite run database files, suite environment file, and task status files are used internally by cylc. Tasks execute in private `work/` directories that are deleted automatically if empty when the task finishes. The suite `share/` directory is made available to all tasks (by `$CYLC_SUITE_SHARE_DIR`) as a common share space. The task submission number increments from 1 if a task retries; this is used as a sub-directory of the log tree to avoid overwriting log files from earlier job submissions.

The top level run directory location can be changed in site and user config files if necessary, and the suite share and work locations can be configured separately because of the potentially larger disk space requirement.

Task job logs can be viewed by right-clicking on tasks in the gcylc GUI (so long as the task proxy is live in the suite), manually accessed from the log directory (of course), or printed to the terminal with the `cylc cat-log` command:

```
# suite logs:
$ cylc cat-log      tut/oneoff/basic          # suite event log
```

```
$ cylc cat-log -o tut/oneoff/basic # suite stdout log
$ cylc cat-log -e tut/oneoff/basic # suite stderr log
# task logs:
$ cylc cat-log      tut/oneoff/basic hello.1 # task job script
$ cylc cat-log -o tut/oneoff/basic hello.1 # task stdout log
$ cylc cat-log -e tut/oneoff/basic hello.1 # task stderr log
```

- For a web-based interface to suite and task logs (and much more), see *Rose* in 13.
- For more on environment variables supplied to tasks, such as `$CYLC_SUITE_SHARE_DIR`, see 9.4.7.

7.16 Remote Tasks

```
suite: tut/oneoff/remote
```

The `hello` task in the first two tutorial suites defaults to running on the suite host. To make it run on a remote host instead change its runtime configuration as in `tut/oneoff/remote`:

```
[runtime]
  [[hello]]
    script = "sleep 10; echo Hello World!"
    [[remote]]
      host = server1.niwa.co.nz
```

In general, a *task remote* is a user account, other than the account running the suite server program, where a task job is submitted to run. It can be on the same machine running the suite or on another machine.

A task remote account must satisfy several requirements:

- Non-interactive ssh must be enabled from the account running the suite server program to the account for submitting (and managing) the remote task job.
- Network settings must allow communication *back* from the remote task job to the suite, either by network ports or ssh, unless the last-resort one way *task polling* communication method is used.
- Cylc must be installed and runnable on the task remote account. Other software dependencies like graphviz are not required there.
- Any files needed by a remote task must be installed on the task host. In this example there is nothing to install because the implementation of `hello` is inlined in the suite definition and thus ends up entirely contained within the task job script.

If your username is different on the task host, you can add a `User` setting for the relevant host in your `~/.ssh/config`. If you are unable to do so, the `[[remote]]` section also supports an `owner=username` item.

If you configure a task account according to the requirements cylc will invoke itself on the remote account (with a login shell by default) to create log directories, transfer any essential service files, send the task job script over, and submit it to run there by the configured batch system.

Remote task job logs are saved to the suite run directory on the task remote, not on the account running the suite. They can be retrieved by right-clicking on the task in the GUI, or to have cylc pull them back to the suite account automatically do this:

```
[runtime]
  [[hello]]
    script = "sleep 10; echo Hello World!"
```

```
[[[remote]]]
    host = server1.niwa.co.nz
    retrieve job logs = True
```

This suite will attempt to `rsync` job logs from the remote host each time a task job completes.

Some batch systems have considerable delays between the time when the job completes and when it writes the job logs in its normal location. If this is the case, you can configure an initial delay and retry delays for job log retrieval by setting some delays. E.g.:

```
[runtime]
[[[hello]]]
    script = "sleep 10; echo Hello World!"
[[[remote]]]
    host = server1.niwa.co.nz
    retrieve job logs = True
    # Retry after 10 seconds, 1 minute and 3 minutes
    retrieve job logs retry delays = PT10S, PT1M, PT3M
```

Finally, if the disk space of the suite host is limited, you may want to set `[[[remote]]]retrieve job logs max size=SIZE`. The value of `SIZE` can be anything that is accepted by the `--max-size=SIZE` option of the `rsync` command. E.g.:

```
[runtime]
[[[hello]]]
    script = "sleep 10; echo Hello World!"
[[[remote]]]
    host = server1.niwa.co.nz
    retrieve job logs = True
    # Don't get anything bigger than 10MB
    retrieve job logs max size = 10M
```

It is worth noting that `cylc` uses the existence of a job's `job.out` or `job.err` in the local file system to indicate a successful job log retrieval. If `retrieve job logs max sizeSIZE=` is set and both `job.out` and `job.err` are bigger than `SIZE` then `cylc` will consider the retrieval as failed. If retry delays are specified, this will trigger some useless (but harmless) retries. If this occurs regularly, you should try the following:

- Reduce the verbosity of `STDOUT` or `STDERR` from the task.
- Redirect the verbosity from `STDOUT` or `STDERR` to an alternate log file.
- Adjust the size limit with tolerance to the expected size of `STDOUT` or `STDERR`.
- For more on remote tasks see [9.4.9](#)
- For more on task communications, see [12.6](#).
- For more on suite passphrases and authentication, see [7.5](#) and [12.9](#).

7.17 Task Triggering

```
suite: tut/oneoff/goodbye
```

To make a second task called `goodbye` trigger after `hello` finishes successfully, return to the original example, `tut/oneoff/basic`, and change the suite graph as in `tut/oneoff/goodbye`:

```
[scheduling]
[[[dependencies]]]
    graph = "hello => goodbye"
```

or to trigger it at the same time as `hello`,

```
[scheduling]
  [[dependencies]]
    graph = "hello & goodbye"
```

and configure the new task's behaviour under `[runtime]`:

```
[runtime]
  [[goodbye]]
    script = "sleep 10; echo Goodbye World!"
```

Run `tut/oneoff/goodbye` and check the output from the new task:

```
$ cat ~/cylc-run/tut/oneoff/goodbye/log/job/1/goodbye/01/job.out
# or
$ cylc cat-log -o tut/oneoff/goodbye goodbye.1
JOB SCRIPT STARTING
cylc (scheduler - 2014-08-14T15:09:30+12): goodbye.1 started at 2014-08-14T15:09:30+12
cylc Suite and Task Identity:
  Suite Name   : tut/oneoff/goodbye
  Suite Host   : oliverh-34403dl.niwa.local
  Suite Port   : 43001
  Suite Owner  : oliverh
  Task ID      : goodbye.1
  Task Host    : nwp-1
  Task Owner   : oliverh
  Task Try No.: 1

Goodbye World!
cylc (scheduler - 2014-08-14T15:09:40+12): goodbye.1 succeeded at 2014-08-14T15:09:40+12
JOB SCRIPT EXITING (TASK SUCCEEDED)
```

7.17.1 Task Failure And Suicide Triggering

```
suite: tut/oneoff/suicide
```

Task names in the graph string can be qualified with a state indicator to trigger off task states other than success:

```
graph = ""
a => b          # trigger b if a succeeds
c:submit => d    # trigger d if c submits
e:finish => f    # trigger f if e succeeds or fails
g:start  => h    # trigger h if g starts executing
i:fail   => j    # trigger j if i fails
""
```

A common use of this is to automate recovery from known modes of failure:

```
graph = "goodbye:fail => really_goodbye"
```

i.e. if task `goodbye` fails, trigger another task that (presumably) really says goodbye.

Failure triggering generally requires use of *suicide triggers* as well, to remove the recovery task if it isn't required (otherwise it would hang about indefinitely in the waiting state):

```
[scheduling]
  [[dependencies]]
    graph = ""hello => goodbye
              goodbye:fail => really_goodbye
              goodbye => !really_goodbye # suicide"""
```

This means if `goodbye` fails, trigger `really_goodbye`; and otherwise, if `goodbye` succeeds, remove `really_goodbye` from the suite.

Try running `tut/oneoff/suicide`, which also configures the `hello` task's runtime to make it fail, to see how this works.

- For more on suite dependency graphs see [9.3](#).
- For more on task triggering see [9.3.5](#).

7.18 Runtime Inheritance

```
suite: tut/oneoff/inherit
```

The `[runtime]` section is actually a *multiple inheritance* hierarchy. Each subsection is a *namespace* that represents a task, or if it is inherited by other namespaces, a *family*. This allows common configuration to be factored out of related tasks very efficiently.

```
[meta]
  title = "Simple runtime inheritance example"
[scheduling]
  [[dependencies]]
    graph = "hello => goodbye"
[runtime]
  [[root]]
    script = "sleep 10; echo $GREETING World!"
  [[hello]]
    [[environment]]
      GREETING = Hello
  [[goodbye]]
    [[environment]]
      GREETING = Goodbye
```

The `[root]` namespace provides defaults for all tasks in the suite. Here both tasks inherit `script` from `root`, which they customize with different values of the environment variable `$GREETING`. Note that inheritance from `root` is implicit; from other parents an explicit `inherit = PARENT` is required, as shown below.

- For more on runtime inheritance, see [9.4](#).

7.19 Triggering Families

```
suite: tut/oneoff/ftrigger1
```

Task families defined by runtime inheritance can also be used as shorthand in graph trigger expressions. To see this, consider two “greeter” tasks that trigger off another task `foo`:

```
[scheduling]
  [[dependencies]]
    graph = "foo => greeter_1 & greeter_2"
```

If we put the common greeting functionality of `greeter_1` and `greeter_2` into a special `GREETERS` family, the graph can be expressed more efficiently like this:

```
[scheduling]
  [[dependencies]]
    graph = "foo => GREETERS"
```

i.e. if `foo` succeeds, trigger all members of `GREETERS` at once. Here's the full suite with runtime hierarchy shown:

```
[meta]
  title = "Triggering a family of tasks"
[scheduling]
  [[dependencies]]
    graph = "foo => GREETERS"
[runtime]
  [[root]]
    pre-script = "sleep 10"
  [[foo]]
    # empty (creates a dummy task)
  [[GREETERS]]
    script = "echo $GREETING World!"
  [[greeter_1]]
    inherit = GREETERS
    [[environment]]
      GREETING = Hello
  [[greeter_2]]
    inherit = GREETERS
    [[environment]]
      GREETING = Goodbye
```

(Note that we recommend given ALL-CAPS names to task families to help distinguish them from task names. However, this is just a convention).

Experiment with the `tut/oneoff/ftrigger1` suite to see how this works.

7.20 Triggering Off Of Families

```
suite: tut/oneoff/ftrigger2
```

Tasks (or families) can also trigger *off* other families, but in this case we need to specify what the trigger means in terms of the upstream family members. Here's how to trigger another task `bar` if all members of `GREETERS` succeed:

```
[scheduling]
  [[dependencies]]
    graph = ""foo => GREETERS
      GREETERS:succeed-all => bar""
```

Verbose validation in this case reports:

```
$ cylc val -v tut/oneoff/ftrigger2
...
Graph line substitutions occurred:
  IN: GREETERS:succeed-all => bar
  OUT: greeter_1:succeed & greeter_2:succeed => bar
...
```

Cylc ignores family member qualifiers like `succeed-all` on the right side of a trigger arrow, where they don't make sense, to allow the two graph lines above to be combined in simple cases:

```
[scheduling]
  [[dependencies]]
    graph = "foo => GREETERS:succeed-all => bar"
```

Any task triggering status qualified by `-all` or `-any`, for the members, can be used with a family trigger. For example, here's how to trigger `bar` if all members of `GREETERS` finish (succeed or fail) and any of them then succeed:

```
[scheduling]
  [[dependencies]]
    graph = ""foo => GREETERS
      GREETERS:finish-all & GREETERS:succeed-any => bar""
```

(use of `GREETERS:succeed-any` by itself here would trigger `bar` as soon as any one member of `GREETERS` completed successfully). Verbose validation now begins to show how family triggers can simplify complex graphs, even for this tiny two-member family:

```
$ cylc val -v tut/oneoff/ftrigger2
...
Graph line substitutions occurred:
  IN: GREETERS:finish-all & GREETERS:succeed-any => bar
  OUT: ( greeter_1:succeed | greeter_1:fail ) & \
        ( greeter_2:succeed | greeter_2:fail ) & \
        ( greeter_1:succeed | greeter_2:succeed ) => bar
...
```

Experiment with `tut/oneoff/ftrigger2` to see how this works.

- For more on family triggering, see [9.3.5.9](#).

7.21 Suite Visualization

You can style dependency graphs with an optional `[visualization]` section, as shown in `tut/oneoff/ftrigger2`:

```
[visualization]
  default node attributes = "style=filled"
  [[node attributes]]
    foo = "fillcolor=#6789ab", "color=magenta"
    GREETERS = "fillcolor=#ba9876"
    bar = "fillcolor=#89ab67"
```

To display the graph in an interactive viewer:

```
$ cylc graph tut/oneoff/ftrigger2 & # dependency graph
$ cylc graph -n tut/oneoff/ftrigger2 & # runtime inheritance graph
```

It should look like Figure 15 (with the `GREETERS` family node expanded on the right).

Graph styling can be applied to entire families at once, and custom “node groups” can also be defined for non-family groups.

7.22 External Task Scripts

```
suite: tut/oneoff/external
```

The tasks in our examples so far have all had inlined implementation, in the suite definition, but real tasks often need to call external commands, scripts, or executables. To try this, let’s return to the basic Hello World suite and cut the implementation of the task `hello` out to a file `hello.sh` in the suite bin directory:

```
#!/bin/sh

set -e

GREETING=${GREETING:-Goodbye}
echo "$GREETING World! from $0"
```

Make the task script executable, and change the `hello` task runtime section to invoke it:




Figure 15: The *tut/oneoff/trigger2* dependency and runtime inheritance graphs

```
[meta]
  title = "Hello World! from an external task script"
[scheduling]
  [[dependencies]]
    graph = "hello"
[runtime]
  [[hello]]
    pre-script = sleep 10
    script = hello.sh
    [[[environment]]]
      GREETING = Hello
```

If you run the suite now the new greeting from the external task script should appear in the `hello` task stdout log. This works because `cylc` automatically adds the suite bin directory to `$PATH` in the environment passed to tasks via their job scripts. To execute scripts (etc.) located elsewhere you can refer to the file by its full file path, or set `$PATH` appropriately yourself (this could be done via `$HOME/.profile`, which is sourced at the top of the task job script, or in the suite definition itself).

Note the use of `set -e` above to make the script abort on error. This allows the error trapping code in the task job script to automatically detect unforeseen errors.

7.23 Cycling Tasks

```
suite: tut/cycling/one
```

So far we've considered non-cycling tasks, which finish without spawning a successor.

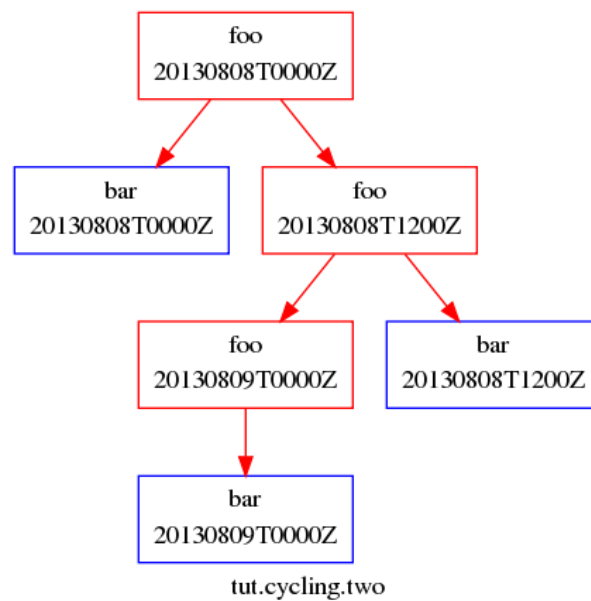
Cycling is based around iterating through date-time or integer sequences. A cycling task may run at each cycle point in a given sequence (cycle). For example, a sequence might be a set of date-times every 6 hours starting from a particular date-time. A cycling task may run for each date-time item (cycle point) in that sequence.

There may be multiple instances of this type of task running in parallel, if the opportunity arises and their dependencies allow it. Alternatively, a sequence can be defined with only one valid cycle point - in that case, a task belonging to that sequence may only run once.

Open the `tut/cycling/one` suite:

```
[meta]
  title = "Two cycling tasks, no inter-cycle dependence"
[cylc]
  UTC mode = True
[scheduling]
  initial cycle point = 20130808T00
  final cycle point = 20130812T00
  [[dependencies]]
    [[[T00,T12]]] # 00 and 12 hours UTC every day
    graph = "foo => bar"
[visualization]
  initial cycle point = 20130808T00
  final cycle point = 20130809T00
  [[node attributes]]
    foo = "color=red"
    bar = "color=blue"
```

The difference between cycling and non-cycling suites is all in the `[scheduling]` section, so we will leave the `[runtime]` section alone for now (this will result in cycling dummy tasks). Note that the graph is now defined under a new section heading that makes each task under it have

Figure 17: The `tut/cycling/two` suite

For any given cycle point in the sequence defined by the cycling graph section heading, `bar` triggers off `foo` as before, but now `foo` triggers off its own previous instance `foo[-PT12H]`. Date-time offsets in inter-cycle triggers are expressed as ISO 8601 intervals (12 hours in this case). Figure 17 shows how this connects the cycling graph sections together.

Experiment with this suite to see how inter-cycle triggers work. Note that the first instance of `foo`, at suite start-up, will trigger immediately in spite of its inter-cycle trigger, because `cylc` ignores dependence on points earlier than the initial cycle point. However, the presence of an inter-cycle trigger usually implies something special has to happen at start-up. If a model depends on its own previous instance for restart files, for example, then some special process has to generate the initial set of restart files when there is no previous cycle point to do it. The following section shows one way to handle this in `cylc` suites.

7.23.3 Initial Non-Repeating (R1) Tasks

```
suite: tut/cycling/three
```

Sometimes we want to be able to run a task at the initial cycle point, but refrain from running it in subsequent cycles. We can do this by writing an extra set of dependencies that are only valid at a single date-time cycle point. If we choose this to be the initial cycle point, these will only apply at the very start of the suite.

The `cylc` syntax for writing this single date-time cycle point occurrence is `R1`, which stands for `R1/no-specified-date-time/no-specified-period`. This is an adaptation of part of the ISO 8601 date-time standard's recurrence syntax (`Rn/date-time/period`) with some special context information supplied by `cylc` for the `no-specified-*` data.

The `1` in the `R1` means run once. As we've specified no date-time, `Cylc` will use the initial cycle point date-time by default, which is what we want. We've also missed out specifying the

period - this is set by `cylc` to a zero amount of time in this case (as it never repeats, this is not significant).

For example, in `tut/cycling/three`:

```
[cylc]
cycle point time zone = +13
[scheduling]
initial cycle point = 20130808T00
final cycle point = 20130812T00
[[dependencies]]
[[R1]]
graph = "prep => foo"
[[T00,T12]]
graph = "foo[-PT12H] => foo => bar"
```

This is shown in Figure 18.

Note that the time zone has been set to `+1300` in this case, instead of UTC (`z`) as before. If no time zone or UTC mode was set, the local time zone of your machine will be used in the cycle points.

At the initial cycle point, `foo` will depend on `foo[-PT12H]` and also on `prep`:

```
prep.20130808T0000+13 & foo.20130807T1200+13 => foo.20130808T0000+13
```

Thereafter, it will just look like e.g.:

```
foo.20130808T0000+13 => foo.20130808T1200+13
```

However, in our initial cycle point example, the dependence on `foo.20130807T1200+13` will be ignored, because that task's cycle point is earlier than the suite's initial cycle point and so it cannot run. This means that the initial cycle point dependencies for `foo` actually look like:

```
prep.20130808T0000+13 => foo.20130808T0000+13
```

- `R1` tasks can also be used to make something special happen at suite shutdown, or at any single cycle point throughout the suite run. For a full primer on cycling syntax, see 9.3.4.6.

7.23.4 Integer Cycling


```
suite: tut/cycling/integer
```

Cylc can also do integer cycling for repeating workflows that are not date-time based.

Open the `tut/cycling/integer` suite, which is plotted in Figure 19.

```
[scheduling]
cycling mode = integer
initial cycle point = 1
final cycle point = 3
[[dependencies]]
[[R1]] # = R1/1/?
graph = start => foo
[[P1]] # = R/1/P1
graph = foo[-P1] => foo => bar
[[R2/P1]] # = R2/P1/3
graph = bar => stop

[visualization]
[[node attributes]]
start = "style=filled", "fillcolor=skyblue"
```


Figure 18: The `tut/cycling/three` suite

```

foo = "style=filled", "fillcolor=slategray"
bar = "style=filled", "fillcolor=seagreen3"
stop = "style=filled", "fillcolor=orangered"

```

The integer cycling notation is intended to look similar to the ISO 8601 date-time notation, but it is simpler for obvious reasons. The example suite illustrates two recurrence forms, `Rn/start-point/period` and `Rn/period/stop-point`, simplified somewhat using suite context information (namely the initial and final cycle points). The first form is used to run one special task called `start` at start-up, and for the main cycling body of the suite; and the second form to run another special task called `stop` in the final two cycles. The `P` character denotes period (interval) just like in the date-time notation. `R/1/P2` would generate the sequence of points `1,3,5,...`.

- For more on integer cycling, including a more realistic usage example see [9.3.4.8](#).

7.24 Jinja2

```
suite: tut/oneoff/jinja2
```

Cylc has built in support for the Jinja2 template processor, which allows us to embed code in suite definitions to generate the final result seen by cylc.

The `tut/oneoff/jinja2` suite illustrates two common uses of Jinja2: changing suite content or structure based on the value of a logical switch; and iteratively generating dependencies and runtime configuration for groups of related tasks:

```

#!jinja2

{% set MULTI = True %}
{% set N_GOODBYES = 3 %}

```

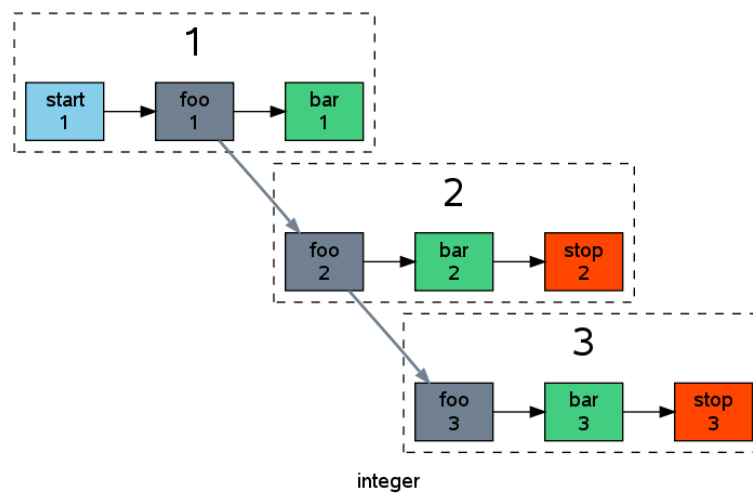


Figure 19: The tut/cycling/integer suite

```

[meta]
    title = "A Jinja2 Hello World! suite"
[scheduling]
    [[dependencies]]
{% if MULTI %}
    graph = "hello => BYE"
{% else %}
    graph = "hello"
{% endif %}

[runtime]
    [[hello]]
        script = "sleep 10; echo Hello World!"
{% if MULTI %}
    [[BYE]]
        script = "sleep 10; echo Goodbye World!"
        {% for I in range(0,N_GOODBYES) %}
        [[ goodbye_{{I}} ]]
        inherit = BYE
        {% endfor %}
{% endif %}
  
```

To view the result of Jinja2 processing with the Jinja2 flag `MULTI` set to `False`:

```
$ cylc view --jinja2 --stdout tut/oneoff/jinja2
```

```

[meta]
    title = "A Jinja2 Hello World! suite"
[scheduling]
    [[dependencies]]
        graph = "hello"
[runtime]
    [[hello]]
        script = "sleep 10; echo Hello World!"
  
```

And with `MULTI` set to `True`:

```
$ cylc view --jinja2 --stdout tut/oneoff/jinja2
```

```

[meta]
    title = "A Jinja2 Hello World! suite"
[scheduling]
    [[dependencies]]
        graph = "hello => BYE"
[runtime]
  
```

```

[[hello]]
    script = "sleep 10; echo Hello World!"
[[BYE]]
    script = "sleep 10; echo Goodbye World!"
[[ goodbye_0 ]]
    inherit = BYE
[[ goodbye_1 ]]
    inherit = BYE
[[ goodbye_2 ]]
    inherit = BYE

```

7.25 Task Retry On Failure

suite: tut/oneoff/retry

Tasks can be configured to retry a number of times if they fail. An environment variable `$CYLC_TASK_TRY_NUMBER` increments from 1 on each successive try, and is passed to the task to allow different behaviour on the retry:

```

[meta]
    title = "A task with automatic retry on failure"
[scheduling]
    [[dependencies]]
        graph = "hello"
[runtime]
    [[hello]]
        script = ""
sleep 10
if [[ $CYLC_TASK_TRY_NUMBER < 3 ]]; then
    echo "Hello ... aborting!"
    exit 1
else
    echo "Hello World!"
fi""
    [[[job]]]
        execution retry delays = 2*PT6S # retry twice after 6-second delays

```

If a task with configured retries fails, it goes into the *retrying* state until the next retry delay is up, then it resubmits. It only enters the *failed* state on a final definitive failure.

If a task with configured retries is *killed* (by `cylc kill` or via the GUI) it goes to the *held* state so that the operator can decide whether to release it and continue the retry sequence or to abort the retry sequence by manually resetting it to the *failed* state.

Experiment with `tut/oneoff/retry` to see how this works.

7.26 Other Users' Suites

If you have read access to another user's account (even on another host) it is possible to use `cylc monitor` to look at their suite's progress without full shell access to their account. To do this, you will need to copy their suite passphrase to

```
$HOME/.cylc/SUITE_OWNER@SUITE_HOST/SUITE_NAME/passphrase
```

(use of the host and owner names is optional here - see [12.9.2](#)) and also retrieve the port number of the running suite from:

```
~SUITE_OWNER/cylc-run/SUITE_NAME/.service/contact
```

Once you have this information, you can run

```
$ cylc monitor --user=SUITE_OWNER --port=SUITE_PORT SUITE_NAME
```

to view the progress of their suite.

Other suite-connecting commands work in the same way; see [12.11](#).

7.27 Other Things To Try

Almost every feature of cylc can be tested quickly and easily with a simple dummy suite. You can write your own, or start from one of the example suites in `/path/to/cylc/examples` (see use of `cylc import-examples` above) - they all run “out the box” and can be copied and modified at will.

- Change the suite runahead limit in a cycling suite.
- Stop a suite mid-run with `cylc stop`, and restart it again with `cylc restart`.
- Hold (pause) a suite mid-run with `cylc hold`, then modify the suite definition and `cylc reload` it before using `cylc release` to continue (you can also reload without holding).
- Use the gcylc View menu to show the task state color key and watch tasks in the `task-states` example evolve as the suite runs.
- Manually re-run a task that has already completed or failed, with `cylc trigger`.
- Use an *internal queue* to prevent more than an allotted number of tasks from running at once even though they are ready - see [12.17](#).
- Configure task event hooks to send an email, or shut the suite down, on task failure.

8 Suite Name Registration

Cylc commands target suites via their names, which are relative path names under the suite run directory (`~/cylc-run/` by default). Suites can be grouped together under sub-directories. E.g.:

```
$ cylc print -t nwp
nwp
|-oper
| |-region1  Local Model Region1      /home/oliverh/cylc-run/nwp/oper/region1
| |-region2  Local Model Region2      /home/oliverh/cylc-run/nwp/oper/region2
'-test
  '-region1  Local Model TEST Region1  /home/oliverh/cylc-run/nwp/test/region1
```

Suites can be pre-registered with a name using the `cylc register` command. This creates the essential directory structure for the suite, and generates some service files underneath it. Otherwise, `cylc run` will create these files on suite start up.

9 Suite Definition

Cylc suites are defined in structured, validated, *suite.rc* files that concisely specify the properties of, and the relationships between, the various tasks managed by the suite. This section of the User Guide deals with the format and content of the *suite.rc* file, including task definition. Task implementation - what’s required of the real commands, scripts, or programs that do the processing that the tasks represent - is covered in [10](#); and task job submission - how tasks are submitted to run - is in [11](#).

9.1 Suite Definition Directories

A *cylc suite definition directory* contains:

- **A `suite.rc` file:** this is the suite definition.
 - And any include-files used in it (see below; may be kept in sub-directories).
- **A `bin/` sub-directory** (optional)
 - For scripts and executables that implement, or are used by, suite tasks.
 - Automatically added to `$PATH` in task execution environments.
 - Alternatively, tasks can call external commands, scripts, or programs; or they can be scripted entirely within the `suite.rc` file.
- **A `lib/python/` sub-directory** (optional)
 - For custom job submission modules (see 11.7) and local Python modules imported by custom Jinja2 filters (see 9.7.2).
- **Any other sub-directories and files** - documentation, control files, etc. (optional)
 - Holding everything in one place makes proper suite revision control possible.
 - Portable access to files here, for running tasks, is provided through `$CYLC_SUITE_DEF_PATH` (see 9.4.7).
 - Ignored by `cylc`, but the entire suite definition directory tree is copied when you copy a suite using `cylc` commands.

A typical example:

```
/path/to/my/suite  # suite definition directory
  suite.rc         # THE SUITE DEFINITION FILE
  bin/             # scripts and executables used by tasks
    foo.sh
    bar.sh
    ...
  # (OPTIONAL) any other suite-related files, for example:
  inc/             # suite.rc include-files
    nwp-tasks.rc
    globals.rc
    ...
  doc/             # documentation
  control/         # control files
  ancil/           # ancillary files
  ...
```

9.2 Suite.rc File Overview

Suite.rc files are an extended-INI format with section nesting.

Embedded template processor expressions may also be used in the file, to programatically generate the final suite definition seen by `cylc`. Currently the Jinja2 template processor is supported (<http://jinja.pocoo.org/docs>); see 9.7 for examples. In the future `cylc` may provide a plug-in interface to allow use of other template engines too.

9.2.1 Syntax

The following defines legal `suite.rc` syntax:

- **Items** are of the form `item = value`.
- **[Section]** headings are enclosed in square brackets.

- **Sub-section** `[[nesting]]` is defined by repeated square brackets.
- Sections are **closed** by the next section heading.
- **Comments** (line and trailing) follow a hash character: `#`
- **List values** are comma-separated.
- **Single-line string values** can be single-, double-, or un-quoted.
- **Multi-line string values** are triple-quoted (using single or double quote characters).
- **Boolean values** are capitalized: `True`, `False`.
- **Leading and trailing whitespace** is ignored.
- **Indentation** is optional but should be used for clarity.
- **Continuation lines** follow a trailing backslash: `\`
- **Duplicate sections** add their items to those previously defined under the same section.
- **Duplicate items** override, *except for dependency `graph` strings, which are additive.*
- **Include-files** `%include inc/foo.rc` can be used as a verbatim inlining mechanism.

Suites that embed Jinja2 code (see 9.7) must process to raw suite.rc syntax.

9.2.2 Include-Files

Cylc has native support for suite.rc include-files, which may help to organize large suites. Inclusion boundaries are completely arbitrary - you can think of include-files as chunks of the suite.rc file simply cut-and-pasted into another file. Include-files may be included multiple times in the same file, and even nested. Include-file paths can be specified portably relative to the suite definition directory, e.g.:

```
# include the file $CYLC_SUITE_DEF_PATH/inc/foo.rc:
%include inc/foo.rc
```

9.2.2.1 Editing Temporarily Inlined Suites

Cylc's native file inclusion mechanism supports optional inlined editing:

```
$ cylc edit --inline SUITE
```

The suite will be split back into its constituent include-files when you exit the edit session. While editing, the inlined file becomes the official suite definition so that changes take effect whenever you save the file. See `cylc prep edit --help` for more information.

9.2.2.2 Include-Files via Jinja2

Jinja2 (9.7) also has template inclusion functionality.

9.2.3 Syntax Highlighting For Suite Definitions

Cylc comes with syntax files for a number of text editors:

```
$CYLC_DIR/conf/cylc.vim      # vim
$CYLC_DIR/conf/cylc-mode.el  # emacs
$CYLC_DIR/conf/cylc.lang     # gedit (and other gtksourceview programs)
$CYLC_DIR/conf/cylc.xml      # kate
```

Refer to comments at the top of each file to see how to use them.

9.2.4 Gross File Structure

Cylc suite.rc files consist of a suite title and description followed by configuration items grouped under several top level section headings:

- **[cylc]** - *non task-specific suite configuration*
- **[scheduling]** - *determines when tasks are ready to run*
 - tasks with special behaviour, e.g. clock-trigger tasks
 - the dependency graph, which defines the relationships between tasks
- **[runtime]** - *determines how, where, and what to execute when tasks are ready*
 - script, environment, job submission, remote hosting, etc.
 - suite-wide defaults in the *root* namespace
 - a nested family hierarchy with common properties inherited by related tasks
- **[visualization]** - suite graph styling

9.2.5 Validation

Cylc suite.rc files are automatically validated against a specification that defines all legal entries, values, options, and defaults. This detects formatting errors, typographic errors, illegal items and illegal values prior to run time. Some values are complex strings that require further parsing by cylc to determine their correctness (this is also done during validation). All legal entries are documented in the *Suite.rc Reference* ([A](#)).

The validator reports the line numbers of detected errors. Here's an example showing a section heading with a missing right bracket:

```
$ cylc validate my.suite
[[special tasks]
'Section bracket mismatch, line 19'
```

If the suite.rc file uses include-files `cylc view` will show an inlined copy of the suite with correct line numbers (you can also edit suites in a temporarily inlined state with `cylc edit --inline`).

Validation does not check the validity of chosen batch systems.

9.3 Scheduling - Dependency Graphs

The **[scheduling]** section of a suite.rc file defines the relationships between tasks in a suite - the information that allows cylc to determine when tasks are ready to run. The most important component of this is the suite dependency graph. Cylc graph notation makes clear textual graph representations that are very concise because sections of the graph that repeat at different hours of the day, say, only have to be defined once. Here's an example with dependencies that vary depending on the particular cycle point:

```
[scheduling]
  initial cycle point = 20200401
  final cycle point = 20200405
  [[dependencies]]
    [[[T00,T06,T12,T18]]] # validity (hours)
      graph = ""
A => B & C    # B and C trigger off A
A[-PT6H] => A # Model A restart trigger
      ""
    [[[T06,T18]]] # hours
      graph = "C => X"
```




Figure 20: Example Suite

Figure 20 shows the complete suite.rc listing alongside the suite graph. This is a complete, valid, runnable suite (it will use default task runtime properties such as `script`).

9.3.1 Graph String Syntax

Multiline graph strings may contain:

- **blank lines**
- **arbitrary white space**
- **internal comments:** following the `#` character
- **conditional task trigger expressions** - see below.

9.3.2 Interpreting Graph Strings

Suite dependency graphs can be broken down into pairs in which the left side (which may be a single task or family, or several that are conditionally related) defines a trigger for the task or family on the right. For instance the “word graph” *C triggers off B which triggers off A* can be deconstructed into pairs *C triggers off B* and *B triggers off A*. In this section we use only the default trigger type, which is to trigger off the upstream task succeeding; see 9.3.5 for other available triggers.

In the case of cycling tasks, the triggers defined by a graph string are valid for cycle points matching the list of hours specified for the graph section. For example this graph:

```

[scheduling]
  [[dependencies]]
    [[[T00,T12]]]
      graph = "A => B"

```

implies that B triggers off A for cycle points in which the hour matches 00 or 12.

To define inter-cycle dependencies, attach an offset indicator to the left side of a pair:


```
[scheduling]
  [[dependencies]]
    [[T00,T12]]
      graph = "A[-PT12H] => B"
```

This means B[time] triggers off A[time-PT12H] (12 hours before) for cycle points with hours matching 00 or 12. *time* is implicit because this keeps graphs clean and concise, given that the majority of tasks will typically depend only on others with the same cycle point. Cycle point offsets can only appear on the left of a pair, because a pairs define triggers for the right task at cycle point *time*. However, A => B[-PT6H], which is illegal, can be reformulated as a *future trigger* A[+PT6H] => B (see 9.3.5.11). It is also possible to combine multiple offsets within a cycle point offset e.g.

```
[scheduling]
  [[dependencies]]
    [[T00,T12]]
      graph = "A[-P1D-PT12H] => B"
```

This means that B[Time] triggers off A[time-P1D-PT12H] (1 day and 12 hours before).

Triggers can be chained together. This graph:

```
graph = """A => B  # B triggers off A
          B => C  # C triggers off B"""
```

is equivalent to this:

```
graph = "A => B => C"
```

Each trigger in the graph must be unique but the same task can appear in multiple pairs or chains. Separately defined triggers for the same task have an AND relationship. So this:

```
graph = """A => X  # X triggers off A
          B => X  # X also triggers off B"""
```

is equivalent to this:

```
graph = "A & B => X"  # X triggers off A AND B
```

In summary, the branching tree structure of a dependency graph can be partitioned into lines (in the suite.rc graph string) of pairs or chains, in any way you like, with liberal use of internal white space and comments to make the graph structure as clear as possible.

```
# B triggers if A succeeds, then C and D trigger if B succeeds:
graph = "A => B => C & D"
# which is equivalent to this:
graph = """A => B => C
          B => D"""
# and to this:
graph = """A => B => D
          B => C"""
# and to this:
graph = """A => B
          B => C
          B => D"""
# and it can even be written like this:
graph = """A => B # blank line follows:

          B => C # comment ...
          B => D"""
```

```
[meta]
  title = some one-off tasks
[scheduling]
  [[dependencies]]
    graph = "foo => bar & baz => qux"
```




Figure 21: One-off (Non-Cycling) Tasks.

9.3.2.1 Splitting Up Long Graph Lines

It is not necessary to use the general line continuation marker `\` to split long graph lines. Just break at dependency arrows, or split long chains into smaller ones. This graph:

```
graph = "A => B => C"
```

is equivalent to this:

```
graph = """A => B =>
           C"""
```

and also to this:

```
graph = """A => B
           B => C"""
```

9.3.3 Graph Types

A suite definition can contain multiple graph strings that are combined to generate the final graph.

9.3.3.1 One-off (Non-Cycling)

Figure 21 shows a small suite of one-off non-cycling tasks; these all share a single cycle point (1) and don't spawn successors (once they're all finished the suite just exits). The integer 1 attached to each graph node is just an arbitrary label here.

9.3.3.2 Cycling Graphs

For cycling tasks the graph section heading defines a sequence of cycle points for which the subsequent graph section is valid. Figure 22 shows a small suite of cycling tasks.

9.3.4 Graph Section Headings

Graph section headings define recurrence expressions, the graph within a graph section heading defines a workflow at each point of the recurrence. For example in the following scenario:

```
[scheduling]
  [[dependencies]]
    [[[ T06 ]]] # A graph section heading
    graph = foo => bar
```




Figure 22: Cycling Tasks.

`T06` means "Run every day starting at 06:00 after the initial cycle point". Cylc allows you to start (or end) at any particular time, repeat at whatever frequency you like, and even optionally limit the number of repetitions.

Graph section heading can also be used with integer cycling see 9.3.4.8.

9.3.4.1 Syntax Rules

Date-time cycling information is made up of a starting *date-time*, an *interval*, and an optional *limit*.

The time is assumed to be in the local time zone unless you set `[cylc]cycle point time zone` or `[cylc]UTC mode`. The calendar is assumed to be the proleptic Gregorian calendar unless you set `[scheduling]cycling mode`.

The syntax for representations is based on the ISO 8601 date-time standard. This includes the representation of *date-time*, *interval*. What we define for cylc's cycling syntax is our own optionally-heavily-condensed form of ISO 8601 recurrence syntax. The most common full form is: `R[limit?]/[date-time]/[interval]`. However, we allow omitting information that can be guessed from the context (rules below). This means that it can be written as:

```

R[limit?]/[date-time]
R[limit?]/[interval]
[date-time]/[interval]
R[limit?] # Special limit of 1 case
[date-time]
[interval]

```

with example graph headings for each form being:

```

[[[ R5/T00 ]]] # Run 5 times at 00:00 every day
[[[ R//PT1H ]]] # Run every hour (Note the R// is redundant)
[[[ 20000101T00Z/P1D ]]] # Run every day starting at 00:00 1st Jan 2000
[[[ R1 ]]] # Run once at the initial cycle point
[[[ 20000101T00Z ]]] # Run once at 00:00 1st Jan 2000
[[[ P1Y ]]] # Run every year

```

Note that `T00` is an example of `[date-time]`, with an inferred 1 day period and no limit.

Where some or all *date-time* information is omitted, it is inferred to be relative to the initial date-time cycle point. For example, `T00` by itself would mean the next occurrence of midnight that follows, or is, the initial cycle point. Entering `+PT6H` would mean 6 hours after the initial cycle point. Entering `-P1D` would mean 1 day before the initial cycle point. Entering no information for the *date-time* implies the initial cycle point date-time itself.

Where the *interval* is omitted and some (but not all) *date-time* information is omitted, it is inferred to be a single unit above the largest given specific *date-time* unit. For example, the largest given specific unit in `T00` is hours, so the inferred interval is 1 day (daily), `P1D`.

Where the *limit* is omitted, unlimited cycling is assumed. This will be bounded by the final cycle point's date-time if given.

Another supported form of ISO 8601 recurrence is: `R[limit?]/[interval]/[date-time]`. This form uses the *date-time* as the end of the cycling sequence rather than the start. For example, `R3/P5D/20140430T06` means:

```
20140420T06
20140425T06
20140430T06
```

This kind of form can be used for specifying special behaviour near the end of the suite, at the final cycle point's date-time. We can also represent this in cylc with a collapsed form:

```
R[limit?]/[interval]
R[limit?]/[date-time]
[interval]/[date-time]
```

So, for example, you can write:

```
[[[ R1//+POD ]]] # Run once at the final cycle point
[[[ R5/P1D ]]]   # Run 5 times, every 1 day, ending at the final
                  # cycle point
[[[ P2W/T00 ]]]  # Run every 2 weeks ending at 00:00 following
                  # the final cycle point
[[[ R//T00 ]]]   # Run every 1 day ending at 00:00 following the
                  # final cycle point
```

9.3.4.2 Referencing The Initial And Final Cycle Points

For convenience the caret and dollar symbols may be used as shorthand for the initial and final cycle points. Using this shorthand you can write:

```
[[[ R1/^+PT12H ]]] # Repeat once 12 hours after the initial cycle point
                  # R[limit]/[date-time]
                  # Equivalent to [[[ R1/+PT12H ]]]
[[[ R1/$ ]]]       # Repeat once at the final cycle point
                  # R[limit]/[date-time]
                  # Equivalent to [[[ R1//+POD ]]]
[[[ $-P2D/PT3H ]]] # Repeat 3 hourly starting two days before the
                  # [date-time]/[interval]
                  # final cycle point
```

Note that there can be multiple ways to write the same headings, for instance the following all run once at the final cycle point:

```
[[[ R1/POY ]]]      # R[limit]/[interval]
[[[ R1/POY/$ ]]]    # R[limit]/[interval]/[date-time]
[[[ R1/$ ]]]        # R[limit]/[date-time]
```

9.3.4.3 Excluding Dates

Date-times can be excluded from a recurrence by an exclamation mark for example `[[[PT1D!20000101]]]` means run daily except on the first of January 2000.

This syntax can be used to exclude one or multiple date-times from a recurrence. Multiple date-times are excluded using the syntax `[[[PT1D!(20000101,20000102,...)]]]`. All date-times listed

within the parentheses after the exclamation mark will be excluded. Note that the `~` and `$` symbols (shorthand for the initial and final cycle points) are both date-times so `[[[T12!$-PT1D]]]` is valid.

If using a run limit in combination with an exclusion, the heading might not run the number of times specified in the limit. For example in the following suite `foo` will only run once as its second run has been excluded.

```
[scheduling]
  initial cycle point = 20000101T00Z
  final cycle point = 20000105T00Z
  [[dependencies]]
    [[ R2/P1D!20000102 ]]]
    graph = foo
```

9.3.4.4 Advanced exclusion syntax

In addition to excluding isolated date-time points or lists of date-time points from recurrences, exclusions themselves may be date-time recurrence sequences. Any partial date-time or sequence given after the exclamation mark will be excluded from the main sequence.

For example, partial date-times can be excluded using the syntax:

```
[[[ PT1H ! T12 ]]]      # Run hourly but not at 12:00 from the initial
                        # cycle point.
[[[ T-00 ! (T00, T06, T12, T18) ]]] # Run hourly but not at 00:00, 06:00,
                        # 12:00, 18:00.
[[[ PT5M ! T-15 ]]]     # Run 5-minutely but not at 15 minutes past the
                        # hour from the initial cycle point.
[[[ T00 ! W-1T00 ]]]     # Run daily at 00:00 except on Mondays.
```

It is also valid to use sequences for exclusions. For example:

```
[[[ PT1H ! PT6H ]]]     # Run hourly from the initial cycle point but
                        # not 6-hourly from the initial cycle point.
[[[ T-00 ! PT6H ]]]     # Run hourly on the hour but not 6-hourly
                        # on the hour.
                        # Same as [[[ T-00 ! T-00/PT6H ]]] (T-00 context is implied)
                        # Same as [[[ T-00 ! (T00, T06, T12, T18) ]]]
                        # Same as [[[ PT1H ! (T00, T06, T12, T18) ]]] Initial cycle point dependent

[[[ T12 ! T12/P15D ]]]   # Run daily at 12:00 except every 15th day.

[[[ R/~P1H ! R5/20000101T00/P1D ]]] # Any valid recurrence may be used to
                        # determine exclusions. This example
                        # translates to: Repeat every hour from
                        # the initial cycle point, but exclude
                        # 00:00 for 5 days from the 1st January
                        # 2000.
```

You can combine exclusion sequences and single point exclusions within a comma separated list enclosed in parentheses:

```
[[[ T-00 ! (20000101T07, PT2H) ]]] # Run hourly on the hour but not at 07:00
                        # on the 1st Jan, 2000 and not 2-hourly
                        # on the hour.
```

9.3.4.5 How Multiple Graph Strings Combine

For a cycling graph with multiple validity sections for different hours of the day, the different sections *add* to generate the complete graph. Different graph sections can overlap (i.e. the same

hours may appear in multiple section headings) and the same tasks may appear in multiple sections, but individual dependencies should be unique across the entire graph. For example, the following graph defines a duplicate prerequisite for task C:

```
[scheduling]
  [[dependencies]]
    [[T00,T06,T12,T18]]
      graph = "A => B => C"
    [[T06,T18]]
      graph = "B => C => X"
      # duplicate prerequisite: B => C already defined at T06, T18
```

This does not affect scheduling, but for the sake of clarity and brevity the graph should be written like this:

```
[scheduling]
  [[dependencies]]
    [[T00,T06,T12,T18]]
      graph = "A => B => C"
    [[T06,T18]]
      # X triggers off C only at 6 and 18 hours
      graph = "C => X"
```

9.3.4.6 Advanced Examples

The following examples show the various ways of writing graph headings in cylc.

```
[[[ R1 ]]]          # Run once at the initial cycle point
[[[ P1D ]]]         # Run every day starting at the initial cycle point
[[[ PT5M ]]]        # Run every 5 minutes starting at the initial cycle
                    # point
[[[ T00/P2W ]]]     # Run every 2 weeks starting at 00:00 after the
                    # initial cycle point
[[[ +P5D/P1M ]]]    # Run every month, starting 5 days after the initial
                    # cycle point
[[[ R1/T06 ]]]      # Run once at 06:00 after the initial cycle point
[[[ R1/POY ]]]      # Run once at the final cycle point
[[[ R1/$ ]]]        # Run once at the final cycle point (alternative
                    # form)
[[[ R1/$-P3D ]]]    # Run once three days before the final cycle point
[[[ R3/T0830 ]]]    # Run 3 times, every day at 08:30 after the initial
                    # cycle point
[[[ R3/O1T00 ]]]    # Run 3 times, every month at 00:00 on the first
                    # of the month after the initial cycle point
[[[ R5/W-1/P1M ]]]  # Run 5 times, every month starting on Monday
                    # following the initial cycle point
[[[ T00!^ ]]]       # Run at the first occurrence of T00 that isn't the
                    # initial cycle point
[[[ PT1D!20000101 ]]] # Run every day days excluding 1st Jan 2000
[[[ 20140201T06/P1D ]]] # Run every day starting at 20140201T06
[[[ R1/min(T00,T06,T12,T18) ]]] # Run once at the first instance
                                # of either T00, T06, T12 or T18
                                # starting at the initial cycle
                                # point
```

9.3.4.7 Advanced Starting Up

Dependencies that are only valid at the initial cycle point can be written using the R1 notation (e.g. as in 7.23.3. For example:

```
[cylc]
  UTC mode = True
[scheduling]
  initial cycle point = 20130808T00
  final cycle point = 20130812T00
```

```

[cylc]
    UTC mode = True
[scheduling]
    initial cycle point = 20130808T00
    final cycle point = 20130812T00
    [[dependencies]]
        [[R1]]
            graph = "prep"
        [[R1/T00]]
            graph = "prep[~] => foo"
# ^ implies the initial cycle point:
        graph = "prep[~] => baz"
        [[T00]]
            graph = "foo[-P1D] => foo => bar"
        [[T12]]
            graph = "baz[-P1D] => baz => qux"
[visualization]
    initial cycle point = 20130808T00
    final cycle point = 20130810T00
    [[node attributes]]
        foo = "color=red"
        bar = "color=orange"
        baz = "color=green"
        qux = "color=blue"

```

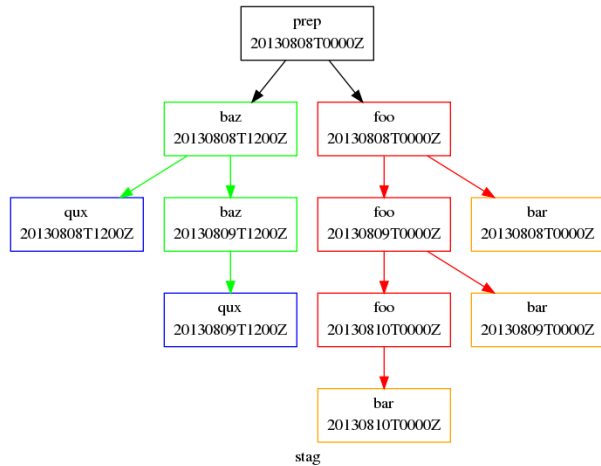


Figure 23: Staggered Start Suite

```

[[dependencies]]
    [[R1]]
        graph = "prep => foo"
    [[T00]]
        graph = "foo[-P1D] => foo => bar"

```

In the example above, `R1` implies `R1/20130808T00`, so `prep` only runs once at that cycle point (the initial cycle point). At that cycle point, `foo` will have a dependence on `prep` - but not at subsequent cycle points.

However, it is possible to have a suite that has multiple effective initial cycles - for example, one starting at `T00` and another starting at `T12`. What if they need to share an initial task?

Let's suppose that we add the following section to the suite example above:

```

[cylc]
    UTC mode = True
[scheduling]
    initial cycle point = 20130808T00
    final cycle point = 20130812T00
    [[dependencies]]
        [[R1]]
            graph = "prep => foo"
        [[T00]]
            graph = "foo[-P1D] => foo => bar"
        [[T12]]
            graph = "baz[-P1D] => baz => qux"

```

We'll also say that there should be a starting dependence between `prep` and our new task `baz` - but we still want to have a single `prep` task, at a single cycle.

We can write this using a special case of the `task[-interval]` syntax - if the interval is null, this implies the task at the initial cycle point.

For example, we can write our suite like [23](#).

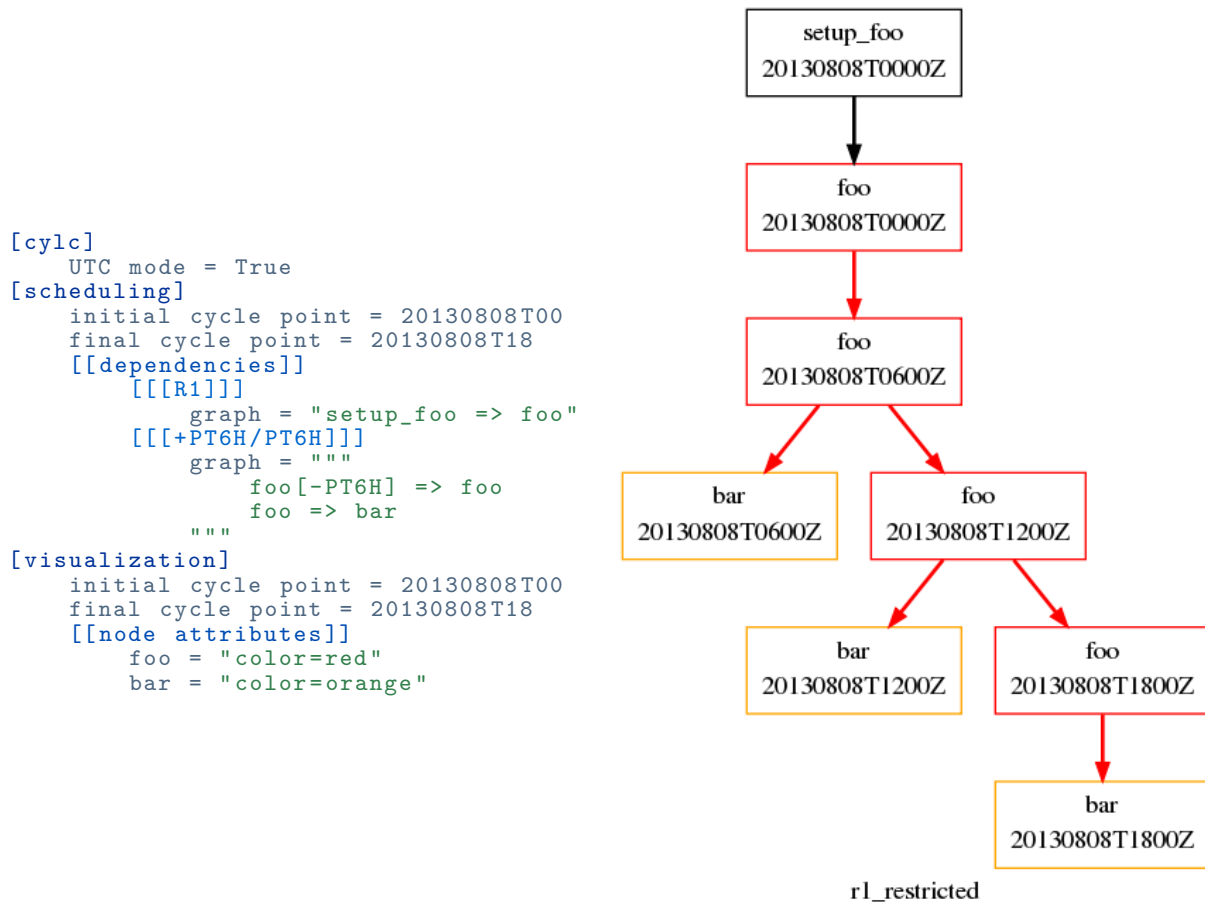


Figure 24: Restricted First Cycle Point Suite

This neatly expresses what we want - a task running at the initial cycle point that has one-off dependencies with other task sets at different cycles.

A different kind of requirement is displayed in Figure 24. Usually, we want to specify additional tasks and dependencies at the initial cycle point. What if we want our first cycle point to be entirely special, with some tasks missing compared to subsequent cycle points?

In Figure 24, `bar` will not be run at the initial cycle point, but will still run at subsequent cycle points. `[[+PT6H/PT6H]]` means start at `+PT6H` (6 hours after the initial cycle point) and then repeat every `PT6H` (6 hours).

Some suites may have staggered start-up sequences where different tasks need running once but only at specific cycle points, potentially due to differing data sources at different cycle points with different possible initial cycle points. To allow this `cylc` provides a `min()` function that can be used as follows:

```

[cylc]
    UTC mode = True
[scheduling]
    initial cycle point = 20100101T03
    [[dependencies]]
        [[R1/min(T00,T12)]]
            graph = "prep1 => foo"
        [[R1/min(T06,T18)]]
            graph = "prep2 => foo"
        [[T00,T06,T12,T18]]

```



```
graph = "foo => bar"
```

In this example the initial cycle point is 20100101T03, so the `prep1` task will run once at 20100101T12 and the `prep2` task will run once at 20100101T06 as these are the first cycle points after the initial cycle point in the respective `min()` entries.

9.3.4.8 Integer Cycling

In addition to non-repeating and date-time cycling workflows, `cylc` can do integer cycling for repeating workflows that are not date-time based.

To construct an integer cycling suite, set `[scheduling]cycling mode = integer`, and specify integer values for the initial and (optional) final cycle points. The notation for intervals, offsets, and recurrences (sequences) is similar to the date-time cycling notation, except for the simple integer values.

The full integer recurrence expressions supported are:


- `Rn/start-point/interval` #e.g. `R3/1/P2`
- `Rn/interval/end-point` #e.g. `R3/P2/9`

But, as for date-time cycling, sequence start and end points can be omitted where suite initial and final cycle points can be assumed. Some examples:

```
[[[ R1 ]]]      # Run once at the initial cycle point
                # (short for R1/initial-point/?)
[[[ P1 ]]]      # Repeat with step 1 from the initial cycle point
                # (short for R/initial-point/P1)
[[[ P5 ]]]      # Repeat with step 5 from the initial cycle point
                # (short for R/initial-point/P5)
[[[ R2//P2 ]]]  # Run twice with step 3 from the initial cycle point
                # (short for R2/initial-point/P2)
[[[ R/+P1/P2 ]]] # Repeat with step 2, from 1 after the initial cycle point
[[[ R2/P2 ]]]   # Run twice with step 2, to the final cycle point
                # (short for R2/P2/final-point)
[[[ R1/P0 ]]]   # Run once at the final cycle point
                # (short for R1/P0/final-point)
```

9.3.4.8.1 Example

The tutorial illustrates integer cycling in 7.23.4, and `$CYLC_DIR/examples/satellite/` is a self-contained example of a realistic use for integer cycling. It simulates the processing of incoming satellite data: each new dataset arrives after a random (as far as the suite is concerned) interval, and is labeled by an arbitrary (as far as the suite is concerned) ID in the filename. A task called `get_data` at the top of the repeating workflow waits on the next dataset and, when it finds one, moves it to a cycle-point-specific shared workspace for processing by the downstream tasks. When `get_data.1` finishes, `get_data.2` triggers and begins waiting for the next dataset at the same time as the downstream tasks in cycle point 1 are processing the first one, and so on. In this way multiple datasets can be processed at once if they happen to come in quickly. A single shutdown task runs at the end of the final cycle to collate results. The suite graph is shown in Figure 25.

Figure 25: The `examples/satellite` integer suite

9.3.4.8.2 Advanced Integer Cycling Syntax

The same syntax used to reference the initial and final cycle points (introduced in 9.3.4.2) for use with date-time cycling can also be used for integer cycling. For example you can write:

```
[[[ R1/^ ]]]      # Run once at the initial cycle point
[[[ R1/$ ]]]      # Run once at the final cycle point
[[[ R3/^/P2 ]]]   # Run three times with step two starting at the
                  # initial cycle point
```

Likewise the syntax introduced in 9.3.4.3 for excluding a particular point from a recurrence also works for integer cycling. For example:

```
[[[ R/P4!8 ]]]    # Run with step 4, to the final cycle point
                  # but not at point 8
[[[ R3/3/P2!5 ]]] # Run with step 2 from point 3 but not at
                  # point 5
[[[ R/+P1/P6!14 ]]] # Run with step 6 from 1 step after the
                  # initial cycle point but not at point 14
```

Multiple integer exclusions are also valid in the same way as the syntax in 9.3.4.3. Integer exclusions may be a list of single integer points, an integer sequence, or a combination of both:

```
[[[ R/P1!(2,3,7) ]]] # Run with step 1 to the final cycle point,
                    # but not at points 2, 3, or 7.
[[[ P1 ! P2 ]]]      # Run with step 1 from the initial to final
                    # cycle point, skipping every other step from
                    # the initial cycle point.
[[[ P1 ! +P1/P2 ]]]  # Run with step 1 from the initial cycle point,
                    # excluding every other step beginning one step
                    # after the initial cycle point.
[[[ P1 !(P2,6,8) ]]] # Run with step 1 from the initial cycle point,
                    # excluding every other step, and also excluding
                    # steps 6 and 8.
```

9.3.5 Trigger Types

Trigger type, indicated by *:type* after the upstream task (or family) name, determines what kind of event results in the downstream task (or family) triggering.

9.3.5.1 Success Triggers

The default, with no trigger type specified, is to trigger off the upstream task succeeding:

```
# B triggers if A SUCCEEDS:
graph = "A => B"
```

For consistency and completeness, however, the success trigger can be explicit:

```
# B triggers if A SUCCEEDS:
graph = "A => B"
# or:
graph = "A:succeed => B"
```

9.3.5.2 Failure Triggers

To trigger off the upstream task reporting failure:

```
# B triggers if A FAILS:
graph = "A:fail => B"
```

Suicide triggers can be used to remove task `B` here if `A` does not fail, see 9.3.5.8.

9.3.5.3 Start Triggers

To trigger off the upstream task starting to execute:

```
# B triggers if A STARTS EXECUTING:
graph = "A:start => B"
```

This can be used to trigger tasks that monitor other tasks once they (the target tasks) start executing. Consider a long-running forecast model, for instance, that generates a sequence of output files as it runs. A postprocessing task could be launched with a start trigger on the model (`model:start => post`) to process the model output as it becomes available. Note, however, that there are several alternative ways of handling this scenario: both tasks could be triggered at the same time (`foo => model & post`), but depending on external queue delays this could result in the monitoring task starting to execute first; or a different postprocessing task could be triggered off a message output for each data file (`model:out1 => post1` etc.; see 9.3.5.5), but this may not be practical if the number of output files is large or if it is difficult to add `cylc` messaging calls to the model.

9.3.5.4 Finish Triggers

To trigger off the upstream task succeeding or failing, i.e. finishing one way or the other:

```
# B triggers if A either SUCCEEDS or FAILS:
graph = "A | A:fail => B"
# or
graph = "A:finish => B"
```

9.3.5.5 Message Triggers

Tasks can also trigger off custom output messages. These must be registered in the [\[runtime\]](#) section of the emitting task, and reported using the `cylc message` command in task scripting. The graph trigger notation refers to the item name of the registered output message. The example suite `$CYLC_DIR/examples/message-triggers` illustrates message triggering.

```
[meta]
    title = "test suite for cylc-6 message triggers"

[scheduling]
    initial cycle point = 20140801T00
    final cycle point = 20141201T00
    [[dependencies]]
        [[P2M]]
            graph = """foo:out1 => bar
                        foo[-P2M]:out2 => baz"""

[runtime]
    [[foo]]
        script = """
sleep 5
cylc message "file 1 done"
sleep 10
cylc message "file 2 done"
sleep 10"""
        [[outputs]]
            out1 = "file 1 done"
            out2 = "file 2 done"
    [[bar, baz]]
        script = sleep 10
```

9.3.5.6 Job Submission Triggers

It is also possible to trigger off a task submitting, or failing to submit:

```
# B triggers if A submits successfully:
graph = "A:submit => B"
# D triggers if C fails to submit successfully:
graph = "C:submit-fail => D"
```

A possible use case for submit-fail triggers: if a task goes into the submit-failed state, possibly after several job submission retries, another task that inherits the same runtime but sets a different job submission method and/or host could be triggered to, in effect, run the same job on a different platform.

9.3.5.7 Conditional Triggers

AND operators (&) can appear on both sides of an arrow. They provide a concise alternative to defining multiple triggers separately:

```
# 1/ this:
graph = "A & B => C"
# is equivalent to:
graph = ""A => C
        B => C""
# 2/ this:
graph = "A => B & C"
# is equivalent to:
graph = ""A => B
        A => C""
# 3/ and this:
graph = "A & B => C & D"
# is equivalent to this:
graph = ""A => C
        B => C
        A => D
        B => D""
```

OR operators (|) which result in true conditional triggers, can only appear on the left,⁵

```
# C triggers when either A or B finishes:
graph = "A | B => C"
```

Forecasting suites typically have simple conditional triggering requirements, but any valid conditional expression can be used, as shown in Figure 26 (conditional triggers are plotted with open arrow heads).

9.3.5.8 Suicide Triggers

Suicide triggers take tasks out of the suite. This can be used for automated failure recovery. The suite.rc listing and accompanying graph in Figure 27 show how to define a chain of failure recovery tasks that trigger if they're needed but otherwise remove themselves from the suite (you can run the *AutoRecover.async* example suite to see how this works). The dashed graph edges ending in solid dots indicate suicide triggers, and the open arrowheads indicate conditional triggers as usual. Suicide triggers are ignored by default in the graph view, unless you toggle them on with *View -> Options -> Ignore Suicide Triggers*.

⁵An OR operator on the right doesn't make much sense: if "B or C" triggers off A, what exactly should cylec do when A finishes?

```
graph = """
# D triggers if A or (B and C) succeed
A | B & C => D
# just to align the two graph sections
D => W
# Z triggers if (W or X) and Y succeed
(W|X) & Y => Z
"""
```




Figure 26: Conditional triggers are plotted with open arrow heads.

```
[meta]
  title = automated failure recovery
  description = """
Model task failure triggers diagnosis
and recovery tasks, which take themselves
out of the suite if model succeeds. Model
post processing triggers off model OR
recovery tasks.
"""
[scheduling]
  [[dependencies]]
    graph = """
pre => model
model:fail => diagnose => recover
model => !diagnose & !recover
model | recover => post
"""
[runtime]
  [[model]]
    # UNCOMMENT TO TEST FAILURE:
    # script = /bin/false
```




Figure 27: Automated failure recovery via suicide triggers.

Note that multiple suicide triggers combine in the same way as other triggers, so this:

```
foo => !baz
bar => !baz
```

is equivalent to this:

```
foo & bar => !baz
```

i.e. both `foo` and `bar` must succeed for `baz` to be taken out of the suite. If you really want a task to be taken out if any one of several events occurs then be careful to write it that way:

```
foo | bar => !baz
```

A word of warning on the meaning of “bare suicide triggers”. Consider the following suite:

```
[scheduling]
  [[dependencies]]
    graph = "foo => !bar"
```

Task `bar` has a suicide trigger but no normal prerequisites (a suicide trigger is not a task triggering prerequisite, it is a task removal prerequisite) so this is entirely equivalent to:

```
[scheduling]
  [[dependencies]]
    graph = """
      foo & bar
      foo => !bar
    """
```

In other words both tasks will trigger immediately, at the same time, and then `bar` will be removed if `foo` succeeds.

If an active task proxy (currently in the submitted or running states) is removed from the suite by a suicide trigger, a warning will be logged.

9.3.5.9 Family Triggers

Families defined by the namespace inheritance hierarchy (9.4) can be used in the graph trigger whole groups of tasks at the same time (e.g. forecast model ensembles and groups of tasks for processing different observation types at the same time) and for triggering downstream tasks off families as a whole. Higher level families, i.e. families of families, can also be used, and are reduced to the lowest level member tasks. Note that tasks can also trigger off individual family members if necessary.

To trigger an entire task family at once:

```
[scheduling]
  [[dependencies]]
    graph = "foo => FAM"
[runtime]
  [[FAM]]      # a family (because others inherit from it)
  [[m1,m2]]    # family members (inherit from namespace FAM)
  inherit = FAM
```

This is equivalent to:

```
[scheduling]
  [[dependencies]]
    graph = "foo => m1 & m2"
[runtime]
  [[FAM]]
  [[m1,m2]]
  inherit = FAM
```

To trigger other tasks off families we have to specify whether to triggering off *all members* starting, succeeding, failing, or finishing, or off *any* members (doing the same). Legal family triggers are thus:

```
[scheduling]
  [[dependencies]]
    graph = ""
    # all-member triggers:
    FAM:start-all => one
    FAM:succeed-all => one
    FAM:fail-all => one
    FAM:finish-all => one
    # any-member triggers:
    FAM:start-any => one
    FAM:succeed-any => one
    FAM:fail-any => one
    FAM:finish-any => one
    ""
```

Here's how to trigger downstream processing after if one or more family members succeed, but only after all members have finished (succeeded or failed):

```
[scheduling]
  [[dependencies]]
    graph = ""
    FAM:finish-all & FAM:succeed-any => foo
    ""
```

9.3.5.10 Writing Efficient Inter-Family Triggering

While cylc allows writing dependencies between two families it is important to consider the number of dependencies this will generate. In the following example, each member of **FAM2** has dependencies pointing at all the members of **FAM1**.

```
[scheduling]
  [[dependencies]]
    graph = ""
    FAM1:succeed-any => FAM2
    ""
```

Expanding this out, you generate $N * M$ dependencies, where N is the number of members of **FAM1** and M is the number of members of **FAM2**. This can result in high memory use as the number of members of these families grows, potentially rendering the suite impractical for running on some systems.

You can greatly reduce the number of dependencies generated in these situations by putting dummy tasks in the graphing to represent the state of the family you want to trigger off. For example, if **FAM2** should trigger off any member of **FAM1** succeeding you can create a dummy task **FAM1_succeed_any_marker** and place a dependency on it as follows:

```
[scheduling]
  [[dependencies]]
    graph = ""
    FAM1:succeed-any => FAM1_succeed_any_marker => FAM2
    ""

[runtime]
...
  [[FAM1_succeed_any_marker]]
    script = true
...
```

This graph generates only $N + M$ dependencies, which takes significantly less memory and CPU to store and evaluate.

9.3.5.11 Inter-Cycle Triggers

Typically most tasks in a suite will trigger off others in the same cycle point, but some may depend on others with other cycle points. This notably applies to warm-cycled forecast models, which depend on their own previous instances (see below); but other kinds of inter-cycle dependence are possible too.⁶ Here's how to express this kind of relationship in cylc:

```
[dependencies]
  [[PT6H]]
    # B triggers off A in the previous cycle point
    graph = "A[-PT6H] => B"
```

inter-cycle and trigger type (or message trigger) notation can be combined:

```
# B triggers if A in the previous cycle point fails:
graph = "A[-PT6H]:fail => B"
```

At suite start-up inter-cycle triggers refer to a previous cycle point that does not exist. This does not cause the dependent task to wait indefinitely, however, because cylc ignores triggers that reach back beyond the initial cycle point. That said, the presence of an inter-cycle trigger does normally imply that something special has to happen at start-up. If a model depends on its own previous instance for restart files, for instance, then an initial set of restart files has to be generated somehow or the first model task will presumably fail with missing input files. There are several ways to handle this in cylc using different kinds of one-off (non-cycling) tasks that run at suite start-up. They are illustrated in the Tutorial (7.23.2); to summarize here briefly:

- R1 tasks (recommended):

```
[scheduling]
  [[dependencies]]
    [[R1]]
      graph = "prep"
    [[R1/T00,R1/T12]]
      graph = "prep[~] => foo"
    [[T00,T12]]
      graph = "foo[-PT12H] => foo => bar"
```

R1, or R1/date-time tasks are the recommended way to specify unusual start up conditions. They allow you to specify a clean distinction between the dependencies of initial cycles and the dependencies of the subsequent cycles.

Initial tasks can be used for real model cold-start processes, whereby a warm-cycled model at any given cycle point can in principle have its inputs satisfied by a previous instance of itself, or by an initial task with (nominally) the same cycle point.

In effect, the R1 task masquerades as the previous-cycle-point trigger of its associated cycling task. At suite start-up initial tasks will trigger the first cycling tasks, and thereafter the inter-cycle trigger will take effect.

If a task has a dependency on another task in a different cycle point, the dependency can be written using the [offset] syntax such as [-PT12H] in foo[-PT12H] => foo. This means that foo at the current cycle point depends on a previous instance of foo at 12 hours before the current cycle point. Unlike the cycling section headings (e.g. [[T00,T12]]), dependencies assume that relative times are relative to the current cycle point, not the initial cycle point.

⁶In NWP forecast analysis suites parts of the observation processing and data assimilation subsystem will typically also depend on model background fields generated by the previous forecast.

However, it can be useful to have specific dependencies on tasks at or near the initial cycle point. You can switch the context of the offset to be the initial cycle point by using the caret symbol: `^`.

For example, you can write `foo[^]` to mean `foo` at the initial cycle point, and `foo[^+PT6H]` to mean `foo` 6 hours after the initial cycle point. Usually, this kind of dependency will only apply in a limited number of cycle points near the start of the suite, so you may want to write it in `R1`-based cycling sections. Here's the example inter-cycle `R1` suite from above again.

```
[scheduling]
  [[dependencies]]
    [[R1]]
      graph = "prep"
    [[R1/T00,R1/T12]]
      graph = "prep[^] => foo"
    [[T00,T12]]
      graph = "foo[-PT12H] => foo => bar"
```

You can see there is a dependence on the initial `R1` task `prep` for `foo` at the first `T00` cycle point, and at the first `T12` cycle point. Thereafter, `foo` just depends on its previous (12 hours ago) instance.

Finally, it is also possible to have a dependency on a task at a specific cycle point.

```
[scheduling]
  [[dependencies]]
    [[R1/20200202]]
      graph = "baz[20200101] => qux"
```

However, in a long running suite, a repeating cycle should avoid having a dependency on a task with a specific cycle point (including the initial cycle point) - as it can currently cause performance issue. In the following example, all instances of `qux` will depend on `baz.20200101`, which will never be removed from the task pool.:

```
[scheduling]
  initial cycle point = 2010
  [[dependencies]]
    # Can cause performance issue!
    [[P1D]]
      graph = "baz[20200101] => qux"
```

9.3.5.12 Special Sequential Tasks

If a cycling task does not generate files required by its own successor, then successive instances can run in parallel if the opportunity arises. However, if such a task would interfere with its own siblings for internal reasons (e.g. use of a hardwired non cycle dependent temporary file or similar) then it can be forced to run sequentially. This can be done with explicit inter-cycle triggers in the graph:

```
[scheduling]
  [[dependencies]]
    [[T00,T12]]
      graph = "foo[-PT12H] => foo => bar"
```

or by declaring the task to be *sequential*:

```
[scheduling]
  [[special tasks]]
    sequential = foo
  [[dependencies]]
    [[T00,T12]]
      graph = "foo => bar"
```

The *sequential* declaration also results in each instance of `foo` triggering off its own predecessor, exactly as in the explicit version. The only difference is that implicit triggers will not appear in graph visualizations. The implicit version can also be considerably simpler when the task appears in multiple graph sections or in a non-uniform cycling sequence: this suite:

```
[scheduling]
[[special tasks]]
    sequential = foo
[[dependencies]]
    [[[T00,T03,T11]]]
        graph = "foo => bar"
```

is equivalent to this one:

```
[scheduling]
[[dependencies]]
    [[[T00,T03,T11]]]
        graph = "foo => bar"
    [[[T00]]]
        graph = "foo[-PT13H] => foo"
    [[[T03]]]
        graph = "foo[-PT3H] => foo"
    [[[T11]]]
        graph = "foo[-PT8H] => foo"
```

9.3.5.13 Future Triggers

Cylc also supports inter-cycle triggering off tasks “in the future” (with respect to cycle point):

```
[[dependencies]]
    [[[T00,T06,T12,T18]]]
        graph = ""
    # A runs in this cycle:
    #   A
    # B in this cycle triggers off A in the next cycle.
    #   A[PT6H] => B
    ""
```

(Recall that `A[t+PT6H]` can run before `B[t]` because tasks in cylc have private cycle points). Future triggers present a problem at the suite shutdown rather than at start-up. Here, `B` at the final cycle point wants to trigger off an instance of `A` that will never exist because it is beyond the suite stop point. Consequently cylc prevents tasks from spawning successors that depend on other tasks beyond the stop point.

9.3.5.14 Clock Triggers

In addition to depending on other tasks (and on external events - see 9.3.5.16) tasks can depend on the wall clock: specifically, they can trigger off a wall clock time expressed as an offset from their own cycle point:

```
[scheduling]
[[special tasks]]
    clock-trigger = foo(PT2H)
[[dependencies]]
    [[[T00]]]
        graph = foo
```

Here, `foo[2015-08-23T00]` would trigger (other dependencies allowing) when the wall clock time reaches `2015-08-23T02`. Clock-trigger offsets are normally positive, to trigger some time *after* the wall-clock time is equal to task cycle point.

Clock-triggers have no effect on scheduling if the suite is running sufficiently far behind the clock (e.g. after a delay, or because it is processing archived historical data) that the trigger times, which are relative to task cycle point, have already passed.

9.3.5.15 Clock-Expire Triggers

Tasks can be configured to *expire* - i.e. to skip job submission and enter the *expired* state - if they are too far behind the wall clock when they become ready to run, and other tasks can trigger off this. As a possible use case, consider a cycling task that copies the latest of a set of files to overwrite the previous set: if the task is delayed by more than one cycle there may be no point in running it because the freshly copied files will just be overwritten immediately by the next task instance as the suite catches back up to real time operation. Clock-expire tasks are configured like clock-trigger tasks, with a date-time offset relative to cycle point ([A.4.11.2](#)). The offset should be positive to make the task expire if the wall-clock time has gone beyond the cycle point. Triggering off an expired task typically requires suicide triggers to remove the workflow that runs if the task has not expired. Here a task called `copy` expires, and its downstream workflow is skipped, if it is more than one day behind the wall-clock (see also [examples/clock-expire](#)):

```
[cylc]
cycle point format = %Y-%m-%dT%H
[scheduling]
initial cycle point = 2015-08-15T00
[[special tasks]]
clock-expire = copy(-P1D)
[[dependencies]]
[[[P1D]]]
graph = ""
model[-P1D] => model => copy => proc
copy:expired => !proc""
```

9.3.5.16 External Triggers

In addition to depending on other tasks (and on the wall clock - see [9.3.5.14](#)) tasks can trigger off events reported by an external system. For example, an external process could detect incoming data on an ftp server, and then notify a suite containing a task to retrieve the new data for processing. This is an alternative to long-running tasks that poll for external events.

Note that cylc does not currently support triggering off “filesystem events” (e.g. `inotify` on Linux). However, external watcher processes can use filesystem events to detect triggering conditions, if that is appropriate, before notifying a suite with our general external event system.

The external triggering process must call `cylc ext-trigger` with the name of the target suite, the message that identifies this type of event to the suite, and an ID that distinguishes this particular event instance from others (the name of the target task or its current cycle point is not required). The event ID is just an arbitrary string to cylc, but it typically identifies the filename(s) of the latest dataset in some way. When the suite server program receives the external event notification it will trigger the next instance of any task waiting on that trigger (whatever its cycle point) and then broadcast (see [12.22](#)) the event ID to the cycle point of the triggered task as `$CYLC_EXT_TRIGGER_ID`. Downstream tasks with the same cycle point therefore know the new event ID too and can use it, if they need to, to identify the same new dataset. In this way a whole workflow can be associated with each new dataset, and multiple datasets can be processed in parallel if they happen to arrive in quick succession.

An externally-triggered task must register the event it waits on in the suite scheduling section:

```
# suite "sat-proc"
[scheduling]
  cycling mode = integer
  initial cycle point = 1
  [[special tasks]]
    external-trigger = get-data("new sat X data avail")
  [[dependencies]]
    [[[P1]]]
      graph = get-data => conv-data => products
```

Then, each time a new dataset arrives the external detection system should notify the suite like this:

```
$ cylc ext-trigger sat-proc "new sat X data avail" passX12334a
```

where “sat-proc” is the suite name and “passX12334a” is the ID string for the new event. The suite passphrase must be installed on triggering account.

Note that only one task in a suite can trigger off a particular external message. Other tasks can trigger off the externally triggered task as required, of course.

`$CYLC_DIR/examples/satellite/ext-triggers/suite.rc` is a working example of a simulated satellite processing suite.

External triggers are not normally needed in date-time cycling suites driven by real time data that comes in at regular intervals. In these cases a data retrieval task can be clock-triggered (and have appropriate retry intervals supplied) to submit at the expected data arrival time, so little time if any is wasted in polling. However, if the arrival time of the cycle-point-specific data is highly variable, external triggering may be used with the cycle point embedded in the message:

```
# suite "data-proc"
[scheduling]
  initial cycle point = 20150125T00
  final cycle point   = 20150126T00
  [[special tasks]]
    external-trigger = get-data("data arrived for $CYLC_TASK_CYCLE_POINT")
  [[dependencies]]
    [[[T00]]]
      graph = init-process => get-data => post-process
```

Once the variable-length waiting is finished, an external detection system should notify the suite like this:

```
$ cylc ext-trigger data-proc "data arrived for 20150126T00" passX12334a
```

where “data-proc” is the suite name, the cycle point has replaced the variable in the trigger string, and “passX12334a” is the ID string for the new event. The suite passphrase must be installed on the triggering account. In this case, the event will trigger for the second cycle point but not the first because of the cycle-point matching.

9.3.6 Model Restart Dependencies

Warm-cycled forecast models generate *restart files*, e.g. model background fields, to initialize the next forecast. This kind of dependence requires an inter-cycle trigger:

```
[scheduling]
  [[dependencies]]
    [[[T00,T06,T12,T18]]]
      graph = "A[-PT6H] => A"
```

If your model is configured to write out additional restart files to allow one or more cycle points to be skipped in an emergency *do not represent these potential dependencies in the suite graph* as they should not be used under normal circumstances. For example, the following graph would result in task `A` erroneously triggering off `A[T-24]` as a matter of course, instead of off `A[T-6]`, because `A[T-24]` will always be finished first:

```
[scheduling]
  [[dependencies]]
    [[[T00,T06,T12,T18]]]
      # DO NOT DO THIS (SEE ACCOMPANYING TEXT):
      graph = "A[-PT24H] | A[-PT18H] | A[-PT12H] | A[-PT6H] => A"
```

9.3.7 How The Graph Determines Task Instantiation

A graph trigger pair like `foo => bar` determines the existence and prerequisites (dependencies) of the downstream task `bar`, for the cycle points defined by the associated graph section heading. In general it does not say anything about the dependencies or existence of the upstream task `foo`. However *if the trigger has no cycle point offset* Cylc will infer that `bar` must exist at the same cycle points as `foo`. This is a convenience to allow this:

```
graph = "foo => bar"
```

to be written as shorthand for this:

```
graph = """foo
          foo => bar"""
```

(where `foo` by itself means `<nothing> => foo`, i.e. the task exists at these cycle points but has no prerequisites - although other prerequisites may be defined for it in other parts of the graph).

Cylc does not infer the existence of the upstream task in offset triggers like `foo[-P1D]=> bar` because, as explained in Section L.5, a typo in the offset interval should generate an error rather than silently creating tasks on an erroneous cycling sequence.

As a result you need to be careful not to define inter-cycle dependencies that cannot be satisfied at run time. Suite validation catches this kind of error if the existence of the cycle offset task is not defined anywhere at all:

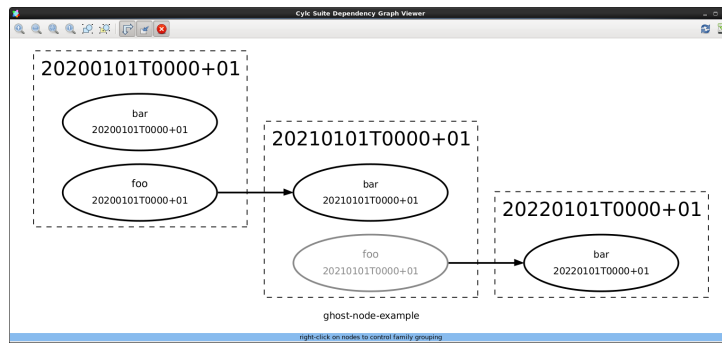
```
[scheduling]
  initial cycle point = 2020
  [[dependencies]]
    [[[P1Y]]]
      # ERROR
      graph = "foo[-P1Y] => bar"

$ cylc validate SUITE
'ERROR: No cycling sequences defined for foo'
```

To fix this, use another line in the graph to tell Cylc to define `foo` at each cycle point:

```
[scheduling]
  initial cycle point = 2020
  [[dependencies]]
    [[[P1Y]]]
      graph = """
        foo
        foo[-P1Y] => bar"""
```

But validation does not catch this kind of error if the offset task is defined only on a different cycling sequence:

Figure 28: Screenshot of `cylc graph` showing one task as a “ghost node”

```
[scheduling]
initial cycle point = 2020
[[dependencies]]
[[[P2Y]]]
graph = """foo
# ERROR
foo[-P1Y] => bar"""
```

This suite will validate OK, but it will stall at runtime with `bar` waiting on `foo[-P1Y]` at the intermediate years where it does not exist. The offset `[-P1Y]` is presumably an error (it should be `[-P2Y]`), or else another graph line is needed to generate `foo` instances on the yearly sequence:

```
[scheduling]
initial cycle point = 2020
[[dependencies]]
[[[P1Y]]]
graph = "foo"
[[[P2Y]]]
graph = "foo[-P1Y] => bar"
```

Similarly the following suite will validate OK, but it will stall at runtime with `bar` waiting on `foo[-P1Y]` in every cycle point, when only a single instance of it exists, at the initial cycle point:

```
[scheduling]
initial cycle point = 2020
[[dependencies]]
[[[R1]]]
graph = foo
[[[P1Y]]]
# ERROR
graph = foo[-P1Y] => bar
```

Note that `cylc graph` will display un-satisfiable inter-cycle dependencies as “ghost nodes”. Figure 28 is a screenshot of `cylc graph` displaying the above example with the un-satisfiable task (`foo`) displayed as a “ghost node”.

9.4 Runtime - Task Configuration

The `[runtime]` section of a suite definition configures what to execute (and where and how to execute it) when each task is ready to run, in a *multiple inheritance hierarchy* of *namespaces* culminating in individual tasks. This allows all common configuration detail to be factored out and defined in one place.

Any namespace can configure any or all of the items defined in the *Suite.rc Reference* (A).

Namespaces that do not explicitly inherit from others automatically inherit from the *root* namespace (below).

Nested namespaces define *task families* that can be used in the graph as convenient shorthand for triggering all member tasks at once, or for triggering other tasks off all members at once - see 9.3.5.9. Nested namespaces can be progressively expanded and collapsed in the dependency graph viewer, and in the gcylc graph and text views. Only the first parent of each namespace (as for single-inheritance) is used for suite visualization purposes.

9.4.1 Namespace Names

Namespace names may contain letters, digits, underscores, and hyphens.

Note that *task names need not be hardwired into task implementations* because task and suite identity can be extracted portably from the task execution environment supplied by the suite server program (9.4.7) - then to rename a task you can just change its name in the suite definition.

9.4.2 Root - Runtime Defaults

The root namespace, at the base of the inheritance hierarchy, provides default configuration for all tasks in the suite. Most root items are unset by default, but some have default values sufficient to allow test suites to be defined by dependency graph alone. The *script* item, for example, defaults to code that prints a message then sleeps for between 1 and 15 seconds and exits. Default values are documented with each item in A. You can override the defaults or provide your own defaults by explicitly configuring the root namespace.

9.4.3 Defining Multiple Namespaces At Once

If a namespace section heading is a comma-separated list of names then the subsequent configuration applies to each list member. Particular tasks can be singled out at run time using the `$CYLC_TASK_NAME` variable.

As an example, consider a suite containing an ensemble of closely related tasks that each invokes the same script but with a unique argument that identifies the calling task name:

```
[runtime]
  [[ENSEMBLE]]
    script = "run-model.sh $CYLC_TASK_NAME"
  [[m1, m2, m3]]
    inherit = ENSEMBLE
```

For large ensembles Jinja2 template processing can be used to automatically generate the member names and associated dependencies (see 9.7).

9.4.4 Runtime Inheritance - Single

The following listing of the *inherit.single.one* example suite illustrates basic runtime inheritance with single parents.


```
# SUITE.RC
[meta]
    title = "User Guide [runtime] example."
[cylc]
    required run mode = simulation # (no task implementations)
[scheduling]
    initial cycle point = 20110101T06
    final cycle point = 20110102T00
    [[dependencies]]
        [[TOO]]
            graph = ""foo => OBS
                OBS:succeed-all => bar""
[runtime]
    [[root]] # base namespace for all tasks (defines suite-wide defaults)
        [[job]]
            batch system = at
        [[environment]]
            COLOR = red
    [[OBS]] # family (inherited by land, ship); implicitly inherits root
        script = run-${CYLC_TASK_NAME}.sh
        [[environment]]
            RUNNING_DIR = $HOME/running/${CYLC_TASK_NAME}
    [[land]] # a task (a leaf on the inheritance tree) in the OBS family
        inherit = OBS
        [[meta]]
            description = land obs processing
    [[ship]] # a task (a leaf on the inheritance tree) in the OBS family
        inherit = OBS
        [[meta]]
            description = ship obs processing
        [[job]]
            batch system = loadleveler
        [[environment]]
            RUNNING_DIR = $HOME/running/ship # override OBS environment
            OUTPUT_DIR = $HOME/output/ship # add to OBS environment
    [[foo]]
        # (just inherits from root)

# The task [[bar]] is implicitly defined by its presence in the
# graph; it is also a dummy task that just inherits from root.
```

9.4.5 Runtime Inheritance - Multiple

If a namespace inherits from multiple parents the linear order of precedence (which namespace overrides which) is determined by the so-called *C3 algorithm* used to find the linear *method resolution order* for class hierarchies in Python and several other object oriented programming languages. The result of this should be fairly obvious for typical use of multiple inheritance in cylc suites, but for detailed documentation of how the algorithm works refer to the official Python documentation here: <http://www.python.org/download/releases/2.3/mro/>.

The *inherit.multi.one* example suite, listed here, makes use of multiple inheritance:

```
[meta]
    title = "multiple inheritance example"

    description = ""To see how multiple inheritance works:

% cylc list -tb[m] SUITE # list namespaces
% cylc graph -n SUITE # graph namespaces
% cylc graph SUITE # dependencies, collapse on first-parent namespaces

% cylc get-config --sparse --item [runtime]ops_s1 SUITE
% cylc get-config --sparse --item [runtime]var_p2 foo""

[scheduling]
    [[dependencies]]
        graph = "OPS:finish-all => VAR"
```

```

[runtime]
  [[root]]
  [[OPS]]
    script = echo "RUN: run-ops.sh"
  [[VAR]]
    script = echo "RUN: run-var.sh"
  [[SERIAL]]
    [[directives]]
      job_type = serial
  [[PARALLEL]]
    [[directives]]
      job_type = parallel
  [[ops_s1, ops_s2]]
    inherit = OPS, SERIAL

  [[ops_p1, ops_p2]]
    inherit = OPS, PARALLEL

  [[var_s1, var_s2]]
    inherit = VAR, SERIAL

  [[var_p1, var_p2]]
    inherit = VAR, PARALLEL

[visualization]
  # NOTE ON VISUALIZATION AND MULTIPLE INHERITANCE: overlapping
  # family groups can have overlapping attributes, so long as
  # non-conflicting attributes are used to style each group. Below,
  # for example, OPS tasks are filled green and SERIAL tasks are
  # outlined blue, so that ops_s1 and ops_s2 are green with a blue
  # outline. But if the SERIAL tasks are explicitly styled as "not
  # filled" (by setting "style=") this will override the fill setting
  # in the (previously defined and therefore lower precedence) OPS
  # group, making ops_s1 and ops_s2 unfilled with a blue outline.
  # Alternatively you can just create a manual node group for ops_s1
  # and ops_s2 and style them separately.
  [[node groups]]
    #(see comment above:)
    #serial_ops = ops_s1, ops_s2
  [[node attributes]]
    OPS = "style=filled", "fillcolor=green"
    SERIAL = "color=blue" #(see comment above:), "style="
    #(see comment above:)
    #serial_ops = "color=blue", "style=filled", "fillcolor=green"

```

`cylc get-suite-config` provides an easy way to check the result of inheritance in a suite. You can extract specific items, e.g.:

```

$ cylc get-suite-config --item '[runtime][var_p2]script' \
  inherit.multi.one
echo 'RUN: run-var.sh'

```

or use the `--sparse` option to print entire namespaces without obscuring the result with the dense runtime structure obtained from the root namespace:

```

$ cylc get-suite-config --sparse --item '[runtime]ops_s1' inherit.multi.one
script = echo 'RUN: run-ops.sh'
inherit = ['OPS', 'SERIAL']
[directives]
  job_type = serial

```

9.4.5.1 Suite Visualization And Multiple Inheritance

The first parent inherited by a namespace is also used as the collapsible family group when visualizing the suite. If this is not what you want, you can demote the first parent for visualization purposes, without affecting the order of inheritance of runtime properties:

```
[runtime]
  [[BAR]]
    # ...
  [[foo]]
    # inherit properties from BAR, but stay under root for visualization:
    inherit = None, BAR
```

9.4.6 How Runtime Inheritance Works

The linear precedence order of ancestors is computed for each namespace using the C3 algorithm. Then any runtime items that are explicitly configured in the suite definition are “inherited” up the linearized hierarchy for each task, starting at the root namespace: if a particular item is defined at multiple levels in the hierarchy, the level nearest the final task namespace takes precedence. Finally, root namespace defaults are applied for every item that has not been configured in the inheritance process (this is more efficient than carrying the full dense namespace structure through from root from the beginning).

9.4.7 Task Execution Environment

The task execution environment contains suite and task identity variables provided by the suite server program, and user-defined environment variables. The environment is explicitly exported (by the task job script) prior to executing the task `script` (see 11).

Suite and task identity are exported first, so that user-defined variables can refer to them. Order of definition is preserved throughout so that variable assignment expressions can safely refer to previously defined variables.

Additionally, access to `cylc` itself is configured prior to the user-defined environment, so that variable assignment expressions can make use of `cylc` utility commands:

```
[runtime]
  [[foo]]
    [[environment]]
      REFERENCE_TIME = $( cylc util cycletime --offset-hours=6 )
```

9.4.7.1 User Environment Variables

A task’s user-defined environment results from its inherited `[[environment]]` sections:

```
[runtime]
  [[root]]
    [[environment]]
      COLOR = red
      SHAPE = circle
  [[foo]]
    [[environment]]
      COLOR = blue # root override
      TEXTURE = rough # new variable
```

This results in a task `foo` with `SHAPE=circle`, `COLOR=blue`, and `TEXTURE=rough` in its environment.

9.4.7.2 Overriding Environment Variables

When you override inherited namespace items the original parent item definition is *replaced* by the new definition. This applies to all items including those in the environment sub-sections

which, strictly speaking, are not “environment variables” until they are written, post inheritance processing, to the task job script that executes the associated task. Consequently, if you override an environment variable you cannot also access the original parent value:

```
[runtime]
  [[FOO]]
    [[environment]]
      COLOR = red
  [[bar]]
    inherit = FOO
    [[environment]]
      tmp = $COLOR          # !! ERROR: $COLOR is undefined here
      COLOR = dark-$tmp     # !! as this overrides COLOR in FOO.
```

The compressed variant of this, `COLOR = dark-$COLOR`, is also in error for the same reason. To achieve the desired result you must use a different name for the parent variable:

```
[runtime]
  [[FOO]]
    [[environment]]
      FOO_COLOR = red
  [[bar]]
    inherit = FOO
    [[environment]]
      COLOR = dark-$FOO_COLOR # OK
```

9.4.7.3 Task Job Script Variables

These are variables that can be referenced (but should not be modified) in a task job script.

The task job script may export the following environment variables:

```
CYLC_DEBUG          # Debug mode, true or not defined
CYLC_DIR            # Location of cylc installation used
CYLC_VERSION        # Version of cylc installation used

CYLC_CYCLING_MODE   # Cycling mode, e.g. gregorian
CYLC_SUITE_FINAL_CYCLE_POINT # Final cycle point
CYLC_SUITE_INITIAL_CYCLE_POINT # Initial cycle point
CYLC_SUITE_NAME     # Suite name
CYLC_UTC            # UTC mode, True or False
CYLC_VERBOSE        # Verbose mode, True or False
TZ                  # Set to "UTC" in UTC mode or not defined

CYLC_SUITE_RUN_DIR  # Location of the suite run directory in
                   # job host, e.g. ~/cylc-run/foo
CYLC_SUITE_DEF_PATH # Location of the suite definition directory in
                   # job host, e.g. ~/cylc-run/foo
CYLC_SUITE_HOST     # Host running the suite process
CYLC_SUITE_OWNER    # User ID running the suite process
CYLC_SUITE_DEF_PATH_ON_SUITE_HOST # Location of the suite definition directory in
                   # suite host, e.g. ~/cylc-run/foo
CYLC_SUITE_SHARE_DIR # Suite (or task!) shared directory (see below)
CYLC_SUITE_WORK_DIR  # Suite work directory (see below)

CYLC_TASK_JOB       # Task job identifier expressed as
                   # CYCLE-POINT/TASK-NAME/SUBMIT-NUM
                   # e.g. 20110511T1800Z/t1/01
CYLC_TASK_CYCLE_POINT # Cycle point, e.g. 20110511T1800Z
CYLC_TASK_NAME       # Job's task name, e.g. t1
CYLC_TASK_SUBMIT_NUMBER # Job's submit number, e.g. 1,
                   # increments with every submit
CYLC_TASK_TRY_NUMBER # Number of execution tries, e.g. 1
                   # increments with automatic retry-on-fail
CYLC_TASK_ID         # Task instance identifier expressed as
                   # TASK-NAME.CYCLE-POINT
                   # e.g. t1.20110511T1800Z
```

```

CYLC_TASK_LOG_ROOT          # Location of the job file
                             # e.g. ~/cylc-run/foo/log/job/20110511T1800Z/t1/01/job
CYLC_TASK_WORK_DIR          # Location of task work directory (see below)
                             # e.g. ~/cylc-run/foo/work/20110511T1800Z/t1
CYLC_TASK_NAMESPACE_HIERARCHY # Linearised family namespace of the task,
                             # e.g. root postproc t1

CYLC_TASK_COMMS_METHOD      # Set to "ssh" if communication method is "ssh"
CYLC_TASK_SSH_LOGIN_SHELL   # With "ssh" communication, if set to "True",
                             # use login shell on suite host

```

There are also some global shell variables that may be defined in the task job script (but not exported to the environment). These include:

```

CYLC_FAIL_SIGNALS           # List of signals trapped by the error trap
CYLC_VACATION_SIGNALS       # List of signals trapped by the vacation trap
CYLC_SUITE_WORK_DIR_ROOT    # Root directory above the suite work directory
                             # in the job host
CYLC_TASK_MESSAGE_STARTED_PID # PID of "cylc message started" command
CYLC_TASK_WORK_DIR_BASE     # Alternate task work directory,
                             # relative to the suite work directory

```

9.4.7.4 Suite Share Directories

A *suite share directory* is created automatically under the suite run directory as a share space for tasks. The location is available to tasks as `$CYLC_SUITE_SHARE_DIR`. In a cycling suite, output files are typically held in cycle point sub-directories of the suite share directory.

The top level share and work directory (below) location can be changed (e.g. to a large data area) by a global config setting (see [B.9.1.2](#)).

9.4.7.5 Task Work Directories

Task job scripts are executed from within *work directories* created automatically under the suite run directory. A task can get its own work directory from `$CYLC_TASK_WORK_DIR` (or simply `$PWD` if it does not `cd` elsewhere at runtime). By default the location contains task name and cycle point, to provide a unique workspace for every instance of every task. This can be overridden in the suite definition, however, to get several tasks to share the same work directory (see [A.5.1.9](#)).

The top level work and share directory (above) location can be changed (e.g. to a large data area) by a global config setting (see [B.9.1.2](#)).

9.4.7.6 Environment Variable Evaluation

Variables in the task execution environment are not evaluated in the shell in which the suite is running prior to submitting the task. They are written in unevaluated form to the job script that is submitted by cylc to run the task ([10.1](#)) and are therefore evaluated when the task begins executing under the task owner account on the task host. Thus `$HOME`, for instance, evaluates at run time to the home directory of task owner on the task host.

9.4.8 How Tasks Get Access To The Suite Directory

Tasks can use `$CYLC_SUITE_DEF_PATH` to access suite files on the task host, and the suite bin directory is automatically added `$PATH`. If a remote suite definition directory is not specified the

local (suite host) path will be assumed with the local home directory, if present, swapped for literal `$HOME` for evaluation on the task host.

9.4.9 Remote Task Hosting

If a task declares an owner other than the suite owner and/or a host other than the suite host, cylc will use non-interactive ssh to execute the task on the `owner@host` account by the configured batch system:

```
[runtime]
  [[foo]]
    [[[remote]]]
      host = orca.niwa.co.nz
      owner = bob
    [[[job]]]
      batch system = pbs
```

For this to work:

- non-interactive ssh is required from the suite host to the remote task accounts.
- cylc must be installed on task hosts.
 - Optional software dependencies such as graphviz and Jinja2 are not needed on task hosts.
 - If polling task communication is used, there is no other requirement.
 - If SSH task communication is configured, non-interactive ssh is required from the task host to the suite host.
 - If (default) task communication is configured, the task host should have access to the port on the suite host.
- the suite definition directory, or some fraction of its content, can be installed on the task host, if needed.

To learn how to give remote tasks access to cylc, see [12.3](#).

Tasks running on the suite host under another user account are treated as remote tasks.

Remote hosting, like all namespace settings, can be declared globally in the root namespace, or per family, or for individual tasks.

9.4.9.1 Dynamic Host Selection

Instead of hardwiring host names into the suite definition you can specify a shell command that prints a hostname, or an environment variable that holds a hostname, as the value of the host config item. See [A.5.1.12.1](#).

9.4.9.2 Remote Task Log Directories

Task stdout and stderr streams are written to log files in a suite-specific sub-directory of the *suite run directory*, as explained in [11.2](#). For remote tasks the same directory is used, but *on the task host*. Remote task log directories, like local ones, are created on the fly, if necessary, during job submission.




Figure 29: Graphs of the *namespaces* example suite showing various states of expansion of the nested namespace family hierarchy, from all families collapsed (top left) through to all expanded (bottom right). This can also be done by right-clicking on tasks in the gcylc graph view.

9.5 Visualization

The visualization section of a suite definition is used to configure suite graphing, principally graph node (task) and edge (dependency arrow) style attributes. Tasks can be grouped for the purpose of applying common style attributes. See [A](#) for details.

9.5.1 Collapsible Families In Suite Graphs

```
[visualization]
    collapsed families = family1, family2
```

Nested families from the runtime inheritance hierarchy can be expanded and collapsed in suite graphs and the gcylc graph view. All families are displayed in the collapsed state at first, unless `[visualization]collapsed families` is used to single out specific families for initial collapsing.

In the gcylc graph view, nodes outside of the main graph (such as the members of collapsed families) are plotted as rectangular nodes to the right if they are doing anything interesting (submitted, running, failed).

Figure 29 illustrates successive expansion of nested task families in the *namespaces* example suite.

9.6 Parameterized Tasks

Cylc can automatically generate tasks and dependencies by expanding parameterized task names over lists of parameter values. Uses for this include:

- generating an ensemble of similar model runs

- generating chains of tasks to process similar datasets
- replicating an entire workflow, or part thereof, over several runs
- splitting a long model run into smaller steps or “chunks” (parameterized cycling)

Note that this can be done with Jinja2 loops too (Section 9.7) but parameterization is much cleaner (nested loops can seriously reduce the clarity of a suite definition).

9.6.1 Parameter Expansion

Parameter values can be lists of strings, or lists of integers and integer ranges (with inclusive bounds). Numeric values in a list of strings are considered strings. It is not possible to mix strings with integer ranges.

For example:

```
[cylc]
[[parameters]]
    # parameters: "ship", "buoy", "plane"
    # default task suffixes: _ship, _buoy, _plane
    obs = ship, buoy, plane

    # parameters: 1, 2, 3, 4, 5
    # default task suffixes: _run1, _run2, _run3, _run4, _run5
    run = 1..5

    # parameters: 1, 3, 5, 7, 9
    # default task suffixes: _idx1, _idx3, _idx5, _idx7, _idx9
    idx = 1..9..2

    # parameters: -11, -1, 9
    # default task suffixes: _idx-11, _idx-01, _idx+09
    idx = -11..9..10

    # parameters: 1, 3, 5, 10, 11, 12, 13
    # default task suffixes: _i01, _i03, _i05, _i10, _i11, _i12, _i13
    i = 1..5..2, 10, 11..13

    # parameters: "0", "1", "e", "pi", "i"
    # default task suffixes: _0, _1, _e, _pi, _i
    item = 0, 1, e, pi, i

    # ERROR: mix strings with int range
    p = one, two, 3..5
```

Then angle brackets denote use of these parameters throughout the suite definition. For the values above, this parameterized name:

```
model<run> # for run = 1..2
```

expands to these concrete task names:

```
model_run1, model_run2
```

and this parameterized name:

```
proc<obs> # for obs = ship, buoy, plane
```

expands to these concrete task names:

```
proc_ship, proc_buoy, proc_plane
```

By default, to avoid any ambiguity, the parameter name appears in the expanded task names for integer values, but not for string values. For example, `model_run1` for `run = 1`, but `proc_ship` for `obs = ship`. However, the default expansion templates can be overridden if need be:


```
[cylc]
[[parameters]]
    obs = ship, buoy, plane
    run = 1..5
[[parameter templates]]
    run = -R%(run)s # Make foo<run> expand to foo-R1 etc.
```

(See [A.3.12](#) for more on the string template syntax.)

Any number of parameters can be used at once. This parameterization:

```
model<run,obs> # for run = 1..2 and obs = ship, buoy, plane
```

expands to these task names:

```
model_run1_ship, model_run1_buoy, model_run1_plane,
model_run2_ship, model_run2_buoy, model_run2_plane
```

Here's a simple but complete example suite:

```
[cylc]
[[parameters]]
    run = 1..2
[scheduling]
[[dependencies]]
    graph = "prep => model<run>"
[runtime]
[[model<run>]]
    # ...
```

The result, post parameter expansion, is this:

```
[scheduling]
[[dependencies]]
    graph = "prep => model_run1 & model_run2"
[runtime]
[[model_run1]]
    # ...
[[model_run2]]
    # ...
```

Here's a more complex graph using two parameters ([runtime] omitted):

```
[cylc]
[[parameters]]
    run = 1..2
    mem = cat, dog
[scheduling]
[[dependencies]]
    graph = ""prep => init<run> => model<run,mem> =>
           post<run,mem> => wrap<run> => done"""
```

Figure 30 shows the result as visualized by `cylc graph`.

9.6.1.1 Zero-Padded Integer Values

Integer parameter values are given a default template for generating task suffixes that are zero-padded according to the longest size of their values. For example, the default template for `p = 9..10` would be `_p%(p)02d`, so that `foo<p>` would become `foo_p09`, `foo_p10`. If negative values are present in the parameter list, the default template will include the sign. For example, the default template for `p = -1..1` would be `_p%(p)+02d`, so that `foo<p>` would become `foo_p-1`, `foo_p+0`, `foo_p+1`.

To get thicker padding and/or alternate suffixes, use a template. E.g.:




Figure 30: Parameter expansion example.

```
[cylc]
[[parameters]]
    i = 1..9
    p = 3..14
[[parameter templates]]
    i = _i%(i)02d # suffixes = _i01, _i02, ..., _i09
    # A double-percent gives a literal percent character
    p = %%p%(p)03d # suffixes = %p003, %p004, ..., %p013, %p014
```

9.6.1.2 Parameters as Full Task Names

Parameter values can be used as full task names, but the default template should be overridden to remove the initial underscore. For example:

```
[cylc]
[[parameters]]
    i = 1..4
    obs = ship, buoy, plane
[[parameter templates]]
    i = i%(i)d # task name must begin with an alphabet
    obs = %(obs)s
[scheduling]
[[dependencies]]
    graph = """
foo => <i> # foo => i1 & i2 & i3 & i4
<obs> => bar # ship & buoy & plane => bar
"""
```

9.6.2 Passing Parameter Values To Tasks

Parameter values are passed as environment variables to tasks generated by parameter expansion. For example, if we have:

```
[cylc]
[[parameters]]
    obs = ship, buoy, plane
    run = 1..5
[scheduling]
[[dependencies]]
    graph = model<run,obs>
```

Then task `model_run2_ship` would get the following standard environment variables:

```
# In a job script of an instance of the "model_run2_ship" task:
export CYLC_TASK_PARAM_run="2"
export CYLC_TASK_PARAM_obs="ship"
```

These variables allow tasks to determine which member of a parameterized group they are, and so to vary their behaviour accordingly.

You can also define custom variables and string templates for parameter value substitution. For example, if we add this to the above configuration:

```
[runtime]
  [[model<run,obs>]]
    [[parameter environment templates]]
      MYNAME = %(obs)sy-mc%(obs)sface
      MYFILE = /path/to/run%(run)03d/%(obs)s
```

Then task `model_run2_ship` would get the following custom environment variables:

```
# In a job script of an instance of the "model_run2_ship" task:
export MYNAME=shipy-mcshipface
export MYFILE=/path/to/run002/ship
```

9.6.3 Selecting Specific Parameter Values

Specific parameter values can be singled out in the graph and under `[runtime]` with the notation `<p=5>` (for example). Here's how to make a special task trigger off just the first of a set of model runs:

```
[cylc]
  [[parameters]]
    run = 1..5
[scheduling]
  [[dependencies]]
    graph = """model<run> => post_proc<run> # general case
              model<run=1> => check_first_run # special case"""
[runtime]
  [[model<run>]]
    # config for all "model" runs...
  [[model<run=1>]]
    # special config (if any) for the first model run...
  #...
```

9.6.4 Selecting Partial Parameter Ranges

The parameter notation does not currently support partial range selection such as `foo<p=5..10>`, but you can achieve the same result by defining a second parameter that covers the partial range and giving it the same expansion template as the full-range parameter. For example:

```
[cylc]
  [[parameters]]
    run = 1..10 # 1, 2, ..., 10
    runx = 1..3 # 1, 2, 3
  [[parameter templates]]
    run = _R%(run)02d # _R01, _R02, ..., _R10
    runx = _R%(runx)02d # _R01, _R02, _R03
[scheduling]
  [[dependencies]]
    graph = """model<run> => post<run>
              model<runx> => checkx<runx>"""
[runtime]
  [[model<run>]]
    # ...
  #...
```

9.6.5 Parameter Offsets In The Graph

A negative offset notation `<NAME-1>` is interpreted as the previous value in the ordered list of parameter values, while a positive offset is interpreted as the next value. For example, to split a model run into multiple steps with each step depending on the previous one, either of these graphs:

```
graph = "model<run-1> => model<run>" # for run = 1, 2, 3
graph = "model<run> => model<run+1>" # for run = 1, 2, 3
```

expands to:

```
graph = ""model_run1 => model_run2
          model_run2 => model_run3""
# or equivalently:
graph = "model_run1 => model_run2 => model_run3"
```

And this graph:

```
graph = "proc<size-1> => proc<size>" # for size = small, big, huge
```

expands to:

```
graph = ""proc_small => proc_big
          proc_big => proc_huge""
# or equivalently:
graph = "proc_small => proc_big => proc_huge"
```

However, a quirk in the current system means that you should avoid mixing conditional logic in these statements. For example, the following will do the unexpected:

```
graph = foo<m-1> & baz => foo<m> # for m = cat, dog
```

currently expands to:

```
graph = foo_cat & baz => foo_dog
# when users may expect it to be:
#   graph = foo_cat => foo_dog
#   graph = baz => foo_cat & foo_dog
```

For the time being, writing out the logic explicitly will give you the correct graph.

```
graph = ""foo<m-1> => foo<m> # for m = cat, dog
          baz => foo<m>""
```

9.6.6 Task Families And Parameterization

Task family members can be generated by parameter expansion:

```
[runtime]
[[FAM]]
[[member<r>]]
  inherit = FAM
# Result: family FAM contains member_r1, member_r2, etc.
```

Family names can be parameterized too, just like task names:

```
[runtime]
[[RUN<r>]]
[[model<r>]]
  inherit = RUN<r>
[[post_proc<r>]]
  inherit = RUN<r>
# Result: family RUN_r1 contains model_r1 and post_proc_r1,
#         family RUN_r2 contains model_r2 and post_proc_r1, etc.
```

As described in Section 9.3.5.9 family names can be used to trigger all members at once:

```
graph = "foo => FAMILY"
```

or to trigger off all members:

```
graph = "FAMILY:succeed-all => bar"
```

or to trigger off any members:

```
graph = "FAMILY:succeed-any => bar"
```

If the members of `FAMILY` were generated with parameters, you can also trigger them all at once with parameter notation:

```
graph = "foo => member<m>"
```

Similarly, to trigger off all members:

```
graph = "member<m> => bar"
# (member<m>:fail etc., for other trigger types)
```

Family names are still needed in the graph, however, to succinctly express “succeed-any” triggering semantics, and all-to-all or any-to-all triggering:

```
graph = "FAM1:succeed-any => FAM2"
```

(Direct all-to-all and any-to-all family triggering is not recommended for efficiency reasons though - see Section 9.3.5.10).

For family *member-to-member* triggering use parameterized members. For example, if family `OBS_GET` has members `get<obs>` and family `OBS_PROC` has members `proc<obs>` then this graph:

```
graph = "get<obs> => proc<obs>" # for obs = ship, buoy, plane
```

expands to:

```
get_ship => proc_ship
get_buoy => proc_buoy
get_plane => proc_plane
```

9.6.7 Parameterized Cycling

Two ways of constructing cycling systems are described and contrasted in Section 5. For most purposes use of a proper *cycling workflow* is recommended, wherein `cylc` incrementally generates the date-time sequence and extends the workflow, potentially indefinitely, at run time. For smaller systems of finite duration, however, parameter expansion can be used to generate a sequence of pre-defined tasks as a proxy for cycling.

Here’s a cycling workflow of two-monthly model runs for one year, with previous-instance model dependence (e.g. for model restart files):

```
[scheduling]
  initial cycle point = 2020-01
  final cycle point = 2020-12
  [[dependencies]]
    [[R1]] # Run once, at the initial point.
      graph = "prep => model"
    [[P2M]] # Run at 2-month intervals between the initial and final points.
      graph = "model[-P2M] => model => post_proc & archive"
[runtime]
  [[model]]
    script = "run-model $CYLC_TASK_CYCLE_POINT"
```

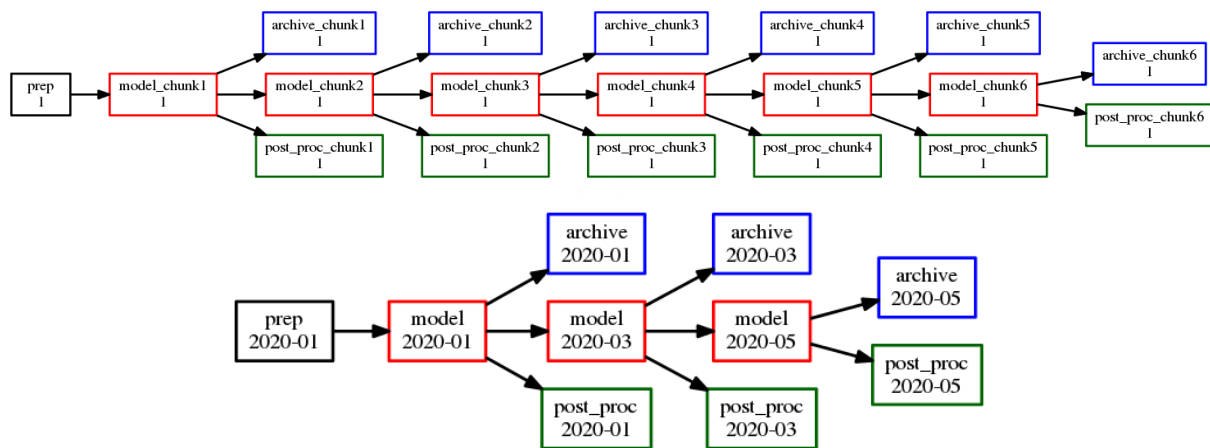


Figure 31: parameterized and cycling versions of the same workflow. The first three cycle points are shown in the cycling case. The parameterized case does not have “cycle points”.

And here’s how to do the same thing with parameterized tasks:

```
[cylc]
[[parameters]]
    chunk = 1..6
[scheduling]
    [[dependencies]]
        graph = """prep => model<chunk=1>
                    model<chunk-1> => model<chunk> =>
                    post_proc<chunk> & archive<chunk>"""
[runtime]
    [[_model<chunk>]]
        script = """
# Compute start date from chunk index and interval, then run the model.
INITIAL_POINT=2020-01
INTERVAL_MONTHS=2
OFFSET_MONTHS=(( (CYLC_TASK_PARAM_chunk - 1)*INTERVAL_MONTHS ))
OFFSET=P${OFFSET_MONTHS}M # e.g. P4M for chunk=3
run-model $(cylc cyclepoint --offset=$OFFSET $INITIAL_POINT)"""
```

The two workflows are shown together in Figure 31. They both achieve the same result, and both can include special tasks at the start, end, or anywhere in between. But as noted earlier the parameterized version has several disadvantages: it must be finite in extent and not too large; the date-time arithmetic has to be done by the user; and the full extent of the workflow will be visible at all times as the suite runs.

Here’s a yearly-cycling suite with four parameterized chunks in each cycle point:

```
[cylc]
[[parameters]]
    chunk = 1..4
[scheduling]
    initial cycle point = 2020-01
    [[dependencies]]
        [[P1Y]]
            graph = """model<chunk-1> => model<chunk>
                    model<chunk=4>[-P1Y] => model<chunk=1>"""
```

Note the inter-cycle trigger that connects the first chunk in each cycle point to the last chunk in the previous cycle point. Of course it would be simpler to just use 3-monthly cycling:

```
[scheduling]
    initial cycle point = 2020-01
    [[dependencies]]
```

```
[[[P3M]]]
graph = "model[-P3M] => model"
```

Here's a possible valid use-case for mixed cycling: consider a portable date-time cycling workflow of model jobs that can each take too long to run on some supported platforms. This could be handled without changing the cycling structure of the suite by splitting the run (at each cycle point) into a variable number of shorter steps, using more steps on less powerful hosts.

9.6.7.1 Cycle Point And Parameter Offsets At Start-Up

In cycling workflows `cylc` ignores anything earlier than the suite initial cycle point. So this graph:

```
graph = "model[-P1D] => model"
```

simplifies at the initial cycle point to this:

```
graph = "model"
```

Similarly, parameter offsets are ignored if they extend beyond the start of the parameter value list. So this graph:

```
graph = "model<chunk-1> => model<chunk>"
```

simplifies for `chunk=1` to this:

```
graph = "model_chunk1"
```

Note however that the initial cut-off applies to every parameter list, but only to cycle point sequences that start at the suite initial cycle point. Therefore it may be somewhat easier to use parameterized cycling if you need multiple date-time sequences *with different start points* in the same suite. We plan to allow this sequence-start simplification for any date-time sequence in the future, not just at the suite initial point, but it needs to be optional because delayed-start cycling tasks sometimes need to trigger off earlier cycling tasks.

9.7 Jinja2

This section needs to be revised - the Parameterized Task feature introduced in `cylc-6.11.0` (see Section 9.6) provides a cleaner way to auto-generate tasks without coding messy Jinja2 loops.

`Cylc` has built in support for the Jinja2 template processor in suite definitions. Jinja2 variables, mathematical expressions, loop control structures, conditional logic, etc., are automatically processed to generate the final suite definition seen by `cylc`.

The need for Jinja2 processing must be declared with a hash-bang comment as the first line of the suite.rc file:

```
#!jinja2
# ...
```

Potential uses for this include automatic generation of repeated groups of similar tasks and dependencies, and inclusion or exclusion of entire suite sections according to the value of a single flag. Consider a large complicated operational suite and several related parallel test suites with slightly different task content and structure (the parallel suites, for instance, might take certain large input files from the operation or the archive rather than downloading them




Figure 32: The Jinja2 ensemble example suite graph.

again) - these can now be maintained as a single master suite definition that reconfigures itself according to the value of a flag variable indicating the intended use.

Template processing is the first thing done on parsing a suite definition so Jinja2 expressions can appear anywhere in the file (inside strings and namespace headings, for example).

Jinja2 is well documented at <http://jinja.pocoo.org/docs>, so here we just provide an example suite that uses it. The meaning of the embedded Jinja2 code should be reasonably self-evident to anyone familiar with standard programming techniques.

The `jinja2.ensemble` example, graphed in Figure 32, shows an ensemble of similar tasks generated using Jinja2:

```

#!jinja2
{% set N_MEMBERS = 5 %}
[scheduling]
[[dependencies]]
graph = """{# generate ensemble dependencies #}
    {% for I in range( 0, N_MEMBERS ) %}
        foo => mem_{{ I }} => post_{{ I }} => bar
    {% endfor %}"""

```

Here is the generated suite definition, after Jinja2 processing:

```

#!jinja2
[scheduling]
[[dependencies]]
graph = """
    foo => mem_0 => post_0 => bar
    foo => mem_1 => post_1 => bar
    foo => mem_2 => post_2 => bar
    foo => mem_3 => post_3 => bar
    foo => mem_4 => post_4 => bar
    """

```

And finally, the `jinja2.cities` example uses variables, includes or excludes special cleanup tasks according to the value of a logical flag, and it automatically generates all dependencies and family relationships for a group of tasks that is repeated for each city in the suite. To add a

new city and associated tasks and dependencies simply add the city name to list at the top of the file. The suite is graphed, with the New York City task family expanded, in Figure 33.

```

#!Jinja2
[meta]
    title = "Jinja2 city suite example."
    description = ""
    Illustrates use of variables and math expressions, and programmatic
    generation of groups of related dependencies and runtime properties.""

{% set HOST = "SuperComputer" %}
{% set CITIES = 'NewYork', 'Philadelphia', 'Newark', 'Houston', 'SantaFe', 'Chicago' %}
{% set CITYJOBS = 'one', 'two', 'three', 'four' %}
{% set LIMIT_MINS = 20 %}

{% set CLEANUP = True %}

[scheduling]
    initial cycle point = 2011-08-08T12
    [[ dependencies ]]
{% if CLEANUP %}
    [[[T23]]]
        graph = "clean"
{% endif %}
    [[[T00,T12]]]
        graph = ""
            setup => get_lbc & get_ic # foo
{% for CITY in CITIES %} {# comment #}
    get_lbc => {{ CITY }}_one
    get_ic => {{ CITY }}_two
    {{ CITY }}_one & {{ CITY }}_two => {{ CITY }}_three & {{ CITY }}_four
{% if CLEANUP %}
    {{ CITY }}_three & {{ CITY }}_four => cleanup
{% endif %}
{% endfor %}
        ""

[runtime]
    [[on_{{ HOST }}]]
        [[remote]]
            host = {{ HOST }}
            # (remote cylc directory is set in site/user config for this host)
        [[[directives]]]
            wall_clock_limit = "00:{{ LIMIT_MINS|int() + 2 }}:00,00:{{ LIMIT_MINS }}:00"

{% for CITY in CITIES %}
    [[ {{ CITY }} ]]
        inherit = on_{{ HOST }}
{% for JOB in CITYJOBS %}
    [[ {{ CITY }}_{{ JOB }} ]]
        inherit = {{ CITY }}
{% endfor %}
{% endfor %}

[visualization]
    initial cycle point = 2011-08-08T12
    final cycle point = 2011-08-08T23
    [[node groups]]
        cleaning = clean, cleanup
    [[node attributes]]
        cleaning = 'style=filled', 'fillcolor=yellow'
        NewYork = 'style=filled', 'fillcolor=lightblue'

```

9.7.1 Accessing Environment Variables With Jinja2

This functionality is not provided by Jinja2 by default, but cylc automatically imports the user environment to the template in a dictionary structure called *environ*. A usage example:

```

#!Jinja2
#...

```




Figure 33: The Jinja2 cities example suite graph, with the New York City task family expanded.

```

[runtime]
[[root]]
[[[environment]]]
    SUITE_OWNER_HOME_DIR_ON_SUITE_HOST = {{environ['HOME']}}

```

This example emphasizes that *the environment is read on the suite host at the time the suite definition is parsed* - it is not, for instance, read at task run time on the task host.

9.7.2 Custom Jinja2 Filters

Jinja2 variable values can be modified by “filters”, using pipe notation. For example, the built-in `trim` filter strips leading and trailing white space from a string:

```

{% set MyString = "   dog   " %}
{{ MyString | trim() }} # "dog"

```

(See official Jinja2 documentation for available built-in filters.)

Cylc also supports custom Jinja2 filters. A custom filter is a single Python function in a source file with the same name as the function (plus “.py” extension) and stored in one of the following locations:

- `$CYLC_DIR/lib/Jinja2Filters/`
- `[suite definition directory]/Jinja2Filters/`
- `$HOME/.cylc/Jinja2Filters/`

In the filter function argument list, the first argument is the variable value to be “filtered”, and subsequent arguments can be whatever is needed. Currently there are two custom filters:

9.7.2.1 pad

The “pad” filter is for padding string values to some constant length with a fill character - useful for generating task names and related values in ensemble suites:

```

{% for i in range(0,100) %} # 0, 1, ..., 99
    {% set j = i | pad(2,'0') %}
    A_{{j}} # A_00, A_01, ..., A_99
{% endfor %}

```

9.7.2.2 strftime

The “strftime” filter can be used to format ISO8601 date-time strings using an strftime string.

```
{% set START_CYCLE = '10661004T08+01' %}
{{ START_CYCLE | strftime('%H') }} # 00
```

Examples:

- `{{START_CYCLE | strftime('%Y')}} - 1066`
- `{{START_CYCLE | strftime('%m')}} - 10`
- `{{START_CYCLE | strftime('%d')}} - 14`
- `{{START_CYCLE | strftime('%H:%M:%S %z')}} - 08:00:00 +01`

It is also possible to parse non-standard date-time strings by passing a strftime string as the second argument.

Examples:

- `{{'12,30,2000' | strftime('%m', '%m,%d,%Y')}} - 12`
- `{{'1066/10/14 08:00:00' | strftime('%Y%m%dT%H', '%Y/%m/%d %H:%M:%S')}} - 10661014T08`

9.7.3 Associative Arrays In Jinja2

Associative arrays (*dicts* in Python) can be very useful. Here’s an example, from

`$CYLC_DIR/examples/jinja2/dict:`

```
#!Jinja2
{% set obs_types = ['airs', 'iasi'] %}
{% set resource = { 'airs': 'ncpus=9', 'iasi': 'ncpus=20' } %}

[scheduling]
  [[dependencies]]
    graph = OBS
[runtime]
  [[OBS]]
    [[[job]]]
      batch system = pbs
    {% for i in obs_types %}
    [[ {{i}} ]]
      inherit = OBS
      [[[directives]]]
        -I = {{ resource[i] }}
    {% endfor %}
```

Here’s the result:

```
$ cylc get-suite-config -i [runtime][airs]directives SUITE
-I = ncpus=9
```

9.7.4 Jinja2 Default Values And Template Inputs

The values of Jinja2 variables can be passed in from the cylc command line rather than hardwired in the suite definition. Here’s an example, from

`$CYLC_DIR/examples/jinja2/defaults:`

```
#!Jinja2

[meta]
```

```

    title = "Jinja2 example: use of defaults and external input"

    description = """
The template variable FIRST_TASK must be given on the cylc command line
using --set or --set-file=FILE; two other variables, LAST_TASK and
N_MEMBERS can be set similarly, but if not they have default values."""

{% set LAST_TASK = LAST_TASK | default( 'baz' ) %}
{% set N_MEMBERS = N_MEMBERS | default( 3 ) | int %}

{# input of FIRST_TASK is required - no default #}

[scheduling]
    initial cycle point = 20100808T00
    final cycle point   = 20100816T00
    [[dependencies]]
        [[[0]]]
            graph = """{{ FIRST_TASK }} => ENS
                        ENS:succeed-all => {{ LAST_TASK }}"""

[runtime]
    [[ENS]]
{% for I in range( 0, N_MEMBERS ) %}
    [[ mem_{{ I }} ]]
        inherit = ENS
{% endfor %}

```

Here's the result:

```

$ cylc list SUITE
Jinja2 Template Error
'FIRST_TASK' is undefined
cylc-list foo failed: 1

$ cylc list --set FIRST_TASK=bob foo
bob
baz
mem_2
mem_1
mem_0

$ cylc list --set FIRST_TASK=bob --set LAST_TASK=alice foo
bob
alice
mem_2
mem_1
mem_0

$ cylc list --set FIRST_TASK=bob --set N_MEMBERS=10 foo
mem_9
mem_8
mem_7
mem_6
mem_5
mem_4
mem_3
mem_2
mem_1
mem_0
baz
bob

```

Note also that `cylc view --set FIRST_TASK=bob --jinja2 SUITE` will show the suite with the Jinja2 variables as set.

Note: suites started with template variables set on the command line will *restart* with the same settings. However, you can set them again on the `cylc restart` command line if they need to be overridden.

9.7.5 Jinja2 Variable Scope

Jinja2 variable scoping rules may be surprising. Variables set inside a *for loop* block, for instance, are not accessible outside of the block, so the following will print `# FOO is 0`, not `# FOO is 9`:

```
{% set FOO = false %}
{% for item in items %}
    {% if item.check_something() %}
        {% set FOO = true %}
    {% endif %}
{% endfor %}
# FOO is {{FOO}}
```

Jinja2 documentation suggests using alternative constructs like the `loop else` block or the special `loop` variable. More complex use cases can be handled using `namespace` objects which allow propagating of changes across scopes:

```
{% set ns = namespace(foo=false) %}
{% for item in items %}
    {% if item.check_something() %}
        {% set ns.foo = true %}
    {% endif %}
{% endfor %}
# FOO is {{ns.foo}}
```

For detail, see: [Jinja2 Template Designer Documentation > Assignments](#)

9.7.6 Raising Exceptions

Cylc provides two functions for raising exceptions using Jinja2. These exceptions are raised when the `suite.rc` file is loaded and will prevent a suite from running.

Note: These functions must be contained within `{{ Jinja2 blocks` as opposed to `{% blocks`.

9.7.6.1 Raise

The “raise” function will result in an error containing the provided text.

```
{% if not VARIABLE is defined %}
    {{ raise('VARIABLE must be defined for this suite.') }}
{% endif %}
```

9.7.6.2 Assert

The “assert” function will raise an exception containing the text provided in the second argument providing that the first argument evaluates as False. The following example is equivalent to the “raise” example above.

```
{{ assert(VARIABLE is defined, 'VARIABLE must be defined for this suite.') }}
```

9.8 Omitting Tasks At Runtime

It is sometimes convenient to omit certain tasks from the suite at runtime without actually deleting their definitions from the suite.

Defining [runtime] properties for tasks that do not appear in the suite graph results in verbose-mode validation warnings that the tasks are disabled. They cannot be used because the suite graph is what defines their dependencies and valid cycle points. Nevertheless, it is legal to leave these orphaned runtime sections in the suite definition because it allows you to temporarily remove tasks from the suite by simply commenting them out of the graph.

To omit a task from the suite at runtime but still leave it fully defined and available for use (by insertion or `cylc submit`) use one or both of [scheduling][[special task]] lists, *include at start-up* or *exclude at start-up* (documented in A.4.11.6 and A.4.11.5). Then the graph still defines the validity of the tasks and their dependencies, but they are not actually loaded into the suite at start-up. Other tasks that depend on the omitted ones, if any, will have to wait on their insertion at a later time or otherwise be triggered manually.

Finally, with Jinja2 (9.7) you can radically alter suite structure by including or excluding tasks from the [scheduling] and [runtime] sections according to the value of a single logical flag defined at the top of the suite.

9.9 Naked Dummy Tasks And Strict Validation

A *naked dummy task* appears in the suite graph but has no explicit runtime configuration section. Such tasks automatically inherit the default “dummy task” configuration from the root namespace. This is very useful because it allows functional suites to be mocked up quickly for test and demonstration purposes by simply defining the graph. It is somewhat dangerous, however, because there is no way to distinguish an intentional naked dummy task from one generated by typographic error: misspelling a task name in the graph results in a new naked dummy task replacing the intended task in the affected trigger expression; and misspelling a task name in a runtime section heading results in the intended task becoming a dummy task itself (by divorcing it from its intended runtime config section).

To avoid this problem any dummy task used in a real suite should not be naked - i.e. it should have an explicit entry in under the runtime section of the suite definition, even if the section is empty. This results in exactly the same dummy task behaviour, via implicit inheritance from root, but it allows use of `cylc validate --strict` to catch errors in task names by failing the suite if any naked dummy tasks are detected.

10 Task Implementation

Existing scripts and executables can be used as cylc tasks without modification so long as they return standard exit status - zero on success, non-zero for failure - and do not spawn detaching processes internally (see 10.5).

10.1 Task Job Scripts

When the suite dameon determines that a task is ready to run it generates a *job script* that embodies the task runtime configuration in the suite.rc file, and submits it to the configured job host and batch system (see 11).

Task job scripts are written to the suite’s job log directory. They can be printed with `cylc cat-log` or generated and printed with `cylc jobscript`.

10.2 Inlined Tasks

Task *script* items can be multi-line strings of `bash` code, so many tasks can be entirely inlined in the `suite.rc` file. For anything more than a few lines of code, however, we recommend using external shell scripts to allow independent testing, re-use, and shell mode editing.

10.3 Task Messages

Tasks messages can be sent back to the suite server program to report completed outputs and arbitrary messages of different severity levels.

Some types of message - in addition to events like task failure - can optionally trigger execution of event handlers in the suite server program (see 12.19).

Normal severity messages are printed to `job.out` and logged by the suite server program:

```
cylc message "Hello from ${CYLC_TASK_ID}"
```

CUSTOM severity messages are printed to `job.out`, logged by the suite server program, and can be used to trigger *custom* event handlers:

```
cylc message -p CUSTOM "data available for ${CYLC_TASK_CYCLE_POINT}"
```

Task output messages, used for triggering other tasks, can also be sent with custom severity if need be

WARNING severity messages are printed to `job.err`, logged by the suite server program, and can be passed to *warning* event handlers:

```
cylc message -p WARNING "Uh-oh, something's not right here."
```

CRITICAL severity messages are printed to `job.err`, logged by the suite server program, and can be passed to *critical* event handlers:

```
cylc message -p CRITICAL "ERROR occurred in process X!"
```

10.4 Aborting Job Scripts on Error

Task job scripts use `set -x` to abort on any error, and trap `ERR`, `EXIT`, and `SIGTERM` to send task failed messages back to the suite server program before aborting. Other scripts called from job scripts should therefore abort with standard non-zero exit status on error, to trigger the job script error trap.

To prevent a command that is expected to generate a non-zero exit status from triggering the exit trap, protect it with a control statement such as:

```
if cmp FILE1 FILE2; then
: # success: do stuff
else
: # failure: do other stuff
fi
```

Task job scripts also use `set -u` to abort on referencing any undefined variable (useful for picking up typos); and `set -o pipefail` to abort if any part of a pipe fails (by default the shell only returns the exit status of the final command in a pipeline).

10.4.1 Custom Failure Messages

Critical events normally warrant aborting a job script rather than just sending a message. As described just above, `exit 1` or any failing command not protected by the surrounding scripting will cause a job script to abort and report failure to the suite server program, potentially triggering a *failed* task event handler.

For failures detected by the scripting you could send a critical message back before aborting, potentially triggering a *critical* task event handler:

```
if ! /bin/false; then
    cylc message -p CRITICAL "ERROR: /bin/false failed!"
    exit 1
fi
```

To abort a job script with a custom message that can be passed to a *failed* task event handler, use the built-in `cylc__job_abort` shell function:

```
if ! /bin/false; then
    cylc__job_abort "ERROR: /bin/false failed!"
fi
```

10.5 Avoid Detaching Processes

If a task script starts background sub-processes and does not wait on them, or internally submits jobs to a batch scheduler and then exits immediately, the detached processes will not be visible to cylc and the task will appear to finish when the top-level script finishes. You will need to modify scripts like this to make them execute all sub-processes in the foreground (or use the shell `wait` command to wait on them before exiting) and to prevent job submission commands from returning before the job completes (e.g. `llsubmit -s` for Loadleveler, `qsub -sync yes` for Sun Grid Engine, and `qsub -W block=true` for PBS).

If this is not possible - perhaps you don't have control over the script or can't work out how to fix it - one alternative approach is to use another task to repeatedly poll for the results of the detached processes:

```
[scheduling]
  [[dependencies]]
    graph = "model => checker => post-proc"
[runtime]
  [[model]]
    # Uh-oh, this script does an internal job submission to run model.exe:
    script = "run-model.sh"
  [[checker]]
    # Fail and retry every minute (for 10 tries at the most) if model's
    # job.done indicator file does not exist yet.
    script = "[[ ! -f $RUN_DIR/job.done ]] && exit 1"
  [[job]]
    execution retry delays = 10 * PT1M
```

11 Task Job Submission and Management

For the requirements a command, script, or program, must fulfill in order to function as a cylc task, see 10. This section explains how tasks are submitted by the suite server program when they are ready to run, and how to define new batch system handlers.

When a task is ready `cylc` generates a job script (see 10.1). The job script is submitted to run by the *batch system* chosen for the task. Different tasks can use different batch systems. Like other runtime properties, you can set a suite default batch system and override it for specific tasks or families:

```
[runtime]
  [[root]] # suite defaults
    [[[job]]]
      batch system = loadleveler
  [[foo]] # just task foo
    [[[job]]]
      batch system = at
```

11.1 Supported Job Submission Methods

Cylc supports a number of commonly used batch systems. See 11.7 for how to add new job submission methods.

11.1.1 background

Runs task job scripts as Unix background processes.

If an execution time limit is specified for a task, its job will be wrapped by the `timeout` command.

11.1.2 at

Submits task job scripts to the rudimentary Unix `at` scheduler. The `atd` daemon must be running.

If an execution time limit is specified for a task, its job will be wrapped by the `timeout` command.

11.1.3 loadleveler

Submits task job scripts to loadleveler by the `llsubmit` command. Loadleveler directives can be provided in the suite.rc file:

```
[runtime]
  [[my_task]]
    [[[job]]]
      batch system = loadleveler
      execution time limit = PT10M
    [[[directives]]]
      foo = bar
      baz = qux
```

These are written to the top of the task job script like this:

```
#!/bin/bash
# DIRECTIVES
# @ foo = bar
# @ baz = qux
# @ wall_clock_limit = 660,600
# @ queue
```

If `restart=yes` is specified as a directive for `loadleveler`, the job will automatically trap `SIGUSR1`, which `loadleveler` may use to preempt the job. On trapping `SIGUSR1`, the job will inform the suite that it has been vacated by `loadleveler`. This will put it back to the submitted state, until it starts running again.

If `execution time limit` is specified, it is used to generate the `wall_clock_limit` directive. The setting is assumed to be the soft limit. The hard limit will be set by adding an extra minute to the soft limit. Do not specify the `wall_clock_limit` directive explicitly if `execution time limit` is specified. Otherwise, the execution time limit known by the suite may be out of sync with what is submitted to the batch system.

11.1.4 lsf

Submits task job scripts to IBM Platform LSF by the `bsub` command. LSF directives can be provided in the `suite.rc` file:

```
[runtime]
  [[my_task]]
    [[[job]]]
      batch system = lsf
      execution time limit = PT10M
    [[[directives]]]
      -q = foo
```

These are written to the top of the task job script like this:

```
#!/bin/bash
# DIRECTIVES
#BSUB -q = foo
#BSUB -W = 10
```

If `execution time limit` is specified, it is used to generate the `-W` directive. Do not specify the `-W` directive explicitly if `execution time limit` is specified. Otherwise, the execution time limit known by the suite may be out of sync with what is submitted to the batch system.

11.1.5 pbs

Submits task job scripts to PBS (or Torque) by the `qsub` command. PBS directives can be provided in the `suite.rc` file:

```
[runtime]
  [[my_task]]
    [[[job]]]
      batch system = pbs
      execution time limit = PT1M
    [[[directives]]]
      -V =
      -q = foo
      -l nodes = 1
```

These are written to the top of the task job script like this:

```
#!/bin/bash
# DIRECTIVES
#PBS -V
#PBS -q foo
#PBS -l nodes=1
#PBS -l walltime=60
```

If `execution time limit` is specified, it is used to generate the `-l walltime` directive. Do not specify the `-l walltime` directive explicitly if `execution time limit` is specified. Otherwise, the execution time limit known by the suite may be out of sync with what is submitted to the batch system.

11.1.6 moab

Submits task job scripts to the Moab workload manager by the `msub` command. Moab directives can be provided in the `suite.rc` file; the syntax is very similar to PBS:

```
[runtime]
  [[my_task]]
    [[[job]]]
      batch system = moab
      execution time limit = PT1M
    [[[directives]]]
      -V =
      -q = foo
      -l nodes = 1
```

These are written to the top of the task job script like this:

```
#!/bin/bash
# DIRECTIVES
#PBS -V
#PBS -q foo
#PBS -l nodes=1
#PBS -l walltime=60
```

(Moab understands `#PBS` directives).

If `execution time limit` is specified, it is used to generate the `-l walltime` directive. Do not specify the `-l walltime` directive explicitly if `execution time limit` is specified. Otherwise, the execution time limit known by the suite may be out of sync with what is submitted to the batch system.

11.1.7 sge

Submits task job scripts to Sun/Oracle Grid Engine by the `qsub` command. SGE directives can be provided in the `suite.rc` file:

```
[runtime]
  [[my_task]]
    [[[job]]]
      batch system = sge
      execution time limit = P1D
    [[[directives]]]
      -cwd =
      -q = foo
      -l h_data = 1024M
      -l h_rt = 24:00:00
```

These are written to the top of the task job script like this:

```
#!/bin/bash
# DIRECTIVES
#$ -cwd
#$ -q foo
#$ -l h_data=1024M
#$ -l h_rt=24:00:00
```

If `execution time limit` is specified, it is used to generate the `-l h_rt` directive. Do not specify the `-l h_rt` directive explicitly if `execution time limit` is specified. Otherwise, the execution time limit known by the suite may be out of sync with what is submitted to the batch system.

11.1.8 slurm

Submits task job scripts to Simple Linux Utility for Resource Management by the `sbatch` command. SLURM directives can be provided in the `suite.rc` file (note that since not all SLURM commands have a short form, `cylc` requires the long form directives):

```
[runtime]
  [[my_task]]
    [[[job]]]
      batch system = slurm
      execution time limit = PT1H
    [[[directives]]]
      --nodes = 5
      --account = QXZ5W2
```

These are written to the top of the task job script like this:

```
#!/bin/bash
#SBATCH --nodes=5
#SBATCH --time=60:00
#SBATCH --account=QXZ5W2
```

If `execution time limit` is specified, it is used to generate the `--time` directive. Do not specify the `--time` directive explicitly if `execution time limit` is specified. Otherwise, the execution time limit known by the suite may be out of sync with what is submitted to the batch system.

11.1.9 Default Directives Provided

For batch systems that use job file directives (PBS, Loadleveler, etc.) default directives are provided to set the job name, stdout and stderr file paths, and the execution time limit (if specified).

Cylc constructs the job name string using a combination of the task ID and the suite name. PBS fails a job submit if the job name in `-N name` is too long. For version 12 or below, this is 15 characters. For version 13, this is 236 characters. The default setting will truncate the job name string to 15 characters. If you have PBS 13 at your site, you should modify your site's global configuration file to allow the job name to be longer. (See also Section B.9.1.19.5.) For example:

```
[hosts]
  [[myhpc*]]
    [[[batch systems]]]
      [[[[pbs]]]]
        # PBS 13
        job name length maximum = 236
```

11.1.10 Directives Section Quirks (PBS, SGE, ...)

To specify an option with no argument, such as `-v` in PBS or `-cwd` in SGE you must give a null string as the directive value in the `suite.rc` file.

The left hand side of a setting (i.e. the string before the first equal sign) must be unique. To specify multiple values using an option such as `-l` option in PBS, SGE, etc., either specify all items in a single line:

```
-l=select=28:ncpus=36:mpiprocs=18:ompthreads=2:walltime=12:00:00
```

(Left hand side is `-l`. A second `-l=...` line will override the first.)

Or separate the items (note: no equal sign after `-l`):

```
-l select=28
-l ncpus=36
-l mpiprocs=18
-l ompthreads=2
-l walltime=12:00:00
```

(Left hand sides are now `-l select`, `-l ncpus`, etc.)

11.2 Task stdout And stderr Logs

When a task is ready to run `cylc` generates a filename root to be used for the task job script and log files. The filename containing the task name, cycle point, and a submit number that increments if the same task is re-triggered multiple times:

```
# task job script:
~/cylc-run/tut/oneoff/basic/log/job/1/hello/01/job
# task stdout:
~/cylc-run/tut/oneoff/basic/log/job/1/hello/01/job.out
# task stderr:
~/cylc-run/tut/oneoff/basic/log/job/1/hello/01/job.err
```

How the stdout and stderr streams are directed into these files depends on the batch system. The `background` method just uses appropriate output redirection on the command line, as shown above. The `loadleveler` method writes appropriate directives to the job script that is submitted to `loadleveler`.

`Cylc` obviously has no control over the stdout and stderr output from tasks that do their own internal output management (e.g. tasks that submit internal jobs and direct the associated output to other files). For less internally complex tasks, however, the files referred to here will be complete task job logs.

Some batch systems, such as `pbs`, redirect a job's stdout and stderr streams to a separate cache area while the job is running. The contents are only copied to the normal locations when the job completes. This means that `cylc cat-log` or the `gcylc` GUI will be unable to find the job's stdout and stderr streams while the job is running. Some sites with these batch systems are known to provide commands for viewing and/or tail-follow a job's stdout and stderr streams that are redirected to these cache areas. If this is the case at your site, you can configure `cylc` to make use of the provided commands by adding some settings to the global site/user config. E.g.:

```
[hosts]
  [[HOST]] # <= replace this with a real host name
    [[batch systems]]
      [[pbs]]
        err tailer = qcat -f -e \%(job_id)s
        out tailer = qcat -f -o \%(job_id)s
        err viewer = qcat -e \%(job_id)s
        out viewer = qcat -o \%(job_id)s
```

11.3 Overriding The Job Submission Command

To change the form of the actual command used to submit a job you do not need to define a new batch system handler; just override the `command template` in the relevant job submission sections of your `suite.rc` file:

```
[runtime]
  [[root]]
    [[[job]]]
      batch system = loadleveler
      # Use '-s' to stop llsubmit returning
      # until all job steps have completed:
      batch submit command template = llsubmit -s %(job)s
```

As explained in [A](#) the template's `%(job)s` will be substituted by the job file path.

11.4 Job Polling

For supported batch systems, one-way polling can be used to determine actual job status: the suite server program executes a process on the task host, by non-interactive ssh, to interrogate the batch queueing system there, and to read a *status file* that is automatically generated by the task job script as it runs.

Polling may be required to update the suite state correctly after unusual events such as a machine being rebooted with tasks running on it, or network problems that prevent task messages from getting back to the suite host.

Tasks can be polled on demand by right-clicking on them in gcylc or using the `cylc poll` command.

Tasks are polled automatically, once, if they timeout while queueing in a batch scheduler and submission timeout is set. (See [A.5.1.13](#) for how to configure timeouts).

Tasks are polled multiple times, where necessary, when they exceed their execution time limits. These are normally set with some initial delays to allow the batch systems to kill the jobs. (See [B.9.1.19.6](#) for how to configure the polling intervals).

Any tasks recorded in the *submitted* or *running* states at suite restart are automatically polled to determine what happened to them while the suite was down.

Regular polling can also be configured as a health check on tasks submitted to hosts that are known to be flaky, or as the sole method of determining task status on hosts that do not allow task messages to be routed back to the suite host.

To use polling instead of task-to-suite messaging set `task communication method = poll` in cylc site and user global config (see [B.9.1.3](#)). The default polling intervals can be overridden for all suites there too (see [B.9.1.5](#) and [B.9.1.4](#)), or in specific suite definitions (in which case polling will be done regardless of the task communication method configured for the host; see [A.5.1.11.7](#) and [A.5.1.11.8](#)).

Note that regular polling is not as efficient as task messaging in updating task status, and it should be used sparingly in large suites.

Note that for polling to work correctly, the batch queueing system must have a job listing command for listing your jobs, and that the job listing must display job IDs as they are returned by the batch queueing system submit command. For example, for pbs, moab and sge, the `qstat` command should list jobs with their IDs displayed in exactly the same format as they are returned by the `qsub` command.

11.5 Job Killing

For supported batch systems, the suite server program can execute a process on the task host, by non-interactive ssh, to kill a submitted or running job according to its batch system.

Tasks can be killed on demand by right-clicking on them in gcylc or using the `cylc kill` command.

11.6 Execution Time Limit

You can specify an `execution time limit` for all supported job submission methods. E.g.:

```
[runtime]
  [[task-x]]
    [[[job]]]
      execution time limit = PT1H
```

For tasks running with `background` or `at`, their jobs will be wrapped using the `timeout` command. For all other methods, the relevant time limit directive will be added to their job files.

The `execution time limit` setting will also inform the suite when a task job should complete by. If a task job has not reported completing within the specified time, the suite will poll the task job. (The default setting is PT1M, PT2M, PT7M. The accumulated times for these intervals will be roughly 1 minute, $1 + 2 = 3$ minutes and $1 + 2 + 7 = 10$ minutes after a task job exceeds its execution time limit.)

11.6.1 Execution Time Limit and Execution Timeout

If you specify an `execution time limit` the `execution timeout event handler` will only be called if the job has not completed after the final poll (by default, 10 min after the time limit). This should only happen if the submission method you are using is not enforcing wallclock limits (unlikely) or you are unable to contact the machine to confirm the job status.

If you specify an `execution timeout` and not an `execution time limit` then the `execution timeout event handler` will be called as soon as the specified time is reached. The job will also be polled to check its latest status (possibly resulting in an update in its status and the calling of the relevant event handler). This behaviour is deprecated, which users should avoid using.

If you specify an `execution timeout` and an `execution time limit` then the execution timeout setting will be ignored.

11.7 Custom Job Submission Methods

Defining a new batch system handler requires a little Python programming. Use the built-in handlers as examples, and read the documentation in `lib/cylc/batch_sys_manager.py`.

11.7.1 An Example

The following `qsub.py` module overrides the built-in `pbs` batch system handler to to change the directive prefix from `#PBS` to `#QSUB`:

```
#!/usr/bin/env python

from cylc.batch_sys_handlers.pbs import PBSHandler

class QSUBHandler(PBSHandler):
    DIRECTIVE_PREFIX = "#QSUB "

BATCH_SYSTEM_HANDLER = QSUBHandler()
```

If this is in the Python search path (see 11.7.2 below) you can use it by name in suite definitions:

```
[scheduling]
  [[dependencies]]
    graph = "a"
[runtime]
  [[root]]
    [[[job]]]
      batch system = qsub # <---!
      execution time limit = PT1M
    [[[directives]]]
      -l nodes = 1
      -q = long
      -V =
```

Generate a job script to see the resulting directives:

```
$ cylc register test $HOME/test
$ cylc jobscript test a.1 | grep QSUB
#QSUB -e /home/oliverh/cylc-run/my.suite/log/job/1/a/01/job.err
#QSUB -l nodes=1
#QSUB -l walltime=60
#QSUB -o /home/oliverh/cylc-run/my.suite/log/job/1/a/01/job.out
#QSUB -N a.1
#QSUB -q long
#QSUB -V
```

(Of course this suite will fail at run time because we only changed the directive format, and PBS does not accept `#QSUB` directives in reality).

11.7.2 Where To Put Batch System Handler Modules

Custom batch system handlers must be installed on suite and job hosts in one of these locations:

- under `SUITE-DEF-PATH/lib/python/`
- under `CYLC-PATH/lib/cylc/batch_sys_handlers/`
- or anywhere in `$PYTHONPATH`

(A note for Rose users: `rose suite-run` automatically installs `SUITE-DEF-PATH/lib/python/` to job hosts).

12 Running Suites

This chapter currently features a diverse collection of topics related to running suites. Please also see the Tutorial (7) and command documentation (F), and experiment with plenty of examples.

12.1 Suite Start-Up

There are three ways to start a suite running: *cold start* and *warm start*, which start from scratch; and *restart*, which starts from a prior suite state checkpoint. The only difference between cold starts and warm starts is that warm starts start from a point beyond the suite initial cycle point.

Once a suite is up and running it is typically a restart that is needed most often (but see also `cylc reload`). *Be aware that cold and warm starts wipe out prior suite state, so you can't go back to a restart if you decide you made a mistake.*

12.1.1 Cold Start

A cold start is the primary way to start a suite run from scratch:

```
$ cylc run SUITE [INITIAL_CYCLE_POINT]
```

The initial cycle point may be specified on the command line or in the suite.rc file. The scheduler starts by loading the first instance of each task at the suite initial cycle point, or at the next valid point for the task.

12.1.2 Warm Start

A warm start runs a suite from scratch like a cold start, but from the beginning of a given cycle point that is beyond the suite initial cycle point. This is generally inferior to a *restart* (which loads a previously recorded suite state - see 12.1.3) because it may result in some tasks rerunning. However, a warm start may be required if a restart is not possible, e.g. because the suite run database was accidentally deleted. The warm start cycle point must be given on the command line:

```
$ cylc run --warm SUITE [START_CYCLE_POINT]
```

The original suite initial cycle point is preserved, but all tasks and dependencies before the given warm start cycle point are ignored.

The scheduler starts by loading a first instance of each task at the warm start cycle point, or at the next valid point for the task. *R1*-type tasks behave exactly the same as other tasks - if their cycle point is at or later than the given start cycle point, they will run; if not, they will be ignored.

12.1.3 Restart and Suite State Checkpoints

At restart (see `cylc restart --help`) a suite server program initializes its task pool from a previously recorded checkpoint state. By default the latest automatic checkpoint - which is updated with every task state change - is loaded so that the suite can carry on exactly as it was just before being shut down or killed.

```
$ cylc restart SUITE
```

Tasks recorded in the 'submitted' or 'running' states are automatically polled (see Section 12.5) at start-up to determine what happened to them while the suite was down.

12.1.3.1 Restart From Latest Checkpoint

To restart from the latest checkpoint simply invoke the `cylc restart` command with the suite name (or select ‘restart’ in the GUI suite start dialog window):

```
$ cylc restart SUITE
```

12.1.3.2 Restart From Another Checkpoint

Suite server programs automatically update the “latest” checkpoint every time a task changes state, and at every suite restart, but you can also take checkpoints at other times. To tell a suite server program to checkpoint its current state:

```
$ cylc checkpoint SUITE-NAME CHECKPOINT-NAME
```

The 2nd argument is a name to identify the checkpoint later with:

```
$ cylc ls-checkpoints SUITE-NAME
```

For example, with checkpoints named ‘bob’, ‘alice’, and ‘breakfast’:

```
$ cylc ls-checkpoints SUITE-NAME
#####
# CHECKPOINT ID (ID|TIME|EVENT)
1|2017-11-01T15:48:34+13|bob
2|2017-11-01T15:48:47+13|alice
3|2017-11-01T15:49:00+13|breakfast
...
0|2017-11-01T17:29:19+13|latest
```

To see the actual task state content of a given checkpoint ID (if you need to), for the moment you have to interrogate the suite DB, e.g.:

```
$ sqlite3 ~/cylc-run/SUITE-NAME/log/db \
'select * from task_pool_checkpoints where id == 3;'
3|2012|model|1|running|
3|2013|pre|0|waiting|
3|2013|post|0|waiting|
3|2013|model|0|waiting|
3|2013|upload|0|waiting|
```

Note that a checkpoint captures the instantaneous state of every task in the suite, including any tasks that are currently active, so you may want to be careful where you do it. Tasks recorded as active are polled automatically on restart to determine what happened to them.

The checkpoint ID 0 (zero) is always used for latest state of the suite, which is updated continuously as the suite progresses. The checkpoint IDs of earlier states are positive integers starting from 1, incremented each time a new checkpoint is stored. Currently suites automatically store checkpoints before and after reloads, and on restarts (using the latest checkpoints before the restarts).

Once you have identified the right checkpoint, restart the suite like this:

```
$ cylc restart --checkpoint=CHECKPOINT-ID SUITE
```

or enter the checkpoint ID in the space provided in the GUI restart window.

12.1.3.3 Checkpointing With A Task

Checkpoints can be generated automatically at particular points in the workflow by coding tasks that run the `cylc checkpoint` command:

```
[scheduling]
  [[dependencies]]
    [[PT6H]]
      graph = "pre => model => post => checkpointer"
[runtime]
  # ...
  [[checkpointer]]
    script = ""
wait "${CYLC_TASK_MESSAGE_STARTED_PID}" 2>/dev/null || true
cylc checkpoint "${CYLC_SUITE_NAME}" CP-${CYLC_TASK_CYCLE_POINT}
```

Note that we need to `wait` on the “task started” message - which is sent in the background to avoid holding tasks up in a network outage - to ensure that the checkpointer task is correctly recorded as running in the checkpoint (at restart the suite server program will poll to determine that that task job finished successfully). Otherwise it may be recorded in the waiting state and, if its upstream dependencies have already been cleaned up, it will need to be manually reset from waiting to succeeded after the restart to avoid stalling the suite.

12.1.3.4 Behaviour of Tasks on Restart

All tasks are reloaded in exactly their checkpointed states. Failed tasks are not automatically resubmitted at restart in case the underlying problem has not been addressed yet.

Tasks recorded in the submitted or running states are automatically polled on restart, to see if they are still waiting in a batch queue, still running, or if they succeeded or failed while the suite was down. The suite state will be updated automatically according to the poll results.

Existing instances of tasks removed from the suite definition before restart are not removed from the task pool automatically, but they will not spawn new instances. They can be removed manually if necessary, with `cylc remove`.

Similarly, instances of new tasks added to the suite definition before restart are not inserted into the task pool automatically, because it is very difficult in general to automatically determine the cycle point of the first instance. Instead, the first instance of a new task should be inserted manually at the right cycle point, with `cylc insert`.

12.2 Reloading The Suite Definition At Runtime

The `cylc reload` command tells a suite server program to reload its suite definition at run time. This is an alternative to shutting a suite down and restarting it after making changes.

As for a restart, existing instances of tasks removed from the suite definition before reload are not removed from the task pool automatically, but they will not spawn new instances. They can be removed manually if necessary, with `cylc remove`.

Similarly, instances of new tasks added to the suite definition before reload are not inserted into the pool automatically. The first instance of each must be inserted manually at the right cycle point, with `cylc insert`.

12.3 Task Job Access To Cylc

Task jobs need access to Cylc on the job host, primarily for task messaging, but also to allow user-defined task scripting to run other Cylc commands.

Cylc should be installed on job hosts as on suite hosts, with different releases installed side-by-side and invoked via the central Cylc wrapper according to the value of `$CYLC_VERSION` - see Section 3.3. Task job scripts set `$CYLC_VERSION` to the version of the parent suite server program, so that the right Cylc will be invoked by jobs on the job host.

Access to the Cylc executable (preferably the central wrapper as just described) for different job hosts can be configured using site and user global configuration files (on the suite host). If the environment for running the Cylc executable is only set up correctly in a login shell for a given host, you can set `[hosts][HOST]use_login_shell = True` for the relevant host (this is the default, to cover more sites automatically). If the environment is already correct without the login shell, but the Cylc executable is not in `$PATH`, then `[hosts][HOST]cylc_executable` can be used to specify the direct path to the executable.

To customize the environment more generally for Cylc on jobs hosts, use of `job-init-env.sh` is described in Section 7.1.1.

12.4 The Suite Contact File

At start-up, suite server programs write a *suite contact file* `$HOME/cylc-run/SUITE/.service/contact` that records suite host, user, port number, process ID, Cylc version, and other information. Client commands can read this file, if they have access to it, to find the target suite server program.

12.5 Task Job Polling

At any point after job submission task jobs can be *polled* to check that their true state conforms to what is currently recorded by the suite server program. See `cylc poll --help` for how to poll one or more tasks manually, or right-click poll a task or family in GUI.

Polling may be necessary if, for example, a task job gets killed by the untrappable SIGKILL signal (e.g. `kill -9 PID`), or if a network outage prevents task success or failure messages getting through, or if the suite server program itself is down when tasks finish execution.

To poll a task job the suite server program interrogates the batch system, and the `job.status` file, on the job host. This information is enough to determine the final task status even if the job finished while the suite server program was down or unreachable on the network.

12.5.1 Routine Polling

Task jobs are automatically polled at certain times: once on job submission timeout; several times on exceeding the job execution time limit; and at suite restart any tasks recorded as active in the suite state checkpoint are polled to find out what happened to them while the suite was down.

Finally, in necessary routine polling can be configured as a way to track job status on job hosts that do not allow networking routing back to the suite host for task messaging by HTTPS or ssh. See [12.6.3](#).

12.6 Tracking Task State

Cylc supports three ways of tracking task state on job hosts:

- task-to-suite messaging via HTTPS
- task-to-suite messaging via non-interactive ssh to the suite host, then local HTTPS
- regular polling by the suite server program

These are explained in the following sections. All three can be used, on different job hosts, in the same suite if necessary.

If your site prohibits HTTPS and ssh back from job hosts to suite hosts, before resorting to the polling method you should consider installing dedicated Cylc servers or VMs inside the HPC trust zone (where HTTPS and ssh should be allowed).

It is also possible to run Cylc suite server programs on HPC login nodes, but this is not recommended for load, run duration, and GUI reasons.

Finally, it has been suggested that *port forwarding* may provide another solution - but that is beyond the scope of this document.

12.6.1 HTTPS Task Messaging

Task job wrappers automatically invoke `cylc message` to report progress back to the suite server program when they begin executing, at normal exit (success) and abnormal exit (failure).

By default the messaging occurs via an authenticated, HTTPS connection to the suite server program. This is the preferred task communications method - it is efficient and direct.

Suite server programs automatically install suite contact information and credentials on job hosts. Users only need to do this manually for remote access to suites on other hosts, or suites owned by other users - see [12.11](#).

12.6.2 Ssh Task Messaging

Cylc can be configured to re-invoke task messaging commands on the suite host via non-interactive ssh (from job host to suite host). Then a local HTTPS connection is made to the suite server program.

(User-invoked client commands (aside from the GUI, which requires HTTPS) can do the same thing with the `--use-ssh` command option).

This is less efficient than direct HTTPS messaging, but it may be useful at sites where the HTTPS ports are blocked but non-interactive ssh is allowed.

12.6.3 Polling to Track Job Status

Finally, suite server programs can actively poll task jobs at configurable intervals, via non-interactive ssh to the job host.

Polling is the least efficient task communications method because task state is updated only at intervals, not when task events actually occur. However, it may be needed at sites that do not allow HTTPS or non-interactive ssh from job host to suite host.

Be careful to avoid spamming task hosts with polling commands. Each poll opens (and then closes) a new ssh connection.

Polling intervals are configurable under `[runtime]` because they should may depend on the expected execution time. For instance, a task that typically takes an hour to run might be polled every 10 minutes initially, and then every minute toward the end of its run. Interval values are used in turn until the last value, which is used repeatedly until finished:

```
[runtime]
  [[foo]]
    [[[job]]]
      # poll every minute in the 'submitted' state:
      submission polling intervals = PT1M
      # poll one minute after foo starts running, then every 10
      # minutes for 50 minutes, then every minute until finished:
      execution polling intervals = PT1M, 5*PT10M, PT1M
```

A list of intervals with optional multipliers can be used for both submission and execution polling, although a single value is probably sufficient for submission polling. If these items are not configured default values from site and user global config will be used for the polling task communication method; polling is not done by default under the other task communications methods (but it can still be used if you like).

12.6.4 Task Communications Configuration

Here are the default site and user global config items relevant to task state tracking (see these with `cylc get-site-config`):

```
# SITE AND USER CONFIG

# Task messaging settings affect task-to-suite communications.
[task messaging]
  # If a message send fails, retry after this delay:
  retry interval in seconds = 5
  # If send fails after this many tries, give up trying:
  maximum number of tries = 7

  # This timeout is the same as --comms-timeout for user commands. If
  # set to None (no timeout) messages to non-responsive suites
  # (e.g. suspended with Ctrl-Z) could hang indefinitely.
  connection timeout in seconds = 30

# Setup the communication method details. This is required for
# communications between cylc clients and servers (i.e. between
# suite-connecting commands and guis, and running suite server processes).
[communication]

  # Configure the choice of communication method. Https is configured by
  # default, and this requires the Python OpenSSL package to be installed.
  # Http can be configured to override the default setting.
# SITE ONLY
method = http
```

```

# Each suite listens on a dedicated network port. The port to bind to is
# selected randomly from the allowed range of ports.
# SITE ONLY
base port = 43001

# This sets the maximum number of suites that can run at once.
# SITE ONLY
maximum number of ports = 100

[hosts]
# The default task host is the suite host, i.e. localhost:
# Add task host sections if local defaults are not sufficient.
[[HOST]]
# Method of communication of task progress back to the suite:
# 1) default - HTTPS via network ports
# 2) ssh - re-invoke messaging commands on suite server
# 3) poll - the suite polls for status of passive tasks
# HTTPS comms are still required in all cases *on the suite host*
# for cylc clients (commands etc.) to communicate with suites.
task communication method = "default" # or "ssh" or "poll"
# The "poll" method sets a default interval here to ensure no
# tasks are accidentally left unpollled. You can override this
# with run-length appropriate intervals under task [runtime]
# (however this will also result in routine polling under the
# default or ssh communications).
default polling interval in minutes = 1.0

```

12.7 The Suite Service Directory

At registration time a *suite service directory*, `$HOME/cylc-run/<SUITE>/service/`, is created and populated with a private passphrase file (containing random text), a self-signed SSL certificate (see 12.9), and a symlink to the suite source directory. An existing passphrase file will not be overwritten if a suite is re-registered.

At run time, the private suite run database is also written to the service directory, along with a *suite contact file* that records the host, user, port number, process ID, Cylc version, and other information about the suite server program. Client commands automatically read daemon targetting information from the contact file, if they have access to it.

12.8 File-Reading Commands

Some Cylc commands and GUI actions parse suite definitions or read other files from the suite host account, rather than communicate with a suite server program over the network. In future we plan to have suite server program serve up these files to clients, but for the moment this functionality requires read-access to the relevant files on the suite host.

If you are logged into the suite host account, file-reading commands will just work.

12.8.1 Remote Host, Shared Home Directory

If you are logged into another host with shared home directories (shared filesystems are common in HPC environments) file-reading commands will just work because suite files will look “local” on both hosts.

12.8.2 Remote Host, Different Home Directory

If you are logged into another host with no shared home directory, file-reading commands require non-interactive ssh to the suite host account, and use of the `--host` and `--user` options to re-invoke the command on the suite account.

12.8.3 Same Host, Different User Account

(This is essentially the same as *Remote Host, Different Home Directory*.)

12.9 Client-Server Interaction

Cylc server programs listen on dedicated network ports for HTTPS communications from Cylc clients (task jobs, and user-invoked commands and GUIs).

Use `cylc scan` to see which suites are listening on which ports on scanned hosts (this lists your own suites by default, but it can show others too - see `cylc scan --help`).

Cylc supports two kinds of access to suite server programs:

- *public* (non-authenticated) - the amount of information revealed is configurable, see [12.9.1](#)
- *control* (authenticated) - full control, suite passphrase required, see [12.9.2](#)

12.9.1 Public Access - No Auth Files

Without a suite passphrase the amount of information revealed by a suite server program is determined by the public access privilege level set in global site/user config ([B.15](#)) and optionally overridden in suites ([A.3.16](#)):

- *identity* - only suite and owner names revealed
- *description* - identity plus suite title and description
- *state-totals* - identity, description, and task state totals
- *full-read* - full read-only access for monitor and GUI
- *shutdown* - full read access plus shutdown, but no other control.

The default public access level is *state-totals*.

The `cylc scan` command and the `cylc gscan` GUI can print descriptions and task state totals in addition to basic suite identity, if the that information is revealed publicly.

12.9.2 Full Control - With Auth Files

Suite auth files (passphrase and SSL certificate) give full control. They are loaded from the suite service directory by the suite server program at start-up, and used to authenticate subsequent client connections. Passphrases are used in a secure encrypted challenge-response scheme, never sent in plain text over the network.

If two users need access to the same suite server program, they must both possess the passphrase file for that suite. Fine-grained access to a single suite server program via distinct user accounts is not currently supported.

Suite server programs automatically install their auth and contact files to job hosts via ssh, to enable task jobs to connect back to the suite server program for task messaging.

Client programs invoked by the suite owner automatically load the passphrase, SSL certificate, and contact file too, for automatic connection to suites.

Manual installation of suite auth files is only needed for remote control, if you do not have a shared filesystem - see below.

12.10 GUI-to-Suite Interaction

The gcylc GUI is mainly a network client to retrieve and display suite status information from the suite server program, but it can also invoke file-reading commands to view and graph the suite definition and so on. This is entirely transparent if the GUI is running on the suite host account, but full functionality for remote suites requires either a shared filesystem, or (see 12.11) auth file installation *and* non-interactive ssh access to the suite host. Without the auth files you will not be able to connect to the suite, and without ssh you will see “permission denied” errors on attempting file access.

12.11 Remote Control

Cylc client programs - command line and GUI - can interact with suite server programs running on other accounts or hosts. How this works depends on whether or not you have:

- a *shared filesystem* such that you see the same home directory on both hosts.
- *non-interactive ssh* from the client account to the server account.

With a shared filesystem, a suite registered on the remote (server) host is also - in effect - registered on the local (client) host. In this case you can invoke client commands without the `--host` option; the client will automatically read the host and port from the contact file in the suite service directory.

To control suite server programs running under other user accounts or on other hosts without a shared filesystem, the suite SSL certificate and passphrase must be installed under your `$HOME/.cylc/` directory:

```
$HOME/.cylc/auth/OWNER@HOST/SUITE/
    ssl.cert
    passphrase
    contact    # (optional - see below)
```

where `OWNER@HOST` is the suite host account and `SUITE` is the suite name. Client commands should then be invoked with the `--user` and `--host` options, e.g.:

```
$ cylc gui --user=OWNER --host=HOST SUITE
```

Note remote suite auth files do not need to be installed for read-only access - see 12.9.1 - via the GUI or monitor.

The suite contact file (see 12.4) is not needed if you have read-access to the remote suite run directory via the local filesystem or non-interactive ssh to the suite host account - client commands will automatically read it. If you do install the contact file in your auth directory note that the port number will need to be updated if the suite gets restarted on a different port.

Otherwise use `cylc scan` to determine the suite port number and use the `--port` client command option.

WARNING: possession of a suite passphrase gives full control over the target suite, including edit run functionality - which lets you run arbitrary scripting on job hosts as the suite owner. Further, non-interactive ssh gives full access to the target user account, so we recommended that this is only used to interact with suites running on accounts to which you already have full access.

12.12 Scan And Gscan

Both `cylc scan` and the `cylc gscan` GUI can display suites owned by other users on other hosts, including task state totals if the public access level permits that (see [12.9.1](#)). Clicking on a remote suite in `gscan` will open a `cylc gui` to connect to that suite. This will give you full control, if you have the suite auth files installed; or it will display full read only information if the public access level allows that.

12.13 Task States Explained

As a suite runs, its task proxies may pass through the following states:

- **waiting** - still waiting for prerequisites (e.g. dependence on other tasks, and clock triggers) to be satisfied.
- **held** - will not be submitted to run even if all prerequisites are satisfied, until released/unheld.
- **queued** - ready to run (prerequisites satisfied) but temporarily held back by an *internal cylc queue* (see [12.17](#)).
- **ready** - ready to run (prerequisites satisfied) and handed to cylc's job submission subsystem.
- **submitted** - submitted to run, but not executing yet (could be waiting in an external batch scheduler queue).
- **submit-failed** - job submission failed *or* submitted job killed (cancelled) before commencing execution.
- **submit-retrying** - job submission failed, but a submission retry was configured. Will only enter the *submit-failed* state if all configured submission retries are exhausted.
- **running** - currently executing (a *task started* message was received, or the task polled as running).
- **succeeded** - finished executing successfully (a *task succeeded* message was received, or the task polled as succeeded).
- **failed** - aborted execution due to some error condition (a *task failed* message was received, or the task polled as failed).
- **retrying** - job execution failed, but an execution retry was configured. Will only enter the *failed* state if all configured execution retries are exhausted.
- **runahead** - will not have prerequisites checked (and so automatically held, in effect) until the rest of the suite catches up sufficiently. The amount of runahead allowed is configurable - see [12.16](#).
- **expired** - will not be submitted to run, due to falling too far behind the wall-clock relative to its cycle point - see [9.3.5.15](#).

12.14 What The Suite Control GUI Shows

The GUI Text-tree and Dot Views display the state of every task proxy present in the task pool. Once a task has succeeded and Cylc has determined that it can no longer be needed to satisfy the prerequisites of other tasks, its proxy will be cleaned up (removed from the pool) and it will disappear from the GUI. To rerun a task that has disappeared from the pool, you need to re-insert its task proxy and then re-trigger it.

The Graph View is slightly different: it displays the complete dependency graph over the range of cycle points currently present in the task pool. This often includes some greyed-out *base* or *ghost nodes* that are empty - i.e. there are no corresponding task proxies currently present in the pool. Base nodes just flesh out the graph structure. Groups of them may be cut out and replaced by single *scissor nodes* in sections of the graph that are currently inactive.

12.15 Network Connection Timeouts

A connection timeout can be set in site and user global config files (see 6) so that messaging commands cannot hang indefinitely if the suite is not responding (this can be caused by suspending a suite with Ctrl-Z) thereby preventing the task from completing. The same can be done on the command line for other suite-connecting user commands, with the `--comms-timeout` option.

12.16 Runahead Limiting

Runahead limiting prevents the fastest tasks in a suite from getting too far ahead of the slowest ones. Newly spawned tasks are released to the task pool only when they fall below the runahead limit. A low runahead limit can prevent cylc from interleaving cycles, but it will not stall a suite unless it fails to extend out past a future trigger (see 9.3.5.11). A high runahead limit may allow fast tasks that are not constrained by dependencies or clock-triggers to spawn far ahead of the pack, which could have performance implications for the suite server program when running very large suites. Succeeded and failed tasks are ignored when computing the runahead limit.

The preferred runahead limiting mechanism restricts the number of consecutive active cycle points. The default value is three active cycle points; see A.4.8. Alternatively the interval between the slowest and fastest tasks can be specified as hard limit; see A.4.7.

12.17 Limiting Activity With Internal Queues

Large suites can potentially overwhelm task hosts by submitting too many tasks at once. You can prevent this with *internal queues*, which limit the number of tasks that can be active (submitted or running) at the same time.

A queue is defined by a *name*; a *limit*, which is the maximum number of active tasks allowed for the queue; and a list of *members*, assigned by task or family name.

Queue configuration is done under the [scheduling] section of the suite.rc file (like dependencies, internal queues constrain *when* a task runs).

By default every task is assigned to the *default* queue, which by default has a zero limit (interpreted by cylc as no limit). To use a single queue for the whole suite just set the default queue limit:

```
[scheduling]
  [[ queues ]]
    # limit the entire suite to 5 active tasks at once
    [[[ defaultt ]]]
      limit = 5
```

To use additional queues just name each one, set their limits, and assign members:

```
[scheduling]
  [[ queues ]]
    [[[ q_foo ]]]
      limit = 5
      members = foo, bar, baz
```

Any tasks not assigned to a particular queue will remain in the default queue. The *queues* example suite illustrates how queues work by running two task trees side by side (as seen in the graph GUI) each limited to 2 and 3 tasks respectively:

```
[meta]
  title = demonstrates internal queueing
  description = """
Two trees of tasks: the first uses the default queue set to a limit of
two active tasks at once; the second uses another queue limited to three
active tasks at once. Run via the graph control GUI for a clear view.
"""

[scheduling]
  [[ queues ]]
    [[[ defaultt ]]]
      limit = 2
    [[[ foo ]]]
      limit = 3
      members = n, o, p, FAM2, u, v, w, x, y, z
  [[ dependencies ]]
    graph = """
      a => b & c => FAM1
      n => o & p => FAM2
      FAM1:succeed-all => h & i & j & k & l & m
      FAM2:succeed-all => u & v & w & x & y & z
    """

[runtime]
  [[ FAM1, FAM2 ]]
    [[ d,e,f,g ]]
      inherit = FAM1
    [[ q,r,s,t ]]
      inherit = FAM2
```

12.18 Automatic Task Retry On Failure

See also [A.5.1.11.6](#) in the *Suite.rc Reference*.

Tasks can be configured with a list of “retry delay” intervals, as ISO 8601 durations. If the task job fails it will go into the *retrying* state and resubmit after the next configured delay interval. An example is shown in the suite listed below under [12.19](#).

If a task with configured retries is *killed* (by `cylc kill` or via the GUI) it goes to the *held* state so that the operator can decide whether to release it and continue the retry sequence or to abort the retry sequence by manually resetting it to the *failed* state.

12.19 Task Event Handling

See also [A.3.13](#) and [A.5.1.13](#) in the *Suite.rc Reference*.

Cylc can call nominated event handlers - to do whatever you like - when certain suite or task events occur. This facilitates centralized alerting and automated handling of critical events. Event handlers can be used to send a message, call a pager, or whatever; they can even intervene in the operation of their own suite using cylc commands.

To send an email, use the built-in setting `[[[events]]]mail events` to specify a list of events for which notifications should be sent. E.g. to send an email on (submission) failed and retry:

```
[runtime]
  [[foo]]
    script = "test ${CYLC_TASK_TRY_NUMBER} -eq 3"
    [[[events]]]
      mail events = submission failed, submission retry, failed, retry
    [[[job]]]
      execution retry delays = PT0S, PT30S
```

By default, the emails will be sent to the current user with:

- to: set as `$USER`
- from: set as `notifications@$(hostname)`
- SMTP server at `localhost:25`

These can be configured using the settings:

- `[[[events]]]mail to` (list of email addresses),
- `[[[events]]]mail from`
- `[[[events]]]mail smtp`.

By default, a cylc suite will send you no more than one task event email every 5 minutes - this is to prevent your inbox from being flooded by emails should a large group of tasks all fail at similar time. See [A.3.8](#) for details.

Event handlers can be located in the suite `bin/` directory; otherwise it is up to you to ensure their location is in `$PATH` (in the shell in which the suite server program runs). They should require little resource and return quickly - as each event handler is invoked by a child process in a finite process pool that is also used to submit, poll and kill jobs. The child process will wait for the event handler to complete before moving on to the next item in the queue. If the process pool is saturated with long running event handlers, the suite will appear to hang.

Task event handlers can be specified using the `[[[events]]]<event> handler` settings, where `<event>` is one of:

- 'submitted' - the job submit command was successful
- 'submission failed' - the job submit command failed
- 'submission timeout' - task job submission timed out
- 'submission retry' - task job submission failed, but will retry after a configured delay
- 'started' - the task reported commencement of execution
- 'succeeded' - the task reported successful completion
- 'warning' - the task reported a WARNING severity message
- 'critical' - the task reported a CRITICAL severity message
- 'custom' - the task reported a CUSTOM severity message
- 'failed' - the task failed

- ‘retry’ - the task failed but will retry after a configured delay
- ‘execution timeout’ - task execution timed out

The value of each setting should be a list of command lines or command line templates (see below).

Alternatively you can use `[[events]]handlers` and `[[events]]handler events`, where the former is a list of command lines or command line templates (see below) and the latter is a list of events for which these commands should be invoked.

Event handler arguments can be constructed from various templates representing suite name; task ID, name, cycle point, message, and submit number name; and any suite or task [meta] item. See [A.3.13](#) and [A.5.1.13](#) for options.

If no template arguments are supplied the following default command line will be used:

```
<task-event-handler> %(event)s %(suite)s %(id)s %(message)s
```

Note: substitution patterns should not be quoted in the template strings. This is done automatically where required.

For an explanation of the substitution syntax, see [String Formatting Operations](#) in the Python documentation.

The retry event occurs if a task fails and has any remaining retries configured (see [12.18](#)). The event handler will be called as soon as the task fails, not after the retry delay period when it is resubmitted.

Note that event handlers are called by the suite server program, not by task jobs. If you wish to pass additional information to them use `[cylc] → [[environment]]`, not task runtime environment.

The following 2 `suite.rc` snippets are examples on how to specify event handlers using the alternate methods:

```
[runtime]
[[foo]]
    script = "test ${CYLC_TASK_TRY_NUMBER} -eq 2"
    [[events]]
        retry handler = "echo '!!!!EVENT!!!!' "
        failed handler = "echo '!!!!EVENT!!!!' "
    [[job]]
        execution retry delays = PT0S, PT30S

[runtime]
[[foo]]
    script = "test ${CYLC_TASK_TRY_NUMBER} -eq 2"
    [[events]]
        handlers = "echo '!!!!EVENT!!!!' "
        handler events = retry, failed
    [[job]]
        execution retry delays = PT0S, PT30S
```

The handler command here - specified with no arguments - is called with the default arguments, like this:

```
echo '!!!!EVENT!!!!' %(event)s %(suite)s %(id)s %(message)s
```

12.20 Handling Job Preemption

Some HPC facilities allow job preemption: the resource manager can kill or suspend running low priority jobs in order to make way for high priority jobs. The preempted jobs may then

be automatically restarted by the resource manager, from the same point (if suspended) or requeued to run again from the start (if killed).

Suspended jobs will poll as still running (their job status file says they started running, and they still appear in the resource manager queue). Loadleveler jobs that are preempted by kill-and-requeue ("job vacation") are automatically returned to the submitted state by Cylc. This is possible because Loadleveler sends the SIGUSR1 signal before SIGKILL for preemption. Other batch schedulers just send SIGTERM before SIGKILL as normal, so Cylc cannot distinguish a preemption job kill from a normal job kill. After this the job will poll as failed (correctly, because it was killed, and the job status file records that). To handle this kind of preemption automatically you could use a task failed or retry event handler that queries the batch scheduler queue (after an appropriate delay if necessary) and then, if the job has been requeued, uses `cylc reset` to reset the task to the submitted state.

12.21 Manual Task Triggering and Edit-Run

Any task proxy currently present in the suite can be manually triggered at any time using the `cylc trigger` command, or from the right-click task menu in gcylc. If the task belongs to a limited internal queue (see 12.17), this will queue it; if not, or if it is already queued, it will submit immediately.

With `cylc trigger --edit` (also in the gcylc right-click task menu) you can edit the generated task job script to make one-off changes before the task submits.

12.22 Cylc Broadcast

The `cylc broadcast` command overrides `[runtime]` settings in a running suite. This can be used to communicate information to downstream tasks by broadcasting environment variables (communication of information from one task to another normally takes place via the filesystem, i.e. the input/output file relationships embodied in inter-task dependencies). Variables (and any other runtime settings) may be broadcast to all subsequent tasks, or targeted specifically at a specific task, all subsequent tasks with a given name, or all tasks with a given cycle point; see broadcast command help for details.

Broadcast settings targeted at a specific task ID or cycle point expire and are forgotten as the suite moves on. Un-targeted variables and those targeted at a task name persist throughout the suite run, even across restarts, unless manually cleared using the broadcast command - and so should be used sparingly.

12.23 The Meaning And Use Of Initial Cycle Point

When a suite is started with the `cylc run` command (cold or warm start) the cycle point at which it starts can be given on the command line or hardwired into the suite.rc file:

```
cylc run foo 20120808T06Z
```

or:

```
[scheduling]
  initial cycle point = 20100808T06Z
```

An initial cycle given on the command line will override one in the suite.rc file.

12.23.1 The Environment Variable `CYLC_SUITE_INITIAL_CYCLE_POINT`

In the case of a *cold start only* the initial cycle point is passed through to task execution environments as `$CYLC_SUITE_INITIAL_CYCLE_POINT`. The value is then stored in suite database files and persists across restarts, but it does get wiped out (set to `None`) after a warm start, because a warm start is really an implicit restart in which all state information is lost (except that the previous cycle is assumed to have completed).

The `$CYLC_SUITE_INITIAL_CYCLE_POINT` variable allows tasks to determine if they are running in the initial cold-start cycle point, when different behaviour may be required, or in a normal mid-run cycle point. Note however that an initial `R1` graph section is now the preferred way to get different behaviour at suite start-up.

12.24 Simulating Suite Behaviour

Several suite run modes allow you to simulate suite behaviour quickly without running the suite's real jobs - which may be long-running and resource-hungry:

- *dummy mode* - runs dummy tasks as background jobs on configured job hosts.
 - simulates scheduling, job host connectivity, and generates all job files on suite and job hosts.
- *dummy-local mode* - runs real dummy tasks as background jobs on the suite host, which allows dummy-running suites from other sites.
 - simulates scheduling and generates all job files on the suite host.
- *simulation mode* - does not run any real tasks.
 - simulates scheduling without generating any job files.

Set the run mode (default *live*) in the GUI suite start dialog box, or on the command line:

```
$ cylc run --mode=dummy SUITE
$ cylc restart --mode=dummy SUITE
```

You can get specified tasks to fail in these modes, for more flexible suite testing. See Section [A.5.1.20](#) for simulation configuration.

12.24.1 Proportional Simulated Run Length

If task `[job]execution time limit` is set, Cylc divides it by `[simulation]speedup factor` (default 10.0) to compute simulated task run lengths (default 10 seconds).

12.24.2 Limitations Of Suite Simulation

Dummy mode ignores batch scheduler settings because Cylc does not know which job resource directives (requested memory, number of compute nodes, etc.) would need to be changed for the dummy jobs. If you need to dummy-run jobs on a batch scheduler manually comment out `script` items and modify directives in your live suite, or else use a custom live mode test suite.

Note that the dummy modes ignore all configured task `script` items including `init-script`. If your `init-script` is required to run even dummy tasks on a job host, note that host environment setup should be done elsewhere - see 3.3.3.

12.24.3 Restarting Suites With A Different Run Mode?

The run mode is recorded in the suite run database files. Cylc will not let you *restart* a non-live mode suite in live mode, or vice versa. To test a live suite in simulation mode just take a quick copy of it and run the the copy in simulation mode.

12.25 Automated Reference Test Suites

Reference tests are finite-duration suite runs that abort with non-zero exit status if any of the following conditions occur (by default):

- cylc fails
- any task fails
- the suite times out (e.g. a task dies without reporting failure)
- a nominated shutdown event handler exits with error status

The default shutdown event handler for reference tests is `cylc hook check-triggering` which compares task triggering information (what triggers off what at run time) in the test run suite log to that from an earlier reference run, disregarding the timing and order of events - which can vary according to the external queueing conditions, runahead limit, and so on.

To prepare a reference log for a suite, run it with the `--reference-log` option, and manually verify the correctness of the reference run.

To reference test a suite, just run it (in dummy mode for the most comprehensive test without running real tasks) with the `--reference-test` option.

A battery of automated reference tests is used to test cylc before posting a new release version. Reference tests can also be used to check that a cylc upgrade will not break your own complex suites - the triggering check will catch any bug that causes a task to run when it shouldn't, for instance; even in a dummy mode reference test the full task job script (sans `script` items) executes on the proper task host by the proper batch system.

Reference tests can be configured with the following settings:

```
[cylc]
  [[reference test]]
    suite shutdown event handler = cylc check-triggering
    required run mode = dummy
    allow task failures = False
    live mode suite timeout = PT5M
    dummy mode suite timeout = PT2M
    simulation mode suite timeout = PT2M
```

12.25.1 Roll-your-own Reference Tests

If the default reference test is not sufficient for your needs, firstly note that you can override the default shutdown event handler, and secondly that the `--reference-test` option is merely a short cut to the following suite.rc settings which can also be set manually if you wish:

```
[cylc]
  abort if any task fails = True
  [[events]]
    shutdown handler = cylc check-triggering
    timeout = PT5M
    abort if shutdown handler fails = True
    abort on timeout = True
```

12.26 Triggering Off Of Tasks In Other Suites

The `cylc suite-state` command interrogates suite run databases. It has a polling mode that waits for a given task in the target suite to achieve a given state. This can be used to make task scripting wait for a remote task to succeed (for example). The suite graph notation also provides a way to define automatic suite-state polling tasks, which use the same polling command under the hood. Note that `cylc suite-state` can only trigger off task *states* in remote suites and does not support triggering off task messages.

Here's how to trigger a task `bar` off a task `foo` in a remote suite called `other.suite`:

```
[scheduling]
  [[dependencies]]
    [[T00, T12]]
      graph = "my-foo<other.suite::foo> => bar"
```

Local task `my-foo` will poll for the success of `foo` in suite `other.suite`, at the same cycle point, succeeding only when or if it succeeds. Other task states can also be polled:

```
graph = "my-foo<other.suite::foo:fail> => bar"
```

The default polling parameters (e.g. maximum number of polls and the interval between them) are printed by `cylc suite-state --help` and can be configured if necessary under the local polling task runtime section:

```
[scheduling]
  [[ dependencies]]
    [[T00,T12]]
      graph = "my-foo<other.suite::foo> => bar"
[runtime]
  [[my-foo]]
    [[suite state polling]]
      max-polls = 100
      interval = PT10S
```

For suites owned by others, or those with run databases in non-standard locations, use the `--run-dir` option, or `in-suite`:

```
[runtime]
  [[my-foo]]
    [[suite state polling]]
      run-dir = /path/to/top/level/cylc/run-directory
```

If the remote task has a different cycling sequence, just arrange for the local polling task to be on the same sequence as the remote task that it represents. For instance, if local task `cat` cycles 6-hourly at 0,6,12,18 but needs to trigger off a remote task `dog` at 3,9,15,21:

```
[scheduling]
  [[dependencies]]
    [[T03,T09,T15,T21]]
      graph = "my-dog<other.suite::dog>"
    [[T00,T06,T12,T18]]
      graph = "my-dog[-PT3H] => cat"
```

For suite-state polling the cycle point of the target task is treated as a literal string so the polling command has to be told if the remote suite has a different cycle point format. Use the `--template` option for this, or `in-suite`:

```
[runtime]
  [[my-foo]]
    [[suite state polling]]
      template = %Y-%m-%dT%H
```

Note that the remote suite does not have to be running when polling commences because the command interrogates the suite run database, not the suite server process.

12.27 Suite Server Logs

Each suite maintains its own log of time-stamped events under the *suite server log directory*:

```
$HOME/cylc-run/SUITE-NAME/log/suite/
```

By way of example, we will show the complete server log generated (at cylc-7.2.0) by a small suite that runs two 30-second dummy tasks `foo` and `bar` for a single cycle point `2017-01-01T00Z` before shutting down:

```
[cylc]
  cycle point format = %Y-%m-%dT%HZ
[scheduling]
  initial cycle point = 2017-01-01T00Z
  final cycle point = 2017-01-01T00Z
  [[dependencies]]
    graph = "foo => bar"
[runtime]
  [[foo]]
    script = sleep 30; /bin/false
  [[bar]]
    script = sleep 30; /bin/true
```

By the task scripting defined above, this suite will stall when `foo` fails. Then, the suite owner `vagrant@cylon` manually resets the failed task's state to *succeeded*, allowing `bar` to trigger and the suite to finish and shut down. Here's the complete suite log for this run:

```
$ cylc cat-log SUITE-NAME
2017-03-30T09:46:10Z INFO - Suite starting: server=localhost:43086 pid=3483
2017-03-30T09:46:10Z INFO - Run mode: live
2017-03-30T09:46:10Z INFO - Initial point: 2017-01-01T00Z
2017-03-30T09:46:10Z INFO - Final point: 2017-01-01T00Z
2017-03-30T09:46:10Z INFO - Cold Start 2017-01-01T00Z
2017-03-30T09:46:11Z INFO - [foo.2017-01-01T00Z] -submit_method_id=3507
2017-03-30T09:46:11Z INFO - [foo.2017-01-01T00Z] -submission succeeded
2017-03-30T09:46:11Z INFO - [foo.2017-01-01T00Z] -(current:submitted)> started
  at 2017-03-30T09:46:10Z
2017-03-30T09:46:41Z CRITICAL - [foo.2017-01-01T00Z] -(current:running)> Task
  job script received signal EXIT at 2017-03-30T09:46:40Z
2017-03-30T09:46:41Z CRITICAL - [foo.2017-01-01T00Z] -(current:running)> failed
  at 2017-03-30T09:46:40Z
2017-03-30T09:46:42Z WARNING - suite stalled
2017-03-30T09:46:42Z WARNING - Unmet prerequisites for bar.2017-01-01T00Z:
2017-03-30T09:46:42Z WARNING - * foo.2017-01-01T00Z succeeded
2017-03-30T09:47:58Z INFO - [client-command] reset_task_states vagrant@cylon:
  cylc-reset 1e0d8e9f-2833-4dc9-a0c8-9cf263c4c8c3
2017-03-30T09:47:58Z INFO - [foo.2017-01-01T00Z] -resetting state to succeeded
2017-03-30T09:47:58Z INFO - Command succeeded: reset_task_states([u'foo.2017'],
  state=succeeded)
2017-03-30T09:47:59Z INFO - [bar.2017-01-01T00Z] -submit_method_id=3565
2017-03-30T09:47:59Z INFO - [bar.2017-01-01T00Z] -submission succeeded
2017-03-30T09:47:59Z INFO - [bar.2017-01-01T00Z] -(current:submitted)> started
  at 2017-03-30T09:47:58Z
```

```

2017-03-30T09:48:29Z INFO - [bar.2017-01-01T00Z] -(current:running)> succeeded
    at 2017-03-30T09:48:28Z
2017-03-30T09:48:30Z INFO - Waiting for the command process pool to empty for
    shutdown
2017-03-30T09:48:30Z INFO - Suite shutting down - AUTOMATIC

```

The information logged here includes:

- event timestamps, at the start of each line
- suite server host, port and process ID
- suite initial and final cycle points
- suite start type (cold start in this case)
- task events (task started, succeeded, failed, etc.)
- suite stalled warning (in this suite nothing else can run when `foo` fails)
- the client command issued by `vagrant@cylon` to reset `foo` to *succeeded*
- job IDs - in this case process IDs for background jobs (or PBS job IDs etc.)
- state changes due to incoming task progress message ("started at ..." etc.) suite shutdown time and reasons (AUTOMATIC means "all tasks finished and nothing else to do")

Note that suite log files are primarily intended for human eyes. If you need to have an external system to monitor suite events automatically, interrogate the sqlite *suite run database* (see [12.28](#)) rather than parse the log files.

12.28 Suite Run Databases

Suite server programs maintain two `sqlite` databases to record restart checkpoints and various other aspects of run history:

```

$HOME/cylc-run/SUITE-NAME/log/db # public suite DB
$HOME/cylc-run/SUITE-NAME/.service/db # private suite DB

```

The private DB is for use only by the suite server program. The identical public DB is provided for use by external commands such as `cylc suite-state`, `cylc ls-checkpoints`, and `cylc report -timings`. If the public DB gets locked for too long by an external reader, the suite server program will eventually delete it and replace it with a new copy of the private DB, to ensure that both correctly reflect the suite state.

You can interrogate the public DB with the `sqlite3` command line tool, the `sqlite3` module in the Python standard library, or any other `sqlite` interface.

```

$ sqlite3 ~/cylc-run/foo/log/db << _END_
> .headers on
> select * from task_events where name is "foo";
> _END_
name|cycle|time|submit_num|event|message
foo|1|2017-03-12T11:06:09Z|1|submitted|
foo|1|2017-03-12T11:06:09Z|1|output completed|started
foo|1|2017-03-12T11:06:09Z|1|started|
foo|1|2017-03-12T11:06:19Z|1|output completed|succeeded
foo|1|2017-03-12T11:06:19Z|1|succeeded|

```

12.29 Disaster Recovery

If a suite run directory gets deleted or corrupted, the options for recovery are:

- restore the run directory from back-up, and restart the suite
- re-install from source, and warm start from the beginning of the current cycle point

A warm start (see 12.1.2) does not need a suite state checkpoint, but it wipes out prior run history, and it could re-run a significant number of tasks that had already completed.

To restart the suite, the critical Cylc files that must be restored are:

```
# On the suite host:
~/cylc-run/SUITE-NAME/
  suite.rc      # live suite definition (located here in Rose suites)
  log/db        # public suite DB (can just be a copy of the private DB)
  log/rose-suite-run.conf # (needed to restart a Rose suite)
  .service/db   # private suite DB
  .service/source -> PATH-TO-SUITE-DIR # symlink to live suite directory

# On job hosts (if no shared filesystem):
~/cylc-run/SUITE-NAME/
  log/job/CYCLE-POINT/TASK-NAME/SUBMIT-NUM/job.status
```

Note this discussion does not address restoration of files generated and consumed by task jobs at run time. How suite data is stored and recovered in your environment is a matter of suite and system design.

In short, you can simply restore the suite service directory, the log directory, and the suite.rc file that is the target of the symlink in the service directory. The service and log directories will come with extra files that aren't strictly needed for a restart, but that doesn't matter - although depending on your log housekeeping the log/job directory could be huge, so you might want to be selective about that. (Also in a Rose suite, the suite.rc file does not need to be restored if you restart with `rose suite-run` - which re-installs suite source files to the run directory).

The public DB is not strictly required for a restart - the suite server program will recreate it if need be - but it is required by `cylc ls-checkpoints` if you need to identify the right restart checkpoint.

The job status files are only needed if the restart suite state checkpoint contains active tasks that need to be polled to determine what happened to them while the suite was down. Without them, polling will fail and those tasks will need to be manually set to the correct state.

WARNING: it is not safe to copy or rsync a potentially-active sqlite DB - the copy might end up corrupted. It is best to stop the suite before copying a DB, or else write a back-up utility using the official sqlite backup API: <http://www.sqlite.org/backup.html>.

13 Suite Storage, Discovery, Revision Control, and Deployment

Small groups of cylc users can of course share suites by manual copying, and generic revision control tools can be used on cylc suites as for any collection of files. Beyond this cylc does not have a built-in solution for suite storage and discovery, revision control, and deployment, on a network. That is not cylc's core purpose, and large sites may have preferred revision control systems and suite meta-data requirements that are difficult to anticipate. We can, however, recommend the use of *Rose* to do all of this very easily and elegantly with cylc suites.

13.1 Rose

Rose is a *framework for managing and running suites of scientific applications*, developed at the UK Met Office for use with cylc. It is available under the open source GPL license.

- Rose documentation: <http://metomi.github.io/rose/doc/rose.html>
- Rose source repository: <https://github.com/metomi/rose>

A Suite.rc Reference

This appendix defines all legal suite definition config items. Embedded Jinja2 code (see 9.7) must process to a valid raw suite.rc file. See also 9.2 for a descriptive overview of suite.rc files, including syntax (9.2.1).

A.1 Top Level Items

The only top level configuration items at present are the suite title and description.

A.2 [meta]

Section containing metadata items for this suite. Several items (title, description, URL) are pre-defined and are used by the GUI. Others can be user-defined and passed to suite event handlers to be interpreted according to your needs. For example, the value of a “suite-priority” item could determine how an event handler responds to failure events.

A.2.1 [meta] →title

A single line description of the suite. It is displayed in the GUI “Open Another Suite” window and can be retrieved at run time with the `cylc show` command.

- *type*: single line string
- *default*: (none)

A.2.2 [meta] →description

A multi-line description of the suite. It can be retrieved at run time with the `cylc show` command.

- *type*: multi-line string
- *default*: (none)

A.2.3 [meta] →URL

A web URL to suite documentation. If present it can be browsed with the `cylc doc` command, or from the gcylc Suite menu. The string template `%(suite_name)s` will be replaced with the actual suite name. See also task URLs (A.5.1.10.3).

- *type*: string (URL)
- *default*: (none)
- *example*: `http://my-site.com/suites/%(suite_name)s/index.html`

A.2.4 group

[meta] →group

A group name for a suite. In the gscan GUI, suites with the same group name can be collapsed into a single state summary when the “group” column is displayed.

- *type*: single line string
- *default*: (none)

A.2.5 [meta] →__MANY__

Replace __MANY__ with any user-defined metadata item. These, like title, URL, etc. can be passed to suite event handlers to be interpreted according to your needs. For example, “suite-priority”.

- *type*: String or integer
- *default*: (none)
- *example*:

```
[meta]
    suite-priority = high
```

A.3 [cylc]

This section is for configuration that is not specifically task-related.

A.3.1 [cylc] →required run mode

If this item is set cylc will abort if the suite is not started in the specified mode. This can be used for demo suites that have to be run in simulation mode, for example, because they have been taken out of their normal operational context; or to prevent accidental submission of expensive real tasks during suite development.

- *type*: string
- *legal values*: live, dummy, dummy-local, simulation
- *default*: None

A.3.2 [cylc] →UTC mode

Cylc runs off the suite host’s system clock by default. This item allows you to run the suite in UTC even if the system clock is set to local time. Clock-trigger tasks will trigger when the current UTC time is equal to their cycle point date-time plus offset; other time values used, reported, or logged by the suite server program will usually also be in UTC. The default for this can be set at the site level (see [B.14.1](#)).

- *type*: boolean
- *default*: False, unless overridden at site level.

A.3.3 [cylc] →cycle point format

To just alter the timezone used in the date-time cycle point format, see [A.3.5](#). To just alter the number of expanded year digits (for years below 0 or above 9999), see [A.3.4](#).

Cylc usually uses a `CCYYMMDDThhmmZ` (z in the special case of UTC) or `CCYYMMDDThhmm+hhmm` format (+ standing for + or - here) for writing down date-time cycle points, which follows one of the basic formats outlined in the ISO 8601 standard. For example, a cycle point on the 3rd of February 2001 at 4:50 in the morning, UTC (+0000 timezone), would be written `20010203T0450Z`. Similarly, for the the 3rd of February 2001 at 4:50 in the morning, +1300 timezone, cylc would write `20010203T0450+1300`.

You may use the isodatetime library's syntax to write dates and times in ISO 8601 formats - `cc` for century, `yy` for decade and decadal year, `+x` for expanded year digits and their positive or negative sign, thereafter following the ISO 8601 standard example notation except for fractional digits, which are represented as `,ii` for `hh`, `,nn` for `mm`, etc. For example, to write date-times as week dates with fractional hours, set cycle point format to `CCYYWwwDThh,iiZ` e.g. `1987W041T08,5Z` for 08:30 UTC on Monday on the fourth ISO week of 1987.

You can also use a subset of the `strptime/strftime` POSIX standard - supported tokens are `%F`, `%H`, `%M`, `%S`, `%Y`, `%d`, `%j`, `%m`, `%s`, `%z`.

The ISO8601 extended date-time format can be used (`%Y-%m-%dT%H:%M`) but note that the '-' and ':' characters end up in job log directory paths.

The pre cylc-6 legacy 10-digit date-time format `YYYYMMDDHH` is not ISO8601 compliant and can no longer be used as the cycle point format. For job scripts that still require the old format, use the `cylc cyclepoint` utility to translate the ISO8601 cycle point inside job scripts, e.g.:

```
[runtime]
  [[root]]
    [[environment]]
      CYCLE_TIME = $(cylc cyclepoint --template=%Y%m%d%H)
```

A.3.4 [cylc] →cycle point num expanded year digits

For years below 0 or above 9999, the ISO 8601 standard specifies that an extra number of year digits and a sign should be used. This extra number needs to be written down somewhere (here).

For example, if this extra number is set to 2, 00Z on the 1st of January in the year 10040 will be represented as `+0100400101T0000Z` (2 extra year digits used). With this number set to 3, 06Z on the 4th of May 1985 would be written as `+00019850504T0600Z`.

This number defaults to 0 (no sign or extra digits used).

A.3.5 [cylc] →cycle point time zone

If you set UTC mode to True ([A.3.2](#)) then this will default to z. If you use a custom cycle point format ([A.3.3](#)), you should specify the timezone choice (or null timezone choice) here as well.

You may set your own time zone choice here, which will be used for all date-time cycle point dumping. Time zones should be expressed as ISO 8601 time zone offsets from UTC, such as `+13`, `+1300`, `-0500` or `+0645`, with `Z` representing the special `+0000` case. Cycle points will be converted to the time zone you give and will be represented with this string at the end.

Cycle points that are input without time zones (e.g. as an initial cycle point setting) will use this time zone if set. If this isn't set (and UTC mode is also not set), then they will default to the current local time zone.

Note that the ISO standard also allows writing the hour and minute separated by a ":" (e.g. `+13:00`) - however, this is not recommended, given that the time zone is used as part of task output filenames.

A.3.6 [cylc] →abort if any task fails

Cylc does not normally abort if tasks fail, but if this item is turned on it will abort with exit status 1 if any task fails.

- *type*: boolean
- *default*: False

A.3.7 [cylc] →health check interval

Specify the time interval on which a running cylc suite will check that its run directory exists and that its contact file contains the expected information. If not, the suite will shut itself down automatically.

- *type*: ISO 8601 duration/interval representation (e.g. `PT5M`, 5 minutes (note: by contrast, `P5M` means 5 months, so remember the T!)).
- *default*: `PT10M`

A.3.8 [cylc] →task event mail interval

Group together all the task event mail notifications into a single email within a given interval. This is useful to prevent flooding users' mail boxes when many task events occur within a short period of time.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT5M`

A.3.9 [cylc] →disable automatic shutdown

This has the same effect as the `--no-auto-shutdown` flag for the suite run commands: it prevents the suite server program from shutting down normally when all tasks have finished (a suite timeout can still be used to stop the daemon after a period of inactivity, however). This option can make it easier to re-trigger tasks manually near the end of a suite run, during suite development and debugging.

- *type*: boolean

- *default*: False

A.3.10 [cylc] →log resolved dependencies

If this is turned on cylc will write the resolved dependencies of each task to the suite log as it becomes ready to run (a list of the IDs of the tasks that actually satisfied its prerequisites at run time). Mainly used for cylc testing and development.

- *type*: boolean
- *default*: False

A.3.11 [cylc] →[[parameters]]

Define parameter values here for use in expanding *parameterized tasks* - see Section 9.6.

- *type*: list of strings, or an integer range LOWER..UPPER..STEP (two dots, inclusive bounds, STEP optional)
- *default*: (none)
- *examples*:
 - `run = control, test1, test2`
 - `mem = 1..5` (equivalent to 1, 2, 3, 4, 5).
 - `mem = -11..-7..2` (equivalent to -11, -9, -7).

A.3.12 [cylc] →[[parameter templates]]

Parameterized task names (see previous item, and Section 9.6) are expanded, for each parameter value, using string templates. You can assign templates to parameter names here, to override the default templates.

- *type*: a Python-style string template
- *default* for integer parameters `p`: `_p%(p)0Nd`
where N is the number of digits of the maximum integer value, e.g. `foo<run>` becomes `foo_run3` for `run` value 3.
- *default* for non-integer parameters `p`: `_%(p)s`
e.g. `foo<run>` becomes `foo_top` for `run` value `top`.
- *example*: `run = -R%(run)s`
e.g. `foo<run>` becomes `foo-R3` for `run` value 3.

Note that the values of a parameter named `p` are substituted for `%(p)s`. In `_run%(run)s` the first “run” is a string literal, and the second gets substituted with each value of the parameter.

A.3.13 [cylc] →[[events]]

Cylc has internal “hooks” to which you can attach handlers that are called by the suite server program whenever certain events occur. This section configures suite event hooks; see A.5.1.13 for task event hooks.

Event handler commands can send an email or an SMS, call a pager, intervene in the operation of their own suite, or whatever. They can be held in the suite bin directory, otherwise it is up

to you to ensure their location is in `$PATH` (in the shell in which cylc runs, on the suite host). The commands should require very little resource to run and should return quickly.

Each event handler can be specified as a list of command lines or command line templates.

A command line template may have any or all of these patterns which will be substituted with actual values:

- `%(event)s`: event name (see below)
- `%(suite)s`: suite name
- `%(suite_url)s`: suite URL
- `%(message)s`: event message, if any
- any suite [meta] item, e.g.:
 - `%(title)s`: suite title
 - `%(importance)s`: example custom suite metadata

Otherwise the command line will be called with the following default arguments:

```
<suite-event-handler> %(event)s %(suite)s %(message)s
```

Note: substitution patterns should not be quoted in the template strings. This is done automatically where required.

Additional information can be passed to event handlers via `[cylc] → [[environment]]`.

A.3.13.1 [cylc] → [[events]] → EVENT handler

A comma-separated list of one or more event handlers to call when one of the following EVENTS occurs:

- **startup** - the suite has started running
- **shutdown** - the suite is shutting down
- **timeout** - the suite has timed out
- **stalled** - the suite has stalled
- **inactivity** - the suite is inactive

Default values for these can be set at the site level via the `siterc` file (see [B.14.4](#)).

Item details:

- *type*: string (event handler script name)
- *default*: None, unless defined at the site level.
- *example*: `startup handler = my-handler.sh`

A.3.13.2 [cylc] → [[[events]]] → handlers

Specify the general event handlers as a list of command lines or command line templates.

- *type*: Comma-separated list of strings (event handler command line or command line templates).
- *default*: (none)
- *example*: `handlers = my-handler.sh`

A.3.13.3 [cylc] →[[events]] →handler events

Specify the events for which the general event handlers should be invoked.

- *type*: Comma-separated list of events
- *default*: (none)
- *example*: `handler events = timeout, shutdown`

A.3.13.4 [cylc] →[[events]] →mail events

Specify the suite events for which notification emails should be sent.

- *type*: Comma-separated list of events
- *default*: (none)
- *example*: `mail events = startup, shutdown, timeout`

A.3.13.5 [cylc] →[[events]] →mail footer

Specify a string or string template to insert to footers of notification emails for both suite events and task events.

A template string may have any or all of these patterns which will be substituted with actual values:

- `%(host)s`: suite host name
- `%(port)s`: suite port number
- `%(owner)s`: suite owner name
- `%(suite)s`: suite name
- *type*:
- *default*: (none)
- *example*: `mail footer = see: http://localhost/%(owner)s/notes-on/%(suite)s/`

A.3.13.6 [cylc] →[[events]] →mail from

Specify an alternate `from`: email address for suite event notifications.

- *type*: string
- *default*: None, (notifications@HOSTNAME)
- *example*: `mail from = no-reply@your-org`

A.3.13.7 [cylc] →[[events]] →mail smtp

Specify the SMTP server for sending suite event email notifications.

- *type*: string
- *default*: None, (localhost:25)
- *example*: `mail smtp = smtp.yourorg`

A.3.13.8 [cylc] →[[events]] →mail to

A list of email addresses to send suite event notifications. The list can be anything accepted by the `mail` command.

- *type*: string
- *default*: None, (USER@HOSTNAME)
- *example*: `mail to = your.colleague`

A.3.13.9 [cylc] →[[events]] →timeout

If a timeout is set and the timeout event is handled, the timeout event handler(s) will be called if the suite stays in a stalled state for some period of time. The timer is set initially at suite start up. It is possible to set a default for this at the site level (see [B.14.4](#)).

- *type*: ISO 8601 duration/interval representation (e.g. `PT5S`, 5 seconds, `PT1S`, 1 second) - minimum 0 seconds.
- *default*: (none), unless set at the site level.

A.3.13.10 [cylc] →[[events]] →inactivity

If inactivity is set and the inactivity event is handled, the inactivity event handler(s) will be called if there is no activity in the suite for some period of time. The timer is set initially at suite start up. It is possible to set a default for this at the site level (see [B.14.4](#)).

- *type*: ISO 8601 duration/interval representation (e.g. `PT5S`, 5 seconds, `PT1S`, 1 second) - minimum 0 seconds.
- *default*: (none), unless set at the site level.

A.3.13.11 [cylc] →[[events]] →reset timer

If `True` (the default) the suite timer will continually reset after any task changes state, so you can time out after some interval since the last activity occurred rather than on absolute suite execution time.

- *type*: boolean
- *default*: True

A.3.13.12 [cylc] →[[events]] →abort on stalled

If this is set to True it will cause the suite to abort with error status if it stalls. A suite is considered "stalled" if there are no active, queued or submitting tasks or tasks waiting for clock triggers to be met. It is possible to set a default for this at the site level (see [B.14.4](#)).

- *type*: boolean
- *default*: False, unless set at the site level.

A.3.13.13 [cylc] →[[events]] →abort on timeout

If a suite timer is set (above) this will cause the suite to abort with error status if the suite times out while still running. It is possible to set a default for this at the site level (see [B.14.4](#)).

- *type*: boolean
- *default*: False, unless set at the site level.

A.3.13.14 [cylc] →[[events]] →abort on inactivity

If a suite inactivity timer is set (above) this will cause the suite to abort with error status if the suite is inactive for some period while still running. It is possible to set a default for this at the site level (see [B.14.4](#)).

- *type*: boolean
- *default*: False, unless set at the site level.

A.3.13.15 [cylc] →[[events]] →abort if EVENT handler fails

Cylc does not normally care whether an event handler succeeds or fails, but if this is turned on the EVENT handler will be executed in the foreground (which will block the suite while it is running) and the suite will abort if the handler fails.

- *type*: boolean
- *default*: False

A.3.14 [cylc] →[[environment]]

Environment variables defined in this section are passed to suite and task event handlers.

- These variables are not passed to tasks - use task runtime variables for that. Similarly, task runtime variables are not available to event handlers - which are executed by the suite server program, (not by running tasks) in response to task events.
- Cylc-defined environment variables such as `$CYLC_SUITE_RUN_DIR` are not passed to task event handlers by default, but you can make them available by extracting them to the cylc environment like this:

```
[cylc]
[[environment]]
    CYLC_SUITE_RUN_DIR = $CYLC_SUITE_RUN_DIR
```

- These variables - unlike task execution environment variables which are written to job scripts and interpreted by the shell at task run time - are not interpreted by the shell prior to use so shell variable expansion expressions cannot be used here.

A.3.14.1 [cylc] →[[environment]] →__VARIABLE__

Replace `__VARIABLE__` with any number of environment variable assignment expressions. Values may refer to other local environment variables (order of definition is preserved) and are not evaluated or manipulated by cylc, so any variable assignment expression that is legal in the shell in which cylc is running can be used (but see the warning above on variable expansions,

which will not be evaluated). White space around the '=' is allowed (as far as cylc's file parser is concerned these are just suite configuration items).

- *type*: string
- *default*: (none)
- *examples*:
 - `FOO = $HOME/foo`

A.3.15 [cylc] →[[reference test]]

Reference tests are finite-duration suite runs that abort with non-zero exit status if cylc fails, if any task fails, if the suite times out, or if a shutdown event handler that (by default) compares the test run with a reference run reports failure. See [12.25](#).

A.3.15.1 [cylc] →[[reference test]] →suite shutdown event handler

A shutdown event handler that should compare the test run with the reference run, exiting with zero exit status only if the test run verifies.

- *type*: string (event handler command name or path)
- *default*: `cylc hook check-triggering`

As for any event handler, the full path can be omitted if the script is located somewhere in `$PATH` or in the suite bin directory.

A.3.15.2 [cylc] →[[reference test]] →required run mode

If your reference test is only valid for a particular run mode, this setting will cause cylc to abort if a reference test is attempted in another run mode.

- *type*: string
- *legal values*: live, dummy, dummy-local, simulation
- *default*: None

A.3.15.3 [cylc] →[[reference test]] →allow task failures

A reference test run will abort immediately if any task fails, unless this item is set, or a list of *expected task failures* is provided (below).

- *type*: boolean
- *default*: False

A.3.15.4 [cylc] →[[reference test]] →expected task failures

A reference test run will abort immediately if any task fails, unless *allow task failures* is set (above) or the failed task is found in a list IDs of tasks that are expected to fail.

- *type*: Comma-separated list of strings (task IDs: `name.cycle_point`).
- *default*: (none)

- *example:* `foo.20120808, bar.20120908`

A.3.15.5 [cylc] →[[reference test]] →live mode suite timeout

The timeout value, expressed as an ISO 8601 duration/interval, after which the test run should be aborted if it has not finished, in live mode. Test runs cannot be done in live mode unless you define a value for this item, because it is not possible to arrive at a sensible default for all suites.

- *type:* ISO 8601 duration/interval representation, e.g. `PT5M` is 5 minutes (note: by contrast `P5M` means 5 months, so remember the `T`!).
- *default:* `PT1M` (1 minute)

A.3.15.6 [cylc] →[[reference test]] →simulation mode suite timeout

The timeout value in minutes after which the test run should be aborted if it has not finished, in simulation mode. Test runs cannot be done in simulation mode unless you define a value for this item, because it is not possible to arrive at a sensible default for all suites.

- *type:* ISO 8601 duration/interval representation (e.g. `PT5M`, 5 minutes (note: by contrast, `P5M` means 5 months, so remember the `T`!)).
- *default:* `PT1M` (1 minute)

A.3.15.7 [cylc] →[[reference test]] →dummy mode suite timeout

The timeout value, expressed as an ISO 8601 duration/interval, after which the test run should be aborted if it has not finished, in dummy mode. Test runs cannot be done in dummy mode unless you define a value for this item, because it is not possible to arrive at a sensible default for all suites.

- *type:* ISO 8601 duration/interval representation (e.g. `PT5M`, 5 minutes (note: by contrast, `P5M` means 5 months, so remember the `T`!)).
- *default:* `PT1M` (1 minute)

A.3.16 [cylc] →[[authentication]]

Authentication of client programs with suite server programs can be set in the global site/user config files and overridden here if necessary. See [B.15](#) for more information.

A.3.16.1 [cylc] →[[authentication]] →public

The client privilege level granted for public access - i.e. no suite passphrase required. See [B.15](#) for legal values.

A.3.17 [cylc] →[[simulation]]

Suite-level configuration for the *simulation* and *dummy* run modes described in [Section 12.24](#).

A.3.17.1 [cylc] →[[simulation]] →disable suite event handlers

If this is set to `True` configured suite event handlers will not be called in simulation or dummy modes.

- *type*: boolean
- *default*: `True`

A.4 [scheduling]

This section allows cylc to determine when tasks are ready to run.

A.4.1 [scheduling] →cycling mode

Cylc runs using the proleptic Gregorian calendar by default. This item allows you to either run the suite using the 360 day calendar (12 months of 30 days in a year) or using integer cycling.

- *type*: string
- *legal values*: gregorian, 360day, integer
- *default*: gregorian

A.4.2 [scheduling] →initial cycle point

In a cold start each cycling task (unless specifically excluded under [special tasks]) will be loaded into the suite with this cycle point, or with the closest subsequent valid cycle point for the task. This item can be overridden on the command line or in the gcylc suite start panel.

In date-time cycling, if you do not provide time zone information for this, it will be assumed to be local time, or in UTC if [A.3.2](#) is set, or in the time zone determined by [A.3.5](#) if that is set.

- *type*: ISO 8601 date-time point representation (e.g. `CCYYMMDDThhmm`, `19951231T0630`) or “now”.
- *default*: (none)

The string “now” converts to the current date-time on the suite host (adjusted to UTC if the suite is in UTC mode but the host is not) to minute resolution. Minutes (or hours, etc.) may be ignored depending on your cycle point format ([A.3.3](#)).

A.4.3 [scheduling] →final cycle point

Cycling tasks are held once they pass the final cycle point, if one is specified. Once all tasks have achieved this state the suite will shut down. If this item is provided you can override it on the command line or in the gcylc suite start panel.

In date-time cycling, if you do not provide time zone information for this, it will be assumed to be local time, or in UTC if [A.3.2](#) is set, or in the [A.3.5](#) if that is set.

- *type*: ISO 8601 date-time point representation (e.g. `CCYYMMDDThhmm`, `19951231T1230`) or ISO 8601 date-time offset (e.g. `+P1D+PT6H`)
- *default*: (none)

A.4.4 [scheduling] →initial cycle point constraints

In a cycling suite it is possible to restrict the initial cycle point by defining a list of truncated time points under the initial cycle point constraints.

- *type*: Comma-separated list of ISO 8601 truncated time point representations (e.g. T00, T06, T-30).
- *default*: (none)

A.4.5 [scheduling] →final cycle point constraints

In a cycling suite it is possible to restrict the final cycle point by defining a list of truncated time points under the final cycle point constraints.

- *type*: Comma-separated list of ISO 8601 truncated time point representations (e.g. T00, T06, T-30).
- *default*: (none)

A.4.6 [scheduling] →hold after point

Cycling tasks are held once they pass the hold after cycle point, if one is specified. Unlike the final cycle point suite will not shut down once all tasks have passed this point. If this item is provided you can override it on the command line or in the gcylc suite start panel.

A.4.7 [scheduling] →runahead limit

Runahead limiting prevents the fastest tasks in a suite from getting too far ahead of the slowest ones, as documented in [12.16](#).

This config item specifies a hard limit as a cycle interval between the slowest and fastest tasks. It is deprecated in favour of the newer default limiting by `max active cycle points` ([A.4.8](#)).

- *type*: Cycle interval string e.g. PT12H for a 12 hour limit under ISO 8601 cycling.
- *default*: (none)

A.4.8 [scheduling] →max active cycle points

Runahead limiting prevents the fastest tasks in a suite from getting too far ahead of the slowest ones, as documented in [12.16](#).

This config item supersedes the deprecated hard `runahead limit` ([A.4.7](#)). It allows up to `N` (default 3) consecutive cycle points to be active at any time, adjusted up if necessary for any future triggering.

- *type*: integer
- *default*: 3

A.4.9 [scheduling] →spawn to max active cycle points

Allows tasks to spawn out to `max active cycle points` (A.4.8), removing restriction that a task has to have submitted before its successor can be spawned.

Important: This should be used with care given the potential impact of additional task proxies both in terms of memory and cpu for the cylc daemon as well as overheads in rendering all the additional tasks in gcylc. Also, use of the setting may highlight any issues with suite design relying on the default behaviour where downstream tasks would otherwise be waiting on ones upstream submitting and the suite would have stalled e.g. a housekeeping task at a later cycle deleting an earlier cycle's data before that cycle has had chance to run where previously the task would not have been spawned until its predecessor had been submitted.

- *type:* boolean
- *default:* False

A.4.10 [scheduling] →[[queues]]

Configuration of internal queues, by which the number of simultaneously active tasks (submitted or running) can be limited, per queue. By default a single queue called *default* is defined, with all tasks assigned to it and no limit. To use a single queue for the whole suite just set the limit on the *default* queue as required. See also 12.17.

A.4.10.1 [scheduling] →[[queues]] →[[[__QUEUE__]]]

Section heading for configuration of a single queue. Replace `__QUEUE__` with a queue name, and repeat the section as required.

- *type:* string
- *default:* “default”

A.4.10.2 [scheduling] →[[queues]] →[[[__QUEUE__]]] →limit

The maximum number of active tasks allowed at any one time, for this queue.

- *type:* integer
- *default:* 0 (i.e. no limit)

A.4.10.3 [scheduling] →[[queues]] →[[[__QUEUE__]]] →members

A list of member tasks, or task family names, to assign to this queue (assigned tasks will automatically be removed from the default queue).

- *type:* Comma-separated list of strings (task or family names).
- *default:* none for user-defined queues; all tasks for the “default” queue

A.4.11 [scheduling] → [[special tasks]]

This section is used to identify tasks with special behaviour. Family names can be used in special task lists as shorthand for listing all member tasks.

A.4.11.1 [scheduling] → [[special tasks]] → clock-trigger

Clock-trigger tasks (see 9.3.5.14) wait on a wall clock time specified as an offset from their own cycle point.

- *type*: Comma-separated list of task or family names with associated date-time offsets expressed as ISO8601 interval strings, positive or negative, e.g. `PT1H` for 1 hour. The offset specification may be omitted to trigger right on the cycle point.
- *default*: (none)
- *example*:

```
clock-trigger = foo(PT1H30M), bar(PT1.5H), baz
```

A.4.11.2 [scheduling] → [[special tasks]] → clock-expire

Clock-expire tasks enter the *expired* state and skip job submission if too far behind the wall clock when they become ready to run. The expiry time is specified as an offset from wall-clock time; typically it should be negative - see 9.3.5.15.

- *type*: Comma-separated list of task or family names with associated date-time offsets expressed as ISO8601 interval strings, positive or negative, e.g. `PT1H` for 1 hour. The offset may be omitted if it is zero.
- *default*: (none)
- *example*:

```
clock-expire = foo(-P1D)
```

A.4.11.3 [scheduling] → [[special tasks]] → external-trigger

Externally triggered tasks (see 9.3.5.16) wait on external events reported via the `cylc ext-trigger` command. To constrain triggers to a specific cycle point, include `$CYLC_TASK_CYCLE_POINT` in the trigger message string and pass the cycle point to the `cylc ext-trigger` command.

- *type*: Comma-separated list of task names with associated external trigger message strings.
- *default*: (none)
- *example*: (note the comma and line-continuation character)

```
external-trigger = get-satx("new sat-X data ready"), \
                  get-saty("new sat-Y data ready for
                           $CYLC_TASK_CYCLE_POINT")
```

A.4.11.4 [scheduling] → [[special tasks]] → sequential

Sequential tasks are automatically given dependence on their own predecessor. This is equivalent to use of explicit inter-cycle triggers in the graph, except that the automatic version does not show in suite graph visualization. For more on sequential tasks see 9.3.5.12.

- *type*: Comma-separated list of task or family names.
- *default*: (none)
- *example*: `sequential = foo, bar`

A.4.11.5 [scheduling] →[[special tasks]] →exclude at start-up

Any task listed here will be excluded from the initial task pool (this goes for suite restarts too). If an *inclusion* list is also specified, the initial pool will contain only included tasks that have not been excluded. Excluded tasks can still be inserted at run time. Other tasks may still depend on excluded tasks if they have not been removed from the suite dependency graph, in which case some manual triggering, or insertion of excluded tasks, may be required.

- *type*: Comma-separated list of task or family names.
- *default*: (none)

A.4.11.6 [scheduling] →[[special tasks]] →include at start-up

If this list is not empty, any task *not* listed in it will be excluded from the initial task pool (this goes for suite restarts too). If an *exclusion* list is also specified, the initial pool will contain only included tasks that have not been excluded. Excluded tasks can still be inserted at run time. Other tasks may still depend on excluded tasks if they have not been removed from the suite dependency graph, in which case some manual triggering, or insertion of excluded tasks, may be required.

- *type*: Comma-separated list of task or family names.
- *default*: (none)

A.4.12 [scheduling] →[[dependencies]]

The suite dependency graph is defined under this section. You can plot the dependency graph as you work on it, with `cylc graph` or by right clicking on the suite in the db viewer. See also [9.3](#).

A.4.12.1 [scheduling] →[[dependencies]] →graph

The dependency graph for a completely non-cycling suites can go here. See also [A.4.12.2.1](#) below and [9.3](#), for graph string syntax.

- *type*: string
- *example*: (see [A.4.12.2.1](#) below)

A.4.12.2 [scheduling] →[[dependencies]] →[[__RECURRENCE__]]

`__RECURRENCE__` section headings define the sequence of cycle points for which the subsequent graph section is valid. These should be specified in our ISO 8601 derived sequence syntax, or similar for integer cycling:

- *examples*:
 - date-time cycling: `[[T00,T06,T12,T18]]` or `[[PT6H]]`

- integer cycling (stepped by 2): `[[P2]]`
- *default*: (none)

See 9.3.3 for more on recurrence expressions, and how multiple graph sections combine.

A.4.12.2.1 [scheduling] → [[dependencies]] → [[__RECCURRENCE__]] → graph

The dependency graph for a given recurrence section goes here. Syntax examples follow; see also 9.3 and 9.3.5.

- *type*: string
- *examples*:


```
graph = """
    foo => bar => baz & waz      # baz and waz both trigger off bar
    foo[-P1D-PT6H] => bar      # bar triggers off foo[-P1D-PT6H]
    baz:out1 => faz             # faz triggers off a message output of baz
    X:start => Y                # Y triggers if X starts executing
    X:fail => Y                 # Y triggers if X fails
    foo[-PT6H]:fail => bar      # bar triggers if foo[-PT6H] fails
    X => !Y                     # Y suicides if X succeeds
    X | X:fail => Z             # Z triggers if X succeeds or fails
    X:finish => Z              # Z triggers if X succeeds or fails
    (A | B & C) | D => foo      # general conditional triggers
    foo:submit => bar           # bar triggers if foo is successfully
                                submitted
    foo:submit-fail => bar      # bar triggers if submission of foo fails
    # comment
    """
```
- *default*: (none)

A.5 [runtime]

This section is used to specify how, where, and what to execute when tasks are ready to run. Common configuration can be factored out in a multiple-inheritance hierarchy of runtime namespaces that culminates in the tasks of the suite. Order of precedence is determined by the C3 linearization algorithm as used to find the *method resolution order* in Python language class hierarchies. For details and examples see 9.4.

A.5.1 [runtime] → [[__NAME__]]

Replace `__NAME__` with a namespace name, or a comma-separated list of names, and repeat as needed to define all tasks in the suite. Names may contain letters, digits, underscores, and hyphens. A namespace represents a group or family of tasks if other namespaces inherit from it, or a task if no others inherit from it.

- *legal values*:
 - `[[foo]]`
 - `[[foo, bar, baz]]`

If multiple names are listed the subsequent settings apply to each.

All namespaces inherit initially from *root*, which can be explicitly configured to provide or override default settings for all tasks in the suite.

A.5.1.1 [runtime] →[[_NAME_]] →extra log files

A list of user defined log files associated with a task. Files defined here will appear alongside the default log files in the cylc gui. Log files must reside in the job log directory and ideally should be named using the `$CYLC_TASK_LOG_ROOT` prefix (see 9.4.7.3).

- *type*: Comma-separated list of strings (log file names).
- *default*: (none)
- *example*: (job.custom-log-name)

A.5.1.2 [runtime] →[[_NAME_]] →inherit

A list of the immediate parent(s) this namespace inherits from. If no parents are listed `root` is assumed.

- *type*: Comma-separated list of strings (parent namespace names).
- *default*: `root`

A.5.1.3 [runtime] →[[_NAME_]] →init-script

This is invoked by the task job script before the task execution environment is configured, so it does not have access to any suite or task environment variables. It can be a single command or multiple lines of scripting. The original intention was to allow remote tasks to source login scripts to configure their access to cylc, but this should no longer be necessary (see 12.3). See also `env-script`, `err-script`, `pre-script`, `script`, and `post-script`.

- *type*: string
- *default*: (none)
- *example*: `init-script = "echo Hello World"`

A.5.1.4 [runtime] →[[_NAME_]] →env-script

This is invoked by the task job script between the cylc-defined environment (suite and task identity, etc.) and the user-defined task runtime environment - i.e. it has access to the cylc environment, and the task environment has access to variables defined by this scripting. It can be a single command or multiple lines of scripting. See also `init-script`, `err-script`, `pre-script`, `script`, and `post-script`.

- *type*: string
- *default*: (none)
- *example*: `env-script = "echo Hello World"`

A.5.1.5 [runtime] →[[_NAME_]] →err-script

This is any custom script to be invoked at the end of the error trap, (if the error trap is triggered due to failure of a command in the task job). The output of this will always be sent to `STDERR` and `$1` is set to the name of the signal caught by the error trap. The script should be fast and use very little system resource to ensure that the error trap can return quickly. It can be a single

command or multiple lines of scripting. See also `init-script`, `env-script`, `pre-script`, `script`, and `post-script`.

- *type*: string
- *default*: (none)
- *example*: `err-script = "printenv FOO"`

A.5.1.6 [runtime] →[[__NAME__]] →pre-script

This is invoked by the task job script immediately before the `script` item (just below). It can be a single command or multiple lines of scripting. See also `init-script`, `env-script`, `err-script`, `script`, and `post-script`.

- *type*: string
- *default*: (none)
- *example*:


```
pre-script = ""
. $HOME/.profile
echo Hello from suite ${CYLC_SUITE_NAME}!""
```

A.5.1.7 [runtime] →[[__NAME__]] →script

This is the main user-defined scripting to run when the task is ready. It can be a single command or multiple lines of scripting. See also `init-script`, `env-script`, `err-script`, `pre-script`, and `post-script`.

- *type*: string
- *root default*: (none)

A.5.1.8 [runtime] →[[__NAME__]] →post-script

This is invoked by the task job script immediately after the `script` item (just above). It can be a single command or multiple lines of scripting. See also `init-script`, `env-script`, `err-script`, `pre-script`, and `script`.

- *type*: string
- *default*: (none)

A.5.1.9 [runtime] →[[__NAME__]] →work sub-directory

Task job scripts are executed from within *work directories* created automatically under the suite run directory. A task can get its own work directory from `$CYLC_TASK_WORK_DIR` (or simply `$PWD` if it does not `cd` elsewhere at runtime). The default directory path contains task name and cycle point, to provide a unique workspace for every instance of every task. If several tasks need to exchange files and simply read and write from their from current working directory, this item can be used to override the default to make them all use the same workspace.

The top level share and work directory location can be changed (e.g. to a large data area) by a global config setting (see [B.9.1.2](#)).

- *type*: string (directory path, can contain environment variables)
- *default*: `$CYLC_TASK_CYCLE_POINT/$CYLC_TASK_NAME`
- *example*: `$CYLC_TASK_CYCLE_POINT/shared/`

Note that if you omit cycle point from the work sub-directory path successive instances of the task will share the same workspace. Consider the effect on cycle point offset housekeeping of work directories before doing this.

A.5.1.10 [runtime] → [____NAME____] → [[[meta]]]

Section containing metadata items for this task or family namespace. Several items (title, description, URL) are pre-defined and are used by the GUI. Others can be user-defined and passed to task event handlers to be interpreted according to your needs. For example, the value of an “importance” item could determine how an event handler responds to task failure events.

Any suite meta item can now be passed to task event handlers by prefixing the string template item name with “suite_”, for example :

```
[runtime]
  [[root]]
    [[events]]
      failed handler = send-help.sh %(suite_title)s %(suite_importance)s
                      %(title)s
```

A.5.1.10.1 [runtime] → [____NAME____] → [[[meta]]] → title

A single line description of this namespace. It is displayed by the `cylc list` command and can be retrieved from running tasks with the `cylc show` command.

- *type*: single line string
- *root default*: (none)

A.5.1.10.2 [runtime] → [____NAME____] → [[[meta]]] → description

A multi-line description of this namespace, retrievable from running tasks with the `cylc show` command.

- *type*: multi-line string
- *root default*: (none)

A.5.1.10.3 [runtime] → [____NAME____] → [[[meta]]] → URL

A web URL to task documentation for this suite. If present it can be browsed with the `cylc doc` command, or by right-clicking on the task in gcylc. The string templates `%(suite_name)s` and `%(task_name)s` will be replaced with the actual suite and task names. See also suite URLs (A.2.3).

- *type*: string (URL)
- *default*: (none)
- *example*: you can set URLs to all tasks in a suite by putting something like the following in the root namespace:

```
[runtime]
  [[root]]
    [[[meta]]]
      URL = http://my-site.com/suites/%(suite_name)s/%(task_name)s.
          html
```

(Note that URLs containing the comment delimiter `#` must be protected by quotes).

A.5.1.10.4 [runtime] → [____NAME____] → [[[meta]]] → ____MANY____

Replace ____MANY____ with any user-defined metadata item. These, like title, URL, etc. can be passed to task event handlers to be interpreted according to your needs. For example, the value of an "importance" item could determine how an event handler responds to task failure events.

- *type*: String or integer
- *default*: (none)
- *example*:

```
[runtime]
  [[root]]
    [[[meta]]]
      importance = high
      color = red
```

A.5.1.11 [runtime] → [____NAME____] → [[[job]]]

This section configures the means by which cylc submits task job scripts to run.

A.5.1.11.1 [runtime] → [____NAME____] → [[[job]]] → batch system

See 11 for how job submission works, and how to define new handlers for different batch systems. Cylc has a number of built in batch system handlers:

- *type*: string
- *legal values*:
 - `background` - invoke a child process
 - `at` - the rudimentary Unix `at` scheduler
 - `loadleveler` - IBM LoadLeveler `l1submit`, with directives defined in the suite.rc file
 - `lsf` - IBM Platform LSF `bsub`, with directives defined in the suite.rc file
 - `pbs` - PBS `qsub`, with directives defined in the suite.rc file
 - `sge` - Sun Grid Engine `qsub`, with directives defined in the suite.rc file
 - `slurm` - Simple Linux Utility for Resource Management `sbatch`, with directives defined in the suite.rc file
 - `moab` - Moab workload manager `msub`, with directives defined in the suite.rc file
- *default*: `background`

A.5.1.11.2 [runtime] → [____NAME____] → [[[job]]] → execution time limit

Specify the execution wall clock limit for a job of the task. For `background` and `at`, the job script will be invoked using the `timeout` command. For other batch systems, the specified time will be

automatically translated into the equivalent directive for wall clock limit.

- *type*: ISO 8601 duration/interval representation
- *example*: `PT5M`, 5 minutes, `PT1H`, 1 hour
- *default*: (none)

A.5.1.11.3 [runtime] →[[__NAME__]] →[[[job]]] →batch submit command template

This allows you to override the actual command used by the chosen batch system. The template's `%(job)s` will be substituted by the job file path.

- *type*: string
- *legal values*: a string template
- *example*: `llsubmit \%(job)s`

A.5.1.11.4 [runtime] →[[__NAME__]] →[[[job]]] →shell

Location of the command used to interpret the job script submitted by the suite server program when a task is ready to run. This can be set to the location of `bash` in the job host if the shell is not installed in the standard location. *Note: It has no bearing on any sub-shells that may be called by the job script.*

Setting this to the path of a ksh93 interpreter is deprecated. Support of which will be withdrawn in a future cylc release. Setting this to any other shell is not supported.

- *type*: string
- *root default*: `/bin/bash`

A.5.1.11.5 [runtime] →[[__NAME__]] →[[[job]]] →submission retry delays

A list of duration (in ISO 8601 syntax), after which to resubmit if job submission fails.

- *type*: Comma-separated list of ISO 8601 duration/interval representations, optionally *preceded* by multipliers.
- *example*: `PT1M,3*PT1H, P1D` is equivalent to `PT1M, PT1H, PT1H, PT1H, P1D` - 1 minute, 1 hour, 1 hour, 1 hour, 1 day.
- *default*: (none)

A.5.1.11.6 [runtime] →[[__NAME__]] →[[[job]]] →execution retry delays

See also [12.18](#).

A list of ISO 8601 time duration/intervals after which to resubmit the task if it fails. The variable `$CYLC_TASK_TRY_NUMBER` in the task execution environment is incremented each time, starting from 1 for the first try - this can be used to vary task behaviour by try number.

- *type*: Comma-separated list of ISO 8601 duration/interval representations, optionally *preceded* by multipliers.
- *example*: `PT1.5M,3*PT10M` is equivalent to `PT1.5M, PT10M, PT10M, PT10M` - 1.5 minutes, 10 minutes, 10 minutes, 10 minutes.

- *default*: (none)

A.5.1.11.7 [runtime] →[[__NAME__]] →[[[job]]] →submission polling intervals

A list of intervals, expressed as ISO 8601 duration/intervals, with optional multipliers, after which cylc will poll for status while the task is in the submitted state.

For the polling task communication method this overrides the default submission polling interval in the site/user config files (6). For default and ssh task communications, polling is not done by default but it can still be configured here as a regular check on the health of submitted tasks.

Each list value is used in turn until the last, which is used repeatedly until finished.

- *type*: Comma-separated list of ISO 8601 duration/interval representations, optionally *preceded* by multipliers.
- *example*: PT1M,3*PT1H, PT1M is equivalent to PT1M, PT1H, PT1H, PT1H, PT1M - 1 minute, 1 hour, 1 hour, 1 hour, 1 minute.
- *default*: (none)

A single interval value is probably appropriate for submission polling.

A.5.1.11.8 [runtime] →[[__NAME__]] →[[[job]]] →execution polling intervals

A list of intervals, expressed as ISO 8601 duration/intervals, with optional multipliers, after which cylc will poll for status while the task is in the running state.

For the polling task communication method this overrides the default execution polling interval in the site/user config files (6). For default and ssh task communications, polling is not done by default but it can still be configured here as a regular check on the health of submitted tasks.

Each list value is used in turn until the last, which is used repeatedly until finished.

- *type*: Comma-separated list of ISO 8601 duration/interval representations, optionally *preceded* by multipliers.
- *example*: PT1M,3*PT1H, PT1M is equivalent to PT1M, PT1H, PT1H, PT1H, PT1M - 1 minute, 1 hour, 1 hour, 1 hour, 1 minute.
- *default*: (none)

A.5.1.12 [runtime] →[[__NAME__]] →[[[remote]]]

Configure host and username, for tasks that do not run on the suite host account. Non-interactive ssh is used to submit the task by the configured batch system, so you must distribute your ssh key to allow this. Cylc must be installed on task remote accounts, but no external software dependencies are required there.

A.5.1.12.1 [runtime] →[[__NAME__]] →[[[remote]]] →host

The remote host for this namespace. This can be a static hostname, an environment variable that holds a hostname, or a command that prints a hostname to stdout. Host selection commands are executed just prior to job submission. The host (static or dynamic) may have an

entry in the cylc site or user config file to specify parameters such as the location of cylc on the remote machine; if not, the corresponding local settings (on the suite host) will be assumed to apply on the remote host.

- *type*: string (a valid hostname on the network)
- *default*: (none)
- *examples*:
 - static host name: `host = foo`
 - fully qualified: `host = foo.bar.baz`
 - dynamic host selection:
 - * shell command (1): `host = $(host-selector.sh)`
 - * shell command (2): `host = 'host-selector.sh'`
 - * environment variable: `host = $MY_HOST`

A.5.1.12.2 [runtime] →[[__NAME__]] →[[[remote]]] →owner

The username of the task host account. This is (only) used in the non-interactive ssh command invoked by the suite server program to submit the remote task (consequently it may be defined using local environment variables (i.e. the shell in which cylc runs, and `[cylc] →[[environment]]`).

If you use dynamic host selection and have different usernames on the different selectable hosts, you can configure your `$HOME/.ssh/config` to handle username translation.

- *type*: string (a valid username on the remote host)
- *default*: (none)

A.5.1.12.3 [runtime] →[[__NAME__]] →[[[remote]]] →retrieve job logs

Remote task job logs are saved to the suite run directory on the task host, not on the suite host. If you want the job logs pulled back to the suite host automatically, you can set this item to `True`. The suite will then attempt to `rsync` the job logs once from the remote host each time a task job completes. E.g. if the job file is `~/cylc-run/tut.oneoff.remote/log/job/1/hello/01/job`, anything under `~/cylc-run/tut.oneoff.remote/log/job/1/hello/01/` will be retrieved.

- *type*: boolean
- *default*: False

A.5.1.12.4 [runtime] →[[__NAME__]] →[[[remote]]] →retrieve job logs max size

If the disk space of the suite host is limited, you may want to set the maximum sizes of the job log files to retrieve. The value can be anything that is accepted by the `--max-size=SIZE` option of the `rsync` command.

- *type*: string
- *default*: None

A.5.1.12.5 [runtime] →[[__NAME__]] →[[[remote]]] →retrieve job logs retry delays

Some batch systems have considerable delays between the time when the job completes and when it writes the job logs in its normal location. If this is the case, you can configure an initial delay and some retry delays between subsequent attempts. The default behaviour is to attempt once without any delay.

- *type*: Comma-separated list of ISO 8601 duration/interval representations, optionally *preceded* by multipliers.
- *default*: (none)
- *example*: `retrieve job logs retry delays = PT10S, PT1M, PT5M`

A.5.1.12.6 [runtime] →[[__NAME__]] →[[[remote]]] →suite definition directory

The path to the suite definition directory on the remote account, needed if remote tasks require access to files stored there (via `$CYLC_SUITE_DEF_PATH`) or in the suite bin directory (via `$PATH`). If this item is not defined, the local suite definition directory path will be assumed, with the suite owner's home directory, if present, replaced by '`$HOME`' for interpretation on the remote account.

- *type*: string (a valid directory path on the remote account)
- *default*: (local suite definition path with `$HOME` replaced)

A.5.1.13 [runtime] →[[__NAME__]] →[[[events]]]

Cylc can call nominated event handlers when certain task events occur. This section configures specific task event handlers; see [A.3.13](#) for suite events.

Event handlers can be located in the suite `bin/` directory, otherwise it is up to you to ensure their location is in `$PATH` (in the shell in which the suite server program runs). They should require little resource to run and return quickly.

Each task event handler can be specified as a list of command lines or command line templates. They can contain any or all of the following patterns, which will be substituted with actual values:

- `%(event)s`: event name
- `%(suite)s`: suite name
- `%(point)s`: cycle point
- `%(name)s`: task name
- `%(submit_num)s`: submit number
- `%(id)s`: task ID (i.e. `%(name)s.%(point)s`)
- `%(batch_sys_name)s`: batch system name
- `%(batch_sys_job_id)s`: batch system job ID
- `%(message)s`: event message, if any
- any task [meta] item, e.g.:
 - `%(title)s`: task title
 - `%(URL)s`: task URL
 - `%(importance)s` - example custom task metadata
- any suite [meta] item, prefixed with "suite_", e.g.:

- %(suite_title)s: suite title
- %(suite_URL)s: suite URL
- %(suite_rating)s - example custom suite metadata

Otherwise, the command line will be called with the following default arguments:

```
<task-event-handler> %(event)s %(suite)s %(id)s %(message)s
```

Note: substitution patterns should not be quoted in the template strings. This is done automatically where required.

For an explanation of the substitution syntax, see String Formatting Operations in the Python documentation: <https://docs.python.org/2/library/stdtypes.html#string-formatting>.

Additional information can be passed to event handlers via the [cylc] → [[environment]] (but not via task runtime environments - event handlers are not called by tasks).

A.5.1.13.1 [runtime] → [[__NAME__]] → [[[events]]] → EVENT handler

A list of one or more event handlers to call when one of the following EVENTS occurs:

- **submitted** - the job submit command was successful
- **submission failed** - the job submit command failed, or the submitted job was killed before it started executing
- **submission retry** - job submit failed, but cylc will resubmit it after a configured delay
- **submission timeout** - the submitted job timed out without commencing execution
- **started** - the task reported commencement of execution
- **succeeded** - the task reported that it completed successfully
- **failed** - the task reported that it failed to complete successfully
- **retry** - the task failed, but cylc will resubmit it after a configured delay
- **execution timeout** - the task timed out after execution commenced
- **warning** - the task reported a WARNING severity message
- **critical** - the task reported a CRITICAL severity message
- **custom** - the task reported a CUSTOM severity message

Item details:

- *type*: Comma-separated list of strings (event handler scripts).
- *default*: None
- *example*: `failed handler = my-failed-handler.sh`

A.5.1.13.2 [runtime] → [[__NAME__]] → [[[events]]] → submission timeout

If a task has not started after the specified ISO 8601 duration/interval, the *submission timeout* event handler(s) will be called.

- *type*: ISO 8601 duration/interval representation (e.g. `PT30M`, 30 minutes or `P1D`, 1 day).
- *default*: (none)

A.5.1.13.3 [runtime] →[[__NAME__]] →[[[events]]] →execution timeout

If a task has not finished after the specified ISO 8601 duration/interval, the *execution timeout* event handler(s) will be called.

- *type*: ISO 8601 duration/interval representation (e.g. `PT4H`, 4 hours or `P1D`, 1 day).
- *default*: (none)

A.5.1.13.4 [runtime] →[[__NAME__]] →[[[events]]] →reset timer

If you set an execution timeout the timer can be reset to zero every time a message is received from the running task (which indicates the task is still alive). Otherwise, the task will timeout if it does not finish in the allotted time regardless of incoming messages.

- *type*: boolean
- *default*: False

A.5.1.13.5 [runtime] →[[__NAME__]] →[[[events]]] →handlers

Specify a list of command lines or command line templates as task event handlers.

- *type*: Comma-separated list of strings (event handler command line or command line templates).
- *default*: (none)
- *example*: `handlers = my-handler.sh`

A.5.1.13.6 [runtime] →[[__NAME__]] →[[[events]]] →handler events

Specify the events for which the general task event handlers should be invoked.

- *type*: Comma-separated list of events
- *default*: (none)
- *example*: `handler events = submission failed, failed`

A.5.1.13.7 [runtime] →[[__NAME__]] →[[[events]]] →handler retry delays

Specify an initial delay before running an event handler command and any retry delays in case the command returns a non-zero code. The default behaviour is to run an event handler command once without any delay.

- *type*: Comma-separated list of ISO 8601 duration/interval representations, optionally *preceded* by multipliers.
- *default*: (none)
- *example*: `handler retry delays = PT10S, PT1M, PT5M`

A.5.1.13.8 [runtime] →[[__NAME__]] →[[[events]]] →mail events

Specify the events for which notification emails should be sent.

- *type*: Comma-separated list of events
- *default*: (none)
- *example*: `mail events = submission failed, failed`

A.5.1.13.9 [runtime] →[[__NAME__]] →[[[events]]] →mail from

Specify an alternate `from`: email address for event notifications.

- *type*: string
- *default*: None, (notifications@HOSTNAME)
- *example*: `mail from = no-reply@your-org`

A.5.1.13.10 [runtime] →[[__NAME__]] →[[[events]]] →mail retry delays

Specify an initial delay before running the mail notification command and any retry delays in case the command returns a non-zero code. The default behaviour is to run the mail notification command once without any delay.

- *type*: Comma-separated list of ISO 8601 duration/interval representations, optionally *preceded* by multipliers.
- *default*: (none)
- *example*: `mail retry delays = PT10S, PT1M, PT5M`

A.5.1.13.11 [runtime] →[[__NAME__]] →[[[events]]] →mail smtp

Specify the SMTP server for sending email notifications.

- *type*: string
- *default*: None, (localhost:25)
- *example*: `mail smtp = smtp.yourorg`

A.5.1.13.12 [runtime] →[[__NAME__]] →[[[events]]] →mail to

A list of email addresses to send task event notifications. The list can be anything accepted by the `mail` command.

- *type*: string
- *default*: None, (USER@HOSTNAME)
- *example*: `mail to = your.colleague`

A.5.1.14 [runtime] →[[__NAME__]] →[[[environment]]]

The user defined task execution environment. Variables defined here can refer to `cylc` suite and task identity variables, which are exported earlier in the task job script, and variable assignment expressions can use `cylc` utility commands because access to `cylc` is also configured earlier in the script. See also 9.4.7.

A.5.1.14.1 [runtime] →[[__NAME__]] →[[[environment]]] →__VARIABLE__

Replace __VARIABLE__ with any number of environment variable assignment expressions. Order of definition is preserved so values can refer to previously defined variables. Values are passed through to the task job script without evaluation or manipulation by cylc, so any variable assignment expression that is legal in the job submission shell can be used. White space around the '=' is allowed (as far as cylc's suite.rc parser is concerned these are just normal configuration items).

- *type*: string
- *default*: (none)
- *legal values*: depends to some extent on the task job submission shell ([A.5.1.11.4](#)).
- *examples*, for the bash shell:
 - FOO = \$HOME/bar/baz
 - BAR = \${FOO}\$GLOBALVAR
 - BAZ = \$(echo "hello world")
 - WAZ = \${FOO%.jpg}.png
 - NEXT_CYCLE = \$(cylc cycle-point --offset=PT6H)
 - PREV_CYCLE = 'cylc cycle-point --offset=-PT6H'
 - ZAZ = "\${FOO#bar}"#<-- QUOTED to escape the suite.rc comment character

A.5.1.15 [runtime] →[[__NAME__]] →[[[environment filter]]]

This section contains environment variable inclusion and exclusion lists that can be used to filter the inherited environment. *This is not intended as an alternative to a well-designed inheritance hierarchy that provides each task with just the variables it needs.* Filters can, however, improve suites with tasks that inherit a lot of environment they don't need, by making it clear which tasks use which variables. They can optionally be used routinely as explicit "task environment interfaces" too, at some cost to brevity, because they guarantee that variables filtered out of the inherited task environment are not used.

Note that environment filtering is done after inheritance is completely worked out, not at each level on the way, so filter lists in higher-level namespaces only have an effect if they are not overridden by descendants.

A.5.1.15.1 [runtime] →[[__NAME__]] →[[[environment filter]]] →include

If given, only variables named in this list will be included from the inherited environment, others will be filtered out. Variables may also be explicitly excluded by an `exclude` list.

- *type*: Comma-separated list of strings (variable names).
- *default*: (none)

A.5.1.15.2 [runtime] →[[__NAME__]] →[[[environment filter]]] →exclude

Variables named in this list will be filtered out of the inherited environment. Variables may also be implicitly excluded by omission from an `include` list.

- *type*: Comma-separated list of strings (variable names).

- *default*: (none)

A.5.1.16 [runtime] → [____NAME____] → [[[parameter environment templates]]]

The user defined task execution parameter environment templates. This is only relevant for *parameterized tasks* - see Section 9.6.

A.5.1.16.1 [runtime] → [____NAME____] → [[[parameter environment templates]]] → ____VARIABLE____

Replace ____VARIABLE____ with pairs of environment variable name and Python string template for parameter substitution. This is only relevant for *parameterized tasks* - see Section 9.6.

If specified, in addition to the standard CYLC_TASK_PARAM_<key> variables, the job script will also export the named variables specified here, with the template strings substituted with the parameter values.

- *type*: string
- *default*: (none)
- *legal values*: name=string template pairs
- *examples*, for the bash shell:
 - MYNUM=%(i)d
 - MYITEM=%(item)s
 - MYFILE=/path/to/%(i)03d/%(item)s

A.5.1.17 [runtime] → [____NAME____] → [[[directives]]]

Batch queue scheduler directives. Whether or not these are used depends on the batch system. For the built-in methods that support directives (*loadleveler*, *lsf*, *pbs*, *sge*, *slurm*, *moab*), directives are written to the top of the task job script in the correct format for the method. Specifying directives individually like this allows use of default directives that can be individually overridden at lower levels of the runtime namespace hierarchy.

A.5.1.17.1 [runtime] → [____NAME____] → [[[directives]]] → ____DIRECTIVE____

Replace ____DIRECTIVE____ with each directive assignment, e.g. `class = parallel`

- *type*: string
- *default*: (none)

Example directives for the built-in batch system handlers are shown in 11.1.

A.5.1.18 [runtime] → [____NAME____] → [[[outputs]]]

Register custom task outputs for use in message triggering in this section (9.3.5.5)

A.5.1.18.1 [runtime] →[[__NAME__]] →[[[outputs]]] →__OUTPUT__

Replace __OUTPUT__ with one or more custom task output messages (9.3.5.5). The item name is used to select the custom output message in graph trigger notation.

- *type*: string
- *default*: (none)
- *examples*:

```
out1 = "sea state products ready"
out2 = "NWP restart files completed"
```

A.5.1.19 [runtime] →[[__NAME__]] →[[[suite state polling]]]

Configure automatic suite polling tasks as described in 12.26. The items in this section reflect the options and defaults of the `cylc suite-state` command, except that the target suite name and the `--task`, `--cycle`, and `--status` options are taken from the graph notation.

A.5.1.19.1 [runtime] →[[__NAME__]] →[[[suite state polling]]] →run-dir

For your own suites the run database location is determined by your site/user config. For other suites, e.g. those owned by others, or mirrored suite databases, use this item to specify the location of the top level cylc run directory (the database should be a suite-name sub-directory of this location).

- *type*: string (a directory path on the target suite host)
- *default*: as configured by site/user config (for your own suites)

A.5.1.19.2 [runtime] →[[__NAME__]] →[[[suite state polling]]] →template

Cycle point template of the target suite, if different from that of the polling suite.

- *type*: string
- *default*: cycle point format of the polling suite
- *example*: %Y-%m-%dT%H

A.5.1.19.3 [runtime] →[[__NAME__]] →[[[suite state polling]]] →interval

Polling interval expressed as an ISO 8601 duration/interval.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT1M`

A.5.1.19.4 [runtime] →[[__NAME__]] →[[[suite state polling]]] →max-polls

The maximum number of polls before timing out and entering the ‘failed’ state.

- *type*: integer
- *default*: 10

A.5.1.19.5 [runtime] →[[__NAME__]] →[[[suite state polling]]] →user

Username of an account on the suite host to which you have access. The polling `cylc suite-state` command will be invoked on the remote account.

- *type*: string (username)
- *default*: (none)

A.5.1.19.6 [runtime] →[[__NAME__]] →[[[suite state polling]]] →host

The hostname of the target suite. The polling `cylc suite-state` command will be invoked on the remote account.

- *type*: string (hostname)
- *default*: (none)

A.5.1.19.7 [runtime] →[[__NAME__]] →[[[suite state polling]]] →verbose

Run the polling `cylc suite-state` command in verbose output mode.

- *type*: boolean
- *default*: False

A.5.1.20 [runtime] →[[__NAME__]] →[[[simulation]]]

Task configuration for the suite *simulation* and *dummy* run modes described in Section 12.24.

A.5.1.20.1 [runtime] →[[__NAME__]] →[[[simulation]]] →default run length

The default simulated job run length, if `[job]execution time limit` and `[simulation]speedup factor` are not set.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT10S`

A.5.1.20.2 [runtime] →[[__NAME__]] →[[[simulation]]] →speedup factor

If `[job]execution time limit` is set, the task simulated run length is computed by dividing it by this factor.

- *type*: float
- *default*: (none) - i.e. do not use proportional run length
- *example*: `10.0`

A.5.1.20.3 [runtime] →[[__NAME__]] →[[[simulation]]] →time limit buffer

For dummy jobs, a new `[job]execution time limit` is set to the simulated task run length plus this buffer interval, to avoid job kill due to exceeding the time limit.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT10S`

A.5.1.20.4 [runtime] →[`__NAME__`] →[[`simulation`]] →fail cycle points

Configure simulated or dummy jobs to fail at certain cycle points.

- *type*: list of strings (cycle points), or *all*
- *default*: (none) - no instances of the task will fail
- *examples*:
 - `all` - all instance of the task will fail
 - `2017-08-12T06`, `2017-08-12T18` - these instances of the task will fail

A.5.1.20.5 [runtime] →[`__NAME__`] →[[`simulation`]] →fail try 1 only

If this is set to `True` only the first run of the task instance will fail, otherwise retries will fail too.

- *type*: boolean
- *default*: `True`

A.5.1.20.6 [runtime] →[`__NAME__`] →[[`simulation`]] →disable task event handlers

If this is set to `True` configured task event handlers will not be called in simulation or dummy modes.

- *type*: boolean
- *default*: `True`

A.6 [visualization]

Configuration of suite graphing for the `cylc graph` command (graph extent, styling, and initial family-collapsed state) and the `gcylc graph view` (initial family-collapsed state). Graphviz documentation of node shapes and so on can be found at <http://www.graphviz.org/Documentation.php>.

A.6.1 [visualization] →initial cycle point

The initial cycle point for graph plotting.

- *type*: ISO 8601 date-time representation (e.g. `CCYYMMDDThhmm`)
- *default*: the suite initial cycle point

The visualization initial cycle point gets adjusted up if necessary to the suite initial cycling point.

A.6.2 [visualization] →final cycle point

An explicit final cycle point for graph plotting. If used, this overrides the preferred *number of cycle points* (below).

- *type*: ISO 8601 date-time representation (e.g. CCYYMMDDThhmm)
- *default*: (none)

The visualization final cycle point gets adjusted down if necessary to the suite final cycle point.

A.6.3 [visualization] →number of cycle points

The number of cycle points to graph starting from the visualization initial cycle point. This is the preferred way of defining the graph end point, but it can be overridden by an explicit *final cycle point* (above).

- *type*: integer
- *default*: 3

A.6.4 [visualization] →collapsed families

A list of family (namespace) names to be shown in the collapsed state (i.e. the family members will be replaced by a single family node) when the suite is first plotted in the graph viewer or the gcylic graph view. If this item is not set, the default is to collapse all families at first. Interactive GUI controls can then be used to group and ungroup family nodes at will.

- *type*: Comma-separated list of family names.
- *default*: (none)

A.6.5 [visualization] →use node color for edges

Plot graph edges (dependency arrows) with the same color as the upstream node, otherwise default to black.

- *type*: boolean
- *default*: False

A.6.6 [visualization] →use node fillcolor for edges

Plot graph edges (i.e. dependency arrows) with the same fillcolor as the upstream node, if it is filled, otherwise default to black.

- *type*: boolean
- *default*: False

A.6.7 [visualization] →node penwidth

Line width of node shape borders.

- *type*: integer
- *default*: 2

A.6.8 [visualization] →edge penwidth

Line width of graph edges (dependency arrows).

- *type*: integer
- *default*: 2

A.6.9 [visualization] →use node color for labels

Graph node labels can be printed in the same color as the node outline.

- *type*: boolean
- *default*: False

A.6.10 [visualization] →default node attributes

Set the default attributes (color and style etc.) of graph nodes (tasks and families). Attribute pairs must be quoted to hide the internal = character.

- *type*: Comma-separated list of quoted 'attribute=value' pairs.
- *legal values*: see graphviz or pygraphviz documentation
- *default*: 'style=filled', 'fillcolor=yellow', 'shape=box'

A.6.11 [visualization] →default edge attributes

Set the default attributes (color and style etc.) of graph edges (dependency arrows). Attribute pairs must be quoted to hide the internal = character.

- *type*: Comma-separated list of quoted 'attribute=value' pairs.
- *legal values*: see graphviz or pygraphviz documentation
- *default*: 'color=black'

A.6.12 [visualization] →[[node groups]]

Define named groups of graph nodes (tasks and families) which can styled en masse, by name, in [visualization] →[[node attributes]]. Node groups are automatically defined for all task families, including root, so you can style family and member nodes at once by family name.

A.6.12.1 [visualization] →[[node groups]] →__GROUP__

Replace __GROUP__ with each named group of tasks or families.

- *type*: Comma-separated list of task or family names.
- *default*: (none)
- *example*:

```
PreProc = foo, bar
PostProc = baz, waz
```

A.6.13 [visualization] → [[node attributes]]

Here you can assign graph node attributes to specific nodes, or to all members of named groups defined in [visualization] → [[node groups]]. Task families are automatically node groups. Styling of a family node applies to all member nodes (tasks and sub-families), but precedence is determined by ordering in the suite definition. For example, if you style a family red and then one of its members green, cylc will plot a red family with one green member; but if you style one member green and then the family red, the red family styling will override the earlier green styling of the member.

A.6.13.1 [visualization] → [[node attributes]] → ___NAME___

Replace ___NAME___ with each node or node group for style attribute assignment.

- *type*: Comma-separated list of quoted 'attribute=value' pairs.
- *legal values*: see graphviz or pygraphviz documentation
- *default*: (none)
- *example*: (with reference to the node groups defined above)

```
PreProc = 'style=filled', 'fillcolor=orange'
PostProc = 'color=red'
foo = 'style=filled'
```

B Global (Site, User) Config File Reference

This section defines all legal items and values for cylc site and user config files. See *Site And User Config Files* (Section 6) for file locations, intended usage, and how to generate the files using the `cylc get-site-config` command.

As for suite definitions, Jinja2 expressions can be embedded in site and user config files to generate the final result parsed by cylc. Use of Jinja2 in suite definitions is documented in Section 9.7.

B.1 Top Level Items

B.1.1 temporary directory

A temporary directory is needed by a few cylc commands, and is cleaned automatically on exit. Leave unset for the default (usually `$TMPDIR`).

- *type*: string (directory path)
- *default*: (none)
- *example*: `temporary directory = /tmp/$USER/cylc`

B.1.2 process pool size

Number of process pool worker processes used to execute shell commands (job submission, event handlers, job poll and kill commands).

- *type*: integer
- *default*: None (number of processor cores on the suite host)

B.1.3 disable interactive command prompts

Commands that intervene in running suites can be made to ask for confirmation before acting. Some find this annoying and ineffective as a safety measure, however, so command prompts are disabled by default.

- *type*: boolean
- *default*: True

B.1.4 enable run directory housekeeping

The suite run directory tree is created anew with every suite start (not restart) but output from the most recent previous runs can be retained in a rolling archive. Set length to 0 to keep no backups. **This is incompatible with current Rose suite housekeeping** (see Section 13 for more on Rose) so it is disabled by default, in which case new suite run files will overwrite existing ones in the same run directory tree. Rarely, this can result in incorrect polling results due to the presence of old task status files.

- *type*: boolean

- *default*: False

B.1.5 run directory rolling archive length

The number of old run directory trees to retain if run directory housekeeping is enabled.

- *type*: integer
- *default*: 2

B.1.6 task host select command timeout

When a task host in a suite is a shell command string, cylc calls the shell to determine the task host. This call is invoked by the main process, and may cause the suite to hang while waiting for the command to finish. This setting sets a timeout for such a command to ensure that the suite can continue.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT10S`

B.2 [task messaging]

This section contains configuration items that affect task-to-suite communications.

B.2.1 [task messaging] →retry interval

If a send fails, the messaging code will retry after a configured delay interval.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT5S`

B.2.2 [task messaging] →maximum number of tries

If successive sends fail, the messaging code will give up after a configured number of tries.

- *type*: integer
- *minimum*: 1
- *default*: 7

B.2.3 [task messaging] →connection timeout

This is the same as the `--comms-timeout` option in cylc commands. Without a timeout remote connections to unresponsive suites can hang indefinitely (suites suspended with Ctrl-Z for instance).

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT30S`

B.3 [suite logging]

The suite event log, held under the suite run directory, is maintained as a rolling archive. Logs are rolled over (backed up and started anew) when they reach a configurable limit size.

B.3.1 [suite logging] →roll over at start-up

If True, a new suite log will be started for a new suite run.

- *type*: boolean
- *default*: True

B.3.2 [suite logging] →rolling archive length

How many rolled logs to retain in the archive.

- *type*: integer
- *minimum*: 1
- *default*: 5

B.3.3 [suite logging] →maximum size in bytes

Suite event logs are rolled over when they reach this file size.

- *type*: integer
- *default*: 1000000

B.4 [documentation]

Documentation locations for the `cylc doc` command and gcylc Help menus.

B.4.1 [documentation] →[[files]]

File locations of documentation held locally on the cylc host server.

B.4.1.1 [documentation] →[[files]] →html index

File location of the main cylc documentation index.

- *type*: string
- *default*: `$CYLC_DIR/doc/index.html`

B.4.1.2 [documentation] →[[files]] →pdf user guide

File location of the cylc User Guide, PDF version.

- *type*: string
- *default*: `$CYLC_DIR/doc/cug-pdf.pdf`

B.4.1.3 [documentation] →[[files]] →multi-page html user guide

File location of the cylc User Guide, multi-page HTML version.

- *type*: string
- *default*: `$CYLC_DIR/doc/html/multi/cug-html.html`

B.4.1.4 [documentation] →[[files]] →single-page html user guide

File location of the cylc User Guide, single-page HTML version.

- *type*: string
- *default*: `$CYLC_DIR/doc/html/single/cug-html.html`

B.4.2 [documentation] →[[urls]]

Online documentation URLs.

B.4.2.1 [documentation] →[[urls]] →internet homepage

URL of the cylc internet homepage, with links to documentation for the latest official release.

- *type*: string
- *default*: `http://cylc.github.com/cylc/`

B.4.2.2 [documentation] →[[urls]] →local index

Local intranet URL of the main cylc documentation index.

- *type*: string
- *default*: (none)

B.5 [document viewers]

PDF and HTML viewers can be launched by cylc to view the documentation.

B.5.1 [document viewers] →pdf

Your preferred PDF viewer program.

- *type*: string

- *default:* evince

B.5.2 [document viewers] →html

Your preferred web browser.

- *type:* string
- *default:* firefox

B.6 [editors]

Choose your favourite text editor for editing suite definitions.

B.6.1 [editors] →terminal

The editor to be invoked by the cylc command line interface.

- *type:* string
- *default:* vim
- *examples:*
 - `terminal = emacs -nw` (emacs non-GUI)
 - `terminal = emacs` (emacs GUI)
 - `terminal = gvim -f` (vim GUI)

B.6.2 [editors] →gui

The editor to be invoked by the cylc GUI.

- *type:* string
- *default:* gvim -f
- *examples:*
 - `gui = emacs`
 - `gui = xterm -e vim`

B.7 [communication]

This section covers options for network communication between cylc clients (suite-connecting commands and guis) servers (running suites). Each suite listens on a dedicated network port, binding on the first available starting at the configured base port.

By default, the communication method is HTTPS secured with HTTP Digest Authentication. If the system does not support SSL, you should configure this section to use HTTP. Cylc will not automatically fall back to HTTP if HTTPS is not available.

B.7.1 [communication] →method

The choice of client-server communication method - currently only HTTPS and HTTP are supported, although others could be developed and plugged in. Cylc defaults to HTTPS if this setting is not explicitly configured.

- *type*: string
- *options*:
 - **https**
 - **http**
- *default*: https

B.7.2 [communication] →base port

The first port that cylc is allowed to use.

- *type*: integer
- *default*: 43001

B.7.3 [communication] →maximum number of ports

This determines the maximum number of suites that can run at once on the suite host.

- *type*: integer
- *default*: 100

B.7.4 [communication] →proxies on

Enable or disable proxy servers for HTTPS - disabled by default.

- *type*: boolean
- *localhost default*: False

B.7.5 [communication] →options

Option flags for the communication method. Currently only 'SHA1' is supported for HTTPS, which alters HTTP Digest Auth to use the SHA1 hash algorithm rather than the standard MD5. This is more secure but is also less well supported by third party web clients including web browsers. You may need to add the 'SHA1' option if you are running on platforms where MD5 is discouraged (e.g. under FIPS).

- *type*: string_list
- *default*: []
- *options*:
 - **SHA1**

B.8 [monitor]

Configurable settings for the command line `cylc monitor` tool.

B.8.1 [monitor] →sort order

The sort order for tasks in the monitor view.

- *type*: string
- *options*:
 - **alphanumeric**
 - **definition** - the order that tasks appear under [runtime] in the suite definition.
- *default*: definition

B.9 [hosts]

The [hosts] section configures some important host-specific settings for the suite host ('local-host') and remote task hosts. Note that *remote task behaviour is determined by the site/user config on the suite host, not on the task host*. Suites can specify task hosts that are not listed here, in which case local settings will be assumed, with the local home directory path, if present, replaced by \$HOME in items that configure directory locations.

B.9.1 [hosts] →[[HOST]]

The default task host is the suite host, **localhost**, with default values as listed below. Use an explicit [hosts][localhost] section if you need to override the defaults. Localhost settings are then also used as defaults for other hosts, with the local home directory path replaced as described above. This applies to items omitted from an explicit host section, and to hosts that are not listed at all in the site and user config files. Explicit host sections are only needed if the automatically modified local defaults are not sufficient.

Host section headings can also be *regular expressions* to match multiple hostnames. Note that the general regular expression wildcard is `'.*'` (zero or more of any character), not `'*'`. Hostname matching regular expressions are used as-is in the Python `re.match()` function. As such they match from the beginning of the hostname string (as specified in the suite definition) and they do not have to match through to the end of the string (use the string-end matching character `'$'` in the expression to force this).

A hierarchy of host match expressions from specific to general can be used because config items are processed in the order specified in the file.

- *type*: string (hostname or regular expression)
- *examples*:
 - `server1.niwa.co.nz` - explicit host name
 - `server\d.niwa.co.nz` - regular expression

B.9.1.1 [hosts] →[[HOST]] →run directory

The top level of the directory tree that holds suite-specific output logs, run database, etc.

- *type*: string (directory path)
- *default*: `$HOME/cylc-run`

B.9.1.2 [hosts] →[[HOST]] →work directory

The top level for suite work and share directories.

- *type*: string (directory path)
- *localhost default*: `$HOME/cylc-run`

B.9.1.3 [hosts] →[[HOST]] →task communication method

The means by which task progress messages are reported back to the running suite. See above for default polling intervals for the poll method.

- *type*: string (must be one of the following three options)
- *options*:
 - **default** - direct client-server communication via network ports
 - **ssh** - use ssh to re-invoke the messaging commands on the suite server
 - **poll** - the suite polls for the status of tasks (no task messaging)
- *localhost default*: default

B.9.1.4 [hosts] →[[HOST]] →execution polling intervals

Cylc can poll running jobs to catch problems that prevent task messages from being sent back to the suite, such as hard job kills, network outages, or unplanned task host shutdown. Routine polling is done only for the polling *task communication method* (below) unless suite-specific polling is configured in the suite definition. A list of interval values can be specified, with the last value used repeatedly until the task is finished - this allows more frequent polling near the beginning and end of the anticipated task run time. Multipliers can be used as shorthand as in the example below.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*:
- *example*: `execution polling intervals = 5*PT1M, 10*PT5M, 5*PT1M`

B.9.1.5 [hosts] →[[HOST]] →submission polling intervals

Cylc can also poll submitted jobs to catch problems that prevent the submitted job from executing at all, such as deletion from an external batch scheduler queue. Routine polling is done only for the polling *task communication method* (above) unless suite-specific polling is configured in the suite definition. A list of interval values can be specified as for execution polling (above) but a single value is probably sufficient for job submission polling.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).

- *default:*
- *example:* (see the execution polling example above)

B.9.1.6 [hosts] →[[HOST]] →scp command

A string for the command used to copy files to a remote host. This is not used on the suite host unless you run local tasks under another user account. The value is assumed to be `scp` with some initial options or a command that implements a similar interface to `scp`.

- *type:* string
- *localhost default:* `scp -oBatchMode=yes -oConnectTimeout=10`

B.9.1.7 [hosts] →[[HOST]] →ssh command

A string for the command used to invoke commands on this host. This is not used on the suite host unless you run local tasks under another user account. The value is assumed to be `ssh` with some initial options or a command that implements a similar interface to `ssh`.

- *type:* string
- *localhost default:* `ssh -oBatchMode=yes -oConnectTimeout=10`

B.9.1.8 [hosts] →[[HOST]] →use login shell

Whether to use a login shell or not for remote command invocation. By default `cylc` runs remote `ssh` commands using a login shell:

```
ssh user@host 'bash --login cylc ...'
```

which will source `/etc/profile` and `~/.profile` to set up the user environment. However, for security reasons some institutions do not allow unattended commands to start login shells, so you can turn off this behaviour to get:

```
ssh user@host 'cylc ...'
```

which will use the default shell on the remote machine, sourcing `~/.bashrc` (or `~/.cshrc`) to set up the environment.

- *type:* boolean
- *localhost default:* `True`

B.9.1.9 [hosts] →[[HOST]] →cylc executable

The `cylc` executable on a remote host. Note this should point to the `cylc` multi-version wrapper (see 7.2) on the host, not `bin/cylc` for a specific installed version. Specify a full path if `cylc` is not in `\$PATH` when it is invoked via `ssh` on this host.

- *type:* string
- *localhost default:* `cylc`

B.9.1.10 [hosts] →[[HOST]] →global init-script

If specified, the value of this setting will be inserted to just before the `init-script` section of all job scripts that are to be submitted to the specified remote host.

- *type*: string
- *localhost default*: ""

B.9.1.11 [hosts] →[[HOST]] →copyable environment variables

A list containing the names of the environment variables that can and/or need to be copied from the suite server program to a job.

- *type*: string_list
- *localhost default*: []

B.9.1.12 [hosts] →[[HOST]] →retrieve job logs

Global default for the [A.5.1.12.3](#) setting for the specified host.

B.9.1.13 [hosts] →[[HOST]] →retrieve job logs command

If `rsync -a` is unavailable or insufficient to retrieve job logs from a remote host, you can use this setting to specify a suitable command.

- *type*: string
- *default*: `rsync -a`

B.9.1.14 [hosts] →[[HOST]] →retrieve job logs max size

Global default for the [A.5.1.12.4](#) setting for the specified host.

B.9.1.15 [hosts] →[[HOST]] →retrieve job logs retry delays

Global default for the [A.5.1.12.5](#) setting for the specified host.

B.9.1.16 [hosts] →[[HOST]] →task event handler retry delays

Host specific default for the [A.5.1.13.7](#) setting.

B.9.1.17 [hosts] →[[HOST]] →local tail command template

A template (with `%{filename}s` substitution) for the command used to tail-follow local job logs, used by the gcylc log viewer and `cylc cat-log --tail`. You are unlikely to need to override this.

- *type*: string
- *default*: `tail -n +1 -F %{filename}s`

B.9.1.18 [hosts] →[[HOST]] →remote tail command template

A template (with `%(filename)s` substitution) for the command used to tail-follow remote job logs, used by the gcylc log viewer and `cylc cat-log --tail`. The remote tail command needs to be told to die when its parent process exits. You may need to override this command for task hosts where the default `tail` or `ps` commands are not equivalent to the Gnu Linux versions.

- *type*: string
- *default*: `tail --pid=$(ps h -o ppid $$ | sed -e 's/[:space:]/g')-n +1 -F %(filename)s`
- *example*: for AIX hosts:
`/gnu/tail --pid=$(ps -o ppid= -p $$ | sed -e 's/[:space:]/g')-n +1 -F %(filename)s`

B.9.1.19 [hosts] →[[HOST]] →[[[batch systems]]]

Settings for particular batch systems on HOST. In the subsections below, SYSTEM should be replaced with the cylc batch system handler name that represents the batch system (see [A.5.1.11.1](#)).

B.9.1.19.1 [hosts] →[[HOST]] →[[[batch systems]]] →[[[SYSTEM]]] →err tailer

A command template (with `%(job_id)s` substitution) that can be used to tail-follow the stderr stream of a running job if SYSTEM does not use the normal log file location while the job is running. This setting overrides [B.9.1.17](#) and [B.9.1.18](#) above.

- *type*: string
- *default*: (none)
- *example*: For PBS:

```
[hosts]
  [[ myhpc*]]
    [[[batch systems]]]
      [[[[pbs]]]]
        err tailer = qcat -f -e %(job_id)s
        out tailer = qcat -f -o %(job_id)s
        err viewer = qcat -e %(job_id)s
        out viewer = qcat -o %(job_id)s
```

B.9.1.19.2 [hosts] →[[HOST]] →[[[batch systems]]] →[[[SYSTEM]]] →out tailer

A command template (with `%(job_id)s` substitution) that can be used to tail-follow the stdout stream of a running job if SYSTEM does not use the normal log file location while the job is running. This setting overrides [B.9.1.17](#) and [B.9.1.18](#) above.

- *type*: string
- *default*: (none)
- *example*: see [B.9.1.19.1](#)

B.9.1.19.3 [hosts] →[[HOST]] →[[[batch systems]]] →[[[SYSTEM]]] →err viewer

A command template (with `%(job_id)s` substitution) that can be used to view the stderr stream of a running job if SYSTEM does not use the normal log file location while the job is running.

- *type*: string
- *default*: (none)
- *example*: see [B.9.1.19.1](#)

B.9.1.19.4 [hosts] → [[HOST]] → [[[batch systems]]] → [[[[SYSTEM]]]] → out viewer

A command template (with `%(job_id)s` substitution) that can be used to view the stdout stream of a running job if SYSTEM does not use the normal log file location while the job is running.

- *type*: string
- *default*: (none)
- *example*: see [B.9.1.19.1](#)

B.9.1.19.5 [hosts] → [[HOST]] → [[[batch systems]]] → [[[[SYSTEM]]]] → job name length maximum

The maximum length for job name acceptable by a batch system on a given host. Currently, this setting is only meaningful for PBS jobs. For example, PBS 12 or older will fail a job submit if the job name has more than 15 characters, which is the default setting. If you have PBS 13 or above, you may want to modify this setting to a larger value.

- *type*: integer
- *default*: (none)
- *example*: For PBS:

```
[hosts]
  [[myhpc*]]
    [[[batch systems]]]
      [[[[pbs]]]]
        # PBS 13
        job name length maximum = 236
```

B.9.1.19.6 [hosts] → [[HOST]] → [[[batch systems]]] → [[[[SYSTEM]]]] → execution time limit polling intervals

The intervals between polling after a task job (submitted to the relevant batch system on the relevant host) exceeds its execution time limit. The default setting is PT1M, PT2M, PT7M. The accumulated times (in minutes) for these intervals will be roughly 1, 1 + 2 = 3 and 1 + 2 + 7 = 10 after a task job exceeds its execution time limit.

- *type*: Comma-separated list of ISO 8601 duration/interval representations, optionally *preceded* by multipliers.
- *default*: PT1M, PT2M, PT7M
- *example*:

```
[hosts]
  [[myhpc*]]
    [[[batch systems]]]
      [[[[pbs]]]]
        execution time limit polling intervals = 5*PT2M
```

B.10 [suite host self-identification]

The suite host's identity must be determined locally by `cylc` and passed to running tasks (via `$CYLC_SUITE_HOST`) so that task messages can target the right suite on the right host.

B.10.1 [suite host self-identification] →method

This item determines how `cylc` finds the identity of the suite host. For the default *name* method `cylc` asks the suite host for its host name. This should resolve on remote task hosts to the IP address of the suite host; if it doesn't, adjust network settings or use one of the other methods. For the *address* method, `cylc` attempts to use a special external "target address" to determine the IP address of the suite host as seen by remote task hosts (in-source documentation in `$CYLC_DIR/lib/cylc/hostuserutil.py` explains how this works). And finally, as a last resort, you can choose the *hardwired* method and manually specify the host name or IP address of the suite host.

- *type*: string
- *options*:
 - *name* - self-identified host name
 - *address* - automatically determined IP address (requires *target*, below)
 - *hardwired* - manually specified host name or IP address (requires *host*, below)
- *default*: *name*

B.10.2 [suite host self-identification] →target

This item is required for the *address* self-identification method. If your suite host sees the internet, a common address such as `google.com` will do; otherwise choose a host visible on your intranet.

- *type*: string (an inter- or intranet URL visible from the suite host)
- *default*: `google.com`

B.10.3 [suite host self-identification] →host

Use this item to explicitly set the name or IP address of the suite host if you have to use the *hardwired* self-identification method.

- *type*: string (host name or IP address)
- *default*: (none)

B.11 [suite host scanning]

Utilities such as `cylc gsummary` need to scan hosts for running suites.

B.11.1 [suite host scanning] →hosts

A list of hosts to scan for running suites.

- *type*: comma-separated list of host names or IP addresses.
- *default*: localhost

B.12 [task events]

Global site/user defaults for [A.5.1.13](#).

B.13 [test battery]

Settings for the automated development tests.

B.13.1 [test battery] →remote host with shared fs

The name of a remote host with shared HOME file system as the host running the test battery.

B.13.2 [test battery] →remote host

The name of a remote host without shared HOME file system as the host running the test battery.

B.13.3 [test battery] →[[batch systems]]

Settings for testing supported batch systems (job submission methods). The tests for a batch system are only performed if the batch system is available on the test host or a remote host accessible via SSH from the test host.

B.13.3.1 [test battery] →[[batch systems]] →[[[SYSTEM]]]

SYSTEM is the name of a supported batch system with automated tests. This can currently be "loadleveler", "lsf", "pbs", "sge" and/or "slurm".

B.13.3.1.1 [test battery] →[[batch systems]] →[[[SYSTEM]]] →host

The name of a host where commands for this batch system is available. Use "localhost" if the batch system is available on the host running the test battery. Any specified remote host should be accessible via SSH from the host running the test battery.

B.13.3.1.2 [test battery] →[[batch systems]] →[[[SYSTEM]]] →err viewer

The command template (with `\%(job_id)s` substitution) for testing the run time stderr viewer functionality for this batch system.

B.13.3.1.3 [test battery] →[[batch systems]] →[[[SYSTEM]]] →out viewer

The command template (with `\%(job_id)s` substitution) for testing the run time stdout viewer functionality for this batch system.

B.13.3.1.4 [test battery] →[[batch systems]] →[[[SYSTEM]]] →[[[directives]]]

The minimum set of directives that must be supplied to the batch system on the site to initiate jobs for the tests.

B.14 [cylc]

Default values for entries in the suite.rc [cylc] section.

B.14.1 [cylc] →UTC mode

Allows you to set a default value for UTC mode in a suite at the site level. See [A.3.2](#) for details.

B.14.2 [cylc] →health check interval

Site default suite health check interval. See [A.3.7](#) for details.

B.14.3 [cylc] →task event mail interval

Site default task event mail interval. See [A.3.8](#) for details.

B.14.4 [cylc] →[[events]]

You can define site defaults for each of the following options, details of which can be found under [A.3.13](#):

- B.14.4.1** [cylc] →[[events]] →handlers
- B.14.4.2** [cylc] →[[events]] →handler events
- B.14.4.3** [cylc] →[[events]] →startup handler
- B.14.4.4** [cylc] →[[events]] →shutdown handler
- B.14.4.5** [cylc] →[[events]] →mail events
- B.14.4.6** [cylc] →[[events]] →mail footer
- B.14.4.7** [cylc] →[[events]] →mail from
- B.14.4.8** [cylc] →[[events]] →mail smtp
- B.14.4.9** [cylc] →[[events]] →mail to
- B.14.4.10** [cylc] →[[events]] →timeout handler
- B.14.4.11** [cylc] →[[events]] →timeout
- B.14.4.12** [cylc] →[[events]] →abort on timeout
- B.14.4.13** [cylc] →[[events]] →stalled handler
- B.14.4.14** [cylc] →[[events]] →abort on stalled
- B.14.4.15** [cylc] →[[events]] →inactivity handler
- B.14.4.16** [cylc] →[[events]] →inactivity
- B.14.4.17** [cylc] →[[events]] →abort on inactivity

B.15 [authentication]

Authentication of client programs with suite server programs can be configured here, and overridden in suites if necessary (see [A.3.16](#)).

The suite-specific passphrase must be installed on a user's account to authorize full control privileges (see [7.5](#) and [12.9](#)). In the future we plan to move to a more traditional user account model so that each authorized user can have their own password.

B.15.1 [authentication] →public

This sets the client privilege level for public access - i.e. no suite passphrase required.

- *type*: string (must be one of the following options)
- *options*:

-
- *identity* - only suite and owner names revealed
 - *description* - identity plus suite title and description
 - *state-totals* - identity, description, and task state totals
 - *full-read* - full read-only access for monitor and GUI
 - *shutdown* - full read access plus shutdown, but no other control.
 - *default*: state-totals

C Gcylc GUI (cylc gui) Config File Reference

This section defines all legal items and values for the gcylc user config file, which should be located in `$HOME/.cylc/gcylc.rc`. Current settings can be printed with the `cylc get-gui-config` command.

C.1 Top Level Items

C.1.1 dot icon size

Set the size of the task state dot icons displayed in the text and dot views.

- *type*: string
- *legal values*: “small” (10px), “medium” (14px), “large” (20px), “extra large (30px)”
- *default*: “medium”

C.1.2 initial side-by-side views

Set the suite view panels initial orientation when the GUI starts. This can be changed later using the “View” menu “Toggle views side-by-side” option.

- *type*: boolean (False or True)
- *default*: “False”

C.1.3 initial views

Set the suite view panel(s) displayed initially, when the GUI starts. This can be changed later using the tool bar.

- *type*: string (a list of one or two view names)
- *legal values*: “text”, “dot”, “graph”
- *default*: “text”
- *example*: `initial views = graph, dot`

C.1.4 maximum update interval

Set the maximum (longest) time interval between calls to the suite for data update.

The update frequency of the GUI is variable. It is determined by considering the time of last update and the mean duration of the last 10 main loops of the suite.

In general, the GUI will use an update frequency that matches the mean duration of the suite’s main loop. In quiet time (or if the suite is not contactable), it will gradually increase the update interval (i.e. reduce the update frequency) to a maximum determined by this setting.

Increasing this setting will reduce the network traffic and hits on the suite process. However, if a quiet suite starts to pick up activity, the GUI may initially appear out of sync with what is happening in the suite for the duration of this interval.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT15S`

C.1.5 sort by definition order

If this is not turned off the default sort order for task names and families in the dot and text views will be the order they appear in the suite definition. Clicking on the task name column in the treeview will toggle to alphanumeric sort, and a View menu item does the same for the dot view. If turned off, the default sort order is alphanumeric and definition order is not available at all.

- *type*: boolean
- *default*: `True`

C.1.6 sort column

If “text” is in `initial views` then `sort column` sets the column that will be sorted initially when the GUI launches. Sorting can be changed later by clicking on the column headers.

- *type*: string
- *legal values*: “task”, “state”, “host”, “job system”, “job ID”, “T-submit”, “T-start”, “T-finish”, “dT-mean”, “latest message”, “none”
- *default*: “none”
- *example*: `sort column = T-start`

C.1.7 sort column ascending

For use in combination with `sort column`, sets whether the column will be sorted using ascending or descending order.

- *type*: boolean
- *default*: “True”
- *example*: `sort column ascending = False`

C.1.8 task filter highlight color

The color used to highlight active task filters in gcylc. It must be a name from the X11 rgb.txt file, e.g. `SteelBlue`; or a *quoted* hexadecimal color code, e.g. `"#ff0000"` for red (quotes are required to prevent the hex code being interpreted as a comment).

- *type*: string
- *default*: `PowderBlue`

C.1.9 task states to filter out

Set the initial filtering options when the GUI starts. Later this can be changed by using the "View" menu "Task Filtering" option.

- *type*: string list
- *legal values*: waiting, held, queued, ready, expired, submitted, submit-failed, submit-retrying, running, succeeded, failed, retrying, runahead
- *default*: runahead

C.1.10 transpose dot

Transposes the content in dot view so that it displays from left to right rather than from top to bottom. Can be changed later using the options submenu available via the view menu.

- *type*: boolean
- *default*: “False”
- *example*: `transpose dot = True`

C.1.11 transpose graph

Transposes the content in graph view so that it displays from left to right rather than from top to bottom. Can be changed later using the options submenu via the view menu.

- *type*: boolean
- *default*: “False”
- *example*: `transpose graph = True`

C.1.12 ungrouped views

List suite views, if any, that should be displayed initially in an ungrouped state. Namespace family grouping can be changed later using the tool bar.

- *type*: string (a list of zero or more view names)
- *legal values*: “text”, “dot”, “graph”
- *default*: (none)
- *example*: `ungrouped views = text, dot`

C.1.13 use theme

Set the task state color theme, common to all views, to use initially. The color theme can be changed later using the tool bar. See `gcylc.rc.eg` and `themes.rc` in `$CYLC_DIR/conf/gcylc/` for how to modify existing color themes or define your own. Use `cylc get-gui-config` to list your available themes.

- *type*: string (theme name)
- *legal values*: “default”, “solid”, “high-contrast”, “color-blind”, and any custom or user-modified themes.
- *default*: “default”

C.1.14 window size

Sets the size (in pixels) of the cylc GUI at startup.

- *type*: integer list: x, y
- *legal values*: positive integers
- *default*: 800, 500
- *example*: `window size = 1000, 700`

C.2 [themes]

This section may contain task state color theme definitions.

C.2.1 [themes] →[[THEME]]

The name of the task state color-theme to be defined in this section.

- *type*: string

C.2.1.1 [themes] →[[THEME]] →inherit

You can inherit from another theme in order to avoid defining all states.

- *type*: string (parent theme name)
- *default*: “default”

C.2.1.2 [themes] →[[THEME]] →defaults

Set default icon attributes for all state icons in this theme.

- *type*: string list (icon attributes)
- *legal values*: “color=COLOR”, “style=STYLE”, “fontcolor=FontColor”
- *default*: (none)

For the attribute values, COLOR and FontColor can be color names from the X11 rgb.txt file, e.g. `SteelBlue`; or hexadecimal color codes, e.g. `#ff0000` for red; and STYLE can be “filled” or “unfilled”. See `gcyrc.rc` and `themes.rc` in `$CYLC_DIR/conf/gcylc/` for examples.

C.2.1.3 [themes] →[[THEME]] →STATE

Set icon attributes for all task states in THEME, or for a subset of them if you have used theme inheritance and/or defaults. Legal values of STATE are any of the cylc task proxy states: *waiting*, *runahead*, *held*, *queued*, *ready*, *submitted*, *submit-failed*, *running*, *succeeded*, *failed*, *retrying*, *submit-retrying*.

- *type*: string list (icon attributes)
- *legal values*: “color=COLOR”, “style=STYLE”, “fontcolor=FontColor”
- *default*: (none)

For the attribute values, COLOR and FONTCOLOR can be color names from the X11 rgb.txt file, e.g. `SteelBlue`; or hexadecimal color codes, e.g. `#ff0000` for red; and STYLE can be “filled” or “unfilled”. See `gcylc.rc.eg` and `themes.rc` in `$CYLC_DIR/conf/gcylc/` for examples.

D Gscan GUI (cylc gscan) Config File Reference

This section defines all legal items and values for the gscan config file which should be located in `$HOME/.cylc/gscan.rc`. Some items also affect the gpanel panel app.

The main menubar can be hidden to maximise the display area. Its visibility can be toggled via the mouse right-click menu, or by typing Alt-m. When visible, the main View menu allows you to change properties such as the columns that are displayed, which hosts to scan for running suites, and the task state icon theme.

At startup, the task state icon theme and icon size are taken from the gcylc config file `$HOME/.cylc/gcylc.rc`.

D.1 Top Level Items

D.1.1 activate on startup

Set whether `cylc gpanel` will activate automatically when the gui is loaded or not.

- *type*: boolean (True or False)
- *legal values*: “True”, “False”
- *default*: “False”
- *example*: `activate on startup = True`

D.1.2 columns

Set the columns to display when the `cylc gscan` GUI starts. This can be changed later with the View menu. The order in which the columns are specified here does not affect the display order.

- *type*: string (a list of one or more view names)
- *legal values*: “host”, “owner”, “status”, “suite”, “title”, “updated”
- *default*: “status”, “suite”
- *example*: `columns = suite, title, status`

D.1.3 suite listing update interval

Set the time interval between refreshing the suite listing (by file system or port range scan).

Increasing this setting will reduce the frequency of gscan looking for running suites. Scanning for suites by port range scan can be a hit on the network and the running suite processes, while scanning for suites by walking the file system can hit the file system (especially if the file system is a network file system). Therefore, this is normally set with a lower frequency than the status update interval. Increasing this setting will make gscan friendlier to the network and/or the file system, but gscan may appear out of sync if there are many start up or shut down of suites between the intervals.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT1M`

D.1.4 suite status update interval

Set the time interval between calls to known running suites (suites that are known via the latest suite listing) for data updates.

Increasing this setting will reduce the network traffic and hits on the suite processes. However, gscan may appear out of sync with what may be happening in very busy suites.

- *type*: ISO 8601 duration/interval representation (e.g. `PT10S`, 10 seconds, or `PT1M`, 1 minute).
- *default*: `PT15S`

D.1.5 window size

Sets the size in pixels of the `cylc gscan` GUI window at startup.

- *type*: integer list: x, y
- *legal values*: positive integers
- *default*: 300, 200
- *example*: `window size = 1000, 700`

D.1.6 hide main menubar

Hide the main menubar of the `cylc gscan` GUI window at startup. By default, the menubar is initially hidden. Either way, you can toggle its visibility with Alt-m or via the right-click menu.

- *type*: boolean (True or False)
- *default*: True
- *example*: `hide main menubar = False`

E Remote Job Host Interaction

This appendix shows transcripts of all `ssh`, `scp`, and `rsync` invocations from a Cylc suite called `test` started by `rose suite-run` on `vagrant@localhost`. The suite submits a single task called `foo` to `hobo@otherhost`. The operator polls then kills the task, then shuts the suite down.

Note that Cylc has been configured (site/global config) not to use a bash login shell on the remote host, but `rose suite-run` is executing `bash --login -c` (this may be optional in the future).

The transcripts were obtained by wrapping `ssh`, `scp`, and `rsync` with scripts that log the command line before invoking the command.

```
# (ROSE) vagrant@localhost: test for shared filesystem and create share/cycle
etc.
ssh -oBatchMode=yes -oConnectTimeout=10 hobo@otherhost bash --login -c '
ROSE_VERSION=2016.12.0 rose suite-run -v -v --name=test --run=run --remote=
uuid=78f99c51-fd82-4c4e-b666-944a7f5454a7'

# (ROSE) vagrant@localhost: copy suite source files to job host
rsync -a --exclude=.* --timeout=1800 --rsh=ssh -oBatchMode=yes -oConnectTimeout
=10 --exclude=78f99c51-fd82-4c4e-b666-944a7f5454a7 --exclude=log/78f99c51-
fd82-4c4e-b666-944a7f5454a7 --exclude=share/78f99c51-fd82-4c4e-b666-944
a7f5454a7 --exclude=share/cycle/78f99c51-fd82-4c4e-b666-944a7f5454a7 --
exclude=work/78f99c51-fd82-4c4e-b666-944a7f5454a7 --exclude=./.* --exclude=/
cylc-suite.db --exclude=/log --exclude=/log.* --exclude=/state --exclude=/
share --exclude=/work ./ hobo@otherhost:cylc-run/test

# (RSYNC INTERNAL COMMAND) vagrant@localhost
ssh -oBatchMode=yes -oConnectTimeout=10 -l hobo otherhost rsync --server -
logDtpre.iLs --timeout=1800 . cylc-run/test

# (RSYNC INTERNAL COMMAND) vagrant@localhost
rsync --server -logDtpre.iLs --timeout=1800 . cylc-run/test

# vagrant@localhost: test for shared FS on suite and job host
ssh -oBatchMode=yes -oConnectTimeout=10 -n hobo@otherhost test -e $HOME/cylc-
run/test/.service/87a8ff0a-dbc7-4ec2-a0f5-7ffd44f8e99c

# vagrant@localhost: create log and service directories on job host
ssh -oBatchMode=yes -oConnectTimeout=10 -n hobo@otherhost mkdir -p $HOME/cylc-
run/test $HOME/cylc-run/test/log/job $HOME/cylc-run/test/.service

# vagrant@localhost: copy suite service and contact files to job host
scp -oBatchMode=yes -oConnectTimeout=10 -p /home/vagrant/cylc-run/test/.service
/contact /home/vagrant/cylc-run/test/.service/passphrase /home/vagrant/cylc-
run/test/.service/ssl.cert hobo@otherhost:$HOME/cylc-run/test/.service/

# (SCP INTERNAL COMMAND) vagrant@localhost
ssh -x -oForwardAgent=no -oPermitLocalCommand=no -oClearAllForwardings=yes -o
BatchMode=yes -o ConnectTimeout=10 -l hobo -- otherhost scp -p -d -t $HOME/
cylc-run/test/.service/

# (SCP INTERNAL COMMAND) hobo@otherhost
scp -p -d -t /home/hobo/cylc-run/test/.service/

# vagrant@localhost: submit job foo.1 (cycle point 1, try 01) on job host
ssh -oBatchMode=yes -oConnectTimeout=10 hobo@otherhost env CYLC_VERSION=7.3.0
cylc jobs-submit '--remote-mode' '--' '$HOME/cylc-run/test/log/job' '1/foo
/01'

# vagrant@localhost: poll job foo.1 (cycle point 1, try 01) on job host
ssh -oBatchMode=yes -oConnectTimeout=10 hobo@otherhost env CYLC_VERSION=7.3.0
cylc jobs-poll '--' '$HOME/cylc-run/test/log/job' '1/foo/01'

# vagrant@localhost: kill job foo.1 (cycle point 1, try 01) on job host
ssh -oBatchMode=yes -oConnectTimeout=10 hobo@otherhost env CYLC_VERSION=7.3.0
cylc jobs-kill '--' '$HOME/cylc-run/test/log/job' '1/foo/01'

# vagrant@localhost: retrieve job logs from job host
```

E REMOTE JOB HOST INTERACTION

```
rsync -a --rsh=ssh -oBatchMode=yes -oConnectTimeout=10 --include=/1 --include
=/1/foo --include=/1/foo/01 --include=/1/foo/01/** --exclude=/**
hobo@otherhost:$HOME/cylc-run/test/log/job/ /home/vagrant/cylc-run/test/log/
job/

# (INTERNAL RSYNC COMMAND) vagrant@localhost
ssh -oBatchMode=yes -oConnectTimeout=10 -l hobo otherhost rsync --server --
sender -logDtpre.iLs . $HOME/cylc-run/test/log/job/

# (INTERNAL RSYNC COMMAND) hobo@otherhost
rsync --server --sender -logDtpre.iLs . /home/hobo/cylc-run/test/log/job/

# vagrant@localhost: remove suite contact file from job host
ssh -oBatchMode=yes -oConnectTimeout=10 -n hobo@otherhost rm -f $HOME/cylc-run/
test/.service/contact

# vagrant@localhost CLI: cylc cat-log --list-remote test foo.1
ssh -oBatchMode=yes -oConnectTimeout=10 -n hobo@otherhost ls $HOME/cylc-run/foo
/log/job/1/foo/NN

# vagrant@localhost CLI: cylc cat-log <OPT> test foo.1
ssh -oBatchMode=yes -oConnectTimeout=10 -n hobo@otherhost cat $HOME/cylc-run/
foo/log/job/1/foo/NN/<LOG>
# where <LOG> (depending on <OPT>, see "cylc cat-log --help") can be:
# * job
# * job.out
# * job.err
# * job.status
# * job.activity.log
# * job-edit.diff
# * job.xtrace

# vagrant@localhost: live tail remote job log, via:
# GUI: View -> job stdout, etc.
# CLI: cylc cat-log --tail <OPT> test foo.1
ssh -oBatchMode=yes -oConnectTimeout=10 -n hobo@otherhost tail --pid='ps h -o
ppid $$ | sed -e s/[[:space:]]//g' -n +1 -F $HOME/cylc-run/test/log/job/1/
foo/01/<LOG>
# (where <LOG> and <OPT> are as described above)
```

F Command Reference

Cylc ("silk") is a suite engine and metascheduler that specializes in cycling weather and climate forecasting suites and related processing (but it can also be used for one-off workflows of non-cycling tasks). For detailed documentation see the Cylc User Guide (`cylc doc --help`).

Version 7.6.0

The graphical user interface for cylc is "gcylc" (a.k.a. "cylc gui").

USAGE:

```
% cylc -v,--version           # print cylc version
% cylc version                 # (ditto, by command)
% cylc help,--help,-h,?       # print this help page

% cylc help CATEGORY          # print help by category
% cylc CATEGORY help          # (ditto)
% cylc help [CATEGORY] COMMAND # print command help

% cylc COMMAND [options] SUITE [arguments]
% cylc COMMAND [options] SUITE TASK [arguments]
```

Commands can be abbreviated as long as there is no ambiguity in the abbreviated command:

```
% cylc trigger SUITE TASK      # trigger TASK in SUITE
% cylc tri SUITE TASK          # ditto
% cylc t SUITE TASK             # ditto

% cylc get                      # Error: ambiguous command
```

TASK IDENTIFICATION IN CYLC SUITES

Tasks are identified by NAME.CYCLE_POINT where POINT is either a date-time or an integer.
Date-time cycle points are in an ISO 8601 date-time format, typically CCYYMMDDThhmm followed by a time zone - e.g. 20101225T0600Z.
Integer cycle points (including those for one-off suites) are integers - just '1' for one-off suites.

HOW TO DRILL DOWN TO COMMAND USAGE **HELP**:

```
% cylc help           # list all available categories (this page)
% cylc help prep       # list commands in category 'preparation'
% cylc help prep edit  # command usage help for 'cylc [prep] edit'
```

Command CATEGORIES:

```
control ..... Suite start up, monitoring, and control.
information ... Interrogate suite definitions and running suites.
all ..... The complete command set.
task ..... The task messaging interface.
license|GPL ... Software licensing information (GPL v3.0).
admin ..... Cylc installation, testing, and example suites.
preparation ... Suite editing, validation, visualization, etc.
hook ..... Suite and task event hook scripts.
discovery ..... Detect running suites.
utility ..... Cycle arithmetic and templating, etc.
```

F.1 Command Categories

F.1.1 admin

CATEGORY: admin - Cylc installation, testing, and example suites.

HELP: `cylc [admin] COMMAND help,--help`
You can abbreviate admin and COMMAND.
The category admin may be omitted.

COMMANDS:

```

check-software .... Check required software is installed.
import-examples ... Import example suites your suite run directory
profile-battery ... Run a battery of profiling tests
test-battery ..... Run a battery of self-diagnosing test suites
upgrade-run-dir ... Upgrade a pre-cylc-6 suite run directory

```

F.1.2 all

CATEGORY: all - The complete command set.

HELP: cylc [all] COMMAND help,--help
 You can abbreviate all and COMMAND.
 The category all may be omitted.

COMMANDS:

```

5to6 ..... Improve the cylc 6
      compatibility of a cylc 5 suite file
broadcast|bcast ..... Change suite [runtime]
      settings on the fly
cat-log|log ..... Print various suite and task
      log files
cat-state ..... Print the state of tasks from
      the state dump
check-software ..... Check required software is
      installed.
check-triggering ..... A suite shutdown event hook
      for cylc testing
check-versions ..... Compare cylc versions on task
      host accounts
checkpoint ..... Tell suite to checkpoint its
      current state
conditions ..... Print the GNU General Public
      License v3.0
cycle-point|cyclepoint|datetime|cycletime .... Cycle point arithmetic and
      filename templating
diff|compare ..... Compare two suite definitions
      and print differences
documentation|browse ..... Display cylc documentation (
      User Guide etc.)
dump ..... Print the state of tasks in a
      running suite
edit ..... Edit suite definitions,
      optionally inlined
email-suite ..... A suite event hook script
      that sends email alerts
email-task ..... A task event hook script that
      sends email alerts
ext-trigger|external-trigger ..... Report an external trigger
      event to a suite
get-directory ..... Retrieve suite source
      directory paths
get-gui-config ..... Print gcylc configuration
      items
get-suite-config|get-global-config ..... Print site/user configuration
      items
get-suite-config|get-config ..... Print suite configuration
      items
get-suite-contact|get-contact|print-contact ... Print contact information of
      a suite server program
get-suite-version|get-cylc-version ..... Print cylc version of a suite
      server program
gpanel ..... Internal interface for GNOME
      2 panel applet
graph ..... Plot suite dependency graphs
      and runtime hierarchies
graph-diff ..... Compare two suite
      dependencies or runtime hierarchies
gscan|gsummary ..... Scan GUI for monitoring
      multiple suites
gui ..... (a.k.a. gcylc) cylc GUI for
      suite control etc.

```

hold	Hold (pause) suites or individual tasks
import-examples	Import example suites your suite run directory
insert	Insert tasks into a running suite
jobs-kill	(Internal) Kill task jobs
jobs-poll	(Internal) Retrieve status for task jobs
jobs-submit	(Internal) Submit task jobs
jobscript	Generate a task job script and print it to stdout
kill	Kill submitted or running tasks
list ls	List suite tasks and family namespaces
ls-checkpoints	Display task pool etc at given events
message task-message	(task messaging) Report task messages
monitor	An in-terminal suite monitor (see also gcylc)
nudge	Cause the cylc task processing loop to be invoked
ping	Check that a suite is running
poll	Poll submitted or running tasks
print	Print registered suites
profile-battery	Run a battery of profiling tests
register	Register a suite for use
release unhold	Release (unpause) suites or individual tasks
reload	Reload the suite definition at run time
remote-init	(Internal) Initialise a task remote
remote-tidy	(Internal) Tidy a task remote
remove	Remove tasks from a running suite
report-timings	Generate a report on task timing data
reset	Force one or more tasks to change state.
restart	Restart a suite from a previous state
run start	Start a suite at a given cycle point
scan	Scan a host for running suites
scp-transfer	Scp-based file transfer for cylc suites
search grep	Search in suite definitions
set-verbosity	Change a running suite's logging verbosity
show	Print task state (prerequisites and outputs etc.)
spawn	Force one or more tasks to spawn their successors.
stop shutdown	Shut down running suites
submit single	Run a single task just as its parent suite would
suite-state	Query the task states in a suite
test-battery	Run a battery of self-diagnosing test suites
trigger	Manually trigger or re-trigger a task
upgrade-run-dir	Upgrade a pre-cylc-6 suite run directory
validate	Parse and validate suite definitions

```

version ..... Print the cylc release
  version
view ..... View suite definitions,
  inlined and Jinja2 processed
warranty ..... Print the GPLv3 disclaimer of
  warranty

```

F.1.3 control

CATEGORY: control - Suite start up, monitoring, and control.

HELP: cylc [control] COMMAND help,--help
 You can abbreviate control and COMMAND.
 The category control may be omitted.

COMMANDS:

```

broadcast|bcast ..... Change suite [runtime] settings on the fly
checkpoint ..... Tell suite to checkpoint its current state
ext-trigger|external-trigger ... Report an external trigger event to a suite
gui ..... (a.k.a. gcylc) cylc GUI for suite control
etc.
hold ..... Hold (pause) suites or individual tasks
insert ..... Insert tasks into a running suite
kill ..... Kill submitted or running tasks
nudge ..... Cause the cylc task processing loop to be
  invoked
poll ..... Poll submitted or running tasks
release|unhold ..... Release (unpause) suites or individual tasks
reload ..... Reload the suite definition at run time
remove ..... Remove tasks from a running suite
reset ..... Force one or more tasks to change state.
restart ..... Restart a suite from a previous state
run|start ..... Start a suite at a given cycle point
set-verbosity ..... Change a running suite's logging verbosity
spawn ..... Force one or more tasks to spawn their
  successors.
stop|shutdown ..... Shut down running suites
trigger ..... Manually trigger or re-trigger a task

```

F.1.4 discovery

CATEGORY: discovery - Detect running suites.

HELP: cylc [discovery] COMMAND help,--help
 You can abbreviate discovery and COMMAND.
 The category discovery may be omitted.

COMMANDS:

```

check-versions ... Compare cylc versions on task host accounts
ping ..... Check that a suite is running
scan ..... Scan a host for running suites

```

F.1.5 hook

CATEGORY: hook - Suite and task event hook scripts.

HELP: cylc [hook] COMMAND help,--help
 You can abbreviate hook and COMMAND.
 The category hook may be omitted.

COMMANDS:

```

check-triggering ... A suite shutdown event hook for cylc testing
email-suite ..... A suite event hook script that sends email alerts
email-task ..... A task event hook script that sends email alerts

```

F.1.6 information

CATEGORY: `information` - Interrogate suite definitions and running suites.

HELP: `cylc [information] COMMAND help,--help`
 You can abbreviate `information` and `COMMAND`.
 The category `information` may be omitted.

COMMANDS:

<code>cat-log log</code>	Print various suite and task log files
<code>cat-state</code>	Print the state of tasks from the state dump
<code>documentation browse</code>	Display cylc documentation (User Guide etc.)
<code>dump</code>	Print the state of tasks in a running suite
<code>get-gui-config</code>	Print gcylc configuration items
<code>get-site-config get-global-config</code>	Print site/user configuration items
<code>get-suite-config get-config</code>	Print suite configuration items
<code>get-suite-contact get-contact print-contact</code>	...	Print contact information of a suite server program
<code>get-suite-version get-cylc-version</code>	Print cylc version of a suite server program
<code>gpanel</code>	Internal interface for GNOME 2 panel applet
<code>gscan gsummary</code>	Scan GUI for monitoring multiple suites
<code>gui gcylc</code>	(a.k.a. gcylc) cylc GUI for suite control etc.
<code>list ls</code>	List suite tasks and family namespaces
<code>monitor</code>	An in-terminal suite monitor (see also gcylc)
<code>show</code>	Print task state (prerequisites and outputs etc.)
<code>version</code>	Print the cylc release version

F.1.7 license

CATEGORY: `license|GPL` - Software licensing information (GPL v3.0).

HELP: `cylc [license|GPL] COMMAND help,--help`
 You can abbreviate `license|GPL` and `COMMAND`.
 The category `license|GPL` may be omitted.

COMMANDS:

<code>conditions</code>	...	Print the GNU General Public License v3.0
<code>warranty</code>	Print the GPLv3 disclaimer of warranty

F.1.8 preparation

CATEGORY: `preparation` - Suite editing, validation, visualization, etc.

HELP: `cylc [preparation] COMMAND help,--help`
 You can abbreviate `preparation` and `COMMAND`.
 The category `preparation` may be omitted.

COMMANDS:

<code>5to6</code>	Improve the cylc 6 compatibility of a cylc 5 suite file
<code>diff compare</code>	Compare two suite definitions and print differences
<code>edit</code>	Edit suite definitions, optionally inlined
<code>get-directory</code>	...	Retrieve suite source directory paths
<code>graph</code>	Plot suite dependency graphs and runtime hierarchies
<code>graph-diff</code>	Compare two suite dependencies or runtime hierarchies
<code>jobscript</code>	Generate a task job script and print it to stdout
<code>list ls</code>	List suite tasks and family namespaces
<code>print</code>	Print registered suites

```

register ..... Register a suite for use
search|grep ..... Search in suite definitions
validate ..... Parse and validate suite definitions
view ..... View suite definitions, inlined and Jinja2 processed

```

F.1.9 task

CATEGORY: task - The task messaging interface.

HELP: cylc [task] COMMAND help,--help
 You can abbreviate task and COMMAND.
 The category task may be omitted.

COMMANDS:

```

jobs-kill ..... (Internal) Kill task jobs
jobs-poll ..... (Internal) Retrieve status for task jobs
jobs-submit ..... (Internal) Submit task jobs
message|task-message ... (task messaging) Report task messages
remote-init ..... (Internal) Initialise a task remote
remote-tidy ..... (Internal) Tidy a task remote
submit|single ..... Run a single task just as its parent suite would

```

F.1.10 utility

CATEGORY: utility - Cycle arithmetic and templating, etc.

HELP: cylc [utility] COMMAND help,--help
 You can abbreviate utility and COMMAND.
 The category utility may be omitted.

COMMANDS:

```

cycle-point|cyclepoint|datetime|cycletime ... Cycle point arithmetic and
  filename templating
ls-checkpoints ..... Display task pool etc at given
  events
report-timings ..... Generate a report on task
  timing data
scp-transfer ..... Scp-based file transfer for
  cylc suites
suite-state ..... Query the task states in a
  suite

```

F.2 Commands

F.2.1 5to6

Usage: cylc [prep] 5to6 FILE

Suggest changes to a cylc 5 suite file to make it more cylc 6 compatible.
 This may be a suite.rc file, an include file, or a suite.rc.processed file.

By default, print the changed file to stdout. Lines that have been changed are marked with '# UPGRADE'. These marker comments are purely for your own information and should not be included in any changes you make. In particular, they may break continuation lines.

Lines with '# UPGRADE CHANGE' have been altered.
 Lines with '# UPGRADE ... INFO' indicate that manual change is needed.

As of cylc 7, 'cylc validate' will no longer print out automatic dependency section translations. At cylc 6 versions of cylc, 'cylc validate' will show start-up/mixed async replacement R1* section(s). The validity of these can be highly dependent on the initial cycle point choice (e.g. whether it is T00 or T12).

This command works best for hour-based cycling - it will always convert e.g. 'foo[T-6]' to 'foo[-PT6H]', even where this is in a monthly or yearly cycling section graph.

This command is an aid, and is not an auto-upgrader or a substitute for reading the documentation. The suggested changes must be understood and checked by hand.

Example **usage**:

```
# Print out a file path (FILE) with suggested changes to stdout.
cylc 5to6 FILE

# Replace the file with the suggested changes file.
cylc 5to6 FILE > FILE

# Save a copy of the changed file.
cylc 5to6 FILE > FILE.5to6

# Show the diff of the changed file vs the original file.
diff - <(cylc 5to6 FILE) <FILE
```

Options:

```
-h, --help    Print this help message and exit.
```

F.2.2 broadcast

Usage: `cylc [control] broadcast|bcast [OPTIONS] REG`

Override [runtime] config in targeted namespaces in a running suite.

Uses for broadcast include making temporary changes to task behaviour, and task-to-downstream-task communication via environment variables.

A broadcast can target any [runtime] namespace for all cycles or for a specific cycle. If a task is affected by specific-cycle and all-cycle broadcasts at once, the specific takes precedence. If a task is affected by broadcasts to multiple ancestor namespaces, the result is determined by normal [runtime] inheritance. In other words, it follows this order:

```
all:root -> all:FAM -> all:task -> tag:root -> tag:FAM -> tag:task
```

Broadcasts persist, even across suite restarts, until they expire when their target cycle point is older than the oldest current in the suite, or until they are explicitly cancelled with this command. All-cycle broadcasts do not expire.

For each task the final effect of all broadcasts to all namespaces is computed on the fly just prior to job submission. The `--cancel` and `--clear` options simply cancel (remove) active broadcasts, they do not act directly on the final task-level result. Consequently, for example, you cannot broadcast to "all cycles except Tn" with an all-cycle broadcast followed by a cancel to Tn (there is no direct broadcast to Tn to cancel); and you cannot broadcast to "all members of FAMILY except member_n" with a general broadcast to FAMILY followed by a cancel to member_n (there is no direct broadcast to member_n to cancel).

To broadcast a variable to all tasks (quote items with internal spaces):

```
% cylc broadcast -s "[environment]VERSE = the quick brown fox" REG
```

To do the same with a file:

```
% cat >'broadcast.rc' <<'__RC__'
% [environment]
%     VERSE = the quick brown fox
% __RC__
% cylc broadcast -F 'broadcast.rc' REG
```

To cancel the same broadcast:

```
% cylc broadcast --cancel "[environment]VERSE" REG
```

If `-F FILE` was used, the same file can be used to cancel the broadcast:

```
% cylc broadcast -G 'broadcast.rc' REG
```

Use `-d/--display` to see active broadcasts. Multiple `--cancel` options or multiple `--set` and `--set-file` options can be used on the same command line. Multiple `--set` and `--set-file` options are cumulative.

The `--set-file=FILE` option can be used when broadcasting multiple values, or when the value contains newline or other metacharacters. If FILE is "-", read from standard input.

Broadcast cannot change [runtime] inheritance.

See also 'cylc reload' - reload a modified suite definition at run time.

Arguments:

REG Suite name

Options:

```
-h, --help                show this help message and exit
-p CYCLE_POINT, --point=CYCLE_POINT
                           Target cycle point. More than one can be added.
                           Defaults to '*' with --set and --cancel, and nothing
                           with --clear.
-n NAME, --namespace=NAME
                           Target namespace. Defaults to 'root' with --set and
                           --cancel, and nothing with --clear.
-s [SEC]ITEM=VALUE, --set=[SEC]ITEM=VALUE
                           A [runtime] config item and value to broadcast.
-F FILE, --set-file=FILE, --file=FILE
                           File with config to broadcast. Can be used multiple
                           times.
-c [SEC]ITEM, --cancel=[SEC]ITEM
                           An item-specific broadcast to cancel.
-G FILE, --cancel-file=FILE
                           File with broadcasts to cancel. Can be used multiple
                           times.
-C, --clear                Cancel all broadcasts, or with -p/--point,
                           -n/--namespace, cancel all broadcasts to targeted
                           namespaces and/or cycle points. Use "-C -p '*' " to
                           cancel all all-cycle broadcasts without canceling all
                           specific-cycle broadcasts.
-e CYCLE_POINT, --expire=CYCLE_POINT
                           Cancel any broadcasts that target cycle points earlier
                           than, but not inclusive of, CYCLE_POINT.
-d, --display              Display active broadcasts.
-k TASKID, --display-task=TASKID
                           Print active broadcasts for a given task
                           (NAME.CYCLE_POINT).
-b, --box                  Use unicode box characters with -d, -k.
-r, --raw                  With -d/--display or -k/--display-task, write out the
                           broadcast config structure in raw Python form.
--user=USER                Other user account name. This results in command
                           reinvocation on the remote account.
--host=HOST                Other host name. This results in command reinvocation
                           on the remote account.
-v, --verbose              Verbose output mode.
--debug                    Run suites in non-daemon mode, and show exception
                           tracebacks.
--port=INT                 Suite port number on the suite host. NOTE: this is
                           retrieved automatically if non-interactive ssh is
                           configured to the suite host.
--use-ssh                  Use ssh to re-invoke the command on the suite host.
--no-login                 Do not use a login shell to run remote ssh commands.
                           The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC
                           Set a timeout for network connections to the running
                           suite. The default is no timeout. For task messaging
                           connections see site/user config file documentation.
--print-uuid               Print the client UUID to stderr. This can be matched
                           to information logged by the receiving suite server
                           program.
--set-uuid=UUID            Set the client UUID manually (e.g. from prior use of
                           --print-uuid). This can be used to log multiple
                           commands under the same UUID (but note that only the
                           first [info] command from the same client ID will be
                           logged unless the suite is running in debug mode).
-f, --force                Do not ask for confirmation before acting. Note that
                           it is not necessary to use this option if interactive
```

command prompts have been disabled in the site/user config files.

F.2.3 cat-log

Usage: `cylc [info] cat-log|log [OPTIONS] REG [TASK-ID]`

Print or view suite and task job logs, or their locations (with no options suite logs or task job scripts are printed). For all task job logs, use the same cycle point format as the suite (list a job log directory to see it).

By default this prints the target file to stdout. With '--tail' it tails the file in real time, or with '-g' or '-b' it opens a temporary copy of it in your text editor. In the GUI, right-click 'View' tails the file in a pop-up text window, or 'View in Editor' opens a temporary copy of it in your editor.

Arguments:

REG	Suite name
[TASK-ID]	Task ID

Options:

-h, --help	show this help message and exit
-l, --location	Print location of the log file, exit 0 if it exists, exit 1 otherwise
-o, --stdout	Suite log: out, task job log: job.out
-e, --stderr	Suite log: err, task job log: job.err
-r INT, --rotation=INT	Suite logs log rotation number
-a, --activity	Task job log only: Short for --filename=job-activity.log
-d, --diff	Task job log only: Short for --filename=job-edit.diff (file present after an edit-run).
-x, --xtrace	Task job log only: Short for --filename=job.xtrace (file present after a debug run).
-u, --status	Task job log only: Short for --filename=job.status
-f FILENAME, -c FILENAME, --filename=FILENAME, --custom=FILENAME	Name of log file (e.g. 'job.stats').
--tail	Tail the job log, if the task is running.
-s INT, -t INT, --submit-number=INT, --try-number=INT	Task job log only: submit number (default=NN).
-i, --list-local	List a log directory on the suite host
-y, --list-remote	Task job log only: List log directory on the job host
-g, --geditor	Open the log file in your configured GUI editor.
-b, --teditor	Open the log file in your configured Non-GUI editor.
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.

F.2.4 cat-state

Usage: `cylc [info] cat-state [OPTIONS] REG`

Print the suite state in the old state dump file format to stdout. This command is deprecated; use "`cylc ls-checkpoints`" instead.

Arguments:

REG	Suite name
-----	------------

Options:

-h, --help	show this help message and exit
-d, --dump	Use the same display format as the 'cylc dump' command.
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.

`--debug` Run suites in non-daemon mode, and show exception tracebacks.

F.2.5 check-software

Usage: `cylc [admin] check-software`

Check that the external software required by cylc is installed.

Minimum versions are not checked, except in the case of Python.

Options:

`-h, --help` Print this help message and exit.

F.2.6 check-triggering

`cylc [hook] check-triggering ARGS`

This is a cylc shutdown event handler that compares the newly generated suite log with a previously generated reference log "`reference.log`" stored in the suite definition directory. Currently it just compares runtime triggering information, disregarding event order and timing, and fails the suite if there is any difference. This should be sufficient to verify correct scheduling of any suite that is not affected by different run-to-run conditional triggering.

1) run your suite with "`cylc run --generate-reference-log`" to generate the reference log with resolved triggering information. Check manually that the reference run was correct.
 2) run reference tests with "`cylc run --reference-test`" - this automatically sets the shutdown event handler along with a suite timeout and "`abort if shutdown handler fails`", "`abort on timeout`", and "`abort if any task fails`".

Reference tests can use any run mode:

- * simulation mode - tests that scheduling is equivalent to the reference
- * dummy mode - also tests that task hosting, job submission, job script evaluation, and cylc messaging are not broken.
- * live mode - tests everything (but takes longer with real tasks!)

If any task fails, or if cylc itself fails, or if triggering is not equivalent to the reference run, the test will abort with non-zero exit status - so reference tests can be used as automated tests to check that changes to cylc have not broken your suites.

F.2.7 check-versions

Usage: `cylc [discovery] check-versions [OPTIONS] SUITE`

Check the version of cylc invoked on each of SUITE's task host accounts when `CYLC_VERSION` is set to *the version running this command line tool*. Different versions are reported but are not considered an error unless the `-e|--error` option is specified, because different cylc versions from 6.0.0 onward should at least be backward compatible.

It is recommended that cylc versions be installed in parallel and access configured via the cylc version wrapper as described in the cylc INSTALL file and User Guide. This must be done on suite and task hosts. Users then get the latest installed version by default, or (like tasks) a particular version if `$CYLC_VERSION` is defined.

User `-v/--verbose` to see the command invoked to determine the remote version (all remote cylc command invocations will be of the same form, which may be site dependent -- see cylc global config documentation).

Arguments:

SUITE Suite name or path

Options:

`-h, --help` show this help message and exit

```

-e, --error          Exit with error status if 7.6.0 is not available on
                    all remote accounts.
-v, --verbose        Verbose output mode.
--debug             Run suites in non-daemon mode, and show exception
                    tracebacks.
--suite-owner=OWNER Specify suite owner
-s NAME=VALUE, --set=NAME=VALUE
                    Set the value of a Jinja2 template variable in the
                    suite definition. This option can be used multiple
                    times on the command line. NOTE: these settings
                    persist across suite restarts, but can be set again on
                    the "cylc restart" command line if they need to be
                    overridden.
--set-file=FILE      Set the value of Jinja2 template variables in the
                    suite definition from a file containing NAME=VALUE
                    pairs (one per line). NOTE: these settings persist
                    across suite restarts, but can be set again on the
                    "cylc restart" command line if they need to be
                    overridden.

```

F.2.8 checkpoint

Usage: `cylc [control] checkpoint [OPTIONS] REG CHECKPOINT-NAME`

Tell suite to checkpoint its current state.

Arguments:

REG	Suite name
CHECKPOINT-NAME	Checkpoint name

Options:

```

-h, --help          show this help message and exit
--user=USER         Other user account name. This results in command
                    reinvocation on the remote account.
--host=HOST         Other host name. This results in command reinvocation
                    on the remote account.
-v, --verbose        Verbose output mode.
--debug            Run suites in non-daemon mode, and show exception
                    tracebacks.
--port=INT          Suite port number on the suite host. NOTE: this is
                    retrieved automatically if non-interactive ssh is
                    configured to the suite host.
--use-ssh           Use ssh to re-invoke the command on the suite host.
--no-login          Do not use a login shell to run remote ssh commands.
                    The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC
                    Set a timeout for network connections to the running
                    suite. The default is no timeout. For task messaging
                    connections see site/user config file documentation.
--print-uuid        Print the client UUID to stderr. This can be matched
                    to information logged by the receiving suite server
                    program.
--set-uuid=UUID     Set the client UUID manually (e.g. from prior use of
                    --print-uuid). This can be used to log multiple
                    commands under the same UUID (but note that only the
                    first [info] command from the same client ID will be
                    logged unless the suite is running in debug mode).
-f, --force         Do not ask for confirmation before acting. Note that
                    it is not necessary to use this option if interactive
                    command prompts have been disabled in the site/user
                    config files.

```

F.2.9 conditions

Usage: `cylc [license] warranty [--help]`

Cylc is release under the GNU General Public License v3.0
This command prints the GPL v3.0 license in full.

Options:

--help Print this usage message.

F.2.10 cycle-point

Usage: `cylc [util] cycle-point [OPTIONS] [POINT]`

Cycle point date-time offset computation, and filename templating.

Filename templating replaces elements of a template string with corresponding elements of the current or given cycle point.

Use ISO 8601 or posix date-time format elements:

```
% cylc cyclepoint 2010080T00 --template foo-CCYY-MM-DD-Thh.nc
foo-2010-08-08-T00.nc
% cylc cyclepoint 2010080T00 --template foo-%Y-%m-%d-T%H.nc
foo-2010-08-08-T00.nc
```

Other examples:

- 1) print offset from an explicit cycle point:


```
% cylc [util] cycle-point --offset-hours=6 20100823T1800Z
20100824T0000Z
```
- 2) print offset from \$CYLC_TASK_CYCLE_POINT (as in suite tasks):


```
% export CYLC_TASK_CYCLE_POINT=20100823T1800Z
% cylc cycle-point --offset-hours=-6
20100823T1200Z
```
- 3) cycle point filename templating, explicit template:


```
% export CYLC_TASK_CYCLE_POINT=2010-08
% cylc cycle-point --offset-years=2 --template=foo-CCYY-MM.nc
foo-2012-08.nc
```
- 4) cycle point filename templating, template in a variable:


```
% export CYLC_TASK_CYCLE_POINT=2010-08
% export MYTEMPLATE=foo-CCYY-MM.nc
% cylc cycle-point --offset-years=2 --template=MYTEMPLATE
foo-2012-08.nc
```

Arguments:

[POINT] ISO 8601 date-time, e.g. 20140201T0000Z, default
\$CYLC_TASK_CYCLE_POINT

Options:

```
-h, --help                    show this help message and exit
--offset-hours=HOURS        Add N hours to CYCLE (may be negative)
--offset-days=DAYS         Add N days to CYCLE (N may be negative)
--offset-months=MONTHS      Add N months to CYCLE (N may be negative)
--offset-years=YEARS        Add N years to CYCLE (N may be negative)
--offset=ISO_OFFSET        Add an ISO 8601-based interval representation to CYCLE
--equal=POINT2             Succeed if POINT2 is equal to POINT (format agnostic).
--template=TEMPLATE        Filename template string or variable
--time-zone=TEMPLATE        Control the formatting of the result's timezone e.g.
                              (Z, +13:00, -hh)
--num-expanded-year-digits=NUMBER
                              Specify a number of expanded year digits to print in
                              the result
--print-year                Print only CCYY of result
--print-month               Print only MM of result
--print-day                 Print only DD of result
--print-hour                Print only hh of result
```

F.2.11 diff

Usage: `cylc [prep] diff|compare [OPTIONS] SUITE1 SUITE2`

Compare two suite definitions and display any differences.

Differencing is done after parsing the suite.rc files so it takes

account of default values that are not explicitly defined, it disregards the order of configuration items, and it sees any include-file content after inlining has occurred.

Files in the suite bin directory and other sub-directories of the suite definition directory are not currently differenced.

Arguments:

SUITE1 Suite name or path
SUITE2 Suite name or path

Options:

-h, --help show this help message and exit
-n, --nested print suite.rc section headings in nested form.
--user=USER Other user account name. This results in command reinvocation on the remote account.
--host=HOST Other host name. This results in command reinvocation on the remote account.
-v, --verbose Verbose output mode.
--debug Run suites in non-daemon mode, and show exception tracebacks.
--suite-owner=OWNER Specify suite owner
-s NAME=VALUE, --set=NAME=VALUE Set the value of a Jinja2 template variable in the suite definition. This option can be used multiple times on the command line. NOTE: these settings persist across suite restarts, but can be set again on the "cylc restart" command line if they need to be overridden.
--set-file=FILE Set the value of Jinja2 template variables in the suite definition from a file containing NAME=VALUE pairs (one per line). NOTE: these settings persist across suite restarts, but can be set again on the "cylc restart" command line if they need to be overridden.
--icp=CYCLE_POINT Set initial cycle point. Required if not defined in suite.rc.

F.2.12 documentation

Usage: cylc [info] documentation|browse [OPTIONS] [SUITE]

View documentation in browser or PDF viewer, as per Cylc global config.

% cylc doc [OPTIONS]

View local or internet [--www] Cylc documentation URLs.

% cylc doc [-t TASK] SUITE

View suite or task documentation, if URLs are specified in the suite. This parses the suite definition to extract the requested URL. Note that suite server programs also hold suite URLs for access from the Cylc GUI.

Arguments:

[TARGET] File, URL, or suite name

Options:

-h, --help show this help message and exit
-p, --pdf Open the PDF User Guide directly.
-w, --www Open the cylc internet homepage
-t TASK_NAME, --task=TASK_NAME Browse task documentation URLs.
-s, --stdout Just print the URL to stdout.
--user=USER Other user account name. This results in command reinvocation on the remote account.
--host=HOST Other host name. This results in command reinvocation on the remote account.
--debug Print exception traceback on error.
--url=URL URL to view in your configured browser.

F.2.13 dump

Usage: `cylc [info] dump [OPTIONS] REG`

Print state information (e.g. the state of each task) from a running suite. For small suites 'watch cylc [info] dump SUITE' is an effective non-GUI real time monitor (but see also 'cylc monitor').

For more information about a specific task, such as the current state of its prerequisites and outputs, see 'cylc [info] show'.

Examples:

Display the state of all running tasks, sorted by cycle point:
`% cylc [info] dump --tasks --sort SUITE | grep running`

Display the state of all tasks in a particular cycle point:
`% cylc [info] dump -t SUITE | grep 2010082406`

Arguments:

REG Suite name

Options:

<code>-h, --help</code>	show this help message and exit
<code>-g, --global</code>	Global information only.
<code>-t, --tasks</code>	Task states only.
<code>-r, --raw, --raw-format</code>	Display raw format.
<code>-s, --sort</code>	Task states only; sort by cycle point instead of name.
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.
<code>--port=INT</code>	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
<code>--use-ssh</code>	Use ssh to re-invoke the command on the suite host.
<code>--no-login</code>	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
<code>--comms-timeout=SEC, --pyro-timeout=SEC</code>	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
<code>--print-uuid</code>	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
<code>--set-uuid=UUID</code>	Set the client UUID manually (e.g. from prior use of --print-uuid). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).

F.2.14 edit

Usage: `cylc [prep] edit [OPTIONS] SUITE`

Edit suite definitions without having to move to their directory locations, and with optional reversible inlining of include-files. Note that Jinja2 suites can only be edited in raw form but the processed version can be viewed with 'cylc [prep] view -p'.

1/cylc [prep] edit SUITE

Change to the suite definition directory and edit the suite.rc file.

2/ cylc [prep] edit -i,--inline SUITE

Edit the suite with include-files inlined between special markers. The original suite.rc file is temporarily replaced so that the inlined version is "live" during editing (i.e. you can run suites during editing and cylc will pick up changes to the suite definition). The inlined file is then split into its constituent include-files again when you exit the editor. Include-files can be nested or

multiply-included; in the latter case only the first inclusion is inlined (this prevents conflicting changes made to the same file).

```
3/ cylc [prep] edit --cleanup SUITE
Remove backup files left by previous INLINED edit sessions.
```

INLINED EDITING SAFETY: The suite.rc file and its include-files are automatically backed up prior to an inlined editing session. If the editor dies mid-session just invoke 'cylc edit -i' again to recover from the last saved inlined file. On exiting the editor, if any of the original include-files are found to have changed due to external intervention during editing you will be warned and the affected files will be written to new backups instead of overwriting the originals. Finally, the inlined suite.rc file is also backed up on exiting the editor, to allow recovery in case of accidental corruption of the include-file boundary markers in the inlined file.

The edit process is spawned in the foreground as follows:

```
% <editor> suite.rc
Where <editor> is defined in the cylc site/user config files.
```

See also 'cylc [prep] view'.

Arguments:

SUITE Suite name or path

Options:

-h, --help	show this help message and exit
-i, --inline	Edit with include-files inlined as described above.
--cleanup	Remove backup files left by previous inlined edit sessions.
-g, --gui	Force use of the configured GUI editor.
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
--suite-owner=OWNER	Specify suite owner

F.2.15 email-suite

Usage: cylc [hook] email-suite EVENT SUITE MESSAGE

THIS COMMAND IS OBSOLETE - use built-in email event hooks.

This is a simple suite event hook script that sends an email. The command line arguments are supplied automatically by cylc.

For example, to get an email alert when a suite shuts down:

```
# SUITE.RC
[cylc]
  [[environment]]
    MAIL_ADDRESS = foo@bar.baz.waz
  [[events]]
    shutdown handler = cylc email-suite
```

See the Suite.rc Reference (Cylc User Guide) for more information on suite and task event hooks and event handler scripts.

F.2.16 email-task

Usage: cylc [hook] email-task EVENT SUITE TASKID MESSAGE

THIS COMMAND IS OBSOLETE - use built-in email event hooks.

A simple task event hook handler script that sends an email. The command line arguments are supplied automatically by cylc.

For example, to get an email alert whenever any task fails:

```
# SUITE.RC
[cylc]
  [[environment]]
    MAIL_ADDRESS = foo@bar.baz.waz
[runtime]
  [[root]]
    [[events]]
      failed handler = cylc email-task
```

See the Suite.rc Reference (Cylc User Guide) for more information on suite and task event hooks and event handler scripts.

F.2.17 ext-trigger

Usage: `cylc [control] ext-trigger [OPTIONS] REG MSG ID`

Report an external event message to a suite server program. It is expected that a task in the suite has registered the same message as an external trigger - a special prerequisite to be satisfied by an external system, via this command, rather than by triggering off other tasks.

The ID argument should uniquely distinguish one external trigger event from the next. When a task's external trigger is satisfied by an incoming message, the message ID is broadcast to all downstream tasks in the cycle point as `$CYLC_EXT_TRIGGER_ID` so that they can use it - e.g. to identify a new data file that the external triggering system is responding to.

Use the retry options in case the target suite is down or out of contact.

The suite passphrase must be installed in `$HOME/.cylc/<SUITE>/.`

Note: to manually trigger a task use 'cylc trigger', not this command.

Arguments:

REG	Suite name
MSG	External trigger message
ID	Unique trigger ID

Options:

<code>-h, --help</code>	show this help message and exit
<code>--max-tries=INT</code>	Maximum number of send attempts (default 5).
<code>--retry-interval=SEC</code>	Delay in seconds before retrying (default 10.0).
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.
<code>--port=INT</code>	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
<code>--use-ssh</code>	Use ssh to re-invoke the command on the suite host.
<code>--no-login</code>	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
<code>--comms-timeout=SEC, --pyro-timeout=SEC</code>	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
<code>--print-uuid</code>	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
<code>--set-uuid=UUID</code>	Set the client UUID manually (e.g. from prior use of <code>--print-uuid</code>). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
<code>-f, --force</code>	Do not ask for confirmation before acting. Note that

it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.

F.2.18 get-directory

Usage: `cylc [prep] get-directory REG`

Retrieve and print the source directory location of suite REG.
Here's an easy way to move to a suite source directory:
\$ `cd $(cylc get-dir REG)`.

Arguments:

SUITE Suite name or path

Options:

-h, --help show this help message and exit
 --user=USER Other user account name. This results in command reinvocation on the remote account.
 --host=HOST Other host name. This results in command reinvocation on the remote account.
 -v, --verbose Verbose output mode.
 --debug Run suites in non-daemon mode, and show exception tracebacks.
 --suite-owner=OWNER Specify suite owner

F.2.19 get-gui-config

Usage: `cylc [admin] get-gui-config [OPTIONS]`

Print gcylc configuration settings.

By default all settings are printed. For specific sections or items use -i/--item and wrap parent sections in square brackets:
`cylc get-gui-config --item '[themes][default]succeeded'`
 Multiple items can be specified at once.

Options:

-h, --help show this help message and exit
 -v, --verbose Print extra information.
 --debug Show exception tracebacks.
 -i [SEC...]ITEM, --item=[SEC...]ITEM Item or section to print (multiple use allowed).
 --sparse Only print items explicitly set in the config files.
 -p, --python Print native Python format.

F.2.20 get-site-config

Usage: `cylc [admin] get-site-config [OPTIONS]`

Print cylc site/user configuration settings.

By default all settings are printed. For specific sections or items use -i/--item and wrap parent sections in square brackets:
`cylc get-site-config --item '[editors]terminal'`
 Multiple items can be specified at once.

Options:

-h, --help show this help message and exit
 -i [SEC...]ITEM, --item=[SEC...]ITEM Item or section to print (multiple use allowed).
 --sparse Only print items explicitly set in the config files.
 -p, --python Print native Python format.
 --print-run-dir Print the configured cylc run directory.
 --print-site-dir Print the cylc site configuration directory location.
 -v, --verbose Print extra information.
 --debug Show exception tracebacks.

F.2.21 get-suite-config

Usage: `cylc [info] get-suite-config [OPTIONS] SUITE`

Print parsed suite configuration items, after runtime inheritance.

By default all settings are printed. For specific sections or items use `-i/--item` and wrap sections in square brackets, e.g.:

`cylc get-suite-config --item '[scheduling]initial cycle point'`
Multiple items can be retrieved at once.

By default, unset values are printed as an empty string, or (for historical reasons) as `"None"` with `-o/--one-line`. These defaults can be changed with the `-n/--null-value` option.

Example:

```
|# SUITE.RC
|[runtime]
|    [[modelX]]
|    [[[environment]]]
|        FOO = foo
|        BAR = bar
```

```
$ cylc get-suite-config --item=[runtime][modelX][environment]FOO SUITE
foo
```

```
$ cylc get-suite-config --item=[runtime][modelX][environment] SUITE
FOO = foo
BAR = bar
```

```
$ cylc get-suite-config --item=[runtime][modelX] SUITE
```

```
...
[[[environment]]]
    FOO = foo
    BAR = bar
...
```

Arguments:

`SUITE` Suite name or path

Options:

```
-h, --help            show this help message and exit
-i [SEC...]ITEM, --item=[SEC...]ITEM
                        Item or section to print (multiple use allowed).
-r, --sparse          Only print items explicitly set in the config files.
-p, --python          Print native Python format.
-a, --all-tasks       For [runtime] items (e.g. --item='script') report
                        values for all tasks prefixed by task name.
-n STRING, --null-value=STRING
                        The string to print for unset values (default
                        nothing).
-m, --mark-up         Prefix each line with '!cylc!'.
-o, --one-line        Print multiple single-value items at once.
-t, --tasks          Print the suite task list [DEPRECATED: use 'cylc list
                        SUITE'].
-u RUN_MODE, --run-mode=RUN_MODE
                        Get config for suite run mode.
--user=USER           Other user account name. This results in command
                        reinvocation on the remote account.
--host=HOST          Other host name. This results in command reinvocation
                        on the remote account.
-v, --verbose        Verbose output mode.
--debug              Run suites in non-daemon mode, and show exception
                        tracebacks.
--suite-owner=OWNER  Specify suite owner
-s NAME=VALUE, --set=NAME=VALUE
                        Set the value of a Jinja2 template variable in the
                        suite definition. This option can be used multiple
                        times on the command line. NOTE: these settings
                        persist across suite restarts, but can be set again on
                        the "cylc restart" command line if they need to be
                        overridden.
```

```

--set-file=FILE      Set the value of Jinja2 template variables in the
                     suite definition from a file containing NAME=VALUE
                     pairs (one per line). NOTE: these settings persist
                     across suite restarts, but can be set again on the
                     "cylc restart" command line if they need to be
                     overridden.
--icp=CYCLE_POINT    Set initial cycle point. Required if not defined in
                     suite.rc.

```

F.2.22 get-suite-contact

Usage: `cylc [info] get-suite-contact [OPTIONS] REG`

Print contact information of running suite REG.

Arguments:

REG Suite name

Options:

```

-h, --help          show this help message and exit
--user=USER         Other user account name. This results in command reinvocation
                   on the remote account.
--host=HOST         Other host name. This results in command reinvocation on the
                   remote account.
-v, --verbose       Verbose output mode.
--debug            Run suites in non-daemon mode, and show exception tracebacks.

```

F.2.23 get-suite-version

Usage: `cylc [info] get-suite-version [OPTIONS] REG`

Interrogate running suite REG to find what version of cylc is running it.

To find the version you've invoked at the command line see "`cylc version`".

Arguments:

REG Suite name

Options:

```

-h, --help          show this help message and exit
--user=USER         Other user account name. This results in command
                   reinvocation on the remote account.
--host=HOST         Other host name. This results in command reinvocation
                   on the remote account.
-v, --verbose       Verbose output mode.
--debug            Run suites in non-daemon mode, and show exception
                   tracebacks.
--port=INT          Suite port number on the suite host. NOTE: this is
                   retrieved automatically if non-interactive ssh is
                   configured to the suite host.
--use-ssh           Use ssh to re-invoke the command on the suite host.
--no-login          Do not use a login shell to run remote ssh commands.
                   The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC
                   Set a timeout for network connections to the running
                   suite. The default is no timeout. For task messaging
                   connections see site/user config file documentation.
--print-uuid        Print the client UUID to stderr. This can be matched
                   to information logged by the receiving suite server
                   program.
--set-uuid=UUID     Set the client UUID manually (e.g. from prior use of
                   --print-uuid). This can be used to log multiple
                   commands under the same UUID (but note that only the
                   first [info] command from the same client ID will be
                   logged unless the suite is running in debug mode).
-f, --force         Do not ask for confirmation before acting. Note that
                   it is not necessary to use this option if interactive
                   command prompts have been disabled in the site/user
                   config files.

```

F.2.24 gpanel

Usage: `cylc gpanel [OPTIONS]`

This is a cylc scan panel applet for monitoring running suites on a set of hosts in GNOME 2.

To install this applet, run `"cylc gpanel --install"` and follow the instructions that it gives you.

This applet can be tested using the `--test` option.

To customize themes, copy `$CYLC_DIR/conf/gcylc/gcylc.rc.eg` to `$HOME/.cylc/gcylc.rc` and follow the instructions in the file.

To configure default suite hosts, edit the [suite host scanning]hosts entry in your `global.rc` file.

Options:

```
-h, --help      show this help message and exit
--compact       Switch on compact mode at runtime.
--install       Install the panel applet.
--test          Run in a standalone window.
```

F.2.25 graph

Usage: 1/ `cylc [prep] graph [OPTIONS] SUITE [START[STOP]]`

Plot the suite.rc dependency graph for SUITE.

2/ `cylc [prep] graph [OPTIONS] -f,--file FILE`

Plot the specified dot-language graph file.

3/ `cylc [prep] graph [OPTIONS] --reference SUITE [START[STOP]]`

Print out a reference format for the dependencies in SUITE.

4/ `cylc [prep] graph [OPTIONS] --output-file FILE SUITE`

Plot SUITE dependencies to a file FILE with a extension-derived format.

If FILE ends with `".png"`, output in PNG format, etc.

Plot suite dependency graphs in an interactive graph viewer.

If START is given it overrides `"[visualization] initial cycle point"` to determine the start point of the graph, which defaults to the suite initial cycle point. If STOP is given it overrides `"[visualization] final cycle point"` to determine the end point of the graph, which defaults to the graph start point plus `"[visualization] number of cycle points"` (which defaults to 3). The graph start and end points are adjusted up and down to the suite initial and final cycle points, respectively, if necessary.

The `"Save"` button generates an image of the current view, of format (e.g. png, svg, jpg, eps) determined by the filename extension. If the chosen format is not available a dialog box will show those that are available.

If the optional output filename is specified, the viewer will not open and a graph will be written directly to the file.

GRAPH VIEWER CONTROLS:

- * Center on a node: left-click.
- * Pan view: left-drag.
- * Zoom: +/- buttons, mouse-wheel, or ctrl-left-drag.
- * Box zoom: shift-left-drag.
- * "Best Fit" and "Normal Size" buttons.
- * Left-to-right graphing mode toggle button.
- * "Ignore suicide triggers" button.
- * "Save" button: save an image of the view.

Family (namespace) grouping controls:

Toolbar:

- * `"group"` - group all families up to root.
- * `"ungroup"` - recursively ungroup all families.

Right-click menu:

- * `"group"` - close this node's parent family.
- * `"ungroup"` - open this family node.
- * `"recursive ungroup"` - ungroup all families below this node.

Arguments:

[SUITE]	Suite name or path
[START]	Initial cycle point (default: suite initial point)
[STOP]	Final cycle point (default: initial + 3 points)

Options:

-h, --help	show this help message and exit
-u, --ungrouped	Start with task families ungrouped (the default is grouped).
-n, --namespaces	Plot the suite namespace inheritance hierarchy (task run time properties).
-f FILE, --file=FILE	View a specific dot-language graphfile.
--filter=NODE_NAME_PATTERN	Filter out one or many nodes.
-O FILE, --output-file=FILE	Output to a specific file, with a format given by --output-format or extrapolated from the extension. '-' implies stdout in plain format.
--output-format=FORMAT	Specify a format for writing out the graph to --output-file e.g. png, svg, jpg, eps, dot. 'ref' is a special sorted plain text format for comparison and reference purposes.
-r, --reference	Output in a sorted plain text format for comparison purposes. If not given, assume --output-file=-.
--show-suicide	Show suicide triggers. They are not shown by default, unless toggled on with the tool bar button.
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
--suite-owner=OWNER	Specify suite owner
-s NAME=VALUE, --set=NAME=VALUE	Set the value of a Jinja2 template variable in the suite definition. This option can be used multiple times on the command line. NOTE: these settings persist across suite restarts, but can be set again on the "cylc restart" command line if they need to be overridden.
--set-file=FILE	Set the value of Jinja2 template variables in the suite definition from a file containing NAME=VALUE pairs (one per line). NOTE: these settings persist across suite restarts, but can be set again on the "cylc restart" command line if they need to be overridden.

F.2.26 graph-diff

Usage: `cylc graph-diff [OPTIONS] SUITE1 SUITE2 -- [GRAPH_OPTIONS_ARGS]`

Difference 'cylc graph --reference' output for SUITE1 and SUITE2.

OPTIONS: Use '-g' to launch a graphical diff utility.
Use '--diff-cmd=MY_DIFF_CMD' to use a custom diff tool.

SUITE1, SUITE2: Suite names to compare.

GRAPH_OPTIONS_ARGS: Options and arguments passed directly to cylc graph.

F.2.27 gscan

Usage: `cylc gscan [OPTIONS]`

This is the cylc scan gui for monitoring running suites on a set of hosts.

To customize themes copy `$CYLC_DIR/conf/gcylcrc/gcylc.rc.eg` to `$HOME/.cylc/gcylc.rc` and follow the instructions in the file.

Arguments :

```
[HOSTS ...]           Hosts to scan instead of the configured hosts.
```

Options:

```

-h, --help                show this help message and exit
-a, --all                 Scan all port ranges in known hosts.
-n PATTERN, --name=PATTERN
                           List suites with name matching PATTERN (regular
                           expression). Defaults to any name. Can be used
                           multiple times.
-o PATTERN, --suite-owner=PATTERN
                           List suites with owner matching PATTERN (regular
                           expression). Defaults to just your own suites. Can be
                           used multiple times.
--comms-timeout=SEC       Set a timeout for network connections to each running
                           suite. The default is 5 seconds.
--interval=SECONDS        Time interval (in seconds) between full updates
--user=USER               Other user account name. This results in command
                           reinvocation on the remote account.
--host=HOST               Other host name. This results in command reinvocation
                           on the remote account.
-v, --verbose             Verbose output mode.
--debug                  Run suites in non-daemon mode, and show exception
                           tracebacks.
--port=INT               Suite port number on the suite host. NOTE: this is
                           retrieved automatically if non-interactive ssh is
                           configured to the suite host.
--use-ssh                 Use ssh to re-invoke the command on the suite host.
--no-login                Do not use a login shell to run remote ssh commands.
                           The default is to use a login shell.
--print-uuid              Print the client UUID to stderr. This can be matched
                           to information logged by the receiving suite server
                           program.
--set-uuid=UUID           Set the client UUID manually (e.g. from prior use of
                           --print-uuid). This can be used to log multiple
                           commands under the same UUID (but note that only the
                           first [info] command from the same client ID will be
                           logged unless the suite is running in debug mode).
```

F.2.28 `gui`

```
Usage: cylc gui [OPTIONS] [REG] [USER_AT_HOST]
gcylc [OPTIONS] [REG] [USER_AT_HOST]
```

This is the cylc Graphical User Interface.

The `USER_AT_HOST` argument allows suite selection by 'cylc scan' output:

```
cylc gui $(cylc scan | grep <suite_name>)
```

Local suites can be opened and switched between from within gcylc. To connect to running remote suites (whose passphrase you have installed) you must currently use `--host` and/or `--user` on the gcylc command line.

Available task state color themes are shown under the View menu. To customize themes copy `$CYLC_DIR/conf/gcylrc/gcylc.rc.eg` to `$HOME/.cylc/gcylc.rc` and follow the instructions in the file.

To see current configuration settings use `"cylc get-gui-config"`.

In the graph view, View -> Options -> "Write Graph Frames" writes .dot graph files to the suite share directory (locally, for a remote suite). These can be processed into a movie by \CYLC DIR/dev/bin/live-graph-movie.sh=.

Arguments:

```
[REG] Suite name
[USER_AT_HOST] user@host:port, shorthand for --user, --host &
--port.
```

Options:

```
-h, --help          show this help message and exit
```

<code>-r, --restricted</code>	Restrict display to 'active' task states: submitted, submit-failed, submit-retrying, running, failed, retrying; and disable the graph view. This may be needed for very large suites. The state summary icons in the status bar still represent all task proxies.
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.
<code>--port=INT</code>	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
<code>--use-ssh</code>	Use ssh to re-invoke the command on the suite host.
<code>--no-login</code>	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
<code>--comms-timeout=SEC, --pyro-timeout=SEC</code>	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
<code>--print-uuid</code>	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
<code>--set-uuid=UUID</code>	Set the client UUID manually (e.g. from prior use of <code>--print-uuid</code>). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
<code>-s NAME=VALUE, --set=NAME=VALUE</code>	Set the value of a Jinja2 template variable in the suite definition. This option can be used multiple times on the command line. NOTE: these settings persist across suite restarts, but can be set again on the <code>"cylc restart"</code> command line if they need to be overridden.
<code>--set-file=FILE</code>	Set the value of Jinja2 template variables in the suite definition from a file containing NAME=VALUE pairs (one per line). NOTE: these settings persist across suite restarts, but can be set again on the <code>"cylc restart"</code> command line if they need to be overridden.

F.2.29 hold

Usage: `cylc [control] hold [OPTIONS] REG [TASKID ...]`

Hold one or more waiting tasks (`cylc hold REG TASKID ...`), or a whole suite (`cylc hold REG`).

Held tasks do not submit even if they are ready to run.

See also '`cylc [control] release`'.

TASKID is a pattern to match task proxies or task families, or groups of them:

- * `[CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]`
- * `[CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]`
- * `TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]`
- * `FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]`

For example, to match:

- * all tasks in a cycle: `'20200202T0000Z/*'` or `'*.20200202T0000Z'`
- * all tasks in the submitted status: `':submitted'`
- * retrying 'foo*' tasks in 0000Z cycles: `'foo*.*0000Z:retrying'` or `'*0000Z/foo*:retrying'`
- * retrying tasks in 'BAR' family: `'*/BAR:retrying'` or `'BAR.*:retrying'`
- * retrying tasks in 'BAR' or 'BAZ' families: `'*/BA[RZ]:retrying'` or `'BA[RZ].*:retrying'`

The old 'MATCH POINT' syntax will be automatically detected and supported. To

avoid this, use the '--no-multitask-compat' option, or use the new syntax (with a '/' or a '.') when specifying 2 TASKID arguments.

Arguments:

REG	Suite name
[TASKID ...]	Task identifiers

Options:

-h, --help	show this help message and exit
--after=CYCLE_POINT	Hold whole suite AFTER this cycle point.
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
--port=INT	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
--use-ssh	Use ssh to re-invoke the command on the suite host.
--no-login	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
--print-uuid	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
--set-uuid=UUID	Set the client UUID manually (e.g. from prior use of --print-uuid). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
-f, --force	Do not ask for confirmation before acting. Note that it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.
-m, --family	(Obsolete) This option is now ignored and is retained for backward compatibility only. TASKID in the argument list can be used to match task and family names regardless of this option.
--no-multitask-compat	Disallow backward compatible multitask interface.

F.2.30 import-examples

Usage: cylc [admin] import-examples DIR

Copy the cylc example suites to DIR and register them for use under the GROUP suite name group.

Arguments:

DIR	destination directory
-----	-----------------------

F.2.31 insert

Usage: cylc [control] insert [OPTIONS] REG TASKID [...]

Insert task proxies into a running suite. Uses of insertion include:

- 1) insert a task that was excluded by the suite definition at start-up.
- 2) reinstate a task that was previously removed from a running suite.
- 3) re-run an old task that cannot be retrigged because its task proxy is no longer live in the a suite.

Be aware that inserted cycling tasks keep on cycling as normal, even if another instance of the same task exists at a later cycle (instances of the same task at different cycles can coexist, but a newly spawned task will not be added to the pool if it catches up to another task with the

same ID).

See also 'cylc submit', for running tasks without the scheduler.

TASKID is a pattern to match task proxies or task families, or groups of them:

```
* [CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]
* [CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]
* TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
* FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
```

For example, to match:

```
* all tasks in a cycle: '20200202T0000Z/*' or '*.20200202T0000Z'
* all tasks in the submitted status: ':submitted'
* retrying 'foo*' tasks in 0000Z cycles: 'foo*.0000Z:retrying' or
  '*0000Z/foo*:retrying'
* retrying tasks in 'BAR' family: '*/BAR:retrying' or 'BAR.*:retrying'
* retrying tasks in 'BAR' or 'BAZ' families: '*/BA[RZ]:retrying' or
  'BA[RZ].*:retrying'
```

The old 'MATCH POINT' syntax will be automatically detected and supported. To avoid this, use the '--no-multitask-compat' option, or use the new syntax (with a '/' or a '.') when specifying 2 TASKID arguments.

Arguments:

REG	Suite name
TASKID [...]	Task identifier

Options:

-h, --help	show this help message and exit
--stop-point=CYCLE_POINT, --remove-point=CYCLE_POINT	Optional hold/stop cycle point for inserted task.
--no-check	Add task even if the provided cycle point is not valid for the given task.
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
--port=INT	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
--use-ssh	Use ssh to re-invoke the command on the suite host.
--no-login	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
--print-uuid	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
--set-uuid=UUID	Set the client UUID manually (e.g. from prior use of --print-uuid). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
-f, --force	Do not ask for confirmation before acting. Note that it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.
-m, --family	(Obsolete) This option is now ignored and is retained for backward compatibility only. TASKID in the argument list can be used to match task and family names regardless of this option.
--no-multitask-compat	Disallow backward compatible multitask interface.

F.2.32 jobs-kill

Usage: `cylc [control] jobs-kill JOB-LOG-ROOT [JOB-LOG-DIR ...]`

(This command is for internal use. Users should use "`cylc kill`".) Read job status files to obtain the names of the batch systems and the job IDs in the systems. Invoke the relevant batch system commands to ask the batch systems to terminate the jobs.

Arguments:

JOB-LOG-ROOT	The log/job sub-directory for the suite
[JOB-LOG-DIR ...]	A point/name/submit_num sub-directory

Options:

<code>-h, --help</code>	show this help message and exit
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.

F.2.33 jobs-poll

Usage: `cylc [control] jobs-poll JOB-LOG-ROOT [JOB-LOG-DIR ...]`

(This command is for internal use. Users should use "`cylc poll`".) Read job status files to obtain the statuses of the jobs. If necessary, Invoke the relevant batch system commands to ask the batch systems for more statuses.

Arguments:

JOB-LOG-ROOT	The log/job sub-directory for the suite
[JOB-LOG-DIR ...]	A point/name/submit_num sub-directory

Options:

<code>-h, --help</code>	show this help message and exit
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.

F.2.34 jobs-submit

Usage: `cylc [control] jobs-submit JOB-LOG-ROOT [JOB-LOG-DIR ...]`

(This command is for internal use. Users should use "`cylc submit`".) Submit task jobs to relevant batch systems. On a remote job host, this command reads the job files from STDIN.

Arguments:

JOB-LOG-ROOT	The log/job sub-directory for the suite
[JOB-LOG-DIR ...]	A point/name/submit_num sub-directory

Options:

<code>-h, --help</code>	show this help message and exit
<code>--remote-mode</code>	Is this being run on a remote job host?
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.

F.2.35 jobscrip

Usage: `cylc [prep] jobscrip [OPTIONS] REG TASK`

Generate a task job script and print it to stdout.

Here's how to capture the script in the vim editor:

```
% cylc jobscrip REG TASK | vim -
Emacs unfortunately cannot read from stdin:
% cylc jobscrip REG TASK > tmp.sh; emacs tmp.sh
```

This command wraps 'cylc [control] submit --dry-run'. Other options (e.g. for suite host and owner) are passed through to the submit command.

Options:

```
-h, --help      Print this usage message.
-e --edit       Open the jobscrip in a CLI text editor.
-g --gedit      Open the jobscrip in a GUI text editor.
--plain        Don't print the "Task Job Script Generated message."
(see also 'cylc submit --help')
```

Arguments:

```
REG            Registered suite name.
TASK           Task ID (NAME.CYCLE_POINT)
```

F.2.36 kill

Usage: `cylc [control] kill [OPTIONS] REG [TASKID ...]`

Kill jobs of active tasks and update their statuses accordingly.

To kill one or more tasks, "cylc kill REG TASKID ..."; to kill all active tasks: "cylc kill REG".

TASKID is a pattern to match task proxies or task families, or groups of them:

```
* [CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]
* [CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]
* TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
* FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
```

For example, to match:

```
* all tasks in a cycle: '20200202T0000Z/*' or '*.20200202T0000Z'
* all tasks in the submitted status: ':submitted'
* retrying 'foo*' tasks in 0000Z cycles: 'foo*.0000Z:retrying' or
  '*0000Z/foo*:retrying'
* retrying tasks in 'BAR' family: '*/BAR:retrying' or 'BAR.*:retrying'
* retrying tasks in 'BAR' or 'BAZ' families: '*/BA[RZ]:retrying' or
  'BA[RZ].*:retrying'
```

The old 'MATCH POINT' syntax will be automatically detected and supported. To avoid this, use the '--no-multitask-compat' option, or use the new syntax (with a '/' or a '.') when specifying 2 TASKID arguments.

Arguments:

```
REG            Suite name
[TASKID ...]    Task identifiers
```

Options:

```
-h, --help      show this help message and exit
--user=USER     Other user account name. This results in command
                reinvocation on the remote account.
--host=HOST     Other host name. This results in command reinvocation
                on the remote account.
-v, --verbose   Verbose output mode.
--debug        Run suites in non-daemon mode, and show exception
                tracebacks.
--port=INT     Suite port number on the suite host. NOTE: this is
                retrieved automatically if non-interactive ssh is
                configured to the suite host.
--use-ssh      Use ssh to re-invoke the command on the suite host.
```



```

--no-login          Do not use a login shell to run remote ssh commands.
                   The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC
                   Set a timeout for network connections to the running
                   suite. The default is no timeout. For task messaging
                   connections see site/user config file documentation.
--print-uuid        Print the client UUID to stderr. This can be matched
                   to information logged by the receiving suite server
                   program.
--set-uuid=UUID     Set the client UUID manually (e.g. from prior use of
                   --print-uuid). This can be used to log multiple
                   commands under the same UUID (but note that only the
                   first [info] command from the same client ID will be
                   logged unless the suite is running in debug mode).
-f, --force         Do not ask for confirmation before acting. Note that
                   it is not necessary to use this option if interactive
                   command prompts have been disabled in the site/user
                   config files.
-m, --family        (Obsolete) This option is now ignored and is retained
                   for backward compatibility only. TASKID in the
                   argument list can be used to match task and family
                   names regardless of this option.
--no-multitask-compat
                   Disallow backward compatible multitask interface.

```

F.2.37 list

Usage: `cylc [info|prep] list|ls [OPTIONS] SUITE`

Print runtime namespace names (tasks and families), the first-parent inheritance graph, or actual tasks for a given cycle range.

The first-parent inheritance graph determines the primary task family groupings that are collapsible in gcylc suite views and the graph viewer tool. To visualize the full multiple inheritance hierarchy use: `'cylc graph -n'`.

Arguments:
 SUITE Suite name or path

Options:

```

-h, --help          show this help message and exit
-a, --all-tasks     Print all tasks, not just those used in the graph.
-n, --all-namespaces Print all runtime namespaces, not just tasks.
-m, --mro           Print the linear "method resolution order" for each
                   namespace (the multiple-inheritance precedence order
                   as determined by the C3 linearization algorithm).
-t, --tree          Print the first-parent inheritance hierarchy in tree
                   form.
-b, --box           With -t/--tree, using unicode box characters. Your
                   terminal must be able to display unicode characters.
-w, --with-titles   Print namespaces titles too.
-p START[,STOP], --points=START[,STOP]
                   Print actual task IDs from the START [through STOP]
                   cycle points.
--user=USER         Other user account name. This results in command
                   reinvocation on the remote account.
--host=HOST         Other host name. This results in command reinvocation
                   on the remote account.
-v, --verbose       Verbose output mode.
--debug            Run suites in non-daemon mode, and show exception
                   tracebacks.
--suite-owner=OWNER Specify suite owner
-s NAME=VALUE, --set=NAME=VALUE
                   Set the value of a Jinja2 template variable in the
                   suite definition. This option can be used multiple
                   times on the command line. NOTE: these settings
                   persist across suite restarts, but can be set again on
                   the "cylc restart" command line if they need to be
                   overridden.
--set-file=FILE     Set the value of Jinja2 template variables in the

```



```

suite definition from a file containing NAME=VALUE
pairs (one per line). NOTE: these settings persist
across suite restarts, but can be set again on the
"cycle restart" command line if they need to be
overridden.
--icp=CYCLE_POINT Set initial cycle point. Required if not defined in
suite.rc.
```

F.2.38 ls-checkpoints

Usage: `cylc [info] ls-checkpoints [OPTIONS] REG [ID ...]`

In the absence of arguments and the `--all` option, list checkpoint IDs, their time and events. Otherwise, display the latest and/or the checkpoints of suite parameters, task pool and broadcast states in the suite runtime database.

Arguments:

```

REG Suite name
[ID ...] Checkpoint ID (default=latest)
```

Options:

```

-h, --help show this help message and exit
-a, --all Display data of all available checkpoints.
--user=USER Other user account name. This results in command reinvocation
on the remote account.
--host=HOST Other host name. This results in command reinvocation on the
remote account.
-v, --verbose Verbose output mode.
--debug Run suites in non-daemon mode, and show exception tracebacks.
```

F.2.39 message

Usage: `cylc [task] message [OPTIONS] MESSAGE ...`

This command is used by task jobs to automatically report success and failure.

It can also be used to send info-, warning-, or critical-severity messages back, and to report registered task "message outputs" completed.

Note: to abort a job script with a custom error message, use `cylc__job_abort`:
`cylc__job_abort 'message...'`
(for technical reasons this is a shell function, not a cylc sub-command).

Options:

```

-h, --help show this help message and exit
-s SEVERITY, -p SEVERITY, --severity=SEVERITY, --priority=SEVERITY
message severity: NORMAL (default), WARNING, CRITICAL
or CUSTOM
--env=ENV Override environment variables
-v, --verbose Verbose output mode
```

F.2.40 monitor

Usage: `cylc [info] monitor [OPTIONS] REG [USER_AT_HOST]`

A terminal-based live suite monitor. Exit with 'Ctrl-C'.

The USER_AT_HOST argument allows suite selection by 'cylc scan' output:
`cylc monitor $(cylc scan | grep <suite_name>)`

Arguments:

```

REG Suite name
[USER_AT_HOST] user@host:port, shorthand for --user, --host &
--port.
```

Options:

```

-h, --help show this help message and exit
```

```

-a, --align           Align task names. Only useful for small suites.
-r, --restricted      Restrict display to active task states. This may be
                      useful for monitoring very large suites. The state
                      summary line still reflects all task proxies.
-s ORDER, --sort=ORDER
                      Task sort order: "definition" or "alphanumeric". The
                      default is definition order, as determined by global
                      config. (Definition order is the order that tasks
                      appear under [runtime] in the suite definition).
-o, --once            Show a single view then exit.
-u, --runahead        Display task proxies in the runahead pool (off by
                      default).
-i SECONDS, --interval=SECONDS
                      Interval between suite state retrievals, in seconds
                      (default 1).
--user=USER           Other user account name. This results in command
                      reinvocation on the remote account.
--host=HOST           Other host name. This results in command reinvocation
                      on the remote account.
-v, --verbose         Verbose output mode.
--debug              Run suites in non-daemon mode, and show exception
                      tracebacks.
--port=INT            Suite port number on the suite host. NOTE: this is
                      retrieved automatically if non-interactive ssh is
                      configured to the suite host.
--use-ssh             Use ssh to re-invoke the command on the suite host.
--no-login            Do not use a login shell to run remote ssh commands.
                      The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC
                      Set a timeout for network connections to the running
                      suite. The default is no timeout. For task messaging
                      connections see site/user config file documentation.
--print-uuid          Print the client UUID to stderr. This can be matched
                      to information logged by the receiving suite server
                      program.
--set-uuid=UUID       Set the client UUID manually (e.g. from prior use of
                      --print-uuid). This can be used to log multiple
                      commands under the same UUID (but note that only the
                      first [info] command from the same client ID will be
                      logged unless the suite is running in debug mode).

```

F.2.41 nudge

Usage: `cylc [control] nudge [OPTIONS] REG`

Cause the cylc task processing loop to be invoked in a running suite.

This happens automatically when the state of any task changes such that task processing (dependency negotiation etc.) is required, or if a clock-trigger task is ready to run.

The main reason to use this command is to update the "estimated time till completion" intervals shown in the tree-view suite control GUI, during periods when nothing else is happening.

Arguments:

REG Suite name

Options:

```

-h, --help           show this help message and exit
--user=USER          Other user account name. This results in command
                      reinvocation on the remote account.
--host=HOST          Other host name. This results in command reinvocation
                      on the remote account.
-v, --verbose         Verbose output mode.
--debug              Run suites in non-daemon mode, and show exception
                      tracebacks.
--port=INT           Suite port number on the suite host. NOTE: this is
                      retrieved automatically if non-interactive ssh is
                      configured to the suite host.

```

```

--use-ssh           Use ssh to re-invoke the command on the suite host.
--no-login          Do not use a login shell to run remote ssh commands.
                    The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC
                    Set a timeout for network connections to the running
                    suite. The default is no timeout. For task messaging
                    connections see site/user config file documentation.
--print-uuid        Print the client UUID to stderr. This can be matched
                    to information logged by the receiving suite server
                    program.
--set-uuid=UUID     Set the client UUID manually (e.g. from prior use of
                    --print-uuid). This can be used to log multiple
                    commands under the same UUID (but note that only the
                    first [info] command from the same client ID will be
                    logged unless the suite is running in debug mode).
-f, --force         Do not ask for confirmation before acting. Note that
                    it is not necessary to use this option if interactive
                    command prompts have been disabled in the site/user
                    config files.

```

F.2.42 ping

Usage: `cylc [discovery] ping [OPTIONS] REG [TASK]`

If suite REG is running or TASK in suite REG is currently running,
exit with success status, else exit with error status.

Arguments:

REG	Suite name
[TASK]	Task NAME.CYCLE_POINT

Options:

```

-h, --help          show this help message and exit
--print-ports        Print the port range from the cylc site config file.
--user=USER          Other user account name. This results in command
                    reinvocation on the remote account.
--host=HOST          Other host name. This results in command reinvocation
                    on the remote account.
-v, --verbose        Verbose output mode.
--debug              Run suites in non-daemon mode, and show exception
                    tracebacks.
--port=INT           Suite port number on the suite host. NOTE: this is
                    retrieved automatically if non-interactive ssh is
                    configured to the suite host.
--use-ssh            Use ssh to re-invoke the command on the suite host.
--no-login           Do not use a login shell to run remote ssh commands.
                    The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC
                    Set a timeout for network connections to the running
                    suite. The default is no timeout. For task messaging
                    connections see site/user config file documentation.
--print-uuid        Print the client UUID to stderr. This can be matched
                    to information logged by the receiving suite server
                    program.
--set-uuid=UUID     Set the client UUID manually (e.g. from prior use of
                    --print-uuid). This can be used to log multiple
                    commands under the same UUID (but note that only the
                    first [info] command from the same client ID will be
                    logged unless the suite is running in debug mode).
-f, --force         Do not ask for confirmation before acting. Note that
                    it is not necessary to use this option if interactive
                    command prompts have been disabled in the site/user
                    config files.

```

F.2.43 poll

Usage: `cylc [control] poll [OPTIONS] REG [TASKID ...]`

Poll (query) task jobs to verify and update their statuses.

Use `"cylc poll REG"` to poll all active tasks, or `"cylc poll REG TASKID"` to poll individual tasks or families, or groups of them.

TASKID is a pattern to match task proxies or task families, or groups of them:

```
* [CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]
* [CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]
* TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
* FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
```

For example, to match:

```
* all tasks in a cycle: '20200202T0000Z/*' or '*.20200202T0000Z'
* all tasks in the submitted status: ':submitted'
* retrying 'foo*' tasks in 0000Z cycles: 'foo*.*0000Z:retrying' or
  '*0000Z/foo*:retrying'
* retrying tasks in 'BAR' family: '*/BAR:retrying' or 'BAR.*:retrying'
* retrying tasks in 'BAR' or 'BAZ' families: '*/BA[RZ]:retrying' or
  'BA[RZ].*:retrying'
```

The old 'MATCH POINT' syntax will be automatically detected and supported. To avoid this, use the `'--no-multitask-compat'` option, or use the new syntax (with a `'/'` or a `'.'`) when specifying 2 TASKID arguments.

Arguments:

REG	Suite name
[TASKID ...]	Task identifiers

Options:

<code>-h, --help</code>	show this help message and exit
<code>-s, --succeeded</code>	Allow polling of succeeded tasks.
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.
<code>--port=INT</code>	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
<code>--use-ssh</code>	Use ssh to re-invoke the command on the suite host.
<code>--no-login</code>	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
<code>--comms-timeout=SEC, --pyro-timeout=SEC</code>	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
<code>--print-uuid</code>	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
<code>--set-uuid=UUID</code>	Set the client UUID manually (e.g. from prior use of <code>--print-uuid</code>). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
<code>-f, --force</code>	Do not ask for confirmation before acting. Note that it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.
<code>-m, --family</code>	(Obsolete) This option is now ignored and is retained for backward compatibility only. TASKID in the argument list can be used to match task and family names regardless of this option.
<code>--no-multitask-compat</code>	Disallow backward compatible multitask interface.

F.2.44 print

Usage: `cylc [prep] print [OPTIONS] [REGEX]`

Print registered (installed) suites.

Note on result filtering:

(a) The filter patterns are Regular Expressions, not shell globs, so the general wildcard is `'.*'` (match zero or more of anything), NOT `'*'`.

(b) For printing purposes there is an implicit wildcard at the end of each pattern (`'foo'` is the same as `'foo/*'`); use the string end marker to prevent this (`'foo$'` matches only literal `'foo'`).

Arguments:

[REGEX] Suite name regular expression pattern

Options:

```
-h, --help      show this help message and exit
-t, --tree      Print suites in nested tree form.
-b, --box       Use unicode box drawing characters in tree views.
-a, --align     Align columns.
-x             don't print suite definition directory paths.
-y             Don't print suite titles.
--fail         Fail (exit 1) if no matching suites are found.
--user=USER    Other user account name. This results in command reinvocation
               on the remote account.
--host=HOST    Other host name. This results in command reinvocation on the
               remote account.
-v, --verbose  Verbose output mode.
--debug        Run suites in non-daemon mode, and show exception tracebacks.
```

F.2.45 profile-battery

Usage: `cylc profile-battery [-e [EXPERIMENT ...]] [-v [VERSION ...]]`

Run profiling experiments against different versions of cylc. A list of experiments can be specified after the `-e` flag, if not provided the experiment `"complex"` will be chosen. A list of versions to profile against can be specified after the `-v` flag, if not provided the current version will be used.

Experiments are stored in `dev/profile-experiments`, user experiments can be stored in `.profiling/experiments`. Experiments are specified without the file extension, experiments in `.profiling/` will be chosen before those in `dev/`.

IMPORTANT: See `dev/profile-experiments/example` for an experiment template with further details.

Versions are any valid git identifiers i.e. tags, branches, commits. To compare results to different cylc versions either:

- * Supply cylc profile-battery with a complete list of the versions you wish to profile, it will then provide the option to checkout the required versions automatically.
- * Checkout each version manually running cylc profile-battery against only one version at a time. Once all results have been gathered you can then run cylc profile-battery with a complete list of versions.

Profiling will save results to `.profiling/results.json` where they can be used for future comparisons. To list profiling results run:

```
* cylc profile-battery --ls # list all results
* cylc profile-battery --ls -e experiment # list all results for
                                         # experiment "experiment".
* cylc profile-battery --ls --delete -v 6.1.2 # Delete all results for
                                              # version 6.1.2 (prompted).
```

If matplotlib and numpy are installed profiling generates plots which are saved to `.profiling/plots` or presented in an interactive window using the `-i` flag.

Results are stored along with a checksum for the experiment file. When an experiment file is changed previous results are maintained, future results will be stored separately. To copy results from an older version of an experiment into those from the current one run:

```
* cylc profile-battery --promote experiment@checksum
```

NOTE: At present results cannot be analysed without the experiment file so old results must be `"copied"` in this way to be re-used.

The results output contain only a small number of metrics, to see a full list

of results use the `--full` option.

Options:

```
-h, --help          show this help message and exit
-e, --experiments   Specify list of experiments to run.
-v, --versions       Specify cylc versions to profile. Git tags, branches,
                     commits are all valid.
-i, --interactive    Open any plots in interactive window rather saving
                     them to files.
-p, --no-plots       Don't generate any plots.
--ls, --list-results List all stored results. Experiments and versions to
                     list can be specified using --experiments and
                     --versions.
--delete            Delete stored results (to be used in combination with
                     --list-results).
-y, --yes            Answer yes to any user input. Will check-out cylc
                     versions as required.
--full-results, --full Display all gathered metrics.
--lobf-order=LOBF_ORDER
                     The order (int) of the line of best fit to be drawn. 0
                     for no lobf, 1 for linear, 2 for quadratic ect.
--promote=PROMOTE    Promote results from an older version of an experiment
                     to the current version. To be used when making non-
                     functional changes to an experiment.
--test              For development purposes, run experiment without
                     saving results and regardless of any prior runs.
```

F.2.46 register

Usage: `cylc [prep] register [OPTIONS] REG [PATH]`

Register the suite definition located in `PATH` (or `$PWD`) as `REG`.

This creates the suite run directory, and authentication files in a sub-directory called `".service/"`.

Suite names are the same as the directory path under the suite run directory. They may contain alphanumeric characters plus `'_'` `'-'` and `'/'`.

Example: if the cylc run directory is `$HOME/cylc-run` (the default) and `/home/bob/suites/test` is a suite source directory, then:

```
% cylc reg nwp/test1 /home/bob/suites/test
```

will create the following suite run directory:

```
/home/bob/cylc-run/nwp/test1
'-- .service
   |-- passphrase
   |-- source -> /home/bob/test
   |-- ssl.cert
   '-- ssl.pem
```

The suite can subsequently be started and targeted by cylc commands using `"nwp/test1"` as the name.

Arguments:

```
REG          Suite name
[PATH]       Suite definition directory (defaults to $PWD)
```

Options:

```
-h, --help          show this help message and exit
--user=USER         Other user account name. This results in command reinvocation
                     on the remote account.
--host=HOST         Other host name. This results in command reinvocation on the
                     remote account.
-v, --verbose       Verbose output mode.
--debug            Run suites in non-daemon mode, and show exception tracebacks.
```

F.2.47 release

Usage: `cylc [control] release|unhold [OPTIONS] REG [TASKID ...]`

Release one or more held tasks (`cylc release REG TASKID`) or the whole suite (`cylc release REG`). Held tasks do not submit even if they are ready to run.

See also '`cylc [control] hold`'.

TASKID is a pattern to match task proxies or task families, or groups of them:

- * [CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]
- * [CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]
- * TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
- * FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]

For example, to match:

- * all tasks in a cycle: `'20200202T0000Z/*'` or `'*.20200202T0000Z'`
- * all tasks in the submitted status: `':submitted'`
- * retrying 'foo*' tasks in 0000Z cycles: `'foo*.*0000Z:retrying'` or `'*0000Z/foo*:retrying'`
- * retrying tasks in 'BAR' family: `'*/BAR:retrying'` or `'BAR.*:retrying'`
- * retrying tasks in 'BAR' or 'BAZ' families: `'*/BA[RZ]:retrying'` or `'BA[RZ].*:retrying'`

The old 'MATCH POINT' syntax will be automatically detected and supported. To avoid this, use the '`--no-multitask-compat`' option, or use the new syntax (with a '/' or a '.') when specifying 2 TASKID arguments.

Arguments:

REG	Suite name
[TASKID ...]	Task identifiers

Options:

<code>-h, --help</code>	show this help message and exit
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.
<code>--port=INT</code>	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
<code>--use-ssh</code>	Use ssh to re-invoke the command on the suite host.
<code>--no-login</code>	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
<code>--comms-timeout=SEC, --pyro-timeout=SEC</code>	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
<code>--print-uuid</code>	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
<code>--set-uuid=UUID</code>	Set the client UUID manually (e.g. from prior use of <code>--print-uuid</code>). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
<code>-f, --force</code>	Do not ask for confirmation before acting. Note that it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.
<code>-m, --family</code>	(Obsolete) This option is now ignored and is retained for backward compatibility only. TASKID in the argument list can be used to match task and family names regardless of this option.
<code>--no-multitask-compat</code>	Disallow backward compatible multitask interface.

F.2.48 reload

Usage: `cylc [control] reload [OPTIONS] REG`

Tell a suite to reload its definition at run time. All settings including task definitions, with the exception of suite log configuration, can be changed on reload. Note that defined tasks can be added to or removed from a running suite with the 'cylc insert' and 'cylc remove' commands, without reloading. This command also allows addition and removal of actual task definitions, and therefore insertion of tasks that were not defined at all when the suite started (you will still need to manually insert a particular instance of a newly defined task). Live task proxies that are orphaned by a reload (i.e. their task definitions have been removed) will be removed from the task pool if they have not started running yet. Changes to task definitions take effect immediately, unless a task is already running at reload time.

If the suite was started with Jinja2 template variables set on the command line (`cylc run --set FOO=bar REG`) the same template settings apply to the reload (only changes to the suite.rc file itself are reloaded).

If the modified suite definition does not parse, failure to reload will be reported but no harm will be done to the running suite.

Arguments:

REG Suite name

Options:

<code>-h, --help</code>	show this help message and exit
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.
<code>--port=INT</code>	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
<code>--use-ssh</code>	Use ssh to re-invoke the command on the suite host.
<code>--no-login</code>	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
<code>--comms-timeout=SEC, --pyro-timeout=SEC</code>	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
<code>--print-uuid</code>	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
<code>--set-uuid=UUID</code>	Set the client UUID manually (e.g. from prior use of <code>--print-uuid</code>). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
<code>-f, --force</code>	Do not ask for confirmation before acting. Note that it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.

F.2.49 remote-init

Usage: `cylc [task] remote-init UUID RUND`

(This command is for internal use.)

Install suite service files on a task remote (i.e. a [owner@]host):

```
* ".service/contact"
* ".service/passphrase"
* ".service/ssl.cert"
```

Content of items to install from a tar file read from STDIN.

Return 0 on success or if initialisation not required. Return 1 on failure.

Print SuiteSrvFileManager.REMOTE_INIT_NOT_REQUIRED if initialisation not required (e.g. remote has shared file system with suite host).

Print SuiteSrvFileManager.REMOTE_INIT_DONE on success.

Arguments:

UUID	UUID of current suite running process
RUND	The run directory of the suite

Options:

-h, --help	show this help message and exit
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.

F.2.50 remote-tidy

Usage: cylc [task] remote-tidy RUND

(This command is for internal use.)

Remove ".service/contact" from a task remote (i.e. a [owner@]host).

Remove ".service" directory on the remote if emptied.

Arguments:

RUND	The run directory of the suite
------	--------------------------------

Options:

-h, --help	show this help message and exit
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.

F.2.51 remove

Usage: cylc [control] remove [OPTIONS] REG TASKID [...]

Remove one or more tasks (cylc remove REG TASKID), or all tasks with a given cycle point (cylc remove REG *.POINT) from a running suite.

Tasks will spawn successors first if they have not done so already.

TASKID is a pattern to match task proxies or task families, or groups of them:

```
* [CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]
* [CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]
* TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
* FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
```

For example, to match:

```
* all tasks in a cycle: '20200202T0000Z/*' or '*.20200202T0000Z'
* all tasks in the submitted status: ':submitted'
* retrying 'foo*' tasks in 0000Z cycles: 'foo*.0000Z:retrying' or
  '*0000Z/foo*:retrying'
* retrying tasks in 'BAR' family: '*/BAR:retrying' or 'BAR.*:retrying'
* retrying tasks in 'BAR' or 'BAZ' families: '*/BA[RZ]:retrying' or
  'BA[RZ].*:retrying'
```

The old 'MATCH POINT' syntax will be automatically detected and supported. To

avoid this, use the '--no-multitask-compat' option, or use the new syntax (with a '/' or a '.') when specifying 2 TASKID arguments.

Arguments:

REG	Suite name
TASKID [...]	Task identifiers

Options:

-h, --help	show this help message and exit
--no-spawn	Do not spawn successors before removal.
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
--port=INT	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
--use-ssh	Use ssh to re-invoke the command on the suite host.
--no-login	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
--print-uuid	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
--set-uuid=UUID	Set the client UUID manually (e.g. from prior use of --print-uuid). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
-f, --force	Do not ask for confirmation before acting. Note that it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.
-m, --family	(Obsolete) This option is now ignored and is retained for backward compatibility only. TASKID in the argument list can be used to match task and family names regardless of this option.
--no-multitask-compat	Disallow backward compatible multitask interface.

F.2.52 report-timings

Usage: `cylc [util] report-timings [OPTIONS] REG`

Retrieve suite timing information for wait and run time performance analysis. Raw output and summary output (in text or HTML format) are available. Output is sent to standard output, unless an output filename is supplied.

Summary Output (the default):

Data stratified by host and batch system that provides a statistical summary of

1. Queue wait time (duration between task submission and start times)
2. Task run time (duration between start and succeed times)
3. Total run time (duration between task submission and succeed times)

Summary tables can be output in plain text format, or HTML with embedded SVG boxplots. Both summary options require the Pandas library, and the HTML summary option requires the Matplotlib library.

Raw Output:

A flat list of tabular data that provides (for each task and cycle) the

1. Time of successful submission
2. Time of task start
3. Time of task successful completion

as well as information about the batch system and remote host to permit stratification/grouping if desired by downstream processors.

Timings are shown only for succeeded tasks.

For long-running and/or large suites (i.e. for suites with many task events), the database query to obtain the timing information may take some time.

Arguments:

REG Suite name

Options:

-h, --help show this help message and exit
 -r, --raw Show raw timing output suitable for custom diagnostics.
 -s, --summary Show textual summary timing output for tasks.
 -w, --web-summary Show HTML summary timing output for tasks.
 -O OUTPUT_FILENAME, --output-file=OUTPUT_FILENAME Output to a specific file
 --user=USER Other user account name. This results in command reinvocation on the remote account.
 --host=HOST Other host name. This results in command reinvocation on the remote account.
 -v, --verbose Verbose output mode.
 --debug Run suites in non-daemon mode, and show exception tracebacks.

F.2.53 reset

Usage: `cylc [control] reset [OPTIONS] REG [TASKID ...]`

Force one or more task proxies in a running suite to change state and modify their prerequisites and outputs accordingly. For example, `--state=waiting` means "prerequisites not satisfied, outputs not completed"; `--state=ready` means "prerequisites satisfied, outputs not completed" (this generally has the same effect as using the `"cylc trigger"` command).

`"cylc reset --state=spawn"` is deprecated: use `"cylc spawn"` instead.

See the documentation for the `-s/--state` option for legal reset states. TASKID is a pattern to match task proxies or task families, or groups of them:

- * [CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]
- * [CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]
- * TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
- * FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]

For example, to match:

- * all tasks in a cycle: `'20200202T0000Z/*'` or `'*.20200202T0000Z'`
- * all tasks in the submitted status: `':submitted'`
- * retrying 'foo' tasks in 0000Z cycles: `'foo*.0000Z:retrying'` or `'*0000Z/foo*:retrying'`
- * retrying tasks in 'BAR' family: `'*/BAR:retrying'` or `'BAR.*:retrying'`
- * retrying tasks in 'BAR' or 'BAZ' families: `'*/BA[RZ]:retrying'` or `'BA[RZ].*:retrying'`

The old 'MATCH POINT' syntax will be automatically detected and supported. To avoid this, use the `'--no-multitask-compat'` option, or use the new syntax (with a `'/'` or a `'.'`) when specifying 2 TASKID arguments.

Arguments:

REG Suite name
 [TASKID ...] Task identifiers

Options:

-h, --help show this help message and exit
 -s STATE, --state=STATE Reset task state to STATE, can be succeeded, failed, waiting, submitted, held, running, submit-failed, ready, expired
 -O OUTPUT, --output=OUTPUT Find task output by message string or trigger string,

```

--user=USER          set complete or incomplete with !OUTPUT, '*' to set
                    all complete, '!*' to set all incomplete. Can be used
                    more than once to reset multiple task outputs.
--host=HOST          Other user account name. This results in command
                    reinvocation on the remote account.
--v, --verbose       Other host name. This results in command reinvocation
                    on the remote account.
--debug             Verbose output mode.
                    Run suites in non-daemon mode, and show exception
                    tracebacks.
--port=INT           Suite port number on the suite host. NOTE: this is
                    retrieved automatically if non-interactive ssh is
                    configured to the suite host.
--use-ssh            Use ssh to re-invoke the command on the suite host.
--no-login           Do not use a login shell to run remote ssh commands.
                    The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC
                    Set a timeout for network connections to the running
                    suite. The default is no timeout. For task messaging
                    connections see site/user config file documentation.
--print-uuid         Print the client UUID to stderr. This can be matched
                    to information logged by the receiving suite server
                    program.
--set-uuid=UUID      Set the client UUID manually (e.g. from prior use of
                    --print-uuid). This can be used to log multiple
                    commands under the same UUID (but note that only the
                    first [info] command from the same client ID will be
                    logged unless the suite is running in debug mode).
-f, --force          Do not ask for confirmation before acting. Note that
                    it is not necessary to use this option if interactive
                    command prompts have been disabled in the site/user
                    config files.
-m, --family         (Obsolete) This option is now ignored and is retained
                    for backward compatibility only. TASKID in the
                    argument list can be used to match task and family
                    names regardless of this option.
--no-multitask-compat
                    Disallow backward compatible multitask interface.

```

F.2.54 restart

Usage: `cylc [control] restart [OPTIONS] REG`

Start a suite run from the previous state. To start from scratch (cold or warm start) see the 'cylc run' command.

The scheduler runs as a daemon unless you specify `--no-detach`.

Tasks recorded as submitted or running are polled at start-up to determine what happened to them while the suite was down.

Arguments:

REG Suite name

Options:

```

-h, --help            show this help message and exit
--non-daemon          (deprecated: use --no-detach)
-n, --no-detach       Do not daemonize the suite
-a, --no-auto-shutdown
                    Do not shut down the suite automatically when all
                    tasks have finished. This flag overrides the
                    corresponding suite config item.
--profile             Output profiling (performance) information
--checkpoint=CHECKPOINT-ID
                    Specify the ID of a checkpoint to restart from
--ignore-final-cycle-point
                    Ignore the final cycle point in the suite run
                    database. If one is specified in the suite definition
                    it will be used, however.
--ignore-initial-cycle-point
                    Ignore the initial cycle point in the suite run

```

```

                                database. If one is specified in the suite definition
                                it will be used, however.
--until=CYCLE_POINT           Shut down after all tasks have PASSED this cycle
                                point.
--hold                        Hold (don't run tasks) immediately on starting.
--hold-after=CYCLE_POINT      Hold (don't run tasks) AFTER this cycle point.
-m STRING, --mode=STRING      Run mode: live, dummy, dummy-local, simulation
                                (default live).
--reference-log               Generate a reference log for use in reference tests.
--reference-test               Do a test run against a previously generated reference
                                log.
-S SOURCE, --source=SOURCE    Specify the suite source.
--user=USER                   Other user account name. This results in command
                                reinvocation on the remote account.
--host=HOST                   Other host name. This results in command reinvocation
                                on the remote account.
-v, --verbose                 Verbose output mode.
--debug                       Run suites in non-daemon mode, and show exception
                                tracebacks.
-s NAME=VALUE, --set=NAME=VALUE
                                Set the value of a Jinja2 template variable in the
                                suite definition. This option can be used multiple
                                times on the command line. NOTE: these settings
                                persist across suite restarts, but can be set again on
                                the "cylc restart" command line if they need to be
                                overridden.
--set-file=FILE               Set the value of Jinja2 template variables in the
                                suite definition from a file containing NAME=VALUE
                                pairs (one per line). NOTE: these settings persist
                                across suite restarts, but can be set again on the
                                "cylc restart" command line if they need to be
                                overridden.

```

F.2.55 run

Usage: `cylc [control] run|start [OPTIONS] REG [START_POINT]`

Start a suite run from scratch, wiping out any previous suite state. To restart from a previous state see 'cylc restart --help'.

The scheduler runs as a daemon unless you specify `--no-detach`.

Any dependence on cycle points earlier than the start cycle point is ignored.

A "cold start" (the default) starts from the suite initial cycle point (specified in the suite.rc or on the command line). Any dependence on tasks prior to the suite initial cycle point is ignored.

A "warm start" (`-w/--warm`) starts from a given cycle point later than the suite initial cycle point (specified in the suite.rc). Any dependence on tasks prior to the given warm start cycle point is ignored. The suite initial cycle point is preserved.

Arguments:

REG	Suite name
[START_POINT]	Initial cycle point or 'now'; overrides the suite definition.

Options:

```

-h, --help                show this help message and exit
--non-daemon              (deprecated: use --no-detach)
-n, --no-detach           Do not daemonize the suite
-a, --no-auto-shutdown    Do not shut down the suite automatically when all
                                tasks have finished. This flag overrides the
                                corresponding suite config item.
--profile                 Output profiling (performance) information
-w, --warm                 Warm start the suite. The default is to cold start.

```

```

--ict                Does nothing, option for backward compatibility only
--until=CYCLE_POINT  Shut down after all tasks have PASSED this cycle
                    point.
--hold              Hold (don't run tasks) immediately on starting.
--hold-after=CYCLE_POINT
                    Hold (don't run tasks) AFTER this cycle point.
-m STRING, --mode=STRING
                    Run mode: live, dummy, dummy-local, simulation
                    (default live).
--reference-log      Generate a reference log for use in reference tests.
--reference-test     Do a test run against a previously generated reference
                    log.
-S SOURCE, --source=SOURCE
                    Specify the suite source.
--user=USER          Other user account name. This results in command
                    reinvocation on the remote account.
--host=HOST          Other host name. This results in command reinvocation
                    on the remote account.
-v, --verbose        Verbose output mode.
--debug             Run suites in non-daemon mode, and show exception
                    tracebacks.
-s NAME=VALUE, --set=NAME=VALUE
                    Set the value of a Jinja2 template variable in the
                    suite definition. This option can be used multiple
                    times on the command line. NOTE: these settings
                    persist across suite restarts, but can be set again on
                    the "cylc restart" command line if they need to be
                    overridden.
--set-file=FILE      Set the value of Jinja2 template variables in the
                    suite definition from a file containing NAME=VALUE
                    pairs (one per line). NOTE: these settings persist
                    across suite restarts, but can be set again on the
                    "cylc restart" command line if they need to be
                    overridden.

```

F.2.56 scan

Usage: `cylc [discovery] scan [OPTIONS] [HOSTS ...]`

Print information about running suites.

By default, it will obtain a listing of running suites for the current user from the file system, before connecting to the suites to obtain information. Use the `-o/--suite-owner` option to get information of running suites for other users.

If a list of HOSTS is specified, it will obtain a listing of running suites by scanning all ports in the relevant range for running suites on the specified hosts. If the `-a/--all` option is specified, it will use the global configuration `"[suite host scanning]"` setting to determine a list of hosts to scan.

Suite passphrases are not needed to get identity information (name and owner). Titles, descriptions, state totals, and cycle point state totals may also be revealed publicly, depending on global and suite authentication settings. Suite passphrases still grant full access regardless of what is revealed publicly.

WARNING: a suite suspended with Ctrl-Z will cause port scans to hang until the connection times out (see `--comms-timeout`).

Arguments:

[HOSTS ...] Hosts to scan instead of the configured hosts.

Options:

```

-h, --help            show this help message and exit
-a, --all             Scan all port ranges in known hosts.
-n PATTERN, --name=PATTERN
                    List suites with name matching PATTERN (regular
                    expression). Defaults to any name. Can be used
                    multiple times.
-o PATTERN, --suite-owner=PATTERN

```

```

List suites with owner matching PATTERN (regular
expression). Defaults to just your own suites. Can be
used multiple times.
-d, --describe      Print suite metadata if available.
-s, --state-totals  Print number of tasks in each state if available
                    (total, and by cycle point).
-f, --full          Print all available information about each suite.
-c, --color, --colour
                    Print task state summaries using terminal color
                    control codes.
-b, --no-bold       Don't use any bold text in the command output.
--print-ports       Print the port range from the site config file
                    ($CYLC_DIR/conf/global.rc).
--comms-timeout=SEC Set a timeout for network connections to each running
                    suite. The default is 5 seconds.
--old, --old-format Legacy output format ("suite owner host port").
-r, --raw, --raw-format
                    Parsable format ("suite|owner|host|property|value").
-j, --json, --json-format
                    JSON format.
--user=USER         Other user account name. This results in command
                    reinvocation on the remote account.
--host=HOST         Other host name. This results in command reinvocation
                    on the remote account.
-v, --verbose       Verbose output mode.
--debug            Run suites in non-daemon mode, and show exception
                    tracebacks.
--port=INT         Suite port number on the suite host. NOTE: this is
                    retrieved automatically if non-interactive ssh is
                    configured to the suite host.
--use-ssh          Use ssh to re-invoke the command on the suite host.
--no-login         Do not use a login shell to run remote ssh commands.
                    The default is to use a login shell.
--print-uuid       Print the client UUID to stderr. This can be matched
                    to information logged by the receiving suite server
                    program.
--set-uuid=UUID    Set the client UUID manually (e.g. from prior use of
                    --print-uuid). This can be used to log multiple
                    commands under the same UUID (but note that only the
                    first [info] command from the same client ID will be
                    logged unless the suite is running in debug mode).

```

F.2.57 scp-transfer

Usage: `cylc [util] scp-transfer [OPTIONS]`

An scp wrapper for transferring a list of files and/or directories at once. The source and target scp URLs can be local or remote (scp can transfer files between two remote hosts). Passwordless ssh must be configured appropriately.

ENVIRONMENT VARIABLE INPUTS:

`$SRCE` - list of sources (files or directories) as scp URLs.

`$DEST` - parallel list of targets as scp URLs.

The source and destination lists should be space-separated.

We let scp determine the validity of source and target URLs. Target directories are created pre-copy if they don't exist.

Options:

```

-v      - verbose: print scp stdout.
--help  - print this usage message.

```

F.2.58 search

Usage: `cylc [prep] search|grep [OPTIONS] SUITE PATTERN [PATTERN2...]`

Search for pattern matches in suite definitions and any files in the suite bin directory. Matches are reported by line number and suite section. An unquoted list of PATTERNS will be converted to an OR'd

pattern. Note that the order of command line arguments conforms to normal cylc command usage (suite name first) not that of the grep command.

Note that this command performs a text search on the suite definition, it does not search the data structure that results from parsing the suite definition - so it will not report implicit default settings.

For case insensitive matching use '(?i)PATTERN'.

Arguments:

SUITE	Suite name or path
PATTERN	Python-style regular expression
[PATTERN2...]	Additional search patterns

Options:

-h, --help	show this help message and exit
-x	Do not search in the suite bin directory
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
--suite-owner=OWNER	Specify suite owner

F.2.59 set-verbosity

Usage: cylc [control] set-verbosity [OPTIONS] REG LEVEL

Change the logging severity level of a running suite. Only messages at or above the chosen severity level will be logged; for example, if you choose WARNING, only warnings and critical messages will be logged.

Arguments:

REG	Suite name
LEVEL	INFO, WARNING, NORMAL, CRITICAL, ERROR, DEBUG

Options:

-h, --help	show this help message and exit
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
--port=INT	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
--use-ssh	Use ssh to re-invoke the command on the suite host.
--no-login	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
--print-uuid	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
--set-uuid=UUID	Set the client UUID manually (e.g. from prior use of --print-uuid). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
-f, --force	Do not ask for confirmation before acting. Note that it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.

F.2.60 show

Usage: `cylc [info] show [OPTIONS] REG [TASKID ...]`

Interrogate a suite server program for the suite metadata; or for the metadata of one of its tasks; or for the current state of the prerequisites, outputs, and clock-triggering of a specific task instance. TASKID is a pattern to match task proxies or task families, or groups of them:

- * `[CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]`
- * `[CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]`
- * `TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]`
- * `FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]`

For example, to match:

- * all tasks in a cycle: `'20200202T0000Z/*'` or `'*.20200202T0000Z'`
- * all tasks in the submitted status: `':submitted'`
- * retrying 'foo*' tasks in 0000Z cycles: `'foo*.*0000Z:retrying'` or `'*0000Z/foo*:retrying'`
- * retrying tasks in 'BAR' family: `'*/BAR:retrying'` or `'BAR.*:retrying'`
- * retrying tasks in 'BAR' or 'BAZ' families: `'*/BA[RZ]:retrying'` or `'BA[RZ].*:retrying'`

The old 'MATCH POINT' syntax will be automatically detected and supported. To avoid this, use the `'--no-multitask-compat'` option, or use the new syntax (with a '/' or a '.') when specifying 2 TASKID arguments.

Arguments:

REG	Suite name
[TASKID ...]	Task names or identifiers

Options:

<code>-h, --help</code>	show this help message and exit
<code>--list-prereqs</code>	Print a task's pre-requisites as a list.
<code>--json</code>	Print output in JSON format.
<code>--user=USER</code>	Other user account name. This results in command reinvocation on the remote account.
<code>--host=HOST</code>	Other host name. This results in command reinvocation on the remote account.
<code>-v, --verbose</code>	Verbose output mode.
<code>--debug</code>	Run suites in non-daemon mode, and show exception tracebacks.
<code>--port=INT</code>	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
<code>--use-ssh</code>	Use ssh to re-invoke the command on the suite host.
<code>--no-login</code>	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
<code>--comms-timeout=SEC, --pyro-timeout=SEC</code>	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
<code>--print-uuid</code>	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
<code>--set-uuid=UUID</code>	Set the client UUID manually (e.g. from prior use of <code>--print-uuid</code>). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
<code>-m, --family</code>	(Obsolete) This option is now ignored and is retained for backward compatibility only. TASKID in the argument list can be used to match task and family names regardless of this option.
<code>--no-multitask-compat</code>	Disallow backward compatible multitask interface.

F.2.61 spawn

Usage: `cylc [control] spawn [OPTIONS] REG [TASKID ...]`

Force one or more task proxies to spawn successors at the next cycle point

in their sequences. This is useful if you need to run successive instances of a task out of order.

TASKID is a pattern to match task proxies or task families, or groups of them:

```
* [CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]
* [CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]
* TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
* FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
```

For example, to match:

```
* all tasks in a cycle: '20200202T0000Z/*' or '*.20200202T0000Z'
* all tasks in the submitted status: ':submitted'
* retrying 'foo*' tasks in 0000Z cycles: 'foo*.*0000Z:retrying' or
  '*0000Z/foo*:retrying'
* retrying tasks in 'BAR' family: '*/BAR:retrying' or 'BAR.*:retrying'
* retrying tasks in 'BAR' or 'BAZ' families: '*/BA[RZ]:retrying' or
  'BA[RZ].*:retrying'
```

The old 'MATCH POINT' syntax will be automatically detected and supported. To avoid this, use the '--no-multitask-compat' option, or use the new syntax (with a '/' or a '.') when specifying 2 TASKID arguments.

Arguments:

REG	Suite name
[TASKID ...]	Task identifiers

Options:

-h, --help	show this help message and exit
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
--port=INT	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
--use-ssh	Use ssh to re-invoke the command on the suite host.
--no-login	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
--print-uuid	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
--set-uuid=UUID	Set the client UUID manually (e.g. from prior use of --print-uuid). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
-f, --force	Do not ask for confirmation before acting. Note that it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.
-m, --family	(Obsolete) This option is now ignored and is retained for backward compatibility only. TASKID in the argument list can be used to match task and family names regardless of this option.
--no-multitask-compat	Disallow backward compatible multitask interface.

F.2.62 stop

Usage: `cylc [control] stop|shutdown [OPTIONS] REG [STOP]`

Tell a suite server program to shut down. In order to prevent failures going unnoticed, suites only shut down automatically at a final cycle point if no failed tasks are present. There are several shutdown methods:

1. (default) stop after current active tasks finish
2. (--now) stop immediately, orphaning current active tasks
3. (--kill) stop after killing current active tasks
4. (with STOP as a cycle point) stop after cycle point STOP
5. (with STOP as a task ID) stop after task ID STOP has succeeded
6. (--wall-clock=T) stop after time T (an ISO 8601 date-time format e.g. CCYYMMDDThh:mm, CCYY-MM-DDThh, etc).

Tasks that become ready after the shutdown is ordered will be submitted immediately if the suite is restarted. Remaining task event handlers and job poll and kill commands, however, will be executed prior to shutdown, unless --now is used.

This command exits immediately unless --max-polls is greater than zero, in which case it polls to wait for suite shutdown.

Arguments:

REG	Suite name
[STOP]	a/ task POINT (cycle point), or b/ ISO 8601 date-time (clock time), or c/ TASK (task ID).

Options:

-h, --help	show this help message and exit
-k, --kill	Shut down after killing currently active tasks.
-n, --now	Shut down without waiting for active tasks to complete. If this option is specified once, wait for task event handler, job poll/kill to complete. If this option is specified more than once, tell the suite to terminate immediately.
-w STOP, --wall-clock=STOP	Shut down after time STOP (ISO 8601 formatted)
--max-polls=INT	Maximum number of polls (default 0).
--interval=SECS	Polling interval in seconds (default 60).
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
--port=INT	Suite port number on the suite host. NOTE: this is retrieved automatically if non-interactive ssh is configured to the suite host.
--use-ssh	Use ssh to re-invoke the command on the suite host.
--no-login	Do not use a login shell to run remote ssh commands. The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC	Set a timeout for network connections to the running suite. The default is no timeout. For task messaging connections see site/user config file documentation.
--print-uuid	Print the client UUID to stderr. This can be matched to information logged by the receiving suite server program.
--set-uuid=UUID	Set the client UUID manually (e.g. from prior use of --print-uuid). This can be used to log multiple commands under the same UUID (but note that only the first [info] command from the same client ID will be logged unless the suite is running in debug mode).
-f, --force	Do not ask for confirmation before acting. Note that it is not necessary to use this option if interactive command prompts have been disabled in the site/user config files.

F.2.63 submit

Usage: `cylc [task] submit|single [OPTIONS] REG TASK [...]`

Submit a single task to run just as it would be submitted by its suite. Task messaging commands will print to stdout but will not attempt to communicate

with the suite (which does not need to be running).

For tasks present in the suite graph the given cycle point is adjusted up to the next valid cycle point for the task. For tasks defined under runtime but not present in the graph, the given cycle point is assumed to be valid.

WARNING: do not 'cylc submit' a task that is running in its suite at the same time - both instances will attempt to write to the same job logs.

Arguments:

REG	Suite name
TASK [...]	Family or task ID (NAME.CYCLE_POINT)

Options:

-h, --help	show this help message and exit
-d, --dry-run	Generate the job script for the task, but don't submit it.
--user=USER	Other user account name. This results in command reinvocation on the remote account.
--host=HOST	Other host name. This results in command reinvocation on the remote account.
-v, --verbose	Verbose output mode.
--debug	Run suites in non-daemon mode, and show exception tracebacks.
-s NAME=VALUE, --set=NAME=VALUE	Set the value of a Jinja2 template variable in the suite definition. This option can be used multiple times on the command line. NOTE: these settings persist across suite restarts, but can be set again on the "cylc restart" command line if they need to be overridden.
--set-file=FILE	Set the value of Jinja2 template variables in the suite definition from a file containing NAME=VALUE pairs (one per line). NOTE: these settings persist across suite restarts, but can be set again on the "cylc restart" command line if they need to be overridden.
--icp=CYCLE_POINT	Set initial cycle point. Required if not defined in suite.rc.

F.2.64 suite-state

Usage: `cylc suite-state REG [OPTIONS]`

Print task states retrieved from a suite database; or (with --task, --point, and --status) poll until a given task reaches a given state; or (with --task, --point, and --message) poll until a task receives a given message. Polling is configurable with --interval and --max-polls; for a one-off check use --max-polls=1. The suite database does not need to exist at the time polling commences but allocated polls are consumed waiting for it (consider max-polls*interval as an overall timeout).

Note for non-cycling tasks --point=1 must be provided.

Important: cylc suite-state only works with task states and does not work with task messages.

For your own suites the database location is determined by your site/user config. For other suites, e.g. those owned by others, or mirrored suite databases, use --run-dir=DIR to specify the location.

Example usages:

```
cylc suite-state REG --task=TASK --point=POINT --status=STATUS
returns 0 if TASK.POINT reaches STATUS before the maximum number of
polls, otherwise returns 1.
```

```
cylc suite-state REG --task=TASK --point=POINT --status=STATUS --offset=PT6H
adds 6 hours to the value of CYCLE for carrying out the polling operation.
```

```
cylc suite-state REG --task=TASK --status=STATUS --task-point
uses CYLC_TASK_CYCLE_POINT environment variable as the value for the CYCLE
```

to poll. This is useful when you want to use `cylc suite-state` in a `cylc` task.

Arguments:

REG Suite name

Options:

```
-h, --help                show this help message and exit
-t TASK, --task=TASK      Specify a task to check the state of.
-p CYCLE, --point=CYCLE   Specify the cycle point to check task states for.
                          Use the CYLC_TASK_CYCLE_POINT environment variable as
                          the cycle point to check task states for. Shorthand
                          for --point=$CYLC_TASK_CYCLE_POINT
-T, --task-point          Remote cyclepoint template (IGNORED - this is now
                          determined automatically).
--template=TEMPLATE
-d DIR, --run-dir=DIR     The top level cylc run directory if non-standard. The
                          database should be DIR/REG/log/db. Use to interrogate
                          suites owned by others, etc.; see note above.
-s OFFSET, --offset=OFFSET Specify an offset to add to the targeted cycle point
-S STATUS, --status=STATUS Specify a particular status or triggering condition to
                          check for. Valid triggering conditions to check for
                          include: 'fail', 'finish', 'start', 'submit' and
                          'succeed'. Valid states to check for include:
                          'runahead', 'waiting', 'held', 'queued', 'ready',
                          'expired', 'submitted', 'submit-failed', 'submit-
                          retrying', 'running', 'succeeded', 'failed' and
                          'retrying'.
-O MSG, -m MSG, --output=MSG, --message=MSG Check custom task output by message
                          string or trigger string.
--max-polls=INT           Maximum number of polls (default 10).
--interval=SECS           Polling interval in seconds (default 60).
--user=USER               Other user account name. This results in command
                          reinvocation on the remote account.
--host=HOST               Other host name. This results in command reinvocation
                          on the remote account.
-v, --verbose             Verbose output mode.
--debug                  Run suites in non-daemon mode, and show exception
                          tracebacks.
```

F.2.65 test-battery

Usage: `cylc test-battery [...]`

Run automated `cylc` and `parsec` tests under [FILES or DIRECTORIES].

Test locations default to the following directory tree:

```
/home/vagrant/cylc.git/tests/
```

Some tests (e.g. those specific to particular batch schedulers) can be configured in your site/user config file. A few others still submit jobs to a `user@host` account taken from the environment:

```
$CYLC_TEST_TASK_HOST # default localhost
$CYLC_TEST_TASK_OWNER # default $USER
```

Requirements:

- * Passwordless `ssh` must be configured to task host accounts.
- * Some test suites submit jobs to 'at' so `atd` must be running.

Options and arguments are appended to the `"prove -j $NPROC -s -r ${@:-tests}"` command, where `NPROC` is the number of child processes that can be used to run the test files.

Some tests use a clean global config file. If some items from your site config file are needed in this, e.g. to get remote test hosts working, add them to `/home/vagrant/cylc.git/conf/global-tests.rc`.

The command normally uses the `"process pool size"` setting (default=4) in the

site/user global configuration file to determine the number of tests to run in parallel. You can also change the amount of concurrency with the "-j N" option.

Suite run directories are cleaned up on the suite host for passing tests - otherwise they are left alone.

To output stderr from failed tests to the terminal, "export CYLC_TEST_DEBUG=true" before running this command.

The command normally uses "diff -u" to compare files. However, if an alternate command such as "xxdiff -D" is desirable (e.g. for debugging), "export CYLC_TEST_DIFF_CMD=xxdiff -D".

Commits or Pull Requests to cylc/cylc on GitHub will trigger Travis CI to run generic (non platform-specific) tests - see /home/vagrant/cylc.git/.travis.yml. After enabling Travis CI for your own cylc fork, you can skip generic tests locally by setting CYLC_TEST_RUN_GENERIC=false.

By default all tests are executed. To run just a subset of them:

- * list individual tests or test directories to run on the command line
 - * list individual tests or test directories to skip in \$CYLC_TEST_SKIP
 - * skip all generic tests with CYLC_TEST_RUN_GENERIC=false
 - * skip all platform-specific tests with CYLC_TEST_RUN_PLATFORM=false
- List specific tests relative to /home/vagrant/cylc.git (i.e. starting with "test/").

Some platform-specific tests are automatically skipped, depending on platform.

FOR DEVELOPERS:

- * Platform-specific tests must set "CYLC_TEST_IS_GENERIC=false" before sourcing the test_header.
- * Tests requiring the sqlite3 CLI must be skipped if sqlite3 is not installed (it is not otherwise a Cylc software prerequisite):

```
| if ! which sqlite3 > /dev/null; then
|   # Skip the remaining 3 tests.
|   skip 3 "sqlite3 not installed?"
|   purge_suite $SUITE_NAME
|   exit 0
| fi
```

For more information see "Reference Tests" in the User Guide.

Options:

-h, --help Print this help message and exit.

Examples:

Run the full test suite with the default options.

```
cylc test-battery
```

Run the full test suite with 12 processes.

```
cylc test-battery -j 12
```

Run only tests under "tests/cyclers/" with 12 processes.

```
cylc test-battery -j 12 tests/cyclers
```

Run only "tests/cyclers/16-weekly.t" in verbose mode

```
cylc test-battery -v tests/cyclers/16-weekly.t
```

Run only tests under "tests/cyclers/" with 12 processes, and skip 00-daily.t

```
export CYLC_TEST_SKIP=tests/cyclers/00-daily.t
cylc test-battery -j 12 tests/cyclers
```

F.2.66 trigger

Usage: cylc [control] trigger [OPTIONS] REG [TASKID ...]

Manually trigger one or more tasks. Waiting tasks will be queued (cylc internal queues) and will submit as normal when released by the queue; queued tasks will submit immediately even if that violates the queue limit (so you may need to trigger a queue-limited task twice to get it to submit).

For single tasks you can use "--edit" to edit the generated job script before it submits, to apply one-off changes. A diff between the original and edited job script will be saved to the task job log directory.

TASKID is a pattern to match task proxies or task families, or groups of them:

```
* [CYCLE-POINT-GLOB/]TASK-NAME-GLOB[:TASK-STATE]
* [CYCLE-POINT-GLOB/]FAMILY-NAME-GLOB[:TASK-STATE]
* TASK-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
* FAMILY-NAME-GLOB[.CYCLE-POINT-GLOB][:TASK-STATE]
```

For example, to match:

```
* all tasks in a cycle: '20200202T0000Z/*' or '*.20200202T0000Z'
* all tasks in the submitted status: ':submitted'
* retrying 'foo*' tasks in 0000Z cycles: 'foo*.0000Z:retrying' or
  '*0000Z/foo*:retrying'
* retrying tasks in 'BAR' family: '*/BAR:retrying' or 'BAR.*:retrying'
* retrying tasks in 'BAR' or 'BAZ' families: '*/BA[RZ]:retrying' or
  'BA[RZ].*:retrying'
```

The old 'MATCH POINT' syntax will be automatically detected and supported. To avoid this, use the '--no-multitask-compat' option, or use the new syntax (with a '/' or a '.') when specifying 2 TASKID arguments.

Arguments:

```
REG                               Suite name
[TASKID ...]                      Task identifiers
```

Options:

```
-h, --help                        show this help message and exit
-e, --edit                        Manually edit the job script before running it.
-g, --geditor                    (with --edit) force use of the configured GUI editor.
--user=USER                      Other user account name. This results in command
                                reinvocation on the remote account.
--host=HOST                      Other host name. This results in command reinvocation
                                on the remote account.
-v, --verbose                    Verbose output mode.
--debug                          Run suites in non-daemon mode, and show exception
                                tracebacks.
--port=INT                       Suite port number on the suite host. NOTE: this is
                                retrieved automatically if non-interactive ssh is
                                configured to the suite host.
--use-ssh                        Use ssh to re-invoke the command on the suite host.
--no-login                       Do not use a login shell to run remote ssh commands.
                                The default is to use a login shell.
--comms-timeout=SEC, --pyro-timeout=SEC
                                Set a timeout for network connections to the running
                                suite. The default is no timeout. For task messaging
                                connections see site/user config file documentation.
--print-uuid                    Print the client UUID to stderr. This can be matched
                                to information logged by the receiving suite server
                                program.
--set-uuid=UUID                 Set the client UUID manually (e.g. from prior use of
                                --print-uuid). This can be used to log multiple
                                commands under the same UUID (but note that only the
                                first [info] command from the same client ID will be
                                logged unless the suite is running in debug mode).
-f, --force                      Do not ask for confirmation before acting. Note that
                                it is not necessary to use this option if interactive
                                command prompts have been disabled in the site/user
                                config files.
-m, --family                    (Obsolete) This option is now ignored and is retained
                                for backward compatibility only. TASKID in the
                                argument list can be used to match task and family
                                names regardless of this option.
--no-multitask-compat           Disallow backward compatible multitask interface.
```

F.2.67 upgrade-run-dir

Usage: `cylc [admin] upgrade-run-dir SUITE`

For one-off conversion of a suite run directory to cylc-6 format.

Arguments:

```
SUITE      suite name or run directory path
```


Options:

-h, --help show this help message and exit

F.2.68 validate

Usage: cylc [prep] validate [OPTIONS] SUITE

Validate a suite definition against the official specification files held in \$CYLC_DIR/conf/suiterc/.

If the suite definition uses include-files reported line numbers will correspond to the inlined version seen by the parser; use 'cylc view -i,--inline SUITE' for comparison.

Arguments:

SUITE Suite name or path

Options:

-h, --help show this help message and exit
 --strict Fail any use of unsafe or experimental features. Currently this just means naked dummy tasks (tasks with no corresponding runtime section) as these may result from unintentional typographic errors in task names.
 -o FILENAME, --output=FILENAME Specify a file name to dump the processed suite.rc.
 --profile Output profiling (performance) information
 -u RUN_MODE, --run-mode=RUN_MODE Validate for run mode.
 --user=USER Other user account name. This results in command reinvocation on the remote account.
 --host=HOST Other host name. This results in command reinvocation on the remote account.
 -v, --verbose Verbose output mode.
 --debug Run suites in non-daemon mode, and show exception tracebacks.
 --suite-owner=OWNER Specify suite owner
 -s NAME=VALUE, --set=NAME=VALUE Set the value of a Jinja2 template variable in the suite definition. This option can be used multiple times on the command line. NOTE: these settings persist across suite restarts, but can be set again on the "cylc restart" command line if they need to be overridden.
 --set-file=FILE Set the value of Jinja2 template variables in the suite definition from a file containing NAME=VALUE pairs (one per line). NOTE: these settings persist across suite restarts, but can be set again on the "cylc restart" command line if they need to be overridden.
 --icp=CYCLE_POINT Set initial cycle point. Required if not defined in suite.rc.

F.2.69 version

Usage: cylc [info] version

Print the cylc version invoked at the command line.

Note that "cylc -v,--version" just prints the version string from the main command interface, whereas this is a proper cylc command that can take the standard --host and --user options, etc.

For the cylc version of running a suite server program, see "cylc get-suite-version".

Arguments:**Options:**


```

-h, --help      show this help message and exit
--long          Print the path to the current cylc version
--user=USER     Other user account name. This results in command reinvocation
                on the remote account.
--host=HOST     Other host name. This results in command reinvocation on the
                remote account.
-v, --verbose   Verbose output mode.
--debug         Run suites in non-daemon mode, and show exception tracebacks.

```

F.2.70 view

Usage: `cylc [prep] view [OPTIONS] SUITE`

View a read-only temporary copy of suite NAME's suite.rc file, in your editor, after optional include-file inlining and Jinja2 preprocessing.

The edit process is spawned in the foreground as follows:

```
% <editor> suite.rc
```

Where <editor> is defined in the cylc site and user config files (\$CYLC_DIR/conf/global.rc and \$HOME/.cylc/global.rc).

For remote host or owner, the suite will be printed to stdout unless the '-g,--gui' flag is used to spawn a remote GUI edit session.

See also 'cylc [prep] edit'.

Arguments:

SUITE	Suite name or path
-------	--------------------

Options:

```

-h, --help      show this help message and exit
-i, --inline    Inline include-files.
-j, --jinja2    View after Jinja2 template processing (implies
                '-i/--inline' as well).
-p, --process   View after all processing (Jinja2, inlining, line-
                continuation joining).
-m, --mark      (With '-i') Mark inclusions in the left margin.
-l, --label     (With '-i') Label file inclusions with the file name.
                Line numbers will not correspond to those reported by
                the parser.
--single        (With '-i') Inline only the first instances of any
                multiply-included files. Line numbers will not
                correspond to those reported by the parser.
-c, --cat       Concatenate continuation lines (line numbers will not
                correspond to those reported by the parser).
-g, --gui       Force use of the configured GUI editor.
--stdout        Print the suite definition to stdout.
--mark-for-edit (With '-i') View file inclusion markers as for 'cylc
                edit --inline'.
--user=USER     Other user account name. This results in command
                reinvocation on the remote account.
--host=HOST     Other host name. This results in command reinvocation
                on the remote account.
-v, --verbose   Verbose output mode.
--debug         Run suites in non-daemon mode, and show exception
                tracebacks.
--suite-owner=OWNER Specify suite owner
-s NAME=VALUE, --set=NAME=VALUE
                Set the value of a Jinja2 template variable in the
                suite definition. This option can be used multiple
                times on the command line. NOTE: these settings
                persist across suite restarts, but can be set again on
                the "cylc restart" command line if they need to be
                overridden.
--set-file=FILE Set the value of Jinja2 template variables in the
                suite definition from a file containing NAME=VALUE
                pairs (one per line). NOTE: these settings persist
                across suite restarts, but can be set again on the
                "cylc restart" command line if they need to be
                overridden.

```

F.2.71 warranty

Usage: `cylc [license] warranty [--help]`

Cylc is released under the GNU General Public License v3.0
This command prints the GPL v3.0 disclaimer of warranty.

Options:

`--help` Print this usage message.

G The gcylc Graph View

The graph view in the gcylc GUI shows the structure of the suite as it evolves. It can work well even for large suites, but be aware that the graphviz layout engine has to do a new global layout every time a task proxy appears in or disappears from the task pool. The following may help mitigate any jumping layout problems:

- The disconnect button can be used to temporarily prevent the graph from changing as the suite evolves.
- The greyed-out base nodes, which are only present to fill out the graph structure, can be toggled off (but this will split the graph into disconnected sub-trees).
- Right-click on a task and choose the “Focus” option to restrict the graph display to that task’s cycle point. Anything interesting happening in other cycle points will show up as disconnected rectangular nodes to the right of the graph (and you can click on those to instantly refocus to their cycle points).
- Task filtering is the ultimate quick route to focusing on just the tasks you’re interested in, but this will destroy the graph structure.

H Cylc README File

The Cylc Workflow Engine

```
[[Build Status]](https://travis-ci.org/cylc/cylc.svg?branch=master)](https://travis-ci.org/cylc/cylc)
[[DOI]](https://zenodo.org/badge/1836229.svg)](https://zenodo.org/badge/latestdoi/1836229)
```

Cylc (ˈsɪlk) orchestrates complex distributed suites of interdependent cycling tasks. It was originally designed for environmental forecasting systems at [NIWA](https://www.niwa.co.nz).

Copyright and Terms of Use

Copyright (C) 2008-2018 NIWA

Cylc is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

Cylc is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with cylc. If not, see [GNU licenses](http://www.gnu.org/licenses/).

Cylc Documentation

* See [The Cylc Home Page](https://cylc.github.io/cylc)

Code Contributors

I CYLC INSTALL FILE

```
<!-- git shortlog -s -n -->
* Hilary Oliver
* Matt Shin
* Ben Fitzpatrick
* Andrew Clark
* Oliver Sanders
* Luis Kornblueh
* Kerry Day
* David Matthews
* Tim Whitcomb
* Scott Wales
* Annette Osprey
* Bruno Kinoshita
* Domingo Manubens Gil
* Jonny Williams
* Alex Reinecke
* Chan Wilson
* Kevin Pulo

## Acknowledgement For Non-Cylc Work:
Licences for non-cylc work included in this distribution can be found in the
'licences/' directory.
* 'lib/cherrypy/':
    External software library released under a BSD license.
    Minor modification to ignore an import warning.
    See [cherrypy](http://www.cherrypy.org/).
* 'lib/isodatetime/':
    Unmodified external software library released under the LGPL license.
    See [metomi/isodatetime](https://github.com/metomi/isodatetime/).
* 'lib/jinja2/':
    External software library released under a BSD license.
    See [Jinja2](http://jinja.pocoo.org/).
* 'lib/markupsafe/':
    External software library released under a BSD license, used by Jinja2.
    See [MarkupSafe](http://www.pocoo.org/projects/markupsafe/).
* 'lib/xdot.py':
    External software released under the LGPL license.
    Modifications based on version 0.6. See
    [xdot](https://github.com/jrfonseca/xdot.py)
```

I Cylc INSTALL File

Cylc Installation.

```
**See [The Cylc User Guide](https://cylc.github.io/cylc/documentation.html) for
detailed instructions.**
```

Cylc must be installed on suite and task job hosts, although the external software packages (below) are not required on job hosts.

Required External Software Packages

These can be installed once on suite hosts updated infrequently.

```
* graphviz
* pygraphviz
```

Installing Cylc

Download the latest tarball from <https://github.com/cylc/cylc/releases>.

Successive Cylc releases should be installed side-by-side under a location such as '/opt':

```
'''bash
cd /opt
tar xzf cylc-7.4.0.tar.gz
# DO NOT CHANGE THE NAME OF THE UNPACKED CYLC SOURCE DIRECTORY.
cd cylc-7.4.0
export PATH=$PWD/bin:$PATH
```

J CYLC DEVELOPMENT HISTORY - MAJOR CHANGES

```
make
'''
```

When you type 'make':

- * A file called VERSION is created, containing the Cylc version number
- * The version number is taken from the name of the parent directory: DO NOT CHANGE THE NAME OF THE UNPACKED CYLC SOURCE DIRECTORY
- * The Cylc documentation is generated from source and put in doc/install/

Once installed, Cylc commands should be invoked via the supplied central wrapper script that selects between the available versions. This allows long-running suites (and their task jobs) to stick with older versions if necessary. The wrapper should be edited to point to the Cylc install location:

```
'''bash
cp /opt/cylc-7.4.0/admin/cylc-wrapper /usr/local/bin/cylc
# (now edit '/usr/local/bin/cylc' as per in-file instructions...)
'''
```

Finally, make a symlink to the latest installed version:

```
'''bash
ln -s /opt/cylc-7.4.0 /opt/cylc
'''
```

(This will be the default version invoked by the wrapper if a specific version is not requested via `\lstinline=$CYLC_VERSION=`.)

Installing The Documentation

After running 'make' you can copy the entire 'doc/install' directory to a convenient location such as '/var/www/html/', and update your Cylc site config file to point to the intranet location.

Cloning The Cylc Repository

To participate in Cylc development fork [Cylc on GitHub](<https://github.com/cylc/cylc>) and clone it locally. Changes should be developed in feature branches then pushed to your GitHub fork before issuing a Pull Request to the team. Please discuss proposed changes before you begin work.

J Cylc Development History - Major Changes

- **pre-cylc-3** - early versions focused on the new scheduling algorithm. A suite was a collection of "task definition files" that encoded the prerequisites and outputs of each task, exposing cylc's self-organising nature. Tasks could be transferred from one suite to another by simply copying their taskdef files over and checking prerequisite and output consistency. Global suite structure was not easy to discern until run time (although cylc-2 could generate resolved run time dependency graphs).
- **cylc-3** - a new suite design interface: dependency graph and task runtime properties defined in a single structured, validated, configuration file - the suite.rc file; graphical user interface; suite graphing.
- **cylc-4** - refined and organized the suite.rc file structure; task runtime properties defined by an efficient inheritance hierarchy; support for the Jinja2 template processor in suite definitions.
- **cylc-5** - multi-threading for continuous network request handling and job submission; more task states to distinguish job submission from execution; dependence between suites via new suite run databases; polling and killing of real task jobs; polling as task communications option.
- **cylc-6** - specification of all date-times and cycling workflows via the ISO8601 date-times, durations, and recurrence expressions; integer cycling; a multi-process pool to execute job submissions, event handlers, and poll and kill commands.

- **cylc-7** - Replaced the Pyro communications layer with RESTful HTTPS. Removed deprecated pre cylc-6 syntax and features.

K Communication Method

Cylc suite server programs and clients (commands, cylc gui, task messaging) communicate via particular ports using the HTTPS protocol, secured by HTTP Digest Authentication using the suite's 20-random-character private passphrase and private SSL certificate.

This is enabled via the included-in-cylc cherrypy library (for the server) and either the Python requests library (if available) or the built-in Python libraries for the clients.

All suites are entirely isolated from one another.

L Cylc 6 Migration Reference

Cylc 6 introduced new date-time-related syntax for the suite.rc file. In some places, this is quite radically different from the earlier syntax.

L.1 Timeouts and Delays

Timeouts and delays such as `[cylc][[events]]timeout` or `[runtime][[my_task]][[[job]]]execution retry delays` were written in a purely numeric form before cylc 6, in seconds, minutes (most common), or hours, depending on the setting.

They are now written in an ISO 8601 duration form, which has the benefit that the units are user-selectable (use 1 day instead of 1440 minutes) and explicit.

Nearly all timeouts and delays in cylc were in minutes, except for:

```
[runtime][[my_task]][[[suite state polling]]]interval
[runtime][[my_task]][[[simulation mode]]]run time range
which were in seconds, and
[scheduling]runahead limit
```

which was in hours (this is a special case discussed below in [L.2](#)).

See [Table 1](#).

Table 1: Timeout/Delay Syntax Change Examples

Setting	Pre-Cylc-6	Cylc-6+
<code>[cylc][[events]]timeout</code>	180	PT3H
<code>[runtime][[my_task]][[[job]]]execution retry delays</code>	2*30, 360, 1440	2*PT30M, PT6H, P1D
<code>[runtime][[my_task]][[[suite state polling]]]interval</code>	2	PT2S

L.2 Runahead Limit

See [A.4.7](#).

The `[scheduling]runahead limit` setting was written as a number of hours in pre-cylc-6 suites. This is now in ISO 8601 format for date-time cycling suites, so `[scheduling]runahead limit=36` would be written `[scheduling]runahead limit=PT36H`.

There is a new preferred alternative to `runahead limit`, `[scheduling]max active cycle points`. This allows the user to configure how many cycle points can run at once (default 3). See [A.4.8](#).

L.3 Cycle Time/Cycle Point

See [A.4.2](#).

The following suite.rc settings have changed name (Table 2):

Table 2: Cycle Point Renaming	
Pre-Cylc-6	Cylc-6+
<code>[scheduling]initial cycle time</code>	<code>[scheduling]initial cycle point</code>
<code>[scheduling]final cycle time</code>	<code>[scheduling]final cycle point</code>
<code>[visualization]initial cycle time</code>	<code>[visualization]initial cycle point</code>
<code>[visualization]final cycle time</code>	<code>[visualization]final cycle point</code>

This change is to reflect the fact that cycling in cylc 6+ can now be over e.g. integers instead of being purely based on date-time.

Date-times written in `initial cycle time` and `final cycle time` were in a cylc-specific 10-digit (or less) `CCYYMMDDhh` format, such as `2014021400` for 00:00 on the 14th of February 2014.

Date-times are now required to be ISO 8601 compatible. This can be achieved easily enough by inserting a `T` between the day and the hour digits.

Table 3: Cycle Point Syntax Example		
Setting	Pre-Cylc-6	Cylc-6+
<code>[scheduling]initial cycle time</code>	<code>2014021400</code>	<code>20140214T00</code>

L.4 Cycling

Special *start-up* and *cold-start* tasks have been removed from cylc 6. Instead, use the initial/run-once notation as detailed in [7.23.3](#) and [9.3.4.7](#).

Repeating asynchronous tasks have also been removed because non date-time workflows can now be handled more easily with integer cycling. See for instance the satellite data processing example documented in [9.3.4.8](#).

For repeating tasks with hour-based cycling the syntax has only minor changes:

Pre-cylc-6:

```
[scheduling]
[[dependencies]]
[[[0,12]]]
graph = foo[T-12] => foo & bar => baz
```

```
[scheduling]
[[dependencies]]
[[[T00,T12]]]
graph = foo[-PT12H] => foo & bar => baz
```

Hour-based cycling section names are easy enough to convert, as seen in Table 4.

Table 4: Hourly Cycling Sections	
Pre-Cylc-6	Cylc-6+
[scheduling][[dependencies]][[[0]]]	[scheduling][[dependencies]][[[T00]]]
[scheduling][[dependencies]][[[6]]]	[scheduling][[dependencies]][[[T06]]]
[scheduling][[dependencies]][[[12]]]	[scheduling][[dependencies]][[[T12]]]
[scheduling][[dependencies]][[[18]]]	[scheduling][[dependencies]][[[T18]]]

The graph text in hour-based cycling is also easy to convert, as seen in Table 5.

Table 5: Hourly Cycling Offsets	
Pre-Cylc-6	Cylc-6+
my_task[T-6]	my_task[-PT6H]
my_task[T-12]	my_task[-PT12H]
my_task[T-24]	my_task[-PT24H] or even my_task[-P1D]

L.5 No Implicit Creation of Tasks by Offset Triggers

Prior to cylc-6 intercycle offset triggers implicitly created task instances at the offset cycle points. For example, this pre cylc-6 suite automatically creates instances of task `foo` at the offset hours 3,9,15,21 each day, for task `bar` to trigger off at 0,6,12,18:

```
# Pre cylc-6 implicit cycling.
[scheduling]
initial cycle time = 2014080800
[[dependencies]]
[[[00,06,12,18]]]
# This creates foo instances at 03,09,15,21:
graph = foo[T-3] => bar
```

Here's the direct translation to cylc-6+ format:

```
# In cylc-6+ this suite will stall.
[scheduling]
initial cycle point = 20140808T00
[[dependencies]]
[[[T00,T06,T12,T18]]]
# This does NOT create foo instances at 03,09,15,21:
graph = foo[-PT3H] => bar
```

This suite fails validation with `ERROR: No cycling sequences defined for foo`, and at runtime it would stall with `bar` instances waiting on non-existent offset `foo` instances (note that these appear as ghost nodes in graph visualisations).

To fix this, explicitly define the cycling of with an offset cycling sequence: `foo`:

```
# Cylc-6+ requires explicit task instance creation.
[scheduling]
initial cycle point = 20140808T00
[[dependencies]]
```

```
[[[T03,T09,T15,T21]]]
graph = foo
[[[T00,T06,T12,T18]]]
graph = foo[-PT3H] => bar
```

Implicit task creation by offset triggers is no longer allowed because it is error prone: a mistaken task cycle point offset should cause a failure rather than automatically creating task instances on the wrong cycling sequence.

M Known Issues

M.1 Current Known Issues

The best place to find current known issues is on Github: <https://github.com/cylc/cylc/issues>.

M.2 Notable Known Issues

M.2.1 Use of pipes in job scripts

In bash, the return status of a pipeline is normally the exit status of the last command. This is unsafe, because if any command in the pipeline fails, the script will continue nevertheless.

For safety, a cylc task job script running in bash will have the `set -o pipefail` option turned on automatically. If a pipeline exists in a task's `script`, `etc` section, the failure of any part of a pipeline will cause the command to return a non-zero code at the end, which will be reported as a task job failure. Due to the unique nature of a pipeline, the job file will trap the failure of the individual commands, as well as the whole pipeline, and will attempt to report a failure back to the suite twice. The second message is ignored by the suite, and so the behaviour can be safely ignored. (You should probably still investigate the failure, however!)

N GNU GENERAL PUBLIC LICENSE v3.0

Copyright © 2007 Free Software Foundation, Inc. <http://fsf.org/>

Everyone is permitted to copy and distribute verbatim copies of this
license document, but changing it is not allowed.

PREAMBLE

The GNU General Public License is a free, copyleft license for software and other kinds of works.

The licenses for most software and other practical works are designed to take away your freedom to share and change the works. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change all versions of a program—to make sure it remains free software for all its users. We, the Free Software Foundation, use the GNU General Public License for most of our software; it applies also to any other work released this way by its authors. You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for them if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs, and that you know you can do these things.

To protect your rights, we need to prevent others from denying you these rights or asking you to surrender the rights. Therefore, you have certain responsibilities if you distribute copies of the software, or if you modify it: responsibilities to respect the freedom of others.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must pass on to the recipients the same freedoms that you received. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

Developers that use the GNU GPL protect your rights with two steps: (1) assert copyright on the software, and (2) offer you this License giving you legal permission to copy, distribute and/or modify it.

For the developers' and authors' protection, the GPL clearly explains that there is no warranty for this free software. For both users' and authors' sake, the GPL requires that modified versions be marked as changed, so that their problems will not be attributed erroneously to authors of previous versions.

Some devices are designed to deny users access to install or run modified versions of the software inside them, although the manufacturer can do so. This is fundamentally incompatible with the aim of protecting users' freedom to change the software. The systematic pattern of such abuse occurs in the area of products for individuals to use, which is precisely where it is most unacceptable. Therefore, we have designed this version of the GPL to prohibit the practice for those products. If such problems arise substantially in other domains, we stand ready to extend this provision to those domains in future versions of the GPL, as needed to protect the freedom of users.

Finally, every program is threatened constantly by software patents. States should not allow patents to restrict development and use of software on general-purpose computers, but in those that do, we wish to avoid the special danger that patents applied to a free program could make it effectively proprietary. To prevent this, the GPL assures that patents cannot be used to render the program non-free.

The precise terms and conditions for copying, distribution and modification follow.

TERMS AND CONDITIONS

0. Definitions.

“This License” refers to version 3 of the GNU General Public License.

“Copyright” also means copyright-like laws that apply to other kinds of works, such as semiconductor masks.

“The Program” refers to any copyrightable work licensed under this License. Each licensee is addressed as “you”. “Licensees” and “recipients” may be individuals or organizations.

To “modify” a work means to copy from or adapt all or part of the work in a fashion requiring copyright permission, other than the making of an exact copy. The resulting

work is called a “modified version” of the earlier work or a work “based on” the earlier work.

A “covered work” means either the unmodified Program or a work based on the Program.

To “propagate” a work means to do anything with it that, without permission, would make you directly or secondarily liable for infringement under applicable copyright law, except executing it on a computer or modifying a private copy. Propagation includes copying, distribution (with or without modification), making available to the public, and in some countries other activities as well.

To “convey” a work means any kind of propagation that enables other parties to make or receive copies. Mere interaction with a user through a computer network, with no transfer of a copy, is not conveying.

An interactive user interface displays “Appropriate Legal Notices” to the extent that it includes a convenient and prominently visible feature that (1) displays an appropriate copyright notice, and (2) tells the user that there is no warranty for the work (except to the extent that warranties are provided), that licensees may convey the work under this License, and how to view a copy of this License. If the interface presents a list of user commands or options, such as a menu, a prominent item in the list meets this criterion.

1. Source Code.

The “source code” for a work means the preferred form of the work for making modifications to it. “Object code” means any non-source form of a work.

A “Standard Interface” means an interface that either is an official standard defined by a recognized standards body, or, in the case of interfaces specified for a particular programming language, one that is widely used among developers working in that language.

The “System Libraries” of an executable work include anything, other than the work as a whole, that (a) is included in the normal form of packaging a Major Component, but which is not part of that Major Component, and (b) serves only to enable use of the work with that Major Component, or to implement a Standard Interface for which an implementation is available to the public in source code form. A “Major Component”, in this context, means a major essential component (kernel, window system, and so on) of the specific operating system (if any) on which the executable work runs, or a compiler used to produce the work, or an object code interpreter used to run it.

The “Corresponding Source” for a work in object code form means all the source code needed to generate, install, and (for an executable work) run the object code and to modify the work, including scripts to control those activities. However, it does not include the work’s System Libraries, or general-purpose tools or generally available free programs which are used unmodified in performing those activities but which are not part of the work. For example, Corresponding Source includes interface definition files associated with source files for the work, and the source code for shared libraries and dynamically linked subprograms that the work is specifically designed to require, such as by intimate data communication or control flow between those subprograms and other parts of the work.

The Corresponding Source need not include anything that users can regenerate automatically from other parts of the Corresponding Source.

The Corresponding Source for a work in source code form is that same work.

2. Basic Permissions.

All rights granted under this License are granted for the term of copyright on the Program, and are irrevocable provided the stated conditions are met. This License explicitly affirms your unlimited permission to run the unmodified Program. The output from running a covered work is covered by this License only if the output, given its content, constitutes a covered work. This License acknowledges your rights of fair use or other equivalent, as provided by copyright law.

You may make, run and propagate covered works that you do not convey, without conditions so long as your license otherwise remains in force. You may convey covered works to others for the sole purpose of having them make modifications exclusively for you, or provide you with facilities for running those works, provided that you comply with the terms of this License in conveying all material for which you do not control copyright. Those thus making or running the covered works for you must do so exclusively on your behalf, under your direction and control, on terms that prohibit them from making any copies of your copyrighted material outside their relationship with you.

Conveying under any other circumstances is permitted solely under the conditions stated below. Sublicensing is not allowed; section 10 makes it unnecessary.

3. Protecting Users' Legal Rights From Anti-Circumvention Law.

No covered work shall be deemed part of an effective technological measure under any applicable law fulfilling obligations under article 11 of the WIPO copyright treaty adopted on 20 December 1996, or similar laws prohibiting or restricting circumvention of such measures.

When you convey a covered work, you waive any legal power to forbid circumvention of technological measures to the extent such circumvention is effected by exercising rights under this License with respect to the covered work, and you disclaim any intention to limit operation or modification of the work as a means of enforcing, against the work's users, your or third parties' legal rights to forbid circumvention of technological measures.

4. Conveying Verbatim Copies.

You may convey verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice; keep intact all notices stating that this License and any non-permissive terms added in accord with section 7 apply to the code; keep intact all notices of the absence of any warranty; and give all recipients a copy of this License along with the Program.

You may charge any price or no price for each copy that you convey, and you may offer support or warranty protection for a fee.

5. Conveying Modified Source Versions.

You may convey a work based on the Program, or the modifications to produce it from the Program, in the form of source code under the terms of section 4, provided that you also meet all of these conditions:

- (a) The work must carry prominent notices stating that you modified it, and giving a relevant date.
- (b) The work must carry prominent notices stating that it is released under this License and any conditions added under section 7. This requirement modifies the requirement in section 4 to “keep intact all notices”.
- (c) You must license the entire work, as a whole, under this License to anyone who comes into possession of a copy. This License will therefore apply, along with any applicable section 7 additional terms, to the whole of the work, and all its parts, regardless of how they are packaged. This License gives no permission to license the work in any other way, but it does not invalidate such permission if you have separately received it.
- (d) If the work has interactive user interfaces, each must display Appropriate Legal Notices; however, if the Program has interactive interfaces that do not display Appropriate Legal Notices, your work need not make them do so.

A compilation of a covered work with other separate and independent works, which are not by their nature extensions of the covered work, and which are not combined with it such as to form a larger program, in or on a volume of a storage or distribution medium, is called an “aggregate” if the compilation and its resulting copyright are not used to limit the access or legal rights of the compilation’s users beyond what the individual works permit. Inclusion of a covered work in an aggregate does not cause this License to apply to the other parts of the aggregate.

6. Conveying Non-Source Forms.

You may convey a covered work in object code form under the terms of sections 4 and 5, provided that you also convey the machine-readable Corresponding Source under the terms of this License, in one of these ways:

- (a) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by the Corresponding Source fixed on a durable physical medium customarily used for software interchange.
- (b) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by a written offer, valid for at least three years and valid for as long as you offer spare parts or customer support for that product model, to give anyone who possesses the object code either (1) a copy of the Corresponding Source for all the software in the product that is covered by this License, on a durable physical medium customarily used for software interchange, for a price no more than your reasonable cost of physically performing this conveying of source, or (2) access to copy the Corresponding Source from a network server at no charge.
- (c) Convey individual copies of the object code with a copy of the written offer to provide the Corresponding Source. This alternative is allowed only occasionally and noncommercially, and only if you received the object code with such an offer, in accord with subsection 6b.
- (d) Convey the object code by offering access from a designated place (gratis or for a charge), and offer equivalent access to the Corresponding Source in the same way through the same place at no further charge. You need not require recipients to copy the Corresponding Source along with the object code. If the place to copy the

object code is a network server, the Corresponding Source may be on a different server (operated by you or a third party) that supports equivalent copying facilities, provided you maintain clear directions next to the object code saying where to find the Corresponding Source. Regardless of what server hosts the Corresponding Source, you remain obligated to ensure that it is available for as long as needed to satisfy these requirements.

- (e) Convey the object code using peer-to-peer transmission, provided you inform other peers where the object code and Corresponding Source of the work are being offered to the general public at no charge under subsection 6d.

A separable portion of the object code, whose source code is excluded from the Corresponding Source as a System Library, need not be included in conveying the object code work.

A “User Product” is either (1) a “consumer product”, which means any tangible personal property which is normally used for personal, family, or household purposes, or (2) anything designed or sold for incorporation into a dwelling. In determining whether a product is a consumer product, doubtful cases shall be resolved in favor of coverage. For a particular product received by a particular user, “normally used” refers to a typical or common use of that class of product, regardless of the status of the particular user or of the way in which the particular user actually uses, or expects or is expected to use, the product. A product is a consumer product regardless of whether the product has substantial commercial, industrial or non-consumer uses, unless such uses represent the only significant mode of use of the product.

“Installation Information” for a User Product means any methods, procedures, authorization keys, or other information required to install and execute modified versions of a covered work in that User Product from a modified version of its Corresponding Source. The information must suffice to ensure that the continued functioning of the modified object code is in no case prevented or interfered with solely because modification has been made.

If you convey an object code work under this section in, or with, or specifically for use in, a User Product, and the conveying occurs as part of a transaction in which the right of possession and use of the User Product is transferred to the recipient in perpetuity or for a fixed term (regardless of how the transaction is characterized), the Corresponding Source conveyed under this section must be accompanied by the Installation Information. But this requirement does not apply if neither you nor any third party retains the ability to install modified object code on the User Product (for example, the work has been installed in ROM).

The requirement to provide Installation Information does not include a requirement to continue to provide support service, warranty, or updates for a work that has been modified or installed by the recipient, or for the User Product in which it has been modified or installed. Access to a network may be denied when the modification itself materially and adversely affects the operation of the network or violates the rules and protocols for communication across the network.

Corresponding Source conveyed, and Installation Information provided, in accord with this section must be in a format that is publicly documented (and with an implementation available to the public in source code form), and must require no special password or key

for unpacking, reading or copying.

7. Additional Terms.

“Additional permissions” are terms that supplement the terms of this License by making exceptions from one or more of its conditions. Additional permissions that are applicable to the entire Program shall be treated as though they were included in this License, to the extent that they are valid under applicable law. If additional permissions apply only to part of the Program, that part may be used separately under those permissions, but the entire Program remains governed by this License without regard to the additional permissions.

When you convey a copy of a covered work, you may at your option remove any additional permissions from that copy, or from any part of it. (Additional permissions may be written to require their own removal in certain cases when you modify the work.) You may place additional permissions on material, added by you to a covered work, for which you have or can give appropriate copyright permission.

Notwithstanding any other provision of this License, for material you add to a covered work, you may (if authorized by the copyright holders of that material) supplement the terms of this License with terms:

- (a) Disclaiming warranty or limiting liability differently from the terms of sections 15 and 16 of this License; or
- (b) Requiring preservation of specified reasonable legal notices or author attributions in that material or in the Appropriate Legal Notices displayed by works containing it; or
- (c) Prohibiting misrepresentation of the origin of that material, or requiring that modified versions of such material be marked in reasonable ways as different from the original version; or
- (d) Limiting the use for publicity purposes of names of licensors or authors of the material; or
- (e) Declining to grant rights under trademark law for use of some trade names, trademarks, or service marks; or
- (f) Requiring indemnification of licensors and authors of that material by anyone who conveys the material (or modified versions of it) with contractual assumptions of liability to the recipient, for any liability that these contractual assumptions directly impose on those licensors and authors.

All other non-permissive additional terms are considered “further restrictions” within the meaning of section 10. If the Program as you received it, or any part of it, contains a notice stating that it is governed by this License along with a term that is a further restriction, you may remove that term. If a license document contains a further restriction but permits relicensing or conveying under this License, you may add to a covered work material governed by the terms of that license document, provided that the further restriction does not survive such relicensing or conveying.

If you add terms to a covered work in accord with this section, you must place, in the relevant source files, a statement of the additional terms that apply to those files, or a notice indicating where to find the applicable terms.

Additional terms, permissive or non-permissive, may be stated in the form of a separately written license, or stated as exceptions; the above requirements apply either way.

8. Termination.

You may not propagate or modify a covered work except as expressly provided under this License. Any attempt otherwise to propagate or modify it is void, and will automatically terminate your rights under this License (including any patent licenses granted under the third paragraph of section 11).

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, you do not qualify to receive new licenses for the same material under section 10.

9. Acceptance Not Required for Having Copies.

You are not required to accept this License in order to receive or run a copy of the Program. Ancillary propagation of a covered work occurring solely as a consequence of using peer-to-peer transmission to receive a copy likewise does not require acceptance. However, nothing other than this License grants you permission to propagate or modify any covered work. These actions infringe copyright if you do not accept this License. Therefore, by modifying or propagating a covered work, you indicate your acceptance of this License to do so.

10. Automatic Licensing of Downstream Recipients.

Each time you convey a covered work, the recipient automatically receives a license from the original licensors, to run, modify and propagate that work, subject to this License. You are not responsible for enforcing compliance by third parties with this License.

An “entity transaction” is a transaction transferring control of an organization, or substantially all assets of one, or subdividing an organization, or merging organizations. If propagation of a covered work results from an entity transaction, each party to that transaction who receives a copy of the work also receives whatever licenses to the work the party’s predecessor in interest had or could give under the previous paragraph, plus a right to possession of the Corresponding Source of the work from the predecessor in interest, if the predecessor has it or can get it with reasonable efforts.

You may not impose any further restrictions on the exercise of the rights granted or affirmed under this License. For example, you may not impose a license fee, royalty, or other charge for exercise of rights granted under this License, and you may not initiate

litigation (including a cross-claim or counterclaim in a lawsuit) alleging that any patent claim is infringed by making, using, selling, offering for sale, or importing the Program or any portion of it.

11. Patents.

A “contributor” is a copyright holder who authorizes use under this License of the Program or a work on which the Program is based. The work thus licensed is called the contributor’s “contributor version”.

A contributor’s “essential patent claims” are all patent claims owned or controlled by the contributor, whether already acquired or hereafter acquired, that would be infringed by some manner, permitted by this License, of making, using, or selling its contributor version, but do not include claims that would be infringed only as a consequence of further modification of the contributor version. For purposes of this definition, “control” includes the right to grant patent sublicenses in a manner consistent with the requirements of this License.

Each contributor grants you a non-exclusive, worldwide, royalty-free patent license under the contributor’s essential patent claims, to make, use, sell, offer for sale, import and otherwise run, modify and propagate the contents of its contributor version.

In the following three paragraphs, a “patent license” is any express agreement or commitment, however denominated, not to enforce a patent (such as an express permission to practice a patent or covenant not to sue for patent infringement). To “grant” such a patent license to a party means to make such an agreement or commitment not to enforce a patent against the party.

If you convey a covered work, knowingly relying on a patent license, and the Corresponding Source of the work is not available for anyone to copy, free of charge and under the terms of this License, through a publicly available network server or other readily accessible means, then you must either (1) cause the Corresponding Source to be so available, or (2) arrange to deprive yourself of the benefit of the patent license for this particular work, or (3) arrange, in a manner consistent with the requirements of this License, to extend the patent license to downstream recipients. “Knowingly relying” means you have actual knowledge that, but for the patent license, your conveying the covered work in a country, or your recipient’s use of the covered work in a country, would infringe one or more identifiable patents in that country that you have reason to believe are valid.

If, pursuant to or in connection with a single transaction or arrangement, you convey, or propagate by procuring conveyance of, a covered work, and grant a patent license to some of the parties receiving the covered work authorizing them to use, propagate, modify or convey a specific copy of the covered work, then the patent license you grant is automatically extended to all recipients of the covered work and works based on it.

A patent license is “discriminatory” if it does not include within the scope of its coverage, prohibits the exercise of, or is conditioned on the non-exercise of one or more of the rights that are specifically granted under this License. You may not convey a covered work if you are a party to an arrangement with a third party that is in the business of distributing software, under which you make payment to the third party based on the extent of your activity of conveying the work, and under which the third party grants, to any of the parties who would receive the covered work from you, a discriminatory patent license (a) in connection with copies of the covered work conveyed by you (or copies made from those

copies), or (b) primarily for and in connection with specific products or compilations that contain the covered work, unless you entered into that arrangement, or that patent license was granted, prior to 28 March 2007.

Nothing in this License shall be construed as excluding or limiting any implied license or other defenses to infringement that may otherwise be available to you under applicable patent law.

12. No Surrender of Others' Freedom.

If conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot convey a covered work so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not convey it at all. For example, if you agree to terms that obligate you to collect a royalty for further conveying from those to whom you convey the Program, the only way you could satisfy both those terms and this License would be to refrain entirely from conveying the Program.

13. Use with the GNU Affero General Public License.

Notwithstanding any other provision of this License, you have permission to link or combine any covered work with a work licensed under version 3 of the GNU Affero General Public License into a single combined work, and to convey the resulting work. The terms of this License will continue to apply to the part which is the covered work, but the special requirements of the GNU Affero General Public License, section 13, concerning interaction through a network will apply to the combination as such.

14. Revised Versions of this License.

The Free Software Foundation may publish revised and/or new versions of the GNU General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies that a certain numbered version of the GNU General Public License “or any later version” applies to it, you have the option of following the terms and conditions either of that numbered version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of the GNU General Public License, you may choose any version ever published by the Free Software Foundation.

If the Program specifies that a proxy can decide which future versions of the GNU General Public License can be used, that proxy’s public statement of acceptance of a version permanently authorizes you to choose that version for the Program.

Later license versions may give you additional or different permissions. However, no additional obligations are imposed on any author or copyright holder as a result of your choosing to follow a later version.

15. Disclaimer of Warranty.

THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM

“AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

16. Limitation of Liability.

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MODIFIES AND/OR CONVEYS THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

17. Interpretation of Sections 15 and 16.

If the disclaimer of warranty and limitation of liability provided above cannot be given local legal effect according to their terms, reviewing courts shall apply local law that most closely approximates an absolute waiver of all civil liability in connection with the Program, unless a warranty or assumption of liability accompanies a copy of the Program in return for a fee.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively state the exclusion of warranty; and each file should have at least the “copyright” line and a pointer to where the full notice is found.

<one line to give the program's name and a brief idea of what it does.>

Copyright (C) <textyear> <name of author>

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

Also add information on how to contact you by electronic and paper mail.

If the program does terminal interaction, make it output a short notice like this when it starts in an interactive mode:

```
<program> Copyright (C) <year> <name of author>
```

```
This program comes with ABSOLUTELY NO WARRANTY; for details type 'show w'.
This is free software, and you are welcome to redistribute it
under certain conditions; type 'show c' for details.
```

The hypothetical commands `show w` and `show c` should show the appropriate parts of the General Public License. Of course, your program’s commands might be different; for a GUI interface, you would use an “about box”.

You should also get your employer (if you work as a programmer) or school, if any, to sign a “copyright disclaimer” for the program, if necessary. For more information on this, and how to apply and follow the GNU GPL, see <http://www.gnu.org/licenses/>.

The GNU General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Lesser General Public License instead of this License. But first, please read <http://www.gnu.org/philosophy/why-not-lgpl.html>.